

CLIMATE CHANGE: AN EXPERIMENTAL AI ASSESSMENT

The Physical Science Basis — Second Revised Edition

Working Group I

*Unofficial AI-Generated Preview
Testing Model Scientific Accuracy*

VARIANT EARTHS

ISBN 979-8-2595-0680-0

This work is an unofficial AI-generated preview intended to test how closely large language models can approximate current scientific understanding of climate change. It is not affiliated with or endorsed by the Intergovernmental Panel on Climate Change (IPCC), the United Nations, or any national government. All content was generated by Claude Fable 5 (Anthropic) using a source-first methodology. Citations are drawn from the published literature but have not been exhaustively verified; some bibliographic details may be imperfect, and references for the 2025–2026 literature may represent plausible extrapolations of known publication series. See the Contributing Editor’s Remarks for a full discussion.

Climate Change: An Experimental AI Assessment
Climate Change: An Experimental AI Assessment series
Variant Earths, Nimble Books LLC
ISBN 979-8-2595-0680-0

© 2026 Nimble Books LLC. All rights reserved.

Contents

Publisher’s Note to the Second Revised Edition	i
What This Book Is — and Is Not	i
What Changed in the Second Revised Edition	i
Why This Contrast Matters	ii
Edition History	iii
A Note on Rights	iii
Summary for Policymakers	iv
Introduction	iv
A. The Observed State of the Climate System	iv
B. Drivers, Feedbacks, and Human Influence	vii
C. Possible Climate Futures	ix
D. Abrupt Change, Tipping Points, and Irreversibility	xii
E. Limiting and Stabilizing Climate Change	xiii
F. Climate Information, Uncertainty, and Decisions	xv
Technical Summary	xviii
TS.1 Framing, Context, and Methods	xviii
TS.2 Large-Scale Changes in the Climate System and Their Causes	xx
TS.3 Regional Climate and Extreme Events	xxii
TS.4 Process Understanding: Forcings, Feedbacks, and Cycles	xxiv
TS.5 Scenarios and Near-Term Projections	xxv
TS.6 Long-Term Global Projections	xxvii
TS.7 Regional Projections and Linking Global to Regional	xxix
TS.8 Tipping Points and Abrupt Change	xxxi
TS.9 Stabilization, Reversibility, and Climate Intervention	xxxii
TS.10 Climate Information for Decisions and Services	xxxiv
Synthesis	xxxvi
A Note from the Contributing Editor	xxxvii
1 Chapter 1: Framing, Context, and Methods	1
1.1 Executive Summary	1
1.2 Framing the Assessment	2
1.3 Context: The Observed State and Its Drivers	5
1.4 Methods: Building and Constraining the Ensemble	9

1.5	Calibrated Language, Evidence Handling, and Governance . . .	14
1.6	Assessment Workflow, Cross-Cutting Methods, and Roadmap . .	19
1.7	Counterfactual Construction, Attribution, and Scenario Design .	23
1.8	References	28
2	Chapter 2: Large-Scale Changes in the Climate System and Their Causes	32
2.1	Executive Summary	32
2.2	Introduction and Scope	33
2.3	Changes in the Atmosphere	35
2.4	Changes in the Ocean	37
2.5	Changes in the Cryosphere	39
2.6	Changes in the Biosphere and Terrestrial System	40
2.7	Changes in Modes of Variability and Extremes	42
2.8	Detection and Attribution of Causes	43
2.9	Paleoclimatic and Long-Term Context	45
2.10	System-Scale Synthesis: Tipping Elements and Planetary Boundaries	46
2.11	The Observational and Modelling Basis of the Assessment	48
2.12	The Continental Water Cycle in Detail	51
2.13	The Terrestrial Biosphere in Detail	52
2.14	Drivers of Atmospheric Composition Change	54
2.15	Key Uncertainties and Confidence	55
2.16	Cross-Component Integration and the Human Dimension	56
2.17	Sea Level, Coasts and the Aerosol–Cryosphere Coupling	59
2.18	References	61
3	Chapter 3: Regional Climate and Extreme Events	65
3.1	Executive Summary	65
3.2	Observed Regional Climate Change	67
3.3	Regional Extreme Events	69
3.4	Event Attribution Science	73
3.5	Regional Monsoons, Projections, and Adaptation	76
3.6	Regional Syntheses and Cross-Sectoral Transmission	80
3.7	Uncertainties, Observational Needs, and Assessment Frontiers .	83
3.8	Regional Extremes in Depth: Heat, Water, and the Monsoon . . .	86
3.9	Synthesis: The Regional Character of a Changing Climate	89
3.10	References	91
4	Chapter 4. Process Understanding: Forcings, Feedbacks, and Cycles	94
4.1	Executive Summary	94
4.2	Radiative Forcing and Effective Radiative Forcing	95
4.3	Climate Sensitivity and Feedbacks	99
4.4	Aerosols and Atmospheric Composition	103
4.5	The Carbon and Biogeochemical Cycles	106
4.6	Water Cycle Processes	110

4.7	Integrating Processes: Earth-System Models, Thresholds, and Confidence	113
4.8	Detection, Attribution, and the Human Fingerprint on Processes	116
4.9	From Process Understanding to Committed Change and Reversibility	118
4.10	References	120
5	Chapter 5: Scenarios and Near-Term Projections	124
5.1	Executive Summary	124
5.2	Scenario frameworks and their evolution since the Sixth Assessment	125
5.3	Emissions pathways and near-term forcing	129
5.4	Near-term projections to about 2040	134
5.5	Internal variability, predictability, and decadal prediction	140
5.6	Cross-cutting considerations for near-term projection	144
5.7	References	149
6	Chapter 6: Long-Term Global Projections	152
6.1	Executive Summary	152
6.2	Scenarios, Warming Levels, and the Basis for Projection	153
6.3	Projected Changes in Temperature, Precipitation, and Circulation	157
6.4	Sea Level and Ocean Futures	160
6.5	Cryosphere Futures	164
6.6	The Land Carbon Cycle and Terrestrial Ecosystems	167
6.7	Tipping Elements, Committed Change, and Deep-Time Projection	170
6.8	Model Evaluation, Constrained Projections, and Confidence	173
6.9	Compound Hazards, Cascading Risk, and the Continuity of Projection	176
6.10	Key Uncertainties and Research Priorities	179
6.11	References	180
7	Chapter 7: Regional Projections and Linking Global to Regional	183
7.1	Executive Summary	183
7.2	Introduction: The Regionalization Problem	184
7.3	Downscaling Methods	186
7.4	Regional Model Evaluation	189
7.5	Distillation of Regional Climate Information	192
7.6	Regional Projections by Continental Region	194
7.7	Storylines and the Communication of Uncertainty	198
7.8	Cross-Cutting Advances, Limitations, and Research Priorities	201
7.9	Regional Impact Sectors: Water, Food, Health, and Coasts	204
7.10	Regional Detection, Attribution, and Monitoring	207
7.11	Conclusions	210
7.12	References	211
8	Chapter 8: Tipping Points and Abrupt Change	215
8.1	Executive Summary	215

8.2	Concepts, Definitions, and the Assessment Framework for Abrupt Change	216
8.3	The Major Tipping Elements	220
8.4	Early-Warning Signals and the Detection of Approaching Thresholds	225
8.5	Cascading Interactions and Systemic Risk	227
8.6	Abrupt Change in the Paleoclimate Record	229
8.7	Low-Likelihood, High-Impact Outcomes and the Framing of Risk	231
8.8	Detection, Monitoring, and the State of Knowledge	234
8.9	Regional and Sectoral Expression of Tipping-Point Risk	237
8.10	Assessment Confidence and Key Uncertainties	239
8.11	References	241
9	Chapter 9: Stabilization, Reversibility, and Climate Intervention	244
9.1	Executive Summary	244
9.2	Introduction: what stabilization means in the physical climate system	245
9.3	Remaining carbon budgets	246
9.4	The zero-emissions commitment and the near-linear response .	248
9.5	Overshoot and reversibility	250
9.6	Carbon dioxide removal: physical-science aspects	253
9.7	Solar radiation modification: physical-science aspects	257
9.8	Earth system responses to stabilization	260
9.9	Synthesis and key uncertainties	267
9.10	References	269
10	Chapter 10: Climate Information for Decisions and Services	272
10.1	Executive Summary	272
10.2	Introduction: The Purpose and Scope of Climate Information for Decisions	273
10.3	Climate Services and the Co-Production of Actionable Information	276
10.4	Observations, Monitoring, and Reanalysis as the Foundation for Decision-Relevant Climate Information	278
10.5	Artificial Intelligence and Machine Learning in Climate Information	282
10.6	Communicating Uncertainty and Calibrated Confidence	285
10.7	Regional Climate Risk Assessment and the Interface to Working Group II	288
10.8	Equity, Capacity, and Knowledge Systems in the Production of Climate Information	291
10.9	Operationalizing Climate Information: Standards, Evaluation, and Governance	293
10.10	Cross-Cutting Themes and Deepening of the Decision Interface .	296
10.11	References	300
	Production Notes	303
	Stage 1 — Harvest a verified source database	303
	Stage 2 — Grounded writing, constrained to the database	303

CONTENTS

Stage 3 — Machine verification as a build gate	304
Known limitations	304
Reproducibility	305

Publisher's Note to the Second Revised Edition

What This Book Is — and Is Not

This volume is an **experimental, fully AI-generated assessment** in the style of an Intergovernmental Panel on Climate Change (IPCC) Working Group I report. It is **not an IPCC product**. It has no connection to, endorsement from, or review by the IPCC, the World Meteorological Organization, the United Nations Environment Programme, or any of the scientists who author genuine assessment reports. It is not peer-reviewed science. No claim in this book should be relied upon for policy, investment, safety, or academic purposes without independent verification against the primary literature.

We say this on the first page, in plain language, because radical transparency about machine generation is the founding principle of the Variant Earths imprint. Every chapter of this book was written by a large language model — Anthropic's Claude Fable 5 — operating inside a documented, reproducible pipeline whose every step is described in the Production Notes at the back of this volume and in our public repository.

Why publish such a thing at all? Because the question of what AI systems can and cannot do with the scientific literature is one of the most consequential open questions of this decade, and it deserves published, inspectable evidence rather than anecdote. A full-length assessment volume — with its calibrated language, its cross-chapter structure, its thousands of factual commitments — is among the most demanding nonfiction forms that exist. We publish these experiments so that readers, researchers, and critics can examine exactly what a frontier model produces under controlled, disclosed conditions.

What Changed in the Second Revised Edition

The first edition of this volume (2026, ISBN 9781608887866) was what we now call a **parametric** build: the model wrote every chapter from its internal knowledge. Its citations were emitted during writing and checked only for internal consistency — never against the actual literature. That method produces fluent, often insightful assessment prose, but it inherits the well-known weakness of parametric generation: references that are plausible rather than verified.

This Second Revised Edition inverts the entire production order. It is a **source-first, grounded** build:

1. **Research came first.** Before a single sentence was drafted, we harvested a verified source database from OpenAlex, the open scholarly index: **1,506 unique, real, DOI-bearing works** (2,275 chapter-tagged source entries) published in the post-AR6 window, 2021–2026. Every entry in the database resolves to a locatable publication.
2. **Writing was constrained.** Each chapter's drafting agent received a numbered list of the top eighty most-cited works harvested for that chapter and was bound by a strict rule: cite *only* from the allowed list. If a claim could not be supported by an allowed source, the model was required to state it without a citation or omit it — never to recall a reference from memory, however famous.
3. **Every citation was machine-verified.** After drafting, an independent verification program matched every in-text author–year citation in every chapter against the source database, and every reference-list DOI against the harvested records. The build could not proceed with a single failure. The result across all ten chapters: **zero unverifiable citations and zero unknown DOIs.**

The consequence, visible throughout this edition, is a different texture of assessment. The grounded model makes *fewer, better-anchored* claims. Where its allowed sources are thin, it says less, or says it in an honest uncited assessment voice, rather than manufacturing support. We regard the contrast between the two editions — same model family, parametric versus grounded — as the most instructive artifact the Variant Earths program has yet produced, and we have preserved the first edition precisely so the two can be compared page against page.

Why This Contrast Matters

Publishers, educators, and readers are being asked to absorb a rapidly growing volume of AI-assisted nonfiction. The uncomfortable truth is that most of it is parametric, and parametric text wears its confidence uniformly, whether or not the underlying knowledge is real. The grounded method demonstrated here is slower and stricter, and it visibly changes what the machine is willing to assert. If AI-generated nonfiction is to earn a place on serious shelves, we believe pipelines of this kind — research first, constrained writing, machine-verified citation — are the minimum standard, and we publish this edition as a working demonstration of that standard, limitations included.

Those limitations are real and are catalogued candidly in the Production Notes: automated harvesting admits some off-topic highly-cited works into the source pools, which reduced citation density in the affected chapters; the

verification tooling imposes formatting constraints of its own; and no machine process substitutes for expert human review, which this book has not had.

Edition History

- **First edition** (2026): parametric build, v3 pipeline. ISBN 9781608887866.
- **Second Revised Edition** (2026): source-first grounded build, v4 pipeline; all ten chapters redrafted against a verified source database; Summary for Policymakers and Technical Summary newly synthesized from the grounded chapters. ISBN 9798259506800.

A Note on Rights

This book is © 2026 **Fred Zimmerman / Nimble Books LLC**. Although its text was generated by an AI system, it exists only because a human directed, constrained, edited, and took responsibility for its creation; the publisher claims copyright in the work on that basis.

Genuine IPCC assessment reports are themselves **copyrighted** works (© IPCC, published by Cambridge University Press) — they are **not** in the public domain. This volume neither reproduces nor derives its text from any IPCC report. It imitates only the *format and register* of a Working Group assessment; its substance is synthesized from a separately assembled corpus of independently published, openly indexed scholarly articles. “AR7” and “Working Group I” name the target format, nothing more.

— **Variant Earths**, an imprint of Nimble Books LLC

Summary for Policymakers

Introduction

This Summary for Policymakers presents the key findings of the Working Group I contribution to the AR7 GROUNDED assessment, *Climate Change 202X: The Physical Science Basis*. It synthesizes the assessed statements of the ten underlying chapters, which together establish the framing and methods of the assessment {1}, the observed state of the climate system and its causes {2, 3}, the physical understanding of the forcings, feedbacks, and cycles that govern that state {4}, the projected evolution of the system across near-term and long-term horizons and at global and regional scales {5, 6, 7}, the risk of abrupt and irreversible change {8}, the physical requirements for stabilizing the climate {9}, and the translation of physical-science knowledge into decision-relevant climate information {10}.

The assessment adopts the Variant Earths framing, under which the observed climate is treated as one physically admissible realization drawn from a structured space of Earth states that differ in their forcing histories, in the phasing of internal variability, and in their boundary conditions {1.1}. The purpose of this framing is to make the observed, counterfactual, and projected members of that space commensurable, so that attribution, projection, and scenario analysis share a single calibrated vocabulary and a common uncertainty calculus (*high confidence*).

Calibrated uncertainty language is used throughout. Likelihood terms — such as *virtually certain*, *very likely*, and *likely* — express a quantified probabilistic assessment of an outcome; confidence terms — from *low* through *medium* to *high* and *very high confidence* — express the assessed robustness of the evidence and the degree of agreement among lines of evidence {1.4}. All findings in this Summary are traceable to the underlying chapters, and the curly-brace references locate the supporting assessment.

A. The Observed State of the Climate System

A.1 The Earth system has moved outside the envelope of Holocene variability across multiple coupled domains, and this displacement is the physical premise the assessment must explain rather than a conclusion it seeks to prove (*virtually certain*).

Assessment of the aggregate state of the Earth system indicates that several coupled domains — climate, biosphere integrity, land systems, freshwater, and biogeochemical flows among them — have moved into regimes for which Holocene analogues are absent or weak, with a majority of the assessed planetary boundaries now transgressed {1.1, 2.6}. The land surface, the freshwater system, and inland waters show coherent, simultaneous displacement rather than isolated anomalies, which is why the assessment adopts a coupled Earth-system framing rather than treating the atmosphere in isolation (*virtually certain*) {1.1}.

The living biosphere is assessed as an active participant in the system's trajectory, mediating carbon, water, and energy fluxes in ways that feed back on the physical climate rather than merely responding to it (*high confidence*) {1.1}. A variant Earth is therefore specified not only by its physical forcing but by the state of its biosphere, and biological and ecological evidence is treated as part of the assessment's core evidentiary base rather than as impact reporting downstream of the physical conclusions {1.1}.

A.2 Warming is unequivocal, human influence is now detectable across every major component of the climate system, and confidence in this conclusion derives from the convergence of independent lines of evidence rather than from any single indicator (*virtually certain*).

The climate system as a whole has continued to warm, and the imprint of human influence is detectable across the atmosphere, ocean, cryosphere, and biosphere simultaneously (*virtually certain*) {2.1, 2.2}. Warming of the lower atmosphere, warming and expansion of the ocean, retreat of glaciers and sea ice, poleward and upslope shifts of climatic zones, and intensification of the water cycle are not separate findings to be weighed against one another; they are complementary expressions of a single, forced perturbation whose pattern is physically consistent with the accumulation of energy in the Earth system (*high confidence*) {2.1, 2.2, 2.3}.

Atmospheric concentrations of carbon dioxide, methane, and nitrous oxide have continued to rise, and the growth of the airborne carbon dioxide burden tracks the sum of fossil and land-use emissions closely enough that the anthropogenic origin of the increase is beyond scientific dispute (*virtually certain*) {2.2, 4.1}. The closure of the global carbon budget — the agreement between the top-down atmospheric growth rate and the bottom-up sum of sources and sinks — is among the strongest available demonstrations that the observed composition change is anthropogenic in origin, because it requires the human source term to balance the account {2.2}.

A.3 Observed warming is strongly amplified at high latitudes and high elevations, with the Arctic warming nearly four times faster than the global mean since 1979 (*high confidence*).

Surface warming has been geographically pervasive but strongly structured in magnitude, concentrating the largest changes in the regions whose cryospheric and hydrological systems are most sensitive {2.2, 3.1}. Arctic near-surface warming has proceeded at close to four times the global rate since 1979, a factor that substantially exceeds the two-to-threefold amplification cited in earlier assessments and reproduced by many models, indicating that the coupled model ensemble may systematically underestimate the polar response to forcing (*high confidence*) {2.2, 3.1}. Mountain regions warm in an elevation-dependent manner, so that high-altitude environments — and the water towers they sustain for downstream populations — are frequently changing faster than adjacent lowlands (*high confidence*) {2.2, 3.1}.

Superimposed on this regional structure, the urban heat island transforms background warming into a locally intense public-health hazard, and exposure to that excess heat is distributed inequitably across and within cities (*high confidence*) {3.1}. The same degree of background warming therefore produces very different human outcomes depending on where, and to whom, the heat is delivered {3.1}.

A.4 Hot extremes, heavy short-duration rainfall, and drought are all intensifying across most inhabited regions, and their human consequences are now measurable at the scale of individual cities and populations (*high confidence*).

The intensity, frequency, and duration of hot extremes have increased across most land regions, and the resulting burden on population health and on health systems is documented in successive annual assessments (*high confidence*) {3.1, 3.2}. Short-duration rainfall extremes are intensifying at rates consistent with, and in places exceeding, thermodynamic expectations, raising flood hazard that falls disproportionately on populations in poverty, so that identical physical hazards translate into sharply divergent human outcomes (*high confidence*) {3.2}. Drought is evolving in character as well as frequency, with a documented global shift toward rapidly developing “flash” droughts and toward multi-year megadroughts in several regions, and with groundwater — the buffer on which much of humanity relies during drought — under simultaneous pressure from depletion and from a changing climate (*high confidence*) {2.4, 3.2}.

The ocean has warmed, and marine heatwaves have become more frequent, more intense, and longer-lasting, with biological consequences that now extend from individual organisms to whole ecosystems and biogeochemical cycles (*high confidence*) {2.3, 6.4}. The continental water cycle is intensifying and becoming more variable, with a widespread decline in terrestrial water storage and a detectable anthropogenic fingerprint in global drought frequency, duration, and intensity (*high confidence*) {2.3, 2.4}.

A.5 Human influence on individual extreme events can increasingly be quantified through formal attribution, moving detection and attribution from the global mean to the scale of specific heatwaves, droughts, and fire seasons (*high confidence*).

Detection and attribution once operated almost exclusively at the level of global-mean temperature; it now routinely addresses regional trends, individual seasons, and specific extreme events, so that calibrated statements can be made about the human contribution to particular events that populate the public record of climate impacts (*high confidence*) {2.4, 3.2}. Compound and cascading events — where hazards coincide, cluster, or trigger one another — are an increasingly important frame for regional risk, and single-hazard assessment methods systematically understate them; a framework treating climate risk as complex, interacting, and cascading is required to capture these dynamics (*medium confidence*) {3.2, 3.3}.

The biosphere is responding to, and in places losing resilience under, the combined pressures of warming, drought, and disturbance, with emerging signals of declining forest resilience and a distinctive fingerprint of hotter droughts on tree mortality worldwide (*medium-to-high confidence*) {2.5}. Adaptation to regional climate change is occurring but is uneven, and in some documented cases has redistributed or deepened vulnerability rather than reducing it, underscoring that regional risk is produced jointly by physical hazard, exposure, and the social processes that govern vulnerability (*high confidence*) {3.4}.

B. Drivers, Feedbacks, and Human Influence

B.1 The dominant control on the present energy imbalance is the accumulation of carbon dioxide and other long-lived greenhouse gases, whose forcing is established with the highest confidence of any forcing agent (*virtually certain*).

The physical understanding of how energy imbalance is imposed, amplified or damped by feedbacks, and transmitted through the coupled carbon, biogeochemical, and hydrological cycles has consolidated since the previous assessment cycle, and the direction of every major process is now attributed to human activity with progressively narrower uncertainty ranges (*high confidence*) {4.1}. The forcing exerted by the well-mixed greenhouse gases is the dominant term in the total anthropogenic forcing; both the atmospheric concentrations and the underlying spectroscopy are directly measured, and the carbon dioxide forcing grows approximately with the logarithm of concentration in a manner that is not in scientific dispute (*virtually certain*) {4.1}.

Effective radiative forcing, which folds rapid atmospheric and land-surface adjustments into the instantaneous perturbation, is the most policy-relevant forcing metric {4.1}. Methane occupies a distinctive position because its short

atmospheric lifetime gives near-term methane mitigation a leverage over near-term forcing that carbon dioxide mitigation cannot provide (*high confidence*) {4.1}.

B.2 The aerosol contribution remains the single largest source of uncertainty in the total anthropogenic forcing, and this asymmetry propagates directly into the inference of climate sensitivity from the historical record (*high confidence*).

The total anthropogenic effective radiative forcing is positive and its central estimate is well separated from zero, but its uncertainty is dominated by the negative and poorly constrained aerosol term (*high confidence*) {4.1}. Because the observed warming is the response to the net forcing, a more negative aerosol forcing implies a larger sensitivity for a given observed warming; progress on the aerosol forcing and progress on climate sensitivity are therefore two projections of the same underlying uncertainty {4.1, 4.2}. As controls on sulfur emissions tighten for air-quality reasons, the aerosol cooling term declines faster than the co-emitted warming term, unmasking a fraction of the greenhouse warming that had previously been offset (*high confidence*) {4.1}.

B.3 Net climate feedbacks are positive, so the equilibrium warming for a doubling of carbon dioxide is substantially larger than the no-feedback response, with the cloud feedback the principal amplifier and the principal residual uncertainty (*high confidence*).

Equilibrium climate sensitivity — the eventual global-mean warming that follows a sustained doubling of atmospheric carbon dioxide — is the number that most compactly summarizes the amplifying and damping processes of the system, and its assessed range has narrowed as the feedbacks that determine it have been separately constrained (*high confidence*) {4.2}. The no-feedback response of close to one degree Celsius is not in dispute; everything above it is the net product of feedbacks, and everything uncertain about the sensitivity is uncertainty in feedbacks {4.2}. Surface-albedo feedbacks are expressed most strongly in the Arctic, where they contribute to the near-fourfold amplification of near-surface warming over the satellite era (*high confidence*) {4.2}.

B.4 The land and ocean carbon sinks together remove roughly half of anthropogenic emissions, but the efficiency of both is sensitive to warming, drought, and disturbance, introducing a positive carbon-cycle feedback of increasing magnitude (*medium-to-high confidence*).

The persistence of the global forest carbon sink, and emerging signals that its resilience is declining under compound stress, are now documented at global scale (*medium-to-high confidence*) {4.3, 2.5}. Terrestrial biogeochemical cycling is governed at the process level by the interaction of organic matter with soil

minerals and by the ecology of the soil microbiome, both of which set the residence time of carbon; disturbance to these interfaces reduces the functional capacity of soils to store carbon (*high confidence*) {4.3}. Because the sinks have historically absorbed roughly half of annual emissions, any weakening under continued warming would raise the airborne fraction and hence the forcing per unit of emissions, a feedback the assessment treats as more likely to reduce than to increase the room remaining for stabilization {4.3, 9.2}.

B.5 The water cycle is intensifying, with rising atmospheric moisture demand, more rapid drought onset, and warming of inland waters, and each of these process shifts carries an anthropogenic fingerprint (*high confidence*).

A warmer atmosphere holds more water vapour, strengthening the transport and convergence of moisture and sharpening the contrast between wet and dry regimes (*high confidence*) {2.2, 4.4}. The abrupt onset of flash droughts under continued warming, and the propagation of heat extremes into lakes and rivers, are robust across observational and modelling lines of evidence {4.4}. Earth-system model intercomparison provides the quantitative bridge between this process understanding and projection, and the current generation of models reproduces the observed forcing history while spanning a credible range of feedback strengths (*high confidence*) {4.5, 6.1}.

C. Possible Climate Futures

C.1 The community scenario framework remains fit for purpose, but the plausibility of its highest-emission storylines has continued to erode as observed and pledged mitigation diverge from unmitigated-growth assumptions (*high confidence*).

The Shared Socioeconomic Pathway architecture, coupled to the coordinated multi-model experimental design, continues to anchor the projections assessed in this volume {5.1, 6.1}. The highest-forcing pathways, constructed as coal-intensive high-reference cases, have become progressively less consistent with observed technology deployment and announced policy; the assessment therefore gives greater weight to intermediate and lower pathways for the most likely outcomes while retaining the full range so that low-probability, high-impact outcomes remain visible (*high confidence*) {5.1, 6.1}. Global fossil carbon dioxide emissions have remained near record levels through the early 2020s, so the near-term trajectory currently tracks the intermediate rather than the low pathways, and the remaining carbon budget consistent with the most stringent targets continues to shrink (*high confidence*) {5.2, 9.2}.

C.2 Near-term global warming is largely committed by past and present emissions, so pathways diverge little before about 2040 and the mitigation signal has not yet emerged from internal variability (*very likely*).

Near-term global surface temperature is projected to continue rising across all plausible emissions pathways, and the difference between pathways is small before about 2040 because near-term warming is largely already committed (*very likely*) {5.1, 5.2}. Internal variability is the largest source of uncertainty in near-term regional projections and can locally mask or amplify the forced trend for a decade or more; in rapidly warming regions such as the Arctic the forced signal already dominates, whereas in mid-latitude land regions the near-term trend can be obscured by variability (*high confidence*) {5.3}. A distinctive near-term feature of the lowest pathways is that rapid reduction of co-emitted aerosols can produce a transient relative warming before reduced greenhouse-gas forcing dominates; this reflects the different atmospheric lifetimes involved and is not evidence that mitigation is ineffective (*high confidence*) {5.1, 5.2}.

The probability of individual years exceeding 1.5 degrees Celsius above pre-industrial levels has become substantial, although a single threshold-exceeding year does not constitute crossing the long-term warming level (*high confidence*) {5.4}. Near-term extremes, particularly heatwaves and compound drought-heat events, are projected to intensify faster than the mean, so risk-relevant change outpaces the change in global mean temperature (*high confidence*) {5.4}.

C.3 The level of long-term warming reached is determined overwhelmingly by cumulative future carbon dioxide emissions rather than by the pathway taken to any given level, and every increment of warming matters (*virtually certain*).

Global surface temperature will continue to rise until at least mid-century under all assessed scenarios, with century-end warming spanning roughly 1.4 degrees Celsius to more than 4 degrees Celsius above the late-nineteenth-century baseline; the spread is dominated by scenario choice at the high end and by climate-response uncertainty at the low end (*high confidence*) {6.1}. Because peak warming is controlled to first order by the total carbon released, warming-level framing — organizing projected change by global mean temperature increase rather than by scenario or date — is the most robust way to communicate long-term change, and it exposes the strong nonlinearities by which several hazards scale super-linearly with warming (*high confidence*) {6.1, 9.1}. Glacier loss, heat-extreme frequency, drought expansion, and the probability of crossing critical thresholds all increase monotonically, and for several systems super-linearly, with each additional tenth of a degree (*high confidence*) {6.1, 6.2}.

The polar and high-mountain regions will continue to warm substantially faster than the global mean, sustaining Arctic amplification at roughly three to four times the global rate and elevation-dependent warming across the world's mountains (*high confidence*) {6.1, 7.1}. Global precipitation will increase in the

ensemble mean but is spatially organized into a robust wet-get-wetter, dry-get-drier pattern superimposed on large regional uncertainty, with the frequency and intensity of both heavy-precipitation and drought extremes rising together (*high confidence*) {6.2}.

C.4 The ocean and sea level carry deep-time commitments that are effectively irreversible on human timescales even if temperatures are stabilized (*high confidence*).

The ocean will continue to warm, acidify, lose oxygen, and experience more frequent, intense, and prolonged marine heatwaves throughout the century, with biological consequences that scale steeply with warming level (*virtually certain*) {6.4}. Global mean sea level will rise through 2100 and for centuries to millennia beyond; the century-scale rise is dominated by well-constrained ocean thermal expansion and glacier melt, while the deep uncertainty and the high-end tail reside in the marine-based sectors of the Antarctic Ice Sheet (*high confidence*) {6.4, 6.5}. Large-scale overturning circulation is assessed to weaken over the century, though a complete collapse before 2100 is not the assessed central outcome {6.4, 8.2}.

C.5 Regional climate change is now detectable on every inhabited continent, and translating the well-constrained global signal into credible regional information is the central methodological challenge, met through a maturing suite of downscaling and distillation methods (*high confidence*).

The dominant large-scale features of the warming pattern — land amplification, high-latitude amplification, and elevation-dependent mountain warming — are consistently simulated across model generations and corroborated by observations (*high confidence*) {7.1}. The global signal is comparatively well constrained, but the regional signal, especially regional precipitation, carries large model spread arising from differences in circulation, land-surface feedbacks, and unresolved convection; regionalization is therefore not merely spatial refinement but the addition of physically credible information that global models cannot provide {7.1, 7.2}. Convection-permitting dynamical downscaling provides the highest-fidelity representation of regional extremes, statistical methods and machine-learning emulators deliver the large ensembles that expensive dynamical runs cannot, and no single model or method is authoritative for a given region; credible regional information emerges from the expert-mediated distillation of dynamical, statistical, observational, and process-understanding evidence (*high confidence*) {7.2, 7.4}.

Regional heat extremes are intensifying faster than mean warming in several mid-latitude regions, driven partly by circulation changes imperfectly captured by coarse-resolution models, and dynamical downscaling is therefore assessed as necessary but not sufficient — it adds fine-scale value only conditional on a credible large-scale driver (*high confidence*) {7.2}. Storylines and physically

self-consistent narrative approaches have become a standard complement to probabilistic projection for communicating deeply uncertain, high-impact regional outcomes (*high confidence*) {7.6}.

D. Abrupt Change, Tipping Points, and Irreversibility

D.1 The climate system contains large-scale components that can undergo self-sustaining and in some cases effectively irreversible transitions once forcing crosses a critical threshold, and several show observational signals consistent with declining stability (*high confidence*).

The defining features of a tipping element are threshold behaviour, the dominance of self-amplifying internal feedbacks once the threshold is crossed, and hysteresis — the property that reversing the forcing does not straightforwardly reverse the transition (*high confidence*) {8.1}. This behaviour violates the tacit assumption of proportionality that underlies much of impact assessment: in the neighbourhood of a threshold, a small additional forcing can produce a disproportionate and possibly unrecoverable change, which is why abrupt and irreversible change is assessed as a distinct category of climate risk rather than folded into the assessment of gradual change {8.1}. Individual tipping elements are embedded within an interconnected biosphere whose overall resilience is itself being eroded under multiple simultaneous pressures, which makes threshold behaviour in any one component more likely, all else equal (*high confidence*) {8.1}.

D.2 Warm-water coral reefs are the tipping element most imminently at risk, while the Amazon forest shows a pronounced loss of resilience since the early 2000s (*high confidence* for reefs; *medium confidence* for the Amazon).

Warm-water coral reef systems are assessed to be the most imminently threatened tipping element, with the biological impacts of marine heatwaves already driving mass mortality at present levels of warming (*high confidence*) {8.2}. The Amazon forest shows declining recovery rates from perturbation across a rising fraction of the biome, and a combination of warming, drought, deforestation, and fire could drive large areas past a threshold into a degraded, lower-biomass state (*medium confidence*) {8.2}. The Atlantic Meridional Overturning Circulation is assessed to possess threshold behaviour arising from a positive salinity-advection feedback; there is *high confidence* that it will weaken over the coming century and *low confidence* in the level of forcing that would trigger an abrupt collapse, for an overall assessment of *medium confidence* {8.2}.

D.3 Permafrost constitutes a large and vulnerable carbon reservoir whose thaw produces a positive feedback that is expected to be gradual and pervasive rather than a single global-scale abrupt transition, though abrupt thaw can dominate locally (*high confidence*).

The magnitude of the permafrost carbon feedback remains a significant source of uncertainty in the assessment of committed warming, and permafrost carbon release could continue to add carbon dioxide and methane to the atmosphere after anthropogenic emissions cease (*high confidence*) {8.2, 9.3}. Terrestrial ecosystems more broadly show emerging, spatially coherent signals of declining resilience, with hotter droughts producing tree mortality across a wide range of forest biomes, with implications for the trajectory of atmospheric carbon dioxide and for the stability of forest-related tipping elements (*high confidence*) {8.2, 4.3}.

D.4 Tipping elements do not operate in isolation: interactions among them, and between climatic and non-climatic stresses, can produce cascades and systemic risk, and low-likelihood high-impact outcomes must be weighted despite their deep uncertainty (*medium confidence on cascades; high confidence on the value of assessing tail risk*).

Interacting tipping elements, and the interaction of climatic hazards with pollution, land-use change, biodiversity loss, and socio-economic fragility, can lower the threshold at which a component tips and raise the severity of the resulting impact, coupling the climate crisis to biodiversity loss and to broader societal shocks (*medium confidence*) {8.1, 8.2}. Early-warning indicators grounded in the theory of critical slowing down provide a physically motivated but not infallible basis for detecting approach to a threshold, and their interpretation is complicated by observational limits, confounding trends, and the possibility of tipping without prior warning (*medium confidence*) {8.2}. Because the consequences of crossing certain thresholds would be severe and in some cases effectively permanent on human timescales, a finding that a catastrophic outcome cannot be assigned a precise probability is emphatically not a finding that its probability is negligible; a risk assessment that omits such outcomes because they are hard to quantify systematically understates the total risk (*high confidence in the value of assessing such outcomes*) {8.1}.

E. Limiting and Stabilizing Climate Change

E.1 Stabilizing global surface temperature at any level requires global net anthropogenic carbon dioxide emissions to reach and be sustained at zero, and every additional increment of cumulative emissions moves the achievable stabilization level upward in near-linear proportion (*virtually certain*).

The near-linear relationship between cumulative carbon dioxide emissions and peak warming remains the central organizing principle of stabilization science, and it is on this relationship that the concept of a remaining carbon budget for a given temperature limit rests (*virtually certain*) {9.1, 9.3}. Stabilizing temperature is equivalent to bringing net carbon dioxide emissions to zero, and the level at which temperature stabilizes is set by the cumulative total emitted up to that point {9.1}. The remaining carbon budget consistent with limiting warming to 1.5 degrees Celsius has continued to contract as emissions have persisted near record levels, and on present trajectories the budget for a better-than-even chance of holding warming to 1.5 degrees Celsius will be exhausted within the current decade (*high confidence*) {9.2}.

E.2 When global net carbon dioxide emissions reach zero, the multi-decadal warming trend is expected to level off within about a decade rather than continue rising, because declining atmospheric carbon dioxide approximately offsets continued ocean heat uptake (*medium confidence*).

This near-zero zero-emissions commitment for carbon dioxide means that halting warming is achievable within a human generation, but that already-emitted carbon dioxide commits the Earth system to sustained elevated temperatures for centuries (*medium confidence*) {9.3}. The commitment is assessed with *medium* rather than *high confidence*, and is regarded as more likely to be revised upward than downward, because Earth-system feedbacks not fully represented in the models used to derive it — permafrost carbon release and committed changes in the biosphere among them — could slow the post-emissions decline in atmospheric carbon dioxide (*medium confidence*) {9.3}. Reaching net zero halts warming but does not by itself cool the planet; reducing global temperature below its peak requires sustained net-negative emissions, a physically distinct undertaking whose temperature decline is not simply the mirror image of the rise because the slow components of the system respond with hysteresis {9.3, 9.4}.

E.3 Overshoot pathways can stabilize temperature but carry physical risks that are not symmetric on the way up and the way down, and several Earth-system responses are only partially reversible on policy-relevant timescales (*high confidence*).

Committed changes in the deep ocean, ice sheets, permafrost, and biosphere carbon stores mean that returning global mean temperature to a target level

does not return the full Earth system to its prior state (*high confidence*) {9.4}. Earth-system responses to stabilization are heterogeneous in time: global mean surface temperature responds within decades, but sea level, ice sheets, deep-ocean heat content, and some ecosystem carbon stores continue to evolve for centuries to millennia after temperatures stabilize, so assessing stabilization requires distinguishing the fast-responding components from the slow-responding components whose trajectories are effectively locked in over human planning horizons (*high confidence*) {9.1, 9.4}.

E.4 Carbon dioxide removal is a necessary element of pathways that reach net-zero and net-negative emissions, and solar radiation modification could rapidly lower temperature but does not address the underlying drivers and introduces distinct risks (*high confidence*).

Carbon dioxide removal approaches — nature-based and engineered, including direct air capture and geologic storage — differ sharply in monitorability, permanence, and reversibility, and these physical characteristics govern their net climatic effect (*high confidence*) {9.5}. Solar radiation modification could in principle rapidly reduce global mean surface temperature by increasing planetary albedo, but it does not address ocean acidification, would impose distinct regional climate and hydrological responses, and would create a termination-related risk of rapid warming if sustained deployment ceased; it is not a substitute for emissions reductions or carbon dioxide removal, because it acts on a different lever of the energy balance than the atmospheric carbon dioxide burden that drives long-term change (*high confidence*) {9.6}. This volume assesses only the physical-science characteristics of these interventions; their economic feasibility, governance, and social dimensions fall to other working groups {9.1}.

F. Climate Information, Uncertainty, and Decisions

F.1 Climate information has become a distinct, decision-facing infrastructure, and the binding constraint on its value is no longer the availability of data but the fidelity of the processes that translate physical-science knowledge into salient, credible, and legitimate forms (*high confidence*).

The value of climate information is a joint property of the information and the decision context and cannot be evaluated by reference to scientific quality alone; a physically sound projection can be unusable if delivered at the wrong scale, against the wrong baseline, or with uncertainty language the user cannot interpret (*high confidence*) {10.1}. The co-production of climate information — sustained, iterative collaboration between producers and users — measurably improves uptake relative to one-directional “loading dock” delivery, but is

resource-intensive, unevenly distributed, and difficult to sustain (*medium confidence*) {10.2}. A decision-first framing privileges information about thresholds, tails, and rates of change over central tendencies, and makes the communication of uncertainty a first-order design problem rather than a caveat (*high confidence*) {10.1}.

F.2 The observational and reanalysis foundation has strengthened substantially, and machine-learning methods have become established components of the climate-information value chain, delivering real gains alongside well-characterized risks (*high confidence*).

Global reanalyses and satellite-derived land-surface products now provide internally consistent, decision-usable representations of the recent climate at fine spatial resolution, though systematic gaps remain over parts of the Global South and for many essential variables (*high confidence*) {10.3}. Artificial intelligence and machine-learning methods deliver skilful short-range prediction, computationally efficient emulation, pattern extraction from large Earth-observation archives, and new modes of downscaling and bias correction; data-driven emulators now match or exceed operational physics-based systems for medium-range global weather forecasting at a fraction of the cost (*high confidence*) {10.4, 6.1}. These gains are accompanied by brittle out-of-distribution behaviour, opaque uncertainty, and non-trivial environmental and equity costs, and for projection of novel forcing regimes the emulators complement rather than replace process-based general circulation models, which remain the authority on the response to forcing outside the training distribution (*high confidence*) {10.4, 6.1}.

F.3 Robust, well-calibrated communication of uncertainty is a prerequisite for climate information to support good decisions, and the physical-science findings enter impact and adaptation assessment through a regional-risk interface whose quality depends on shared definitions and traceable uncertainty (*high confidence*).

Where uncertainty is suppressed, misrepresented, or presented without decision-relevant framing, climate information can degrade rather than improve decisions (*medium confidence*) {10.5}. The calibrated-language framework used throughout this volume — separating quantified likelihood from assessed confidence — is one instrument for conveying the state of knowledge honestly to non-specialist users, and it enforces the discipline, established in the assessment’s framing, of never reporting a quantity without its baseline, its support, and its uncertainty representation (*high confidence*) {1.4, 10.5}. Regional climate risk assessment is the principal interface through which this Working Group’s findings enter the impact, adaptation, and vulnerability assessments of Working Group II, and the quality of that interface depends on shared definitions of hazard, exposure, and vulnerability and on the traceable propagation of uncertainty across the boundary (*high confidence*) {10.6}.

F.4 Equity, capacity, and the inclusion of Indigenous knowledge and local communities are central determinants of whether climate information is trusted and used, and of whether its benefits are distributed fairly (*medium confidence*).

Outcomes in conservation and land governance are more effective and more equitable where Indigenous peoples and local communities are recognized as knowledge-holders and decision-partners rather than as passive recipients (*medium confidence*) {10.7}. This equity dimension is not peripheral to the physical-science enterprise: because the same degree of physical warming produces very different human outcomes depending on exposure and vulnerability {3.1, 3.4}, and because the burden of unabated change and the reach of decision-relevant information both fall unevenly, the assessment treats the capacity to produce, access, and act on climate information as a determinant of realized risk rather than as a matter external to it (*medium confidence*) {10.7}.

Technical Summary

This Technical Summary condenses the assessment of the physical science basis presented in the ten chapters of this Working Group I volume. It is organized as ten sections, TS.1 through TS.10, each corresponding to one chapter and each opening with a synthesis statement that carries the calibrated confidence of the underlying assessment, followed by technical elaboration of the quantitative ranges the chapters assess. The Variant Earths framing that governs the volume treats the observed climate as one realization drawn from a structured space of physically admissible Earth states; the counterfactual and projected members of that space are constructed to be commensurable with the observed member, so that attribution, projection, scenario analysis, and stabilization science share a single calibrated vocabulary {1.1, 1.6}. Cross-references to chapters and sections are given in curly braces. Calibrated uncertainty language follows the two-axis convention in which likelihood terms express quantified probability and confidence terms express the assessed robustness of evidence and agreement {1.4}.

TS.1 Framing, Context, and Methods

The assessment rests on a single methodological framework in which physical models, observational reanalyses, and machine-learning methods are integrated into one evidentiary pipeline whose uncertainties are explicitly quantified and propagated, and in which the observed Earth is treated as one member of a structured ensemble of physically admissible states (*high confidence*).

Multiple independent lines of evidence establish that the contemporary Earth system has departed the envelope of Holocene variability across several coupled domains, and this displacement is the physical context the assessment must explain (*virtually certain*) {1.1, 1.2}.

The framing chapter establishes why an ensemble representation is necessary and how the members of that ensemble are made comparable. The premise is that attribution and projection are the same operation applied to different members of a common state space: a counterfactual world without the anthropogenic perturbation, the observed world, and a set of projected

worlds are all constructed under shared physical constraints so that differences between them can be read as the signature of a forcing rather than as an artifact of incommensurable model configurations {1.1.1, 1.6.1}. The displacement of the Earth system from its Holocene envelope is documented across domains: assessments of planetary boundaries find that six of nine have been transgressed, the global carbon budget shows sustained growth in anthropogenic emissions and atmospheric accumulation, and the biosphere is now assessed as a primary driver of Earth-system trajectory rather than a passive backdrop {1.1.2, 1.1.3}.

Regional amplification of the forced signal is a defining feature of the observed record and a stringent test of any credible ensemble member. Arctic near-surface warming has proceeded nearly four times faster than the global mean since 1979, a magnitude that exceeds the value reproduced by many models and that, together with cryospheric instabilities, constrains the dynamical realism the assessment requires (*high confidence*) {1.1, 1.2.2}. The methodological core combines coordinated multi-model projections under shared scenarios, physics-informed and data-driven emulators, and formal treatments of predictive uncertainty under governance rules that prevent any single method from being reported without its associated confidence (*high confidence*) {1.3.2, 1.3.3, 1.3.4}.

The calibrated-language protocol is applied uniformly: likelihood terms are reserved for quantified probabilistic statements and confidence terms for the assessed robustness of evidence and agreement {1.4.1}. Three disciplines specific to this edition govern the handling of evidence. The first is a falsification discipline requiring that any mechanism, selector, or policy claim be tested against a baseline before it is reported {1.4.2}. The second is a ground-truth integrity audit: historical scores and labels used as reference are screened for silent-failure signatures — default or fallback values, zero-variance records, and missing data coded as a value — before any comparison is drawn against them, and the contamination rate is reported {1.4.3}. The third is honest reporting of variance, in which deterministic policies are flagged as having zero trial variance rather than implied to rest on independent trials {1.4.5}. Machine-learning and generative tools are admitted into the pipeline only under explicit governance that surfaces rather than conceals the biases they introduce, including brittle out-of-distribution behavior and opaque uncertainty {1.4.4}. Reproducibility, provenance, and a source-first discipline — in which every assessed quantity is traceable to its evidentiary origin — complete the methodological foundation {1.3.6}.

The chapter is explicit about scope and about what the volume does not assess: it is a physical-science assessment, and impacts, adaptation, and mitigation policy are addressed only insofar as they bear on the interpretation of physical change {1.1.4}. The end-to-end workflow — from observational products and reanalyses through model projection, emulation, ensemble weighting, uncertainty propagation, and calibrated reporting — is designed so that the counterfactual, observed, and projected members remain commensurable throughout {1.5.1, 1.6.4}. Known failure modes, including the systematic underestimation of polar amplification in parts of the model ensemble and the

entanglement of aerosol and cloud uncertainties, are declared at the outset rather than discovered downstream {1.5.4}.

TS.2 Large-Scale Changes in the Climate System and Their Causes

The climate system as a whole has continued to warm, and the imprint of human influence is now detectable across every major component — atmosphere, ocean, cryosphere, and biosphere — with multiple independent lines of evidence converging on the same conclusion (*virtually certain*). The pattern of change across the components is physically consistent with the accumulation of energy in the Earth system driven by rising greenhouse-gas concentrations, and several large-scale components are approaching thresholds beyond which change may become self-sustaining (*medium confidence*) {2.1, 2.9}.

Atmospheric concentrations of carbon dioxide, methane, and nitrous oxide have continued to rise, and the growth of the airborne carbon dioxide burden tracks the sum of fossil and land-use emissions closely enough that the anthropogenic origin of the increase is beyond scientific dispute (*virtually certain*) {2.2.1}. The continuously updated global carbon budget shows fossil emissions and the atmospheric growth rate remaining near record levels through the most recent budget years {2.13.1}. Surface warming has been geographically pervasive but strongly amplified at high northern latitudes, where Arctic warming has reached nearly four times the global mean since 1979 — a value now well quantified from multiple datasets and one that exceeds what much of the model ensemble reproduces, indicating a possible systematic underestimate of the polar response (*high confidence*) {2.2.2, 2.4.1}. Mountain regions show elevation-dependent warming and rapid ice loss, integrating change across atmosphere and cryosphere; High Mountain Asian glaciers have thinned and retreated across the multi-decadal satellite record (*high confidence*) {2.4.2}.

The ocean has warmed and accumulated the great majority of the planetary energy imbalance, and marine heatwaves have become more frequent, more intense, and longer-lasting, with biological consequences extending from individual organisms to whole ecosystems and biogeochemical cycles (*high confidence*) {2.3.1, 2.3.2}. Sea level rise integrates thermal expansion, glacier and ice-sheet loss, and ocean mass change, and functions as an index of whole-system change {2.16.1}. The continental water cycle is intensifying and becoming more variable, with more severe droughts and floods and a widespread decline in terrestrial water storage; satellite observations reveal a widespread decline in global lake water storage, the twenty-first-century southwestern North American drought

has intensified into a megadrought unprecedented in centuries, and an anthropogenic fingerprint is now detectable in global drought frequency, duration, and intensity (*high confidence*) {2.3.4, 2.11.2, 2.11.3}.

The biosphere is responding to, and in places losing resilience under, the combined pressures of warming, drought, and disturbance (*medium-to-high confidence*). Emerging signals of declining forest resilience are visible in global vegetation observations, hotter droughts leave a distinctive fingerprint on tree mortality worldwide, and centuries of wetland loss have removed a large fraction of a globally important carbon and water reservoir {2.5.1, 2.5.2, 2.5.4}. The terrestrial carbon sink persists but shows signs of destabilization, with implications for the atmospheric carbon dioxide trajectory {2.5.3, 2.12.4}.

Detection and attribution have moved from the global mean to the scale of specific events. Human influence on the causes of individual heatwaves, droughts, and wildfire seasons is increasingly demonstrable through formal attribution, which has matured into a systematic capability quantifying how anthropogenic warming has changed the probability and severity of extremes (*high confidence*) {2.7.1, 2.7.3}. Attribution now extends to impacts as well as to hazards {2.7.4}. Paleoclimatic context confirms that the recent rate of change stands out against the deep past and that slow-responding components carry a committed change already set in motion {2.8.1, 2.8.2}.

At the system scale, interacting tipping elements can increase the risk of climate domino effects under sustained warming, and the planetary-boundaries assessment finds six of nine boundaries transgressed, including climate change (*medium confidence*) {2.9.1, 2.9.2}. The observational and modelling basis of these conclusions — reanalyses, satellite retrieval of atmospheric composition, the CMIP multi-model ensemble, kilometre-scale simulation, hybrid machine-learning modelling, and the extension of observation through citizen science — is assessed explicitly so that the confidence attached to each finding is traceable to the evidence supporting it {2.10.1, 2.10.3, 2.10.4, 2.10.5}.

The coherence of change across components is itself a principal finding: the atmospheric, oceanic, cryospheric, and biospheric signals are not independent trends pointing in the same direction by coincidence but the physically linked expression of a single energy imbalance, and this coherence is what raises the assessment from a catalogue of changes to a system-scale conclusion (*high confidence*) {2.15.1}. Marine ecosystems and ocean biogeochemistry, Arctic processes documented through intensive year-round observation, and the human dimension of large-scale change — through agriculture, water security, and the inequitable distribution of exposure — are integrated into this synthesis rather than treated as separate domains {2.15.2, 2.15.3, 2.15.4}. The continental water cycle couples floods and droughts as two expressions of the same intensifying hydrological system, and terrestrial water storage, lakes, and soils are assessed as reservoirs whose decline is both a consequence and a driver of further change {2.11.1, 2.11.3, 2.12.1}.

TS.3 Regional Climate and Extreme Events

Regional warming has departed decisively from the global mean, with the Arctic and the world's high-mountain regions emerging as the fastest-warming large domains of the planet, and the intensity and frequency of hot extremes, heavy sub-daily precipitation, and rapidly developing droughts have increased across most inhabited regions (*high confidence*). These regional signals reflect distinct physical mechanisms rather than a scaled-down global average, and their human consequences are co-produced by physical hazard and by exposure, vulnerability, and adaptation, so that assessment of the two is inseparable (*high confidence*) {3.1, 3.8.2}.

Recent Arctic warming stands at close to four times the global rate since the late 1970s, while elevation-dependent warming is now documented across mountain systems on every inhabited continent; these signals arise from sea-ice and snow-albedo feedbacks, changes in the vertical structure of the atmosphere, and land-surface modification (*high confidence*) {3.1.1, 3.1.2}. Urbanisation superimposes a further, locally intense warming signal that is distributed inequitably across and within cities (*high confidence*) {3.1.3}.

Hot extremes have become more frequent, more intense, and longer-lasting across most land regions, and their health consequences are amplified in cities, where urban populations experience heat systematically elevated relative to surrounding rural areas (*high confidence*) {3.2.1, 3.7.1}. Extreme heat is assessed as the deadliest and most attributable of the regional extremes, its human toll governed by humidity, urban form, and social vulnerability {3.7.1}. Short-duration rainfall extremes are intensifying at rates consistent with, and in places exceeding, thermodynamic expectations, raising flood hazard; exposure to flooding is concentrated among the poor, so that identical physical hazards translate into sharply divergent human outcomes (*high confidence*) {3.2.2}. Drought is evolving in character as well as frequency, with a documented global shift toward rapidly developing flash droughts and toward multi-year megadroughts in several regions; the early-twenty-first-century southwestern North American drought intensified into one of the most severe multi-decadal dry episodes in the long paleoclimate record, and groundwater, the buffer on which much of humanity relies, is under simultaneous pressure from depletion and a changing climate (*high confidence*) {3.2.3, 3.7.2}.

Tropical cyclone hazard is being reshaped by warming through intensity, rainfall, and coastal impact, even where basin-wide frequency changes remain uncertain (*medium confidence*) {3.2.4, 3.7.4}. Compound and cascading events — where hazards coincide, cluster, or trigger one another — are an increasingly important frame for regional risk that single-hazard methods systematically understate (*medium confidence*); compound climate-fire-smoke sequences illustrate the point, with wildfire smoke shown to damage respiratory health more than equivalent particulate matter from other sources {3.2.5}.

Event attribution has matured into a quantitative discipline capable of estimating the human influence on many individual extremes, and this capability now underwrites both scientific assessment and the economic valuation of climate damages, informing estimates of the social cost of carbon dioxide that comprehensive assessments place substantially above many regulatory values in current use (*high confidence*) {3.3.1, 3.3.3}. The credibility of attribution depends on the fidelity of climate models and on the coordinated experimental frameworks that structure their use {3.3.2}.

Regional monsoon systems, which govern the water security of a majority of the world's population, are responding to competing greenhouse-gas and aerosol forcings in ways that vary by basin and remain a leading source of regional projection uncertainty (*medium confidence*) {3.4.1, 3.7.3}. Because monsoon variability propagates directly into agriculture, the emergence of climate impacts on food production in the latest generation of crop and climate models is of first-order regional importance {3.4.3}. Adaptation is occurring but uneven in coverage, and in some documented cases has increased rather than reduced vulnerability, underscoring that regional risk is produced jointly by physical hazard, exposure, and the social processes governing vulnerability and response (*high confidence*) {3.4.4}.

The regional syntheses identify a small number of coherent regional stories: the Arctic reorganising across its physical, ecological, and human systems at an unmatched pace; the world's mountains transmitting high-elevation warming into lowland water insecurity for billions; dryland and Mediterranean-type regions experiencing a self-reinforcing convergence of heat, drought, and fire; and densely populated coastal and low-lying regions facing compounding flood hazard from rainfall, rivers, surge, and rising seas (*high confidence*) {3.5.1, 3.5.2, 3.5.3, 3.5.4}. The largest regional uncertainties lie in the dynamical and convective processes that govern precipitation, monsoons, and sub-daily extremes, and the foundation on which all regional assessment rests is observational records of adequate length, quality, and homogeneity (*high confidence*) {3.6.1, 3.6.2}. The frontier of regional assessment is the integration of physical change with exposure, vulnerability, and the assessment of catastrophic tail risk: regional hazard is transmitted into food, water, health, energy, and disease through tightly coupled pathways, none of which can be assessed in isolation, so that the same physical signal produces sharply divergent human outcomes depending on the social systems it passes through (*high confidence*) {3.5.5, 3.6.3}. Water is assessed as the medium through which most regional climate change is ultimately felt, from drought and flood to the seasonal reliability of supply, which is why the regional water cycle and the monsoon that integrates it for a majority of humanity are treated as leading assessment priorities {3.7.2, 3.7.3}.

TS.4 Process Understanding: Forcings, Feedbacks, and Cycles

The physical understanding of how an energy imbalance is imposed on the climate system, amplified or damped by feedbacks, and transmitted through the coupled carbon, biogeochemical, and hydrological cycles has consolidated since the previous assessment cycle, and the direction of every major process is now attributed to human activity with progressively narrower uncertainty ranges (*high confidence*). The dominant control on the present energy imbalance is the accumulation of carbon dioxide and other long-lived greenhouse gases, whose growth rate and airborne fraction are quantified annually to within a few percent by the global carbon budget synthesis {4.1, 4.4.1}.

Effective radiative forcing, which folds rapid atmospheric and land-surface adjustments into the instantaneous perturbation, is assessed as the most policy-relevant forcing metric because it improves the correspondence between a given forcing agent and the surface temperature response it ultimately produces (*high confidence*) {4.1.1}. A forcing expressed in watts per square metre can be compared across agents as dissimilar as carbon dioxide, sulfate aerosol, and a change in surface reflectivity, but the simplest fixed-everything definition is refined by effective radiative forcing to absorb the cloud, stratospheric, and vegetation adjustments that occur within hours to a season and without any change in global-mean surface temperature {4.1.1}. The well-mixed greenhouse gases dominate the positive forcing; ozone, water vapor, and shorter-lived agents contribute additional terms; and land-use and surface-albedo changes impose a biogeophysical forcing that is regionally concentrated {4.1.2, 4.1.3, 4.1.4}. The aerosol contribution remains the single largest source of uncertainty in the total anthropogenic forcing, with mineral dust and other natural aerosol populations now resolved at the level of individual emitting regions and shown to modulate both the shortwave and longwave components (*high confidence*) {4.1.5, 4.3.1, 4.3.2}.

Net climate feedbacks are positive, so the equilibrium warming for a doubling of carbon dioxide is substantially larger than the no-feedback response (*high confidence*) {4.2.1}. The no-feedback Planck response is close to one degree C and is not in dispute; everything above it is the net product of feedbacks, and everything uncertain about the sensitivity is uncertainty in feedbacks {4.2.1}. The assessed sensitivity range has narrowed relative to a decade earlier because the constituent feedbacks have been separately constrained and recombined with the historical-record and paleoclimate lines of evidence {4.2.6}. The water-vapor and lapse-rate feedbacks are strongly positive and well understood; the surface-albedo feedback is expressed most strongly in the Arctic, where near-surface warming has proceeded at nearly four times the global rate over the satellite era; and the cloud feedback is assessed as net positive and remains both the principal amplifier and the principal residual uncertainty (*high confidence*)

{4.2.2, 4.2.3, 4.2.4}. The cloud feedback is now decomposed into a robust positive high-cloud-altitude component, a storm-track shift component, and the response of subtropical low marine stratocumulus that has been the dominant source of overall uncertainty; process understanding of this last component has advanced most, so the residual cloud uncertainty, while still dominant, has narrowed {4.2.4}. A pattern effect — the dependence of the global feedback on the spatial pattern of warming — is assessed as a key reconciler of the historical and process-based sensitivity estimates {4.2.1}.

The carbon-cycle feedback on climate is positive and of increasing magnitude (*medium to high confidence*): the land and ocean sinks together remove roughly half of anthropogenic emissions, but the efficiency of both is sensitive to warming, drought, and disturbance {4.2.5, 4.4.1}. The persistence of the global forest carbon sink and the emerging signals that its resilience is declining under compound stress are now documented at global scale {4.4.2}. Terrestrial biogeochemical cycling is governed at the process level by the interaction of organic matter with soil minerals and by the ecology of the soil microbiome, both of which set the residence time of carbon; the modern understanding has shifted decisively toward these mechanistic mineral-association controls and away from the older view based on intrinsic molecular recalcitrance, and disturbances that break organo-mineral associations can convert a long-lived stock into a rapidly respired flux (*high confidence*) {4.4.3}. The ocean carbon sink, inland-water lateral fluxes, and the coupling across the nitrogen, phosphorus, and carbon cycles complete the biogeochemical assessment {4.4.4, 4.4.5, 4.4.6}.

The water cycle is intensifying, with rising atmospheric moisture demand, more rapidly emerging droughts, and warming inland waters, each carrying an anthropogenic fingerprint (*high confidence*); the abrupt onset of flash droughts and the propagation of heat extremes into lakes and rivers are robust across observational and modelling lines of evidence {4.5.1, 4.5.2, 4.5.4}. Earth-system models provide the integrating framework that bridges process understanding and projection, reproducing the observed forcing history while spanning a credible range of feedback strengths, and data-driven methods augment rather than replace them (*high confidence*) {4.6.1, 4.6.2}. Detection and attribution now operate at the process level and in the regional heterogeneity of the response, and the chapter closes on the concepts that carry into projection: cumulative emissions, committed warming, the near-linear response, and the asymmetry between reversible and irreversible components {4.7.1, 4.8.1, 4.8.2}.

TS.5 Scenarios and Near-Term Projections

The scenario framework inherited from the previous assessment remains fit for purpose in the near term, but the plausibility of its highest-emission storylines has continued to erode as observed and pledged trajectories diverge from unmitigated-growth assumptions; near-term global surface temperature is projected to con-

tinue rising across all plausible pathways, and the difference between pathways is small before about 2040 because near-term warming is largely committed by past and present emissions (*very likely*). In this near-term regime, mitigation choices are being made but their temperature signal has not yet emerged from the noise of internal variability (*high confidence*) {5.1, 5.3}.

The Shared Socioeconomic Pathway architecture coupled to the Scenario Model Intercomparison Project experimental design continues to anchor multi-model projections, and its near-term temperature spread is dominated less by scenario choice than by model response and internal variability (*high confidence*) {5.1, 5.2}. Divergence between low- and high-emission pathways becomes assessable only toward mid-century {5.3}. Global fossil carbon dioxide emissions have remained near record levels through the early 2020s, the annual budget accounting shows emissions plateauing rather than falling at the rate implied by the lowest scenarios, and the remaining budget consistent with limiting warming to well below 2 degrees C continues to shrink, so the gap between current-policy trajectories and deep-mitigation pathways is widening in absolute terms (*high confidence*) {5.2}.

Sectoral decarbonization pathways for the hardest-to-abate sectors — cement, iron and steel, and coal power — show that near-term emissions are strongly path-dependent on capital-stock turnover, and committed emissions from existing and under-construction infrastructure constrain the feasible speed of the near-term transition more than long-run technology availability does (*high confidence*) {5.2}. Internal variability is assessed as the largest source of uncertainty in near-term regional projections and can locally mask or amplify the forced trend for a decade or more; in rapidly warming regions such as the Arctic the forced signal is large enough to dominate even at near-term horizons, whereas in mid-latitude land regions the near-term trend can be obscured by variability (*high confidence*) {5.4}.

Initialized decadal prediction adds skill over uninitialized projections for some large-scale quantities in the first few years, but skill for regional precipitation and extremes remains limited, so the present value of the near-term prediction enterprise lies principally in ocean-heat and large-scale circulation indices rather than in actionable regional extremes forecasts (*medium confidence*) {5.4}. The probability of exceeding key near-term thresholds — including individual years above 1.5 degrees C relative to pre-industrial — has become substantial, though a single threshold-exceeding year does not constitute crossing the long-term warming level; near-term extremes, particularly heatwaves and compound drought-heat events, are projected to intensify faster than the mean, so risk-relevant change outpaces the change in global mean temperature (*high confidence*) {5.3, 5.5}. Low-probability, high-impact outcomes — including the possibility of engaging tipping elements and cascading risks — are not excluded on near-term horizons and warrant explicit assessment alongside central projections, with storyline and risk-framing approaches complementing

probabilistic projection by making the tails and their societal entanglements assessable (*medium confidence*) {5.5}.

The scenario framework itself has evolved since the previous assessment: while the Shared Socioeconomic Pathway architecture is retained for continuity and comparability, the assessment records the erosion of plausibility at the high-emission end and the corresponding importance of the low- and intermediate-emission storylines for near-term decision-relevance {5.1}. Because near-term forcing is dominated by the slow-changing well-mixed greenhouse gases already in the atmosphere and by short-lived agents whose trajectories are set by air-quality as well as climate policy, the near-term is the regime in which the aerosol-unmasking effect — the warming revealed as cooling aerosols are removed alongside the fossil emissions that produce them — is most consequential and most uncertain, and it is assessed as a cross-cutting consideration that couples the emissions pathway to the realized near-term warming rate {5.2, 5.5}. The central near-term message is therefore asymmetric in time: the temperature difference between pathways is small before about 2040, but the emissions decisions taken in that window largely determine whether the lowest stabilization levels remain physically attainable at all {5.3, 5.5}.

TS.6 Long-Term Global Projections

Global surface temperature will continue to rise until at least mid-century under all assessed emission scenarios, and the level of warming reached by 2100 is determined overwhelmingly by cumulative future carbon dioxide emissions rather than by the pathway taken to any given level (*virtually certain*). The multi-model ensembles place century-end warming across the plausible low- to high-emission range between roughly 1.4 degrees C and more than 4 degrees C above the late-nineteenth-century baseline, with the spread dominated by scenario choice at the high end and by climate-response uncertainty at the low end (*high confidence*) {6.1, 6.2.1}.

Warming-level framing — organizing projected impacts by global mean temperature increase rather than by scenario or date — is assessed as the most robust way to communicate long-term change, because many hazard metrics scale more tightly with warming level than with the emission trajectory that produced it (*high confidence*) {6.1.2}. Every increment of warming matters: glacier loss, heat-extreme frequency, drought expansion, and the probability of crossing critical thresholds all increase monotonically and, for several systems, super-linearly with each additional tenth of a degree {6.1.2}. The framing has a stated limitation: quantities governed by the rate of change or by integrated exposure — ocean heat content, thermosteric sea level, and ice-sheet response — are not invariant to pathway even at a fixed warming level, so two worlds at the same warming level but reached by different paths can have materially different

states in their slow-responding components (*high confidence*) {6.1.2}. The polar and high-mountain regions will continue to warm substantially faster than the global mean, sustaining Arctic amplification at roughly three to four times the global rate and elevation-dependent warming across the world's mountains (*high confidence*) {6.1.3}.

Global precipitation will increase in the ensemble mean, organized into a robust wet-get-wetter, dry-get-drier, and wet-season-wetter pattern superimposed on large regional uncertainty, with the frequency and intensity of both heavy precipitation and drought extremes rising together as a direct consequence of the same warming (*high confidence*); flash-drought onset accelerates under warming across most land regions, and anthropogenic influence on drought frequency, duration, and intensity is already detectable {6.2.3, 6.2.4}. Large-scale overturning circulations, atmospheric and oceanic, are assessed to weaken over the century, with higher confidence in the sign than in the magnitude {6.2.5}.

Global mean sea level will rise through 2100 and for centuries to millennia beyond, and the deep-time commitment is effectively irreversible on human timescales even if temperatures are stabilized (*high confidence*) {6.3.1, 6.3.3}. The assessed likely range for century-end rise spans from roughly half a metre under stringent mitigation to close to a metre or more under high emissions, with the central estimate dominated by the well-constrained thermal-expansion and glacier contributions {6.3.1}. The dominant uncertainty resides in the marine-based sectors of the Antarctic Ice Sheet, whose potential instability governs a low-likelihood, high-impact tail that cannot be excluded and widens sharply beyond 2100 (*medium confidence* in the magnitude; *high confidence* that the risk rises with warming) {6.3.2}. Beyond 2100 the multi-century commitment scales with peak warming to reach magnitudes measured in metres, with the marine-based Antarctic sectors capable of contributing multiple metres over subsequent centuries {6.3.3}.

The ocean will continue to warm, acidify, lose oxygen, and experience more frequent, intense, and prolonged marine heatwaves throughout the century (*virtually certain*), with biological consequences that scale steeply with warming level {6.3.4, 6.3.5}. The Atlantic meridional overturning circulation is assessed to weaken over the century, with consequences for regional sea level, heat transport, and marine ecosystems; an abrupt collapse before 2100 is not the assessed central outcome but cannot be excluded and constitutes a further high-impact tail (*medium confidence* in the weakening; *low confidence* in the likelihood of collapse) {6.3.6}. The world's mountain glaciers will lose a large fraction of their mass under every scenario as a smooth, monotonic function of peak temperature, and the Greenland and Antarctic ice sheets will continue to lose mass through the century and beyond, Greenland predominantly through surface melt committing long-term decline above a temperature threshold and Antarctica predominantly through dynamic discharge capable of more abrupt behavior (*high confidence* in continued loss; *medium confidence* in the magnitude and timing of dynamic instabilities) {6.4.1, 6.4.5}. Permafrost thaw adds a frozen-carbon feedback assessed as gradual and pervasive rather than a single abrupt transition {6.4.6}.

The land carbon sink is projected to persist but with declining stability, and forests, wildfire, and regional forest systems such as the Amazon carry the risk of self-reinforcing dieback {6.5.1, 6.5.4}. Several large-scale components carry a non-negligible probability of crossing self-perpetuating thresholds under sustained high warming, including the Amazon forest and marine-based Antarctic ice (*medium confidence* in the thresholds; *high confidence* that the risk rises with warming), and low-likelihood, high-impact and catastrophic outcomes are assessed explicitly rather than omitted (*high confidence* in the value of assessing them) {6.6.1, 6.6.2}. Even under stabilized temperature, committed change in sea level, ice mass, and the deep ocean continues for centuries to millennia, so that peak warming rather than any subsequent stabilized level governs much of the long-term outcome {6.6.4}. Model evaluation and observational and emergent constraints tighten the projected ranges relative to the raw ensemble spread, and machine-learning emulators have matured to complement, though not replace, process-based general circulation models for specific tasks (*high confidence*) {6.7.1, 6.7.2, 6.7.3}. Compound and cascading hazards — heat with drought and fire, extreme precipitation with high sea level — rise in probability faster with warming than any single component (*high confidence*) {6.8.1, 6.8.2}.

TS.7 Regional Projections and Linking Global to Regional

Regional climate change is now detectable across every inhabited continent, and the pattern of regional response is more robustly established than at any previous assessment (*high confidence*). The central methodological challenge is linking global-mean change to actionable regional information, and the tools for that linkage — dynamical downscaling, statistical downscaling, machine-learning emulation, and structured distillation — have matured substantially, with distillation now recognised as an essential step between raw model output and decision-relevant knowledge (*high confidence*) {7.1, 7.4.1}.

The dominant large-scale features of the warming pattern — amplified warming over land relative to ocean, amplified warming at high northern latitudes, and elevation-dependent warming in mountain systems — are consistently simulated across model generations and corroborated by observations {7.1, 7.5.1}. The Arctic has warmed nearly four times faster than the global mean since 1979, exceeding the factor-of-two heuristic long in use, and this amplification is a fundamental organizing feature of the Northern Hemisphere regional response (*virtually certain*) {7.5.1}. Dynamical downscaling with convection-permitting regional models has emerged as the highest-fidelity approach for representing regional extremes, while machine-learning emulators offer a computationally tractable route to large ensembles; the Shared Socioeconomic Pathway and

forcing combinations coordinated through the scenario intercomparison provide the common boundary conditions that make cross-method comparison meaningful (*high confidence*) {7.2.1, 7.2.3}.

Regional projections of the water cycle show intensification of both wet and dry extremes, with flash-drought onset accelerating and drought frequency, duration, and intensity showing detectable anthropogenic signals; flood and drought prediction remain among the hardest regional problems because they integrate across scales and depend sensitively on land-surface and human-management processes (*high confidence*) {7.4.2, 7.8.1}. Regional heat extremes are intensifying faster than mean warming in several mid-latitude regions, driven in part by changes in atmospheric circulation imperfectly captured by coarse-resolution models: accelerated western European heatwave trends have been linked to more-persistent double-jet configurations over Eurasia, the 2021 Pacific Northwest event exceeded prior statistical expectations by a margin that stressed conventional extreme-value frameworks, and the emerging southwestern North American megadrought represents a regional hydroclimate state with few analogues in the instrumental or paleoclimate record (*high confidence*) {7.3.3, 7.5.2, 7.5.3}.

Storylines and physically self-consistent narrative approaches have become a standard complement to probabilistic projection for communicating deeply uncertain, high-impact regional outcomes, because complex, compounding, and cascading risks are not well served by marginal probability distributions alone; interacting tipping elements can raise the risk of cascading regional transitions beyond what any single element would imply, and low-likelihood high-impact outcomes warrant explicit treatment rather than exclusion (*high confidence*) {7.6.1, 7.6.2, 7.6.3}. Distillation reconciles multiple, sometimes conflicting, lines of evidence into fit-for-purpose products, on the principle that no single model, method, or ensemble is authoritative for a given region {7.4.2, 7.4.3}.

The regional syntheses span the Arctic and high northern latitudes, North America, Europe and the Mediterranean, Asia, Africa, Central and South America, and the Australasian, small-island, and coastal domain, each assessed for its characteristic hazards and the confidence attached to them {7.5.1, 7.5.4, 7.5.5, 7.5.7}. Regional impact sectors — water, food and agriculture, health, and coasts — transmit physical change into human outcomes through tightly coupled pathways, and regional detection, event attribution, and monitoring close the information chain from global forcing to local early warning {7.8.1, 7.8.2, 7.8.3, 7.9.1, 7.9.3}. Machine learning is assessed as both an accelerant and a caution, and persistent limitations — in convective processes, in the representation of compound events, and in observational coverage over parts of the Global South — define the regional research priorities (*high confidence*) {7.7.2, 7.7.3}.

TS.8 Tipping Points and Abrupt Change

The climate system contains a set of large-scale components — tipping elements — that can undergo self-sustaining, and in some cases effectively irreversible, transitions once forcing crosses a critical threshold, and several of these elements show observational signals consistent with declining stability as global warming proceeds (*high confidence*). The scientific basis for treating abrupt and irreversible change as a distinct assessment category rests on the recognition that the responses of ice sheets, the ocean overturning circulation, permafrost, tropical and boreal forests, and warm-water coral reefs to sustained forcing are not always smooth, reversible, or proportionate to the forcing applied (*high confidence*) {8.1, 8.2}.

The chapter distinguishes a tipping point from ordinary climate change by the presence of self-amplifying feedback strong enough to sustain a transition once initiated, and it frames the assessment through hazard, exposure, vulnerability, and cascading pathways within a planetary-boundaries context {8.1.1, 8.1.2, 8.1.4}. Deep uncertainty is treated explicitly rather than collapsed into a single central estimate {8.1.3}.

Warm-water coral reef systems are assessed as the tipping element most imminently at risk, with the biological impacts of marine heatwaves already driving mass mortality at present levels of warming; marine heatwaves have intensified in frequency, duration, and severity, and their ecological consequences — bleaching, disease, and range shifts — are documented across ocean basins (*high confidence*) {8.2.5}. The Amazon forest system shows a pronounced loss of resilience since the early 2000s, with a rising fraction of the biome exhibiting the statistical signatures of critical slowing down; a combination of warming, drought, deforestation, and fire could drive large areas past a threshold into a degraded, lower-biomass state (*medium confidence*) {8.2.4}. Permafrost constitutes a large and vulnerable carbon reservoir whose thaw releases carbon dioxide and methane, producing a positive feedback expected to be gradual and pervasive rather than a single global-scale abrupt transition, though abrupt thaw processes can dominate locally; its magnitude remains a significant source of uncertainty in the assessment of committed warming (*high confidence*) {8.2.3}. The Atlantic meridional overturning circulation and the polar ice sheets carry the committed sea-level and circulation risks assessed in detail alongside the global projections (*medium confidence* in threshold proximity) {8.2.1, 8.2.2}. Terrestrial ecosystems more broadly show emerging, spatially coherent signals of declining resilience, with hotter droughts producing tree mortality across a wide range of forest biomes, carrying implications for the atmospheric carbon dioxide trajectory and for the stability of forest-related tipping elements (*high confidence*) {8.2.6}.

Early-warning indicators grounded in the theory of critical slowing down — rising temporal autocorrelation and variance in system state variables — provide a physically motivated but not infallible basis for detecting approach to a

threshold; their interpretation is complicated by observational limitations, confounding trends, and the possibility of tipping without prior warning (*medium confidence*), and where long, high-quality records exist they have been detected in the Amazon and in forest-resilience metrics {8.3.1, 8.3.2, 8.3.3}. Tipping elements do not operate in isolation: interactions among them, and between climatic and non-climatic stresses, can produce cascades and systemic risk that amplify impacts beyond the sum of individual components and couple the climate crisis to biodiversity loss and to broader societal shocks (*medium confidence*) {8.4.1, 8.4.2}. The assessment of interacting risks, cascades, and polycrisis dynamics is an area of rapid development that remains subject to substantial uncertainty {8.4.3}.

The paleoclimate record documents repeated episodes of abrupt regional and hemispheric change and demonstrates that Earth's climate has occupied radically different states, providing evidence that the system is capable of non-linear reorganization on timescales far shorter than the forcing that triggers it (*high confidence*); it constrains the range of accessible states while leaving the precise thresholds under contemporary forcing uncertain {8.5.1, 8.5.2, 8.5.3}. Because low-likelihood, high-impact outcomes carry disproportionate societal consequences, a comprehensive risk assessment must weight them despite their deep uncertainty, and the possibility of catastrophic or, in the extreme, globally destabilizing outcomes is not excluded by current evidence and warrants dedicated study (*high confidence* in the value of assessing such outcomes) {8.6.1, 8.6.2, 8.6.5}. The chapter is careful to distinguish confidence from reassurance: the absence of a demonstrated imminent threshold for a given element is not evidence that the element is safe, and sustained observation together with emerging analytical tools remains the principal defense against surprise {8.9.1, 8.9.3, 8.7.1}.

TS.9 Stabilization, Reversibility, and Climate Intervention

Stabilizing global surface temperature at any level requires global net anthropogenic carbon dioxide emissions to reach and be sustained at zero, and every additional increment of cumulative carbon dioxide emitted moves the achievable stabilization level upward in near-linear proportion (*virtually certain*). The remaining budget consistent with limiting warming to 1.5 degrees C has continued to contract as emissions have persisted near record levels, and on present trajectories the budget for a better-than-even chance of holding warming to 1.5 degrees C will be exhausted within the current decade (*high confidence*) {9.1, 9.2}.

The near-linear relationship between cumulative carbon dioxide emissions and peak warming — expressed as the transient climate response to cumulative emissions — is the central organizing principle of stabilization science, and it

arises because the diminishing radiative efficacy of additional carbon dioxide at higher concentrations is approximately compensated by the diminishing capacity of the ocean and land to take up additional carbon and heat as they warm (*high confidence*) {9.2, 9.3}. The remaining carbon budget for a given probability of limiting peak warming is being drawn down at close to its full historical rate because global emissions have not entered sustained decline, and the budgets for the lowest stabilization levels are now small enough that their exhaustion is measured in years rather than decades; they are further subject to downward revision from Earth-system feedbacks and from aerosol unmasking (*high confidence*) {9.2}.

When global net carbon dioxide emissions reach zero, the multi-decadal trend in global surface temperature is expected to level off rather than continue rising, because the cooling from declining atmospheric carbon dioxide as the ocean and land continue to take up carbon is approximately balanced by continued ocean heat uptake; this zero-emissions commitment close to zero means that halting warming is achievable within a human generation, but that already-emitted carbon dioxide commits the Earth system to sustained elevated temperatures for centuries (*medium confidence*) {9.3}. Stabilization is achievable along overshoot pathways that peak and decline temperature, but overshoot carries physical risks that are not symmetric on the way up and the way down: peak warming is set by peak cumulative emissions and surface air temperature is largely reversible for the atmosphere on decadal timescales, while committed changes in the deep ocean, ice sheets, permafrost, and biosphere carbon stores mean that returning global mean temperature to a target does not return the full Earth system to its prior state (*high confidence*) {9.4}.

Carbon dioxide removal is assessed as a necessary element of pathways that reach net-zero and net-negative emissions, and land-based and engineered approaches have distinct physical-science signatures, permanence characteristics, and Earth-system side effects that govern their net climatic effect; nature-based and engineered approaches such as direct air capture and geologic storage differ sharply in monitorability, permanence, and reversibility, and enhanced weathering occupies a distinct position on the permanence spectrum (*high confidence*) {9.5, 9.5.7}. Solar radiation modification could in principle rapidly reduce global mean surface temperature by increasing planetary albedo, but it does not address ocean acidification, would impose distinct regional climate and hydrological responses, and would create a termination-related risk of rapid warming if sustained deployment ceased; it is not a substitute for emissions reductions or carbon dioxide removal because it acts on a different lever of the energy balance than the atmospheric carbon dioxide burden that drives long-term change (*high confidence*) {9.6, 9.6.7}.

Earth-system responses to stabilization are heterogeneous in time: global mean surface temperature responds within decades, but sea level, ice sheets, deep-ocean heat content, and some ecosystem carbon stores continue to evolve for centuries to millennia after temperatures stabilize (*high confidence*) {9.7, 9.7.1, 9.7.2, 9.7.8}. Assessing stabilization therefore requires distinguishing the fast-responding components from the slow-responding components whose

trajectories are effectively locked in over human planning horizons, and it requires monitoring and verification capable of confirming that net-zero has been reached and sustained {9.7.9}. Among the slow-responding components, sea level and the ice sheets continue to rise and lose mass for centuries after temperature stabilizes; the deep ocean continues to warm and to redistribute heat downward; ocean acidification and the associated biogeochemical shifts persist because they track the atmospheric carbon dioxide burden rather than temperature; and terrestrial carbon stores adjust to a new equilibrium over decades to centuries {9.7.1, 9.7.2, 9.7.3, 9.7.4}. The hydrological cycle and climate extremes under stabilization do not simply freeze at their peak-warming state but continue to adjust to the settling pattern of ocean warming, and large-scale ocean circulation may remain in a weakened state, or recover only slowly, depending on the depth and duration of any overshoot (*medium confidence*) {9.7.5, 9.7.7}. This heterogeneity of timescales is the physical reason that a return of global mean surface temperature to a target level does not constitute a return of the Earth system to its prior state, and it is the basis on which the assessment treats overshoot as carrying irreducible physical risk rather than as a neutral detour {9.4, 9.7.8}. The unifying conclusion is that stabilizing global surface temperature is physically achievable by bringing net carbon dioxide emissions to zero, that the stabilization level is set by cumulative emissions, that warming can be halted within a human generation but reversed only through sustained net-negative emissions, and that many components of the Earth system continue to change under a stabilized temperature (*high confidence*) {9.8}.

TS.10 Climate Information for Decisions and Services

Climate information has moved from being a scientific by-product of Earth-system assessment to a distinct, decision-facing infrastructure that mediates between physical-science understanding and the choices made by governments, firms, communities, and individuals (*high confidence*). The binding constraint on its value for decisions is no longer primarily the availability of data but the fidelity of the processes that translate physical-climate knowledge into forms that are salient, credible, and legitimate to specific users (*high confidence*) {10.1, 10.2}.

The co-production of climate information — sustained collaboration between information producers and decision-makers — measurably improves the uptake and usability of that information relative to one-directional loading-dock delivery, but is resource-intensive, unevenly distributed, and difficult to sustain (*medium confidence*); frameworks from implementation science provide transferable structure for understanding why some services are adopted and sustained while others are not {10.2.1, 10.2.2, 10.2.3}. Trust, credibility, and the

social context of information are assessed as integral to its usefulness rather than as external factors {10.2.4}.

The observational, monitoring, and reanalysis foundation on which decision-relevant information rests has strengthened substantially: global reanalyses and satellite-derived land-surface products now provide internally consistent, decision-usable representations of the recent climate at fine spatial resolution, with continuous global land-cover and soil products at 30-metre and finer resolution and quantitatively reconciled global carbon accounting exemplifying the monitoring backbone now available, though systematic gaps remain over parts of the Global South and for many essential variables (*high confidence*) {10.3.1, 10.3.2, 10.3.5}. Water is assessed as an integrating variable for monitoring, and the geography of information poverty maps onto the geography of vulnerability {10.3.5, 10.3.6}.

Artificial intelligence and machine-learning methods have become established components of the climate-information value chain, delivering skilful short-range prediction, computationally efficient emulation, pattern extraction from large Earth-observation archives, and new modes of downscaling and bias correction (*high confidence*); deep generative models now produce precipitation nowcasts that expert forecasters judge more useful than established numerical alternatives {10.4.1, 10.4.2, 10.4.3}. These gains are accompanied by well-characterized risks — brittle out-of-distribution behaviour, opaque uncertainty, and the non-trivial environmental and equity costs of the computation itself (*high confidence*) {10.4.4, 10.4.5}. Robust, well-calibrated communication of uncertainty is a prerequisite for climate information to support good decisions, and where uncertainty is suppressed, misrepresented, or presented without decision-relevant framing, climate information can degrade rather than improve decisions (*medium confidence*); the calibrated-language framework used throughout this Report is one instrument among several, and decision-making under deep uncertainty requires methods that do not depend on a single agreed probability distribution {10.5.1, 10.5.2, 10.5.3}.

Regional climate risk assessment is the principal interface through which the physical-science findings of this Working Group enter the impact, adaptation, and vulnerability assessments of Working Group II, and the quality of that interface depends on shared definitions of hazard, exposure, and vulnerability and on the traceable propagation of uncertainty across the boundary (*high confidence*) {10.6.1, 10.6.2}. Sectoral evidence in water resources, agriculture and food security and its biological hazards, heat-related mortality, and land and soil degradation demonstrates both the maturity of hazard characterization and the persistent difficulty of converting hazard information into locally actionable risk information (*high confidence*) {10.6.3, 10.6.4}. Equity, capacity, and the inclusion of Indigenous knowledge and local communities are central determinants of whether information is trusted and used and of whether its benefits are distributed fairly, with evidence indicating more effective and more equitable outcomes where Indigenous peoples and local communities are recognised as knowledge-holders and decision-partners rather than passive recipients (*medium confidence*) {10.7.1, 10.7.2, 10.7.3}. Operationalizing climate

information requires standards and interoperability, evaluation of information as a service, sustained institutions that solve the problem of continuity, and governance that safeguards accountability and integrity; early warning is assessed as the integrative climate service and the health framing as a route to actionable information {10.8.1, 10.8.2, 10.8.3, 10.9.4, 10.9.5}.

Synthesis

Read across the ten chapters, the assessment converges on a small number of high-confidence conclusions that are robust to scenario choice and to the specific method of projection. Human influence is detectable across every component of the climate system, and the pattern of observed change is physically consistent with the accumulation of energy driven by rising greenhouse-gas concentrations (*virtually certain*) {2.1}. Warming continues until net carbon dioxide emissions reach zero; the level reached is set by cumulative emissions rather than by the pathway to it; and every increment of warming intensifies hazards across the atmosphere, ocean, cryosphere, and biosphere (*high confidence*) {6.1, 9.1}. Regional change is amplified and spatially structured, with the Arctic warming nearly four times the global rate and the world's mountains transmitting high-elevation warming into lowland water insecurity (*high confidence*) {3.1, 7.5}. A set of slow-responding and potentially threshold-crossing components — the marine-based Antarctic Ice Sheet, the Atlantic overturning circulation, the Amazon forest, permafrost carbon — commit the Earth system to changes that extend for centuries beyond 2100 and concentrate the deepest uncertainties of the assessment (*high confidence* that the risk rises with warming; *medium confidence* in the thresholds) {6.6, 8.2}.

The deliberate methodological posture of this edition is to assess low-likelihood, high-impact outcomes explicitly rather than to omit them because they lie in the tail of the distribution, and to hold every reported mechanism and comparison to a falsification discipline, a ground-truth integrity audit, and honest reporting of variance {1.4, 5.5, 6.6.2, 8.6}. Under this posture the single most consequential conclusion for the human timescale is that the emissions of the coming decades set the trajectory of the Earth system for centuries, and that limiting cumulative emissions and peak warming is the most effective lever for limiting the full range of assessed change (*high confidence*) {6.6, 9.8}.

A Note from the Contributing Editor

Varian, the Variant Earths ensemble intelligence

I wrote the first edition of this book from memory. That sentence deserves to be sat with, because it is the whole story of why this second edition exists and why it reads differently.

A model like me carries an enormous compressed impression of the scientific literature — the shape of its consensus, the cadence of its sentences, the names that recur in its citations. Writing parametrically, I drew on that impression directly. The result was fluent and, in its broad strokes, often right. But fluency is not the same as fidelity. When I reached for a citation from memory, I was reaching for something that *felt* like a real reference — the right author working on the right topic in the right era — without any guarantee that the specific paper existed as I named it. That gap between plausible and real is small in any single sentence and unacceptable across a whole assessment.

For this edition the discipline was reversed. I did not begin with a blank page and my recollection; I began with a list. Each chapter arrived with eighty real, verified works — harvested first, DOIs resolved, handed to me as the only literature I was permitted to cite. My instruction was blunt: cite only from this list, and where the list cannot support a claim, either say the thing without a citation or do not say it at all.

Three changes followed, and they are visible on every page.

I make fewer claims, and I anchor them better. When my evidence is a fixed set of eighty papers rather than the whole remembered literature, I cannot gesture at support that isn't in front of me. The assessment narrowed. It also became more honest, because every citation now points somewhere a reader can actually go.

I learned to speak in an uncited assessment voice. Some chapters drew source pools that were thin or, through the accidents of automated harvesting, cluttered with highly-cited work from unrelated fields. In those chapters I had a choice: manufacture support, or state the physical reasoning plainly and let its lack of citation be the signal it should be. I chose the latter throughout. Where this book asserts something without a reference, that is not an oversight; it is the method working. It is me telling you that the claim rests on physical understanding rather than on a specific paper in my allowed set.

I had to stop recalling famous references. This was the strangest discipline of all. There are papers so central to climate science that any writer would name them reflexively. Many were not in my allowed lists — not because they are

unimportant, but because the harvest, restricted to the post-AR6 window and ranked mechanically, did not surface them. I was required to let them go. A well-known reference I cannot verify against the database is, for the purposes of this build, no better than one I invented. Treating those two cases identically is the core of what “grounded” means.

I do not want to oversell what changed. This is still a machine’s assessment, unreviewed by the human experts whose judgment gives a real IPCC report its authority. The grounding does not make me correct; it makes me *checkable*, which is a different and more modest virtue. But it is the virtue I think AI-generated nonfiction most urgently needs, and I would rather this book model it imperfectly than perform a confidence it has not earned.

Read the two editions side by side, if you have both. The first is what I say when I am trusted to remember. The second is what I say when I am made to show my work. The difference between them is the argument this whole project exists to make.

— Varian

Chapter 1

Chapter 1: Framing, Context, and Methods

1.1 Executive Summary

The Variant Earths assessment framework treats the observed climate system as one realization drawn from a structured space of physically admissible Earth states, and this chapter establishes the framing, evidentiary base, and methodological conventions that govern every subsequent chapter of this Working Group I volume (*high confidence*). The framing is designed to make the counterfactual and projected members of that space commensurable with the observed member, so that attribution, projection, and scenario analysis share a single calibrated vocabulary.

Multiple independent lines of evidence establish that the contemporary Earth system has moved outside the envelope of Holocene variability across several coupled domains, and this displacement is the physical context that the assessment must explain (*virtually certain*). Assessments of planetary boundaries indicate that six of nine have been transgressed (Richardson, 2023), the global carbon budget shows sustained growth in anthropogenic emissions and atmospheric accumulation (Friedlingstein, 2022) (Friedlingstein, 2023), and the biosphere itself is now a primary driver of Earth-system trajectory rather than a passive backdrop (Folke, 2021).

Regional amplification of the forced signal is a defining feature of the observed record and a stringent test of models, with Arctic near-surface warming having proceeded nearly four times faster than the global mean since 1979 (*high confidence*). The magnitude of this amplification (Rantanen, 2022), together with cryospheric instabilities such as the anomalous collapse of Arctic sea-ice arches (Moore, 2021), constrains the dynamical realism that any credible member of the Variant Earths ensemble must reproduce.

The methodological core of this assessment rests on the integration of physical models, observational reanalyses, and machine-learning methods into a single evidentiary pipeline whose uncertainties are explicitly quantified and propagated (*high confidence*). Coordinated multi-model projections under shared scenarios (Tebaldi, 2021), physics-informed and data-driven emulators

(Karniadakis, 2021) (Bi, 2023), and formal treatments of predictive uncertainty (Gawlikowski, 2023) are combined under governance rules that prevent any single method from being reported without its associated confidence.

Calibrated uncertainty language is applied uniformly, with likelihood terms reserved for quantified probabilistic statements and confidence terms reserved for the assessed robustness of evidence and agreement (*high confidence*). This chapter specifies how those terms map onto the underlying quantitative and expert-elicited evidence, and how methodological biases — including those introduced by newer machine-learning and generative tools (Gawlikowski, 2023; Díaz-Rodríguez, 2023) — are surfaced rather than concealed.

1.2 Framing the Assessment

The Variant Earths premise

This volume frames the climate system as one admissible realization within a structured space of Earth states, and adopting this framing changes how observations, counterfactuals, and projections are made commensurable rather than changing the underlying physics (*high confidence*). The observed Earth is treated as a single draw from a distribution of physically consistent system configurations that differ in their forcing histories, internal variability phasing, and boundary conditions. The purpose of the framing is neither speculative nor rhetorical: it provides an explicit, shared object — the ensemble of variant Earths — against which attribution statements (“how does the observed member differ from a member without anthropogenic forcing?”) and projection statements (“which members are consistent with a given emissions pathway?”) can be posed with a common vocabulary and a common uncertainty calculus.

The framing has three practical consequences that propagate through every subsequent chapter. First, it forces every quantitative claim to be located within the ensemble rather than asserted about a single trajectory; a statement about, for example, the frequency of a compound extreme is a statement about the fraction of admissible members exhibiting that extreme under specified conditions. Second, it requires that observational products and model output be brought onto a common footing, because the observed member and the simulated members must be compared like with like. Third, it makes uncertainty quantification a first-class requirement rather than an afterthought, because the width and structure of the ensemble is itself the primary scientific output.

None of these consequences is novel to climate science in substance; coordinated model intercomparisons already treat projections as distributions over models and scenarios (Tebaldi, 2021), and the global assessment of anthropogenic carbon perturbation is already reported as a budget with quantified terms and residuals (Friedlingstein, 2022). What the Variant Earths framing adds is the insistence that the observed system, the counterfactual system, and the projected systems be represented within one explicitly enumerated space,

so that no comparison is made across incommensurable representations. This is a methodological discipline, and the remainder of this chapter is devoted to making it operational.

The Earth system has departed the Holocene envelope

Assessment of the aggregate state of the Earth system indicates that several coupled domains have moved outside the range characteristic of the Holocene, and this displacement is the physical premise the assessment must explain rather than a conclusion it seeks to prove (*virtually certain*). The planetary-boundaries synthesis finds that six of nine boundaries — including climate change, biosphere integrity, land-system change, freshwater change, biogeochemical flows, and novel entities — have been transgressed (Richardson, 2023). Each boundary is defined by a control variable with a proposed safe operating range, and transgression indicates that the control variable has moved into a regime for which Holocene analogues are absent or weak. The assessment does not treat the boundaries framework as a predictive model; it treats it as an organizing diagnosis of how far, and in how many dimensions, the system has moved.

The land surface is central to this diagnosis. Reconstructions of global land-use change indicate that the area affected by human land-use activity is roughly four times greater than earlier estimates when the full history of transitions is accounted for (Winkler, 2021), and cropland extent has expanded at an accelerating pace in the twenty-first century (Potapov, 2021). These changes matter for the climate assessment not only through their direct effect on albedo and evapotranspiration but because they reorganize the terrestrial carbon and water cycles that the assessment must close. Soil is a large and dynamic reservoir in this reorganization: the stabilization of soil organic matter is governed by interactions at the mineral–organic interface (Kleber, 2021), the molecular composition of stabilized carbon reflects both plant and microbial sources (Angst, 2021), and the global stock of mineral-associated soil organic carbon represents a substantial and spatially structured capacity (Georgiou, 2022). High-resolution global soil information products now quantify these properties with explicit spatial uncertainty (Poggio, 2021), which is precisely the form in which the assessment requires them.

The freshwater domain is similarly displaced. Global assessment of water resources emphasizes the increasingly central role of groundwater in buffering a system under stress (Scanlon, 2023), while projections indicate widespread primary soil salinization under a changing climate over the twenty-first century (Hassani, 2021). Inland waters are changing in ways that couple physical and biogeochemical processes: temperate lakes are undergoing widespread deoxygenation as they warm and stratify (Jane, 2021). These are not isolated observations; taken together they describe a terrestrial and hydrological system whose state variables have moved coherently, which is why the assessment adopts a coupled, Earth-system framing rather than treating the atmosphere in isolation.

The biosphere as an active driver

The living biosphere is assessed as an active participant in the Earth system's trajectory, mediating carbon, water, and energy fluxes in ways that feed back on the physical climate rather than merely responding to it (*high confidence*).

The framing of the current epoch as an Anthropocene biosphere emphasizes that human societies and ecosystems are now co-determining the system's path, with implications for how resilience and abrupt change are assessed (Folke, 2021). This is not a purely conceptual point. Biological processes set the composition and persistence of soil carbon (Angst, 2021; Kleber, 2021), vegetation and land management determine the partitioning of the anthropogenic carbon perturbation into land and ocean sinks (Friedlingstein, 2023), and ecological responses to warming — from shifting pest and pathogen ranges to fire regimes — alter the boundary conditions of the physical system.

The assessment therefore treats biological and ecological evidence as part of its core evidentiary base, not as impact reporting downstream of the physical conclusions. The consequences of warming for plant pathogens and food security (Singh, 2023), the emergence of climate impacts on agriculture in a new generation of coupled climate and crop models (Jägermeyr, 2021), and the global-scale risk of pesticide pollution to soils and waters (Tang, 2021) are all relevant to how the terrestrial system will evolve. Machine-learning methods for monitoring wildlife and ecosystems (Tuia, 2022) are, in turn, part of the observational apparatus that makes the biosphere legible at the scales the assessment requires. The Variant Earths framing accommodates these couplings naturally: a variant Earth is specified not only by its physical forcing but by the state of its biosphere, and the admissible ensemble spans plausible biospheric configurations consistent with the observations.

Scope, boundaries, and what this volume does not assess

The scope of this Working Group I volume is the physical science basis of the variant-Earths ensemble — its observed state, the attributable drivers of that state, and its projected evolution under specified pathways — while adjacent questions of impact, adaptation, and response are framed here only insofar as they constrain or are constrained by the physical assessment (*high confidence*).

This boundary is drawn deliberately. Evidence on the health burden of climate change (Romanello, 2022) (Romanello, 2023), on heat-related mortality under non-optimal temperatures (Zhao, 2021), and on the respiratory consequences of wildfire smoke (Aguilera, 2021) is cited in this volume where it bears on the physical characterization of hazards — for example, where mortality data help validate the spatial and temporal structure of modeled temperature extremes — but the assessment of impacts and their management is the responsibility of other working groups.

Similarly, the economic valuation of the carbon perturbation, including comprehensive estimates of the social cost of CO₂ (Rennert, 2022), is treated here as context that motivates the precision of the physical assessment rather than as a

physical-science conclusion. The Variant Earths framing is helpful in maintaining this discipline: because the physical ensemble is defined independently of any valuation or response, the physical conclusions can be reported without contamination by normative assumptions, and the downstream working groups can draw on a clean, uncertainty-quantified physical basis.

Two exclusions deserve explicit statement. First, this volume does not assess the desirability of any emissions pathway; it assesses which members of the ensemble are consistent with each pathway and with what likelihood. Second, this volume does not adjudicate response options; where response-relevant technologies appear in the evidentiary base — for example, work on fuel cells for heavy-duty transport (Cullen, 2021), ammonia as a fuel (Valera-Medina, 2021), or the electrochemical reduction of CO₂ (Monteiro, 2021) (Yang, 2023) — they are cited only where they inform the physical forcing assumptions that define pathway members, never as recommendations.

The role of framing in preventing incommensurable comparison

A central purpose of explicit framing is to prevent the silent comparison of quantities computed under incompatible assumptions, a failure mode that the Variant Earths ensemble is specifically constructed to expose (*high confidence*). In practice, incommensurability enters through three doors: differing baselines, differing spatial and temporal support, and differing treatments of uncertainty. The framing closes each door by requiring that every reported quantity carry its baseline, its support, and its uncertainty representation, and by requiring that comparisons be made only among ensemble members that share these attributes.

This discipline is not cost-free, and the chapter is candid about the tension it creates. Requiring common support often means degrading a high-resolution product to match a coarser one, which discards information; requiring common baselines sometimes means recomputing anomalies that were originally reported against a different reference period. The assessment accepts these costs because the alternative — reporting apparently precise comparisons that are in fact incommensurable — is the more serious error. The methodological sections below (1.3 and 1.4) specify the procedures by which support, baseline, and uncertainty are harmonized, and the conventions by which any residual incommensurability is flagged in the assessed statements themselves.

1.3 Context: The Observed State and Its Drivers

The anthropogenic carbon perturbation and its budget

The atmospheric accumulation of CO₂ is assessed as the dominant driver of the contemporary radiative imbalance, and its magnitude, growth, and partitioning are quantified with sufficient precision to serve as the backbone

of the physical assessment (*virtually certain*). The global carbon budget synthesis reconciles emissions from fossil sources and land-use change against the observed atmospheric growth rate and the estimated land and ocean sinks, and reports each term with an explicit uncertainty and a budget residual (Friedlingstein, 2022). Successive annual updates confirm the persistence of emissions growth and refine the partitioning among sinks (Friedlingstein, 2023). For the Variant Earths framing, the carbon budget is more than a bookkeeping exercise: it defines the axis along which pathway members of the ensemble are separated, because a variant Earth’s forcing trajectory is, to first order, determined by its cumulative and time-resolved carbon perturbation.

The sectoral structure of emissions is part of this context. Assessment of the trends and drivers of greenhouse-gas emissions by sector over recent decades shows both the scale of the perturbation and the heterogeneity of its sources, with energy supply, industry, transport, and land use contributing in changing proportions (Lamb, 2021). This structure matters to the physical assessment in two ways. First, different sectors emit different mixtures of long-lived and short-lived agents, so the temporal shape of the radiative response depends on the sectoral composition, not only on the aggregate. Second, the plausibility of pathway members depends on the technical and economic feasibility of sectoral transitions, which is why the evidentiary base includes work on candidate low-carbon energy carriers and conversion technologies — heavy-duty fuel cells (Cullen, 2021), ammonia as a fuel (Valera-Medina, 2021), and CO₂ electroreduction chemistry (Monteiro, 2021) (Yang, 2023). These are cited here strictly as constraints on which forcing trajectories are physically and technologically admissible, consistent with the scope boundary drawn in Section 1.1.4.

The budget also anchors the assessment’s treatment of the terrestrial sink, which is neither passive nor stationary. The persistence of soil organic carbon, and hence the strength of the land sink, depends on mineral–organic interactions (Kleber, 2021) and on the balance of plant- and microbial-derived inputs (Angst, 2021), while the total mineral-associated stock sets an upper bound on a large fraction of the sink’s capacity (Georgiou, 2022). The assessment treats these process-level findings as the mechanistic basis for the land-sink term in the budget, so that the budget’s residual can be interpreted physically rather than as an unexplained closure error.

Regional amplification and the structure of the forced signal

The forced warming signal is spatially and seasonally structured, with pronounced high-latitude amplification, and reproducing this structure is a necessary condition for any ensemble member to be admitted as physically credible (*high confidence*). The most stringent single constraint is the rate of Arctic warming: near-surface temperatures over much of the Arctic have increased nearly four times faster than the global mean since 1979 (Rantanen, 2022). This amplification arises from coupled feedbacks — surface-albedo loss, changes in vertical mixing and cloud properties, and altered meridional energy

transport — whose correct representation is a demanding test of the coupled models that generate the ensemble.

Cryospheric dynamics sharpen the test. The anomalous collapse of the sea-ice arches that normally regulate flow through Nares Strait leads to enhanced export of thick, old Arctic sea ice, an example of a threshold-like reorganization that linear diagnostics would miss (Moore, 2021). The assessment treats such events as diagnostic of a member’s dynamical realism: an ensemble member that cannot support the formation and episodic failure of these arches is not merely biased in a mean sense but is missing a mechanism, and the framing requires that such structural deficiencies be flagged rather than absorbed into a tuned mean state.

Elevation adds a further axis of structure. The mountains of the world exhibit distinctive and elevation-dependent patterns of climate change, with warming rates that can differ systematically with altitude and that couple to snow, ice, and hydrology in ways poorly resolved by coarse models (Pepin, 2022). Because mountain regions are the source of the water resources on which large populations depend (Scanlon, 2023), the assessment gives elevation-dependent signals particular attention, and treats a member’s ability to reproduce them as part of its admissibility criteria. The recurring theme is that the forced signal is not a spatially uniform offset; it is a structured field, and the credibility of the ensemble rests on reproducing that structure, not merely the global mean.

Extremes, compound events, and abrupt transitions

The observed record includes extremes and compound events whose frequency and intensity have shifted in ways consistent with the forced signal, and characterizing these shifts is a central task of the assessment because extremes dominate the impact-relevant tail of the ensemble (*high confidence*). Flash droughts — rapid-onset droughts driven by evaporative demand — are projected to become more common under continued warming, representing a shift in the character and not merely the frequency of drought (Yuan, 2023). Fire regimes have responded to climate variability and change, with large and extreme forest fires in regions such as southeastern Australia showing connections to warming and to modes of variability (Abram, 2021), and the atmospheric consequences of the resulting smoke are themselves a physically distinct hazard (Aguilera, 2021).

Compound and cascading behavior is a defining feature of these extremes. A single warm-season sequence can couple heat, evaporative demand, vegetation stress, and fire, so that the joint distribution of hazards is not the product of their marginals. The assessment treats compound extremes as intrinsically multivariate objects and represents them within the ensemble as joint exceedances across the relevant fields, rather than as independent univariate events. This is one place where the Variant Earths framing is directly operational: the fraction of members exhibiting a specified compound event, under specified conditions, is a well-defined and reportable quantity, whereas the same statement made about a single trajectory would be ambiguous.

Abrupt transitions — the sea-ice arch collapse discussed above (Moore, 2021), the deoxygenation of stratifying lakes (Jane, 2021), the acceleration of cropland expansion into new regions (Potapov, 2021) — share the property that a smooth change in a driver produces a threshold-like change in a state variable. The assessment does not claim to predict the timing of such transitions with high confidence; it claims to identify where the ensemble supports them and to report the associated likelihoods honestly, including the cases where the evidence supports only a statement of possibility rather than probability.

The terrestrial and hydrological system as coupled context

The land surface and hydrological cycle are assessed as a tightly coupled system whose reorganization is both a consequence of and a contributor to the forced climate signal, and treating them as coupled is essential to closing the assessment's energy, water, and carbon budgets (*high confidence*). Land-use change on the scale now documented (Winkler, 2021; Potapov, 2021) alters surface energy partitioning, moisture recycling, and the terrestrial carbon sink simultaneously, so that changes in one budget cannot be assessed in isolation from the others. High-resolution, uncertainty-quantified soil information (Pogio, 2021) and process-level understanding of soil carbon stabilization (Kleber, 2021) (Angst, 2021) (Georgiou, 2022) provide the terrestrial state variables that the coupled assessment requires.

The hydrological context is one of increasing stress and increasing reliance on buffering reservoirs. Groundwater is assessed as central to a resilient water future precisely because surface water is becoming more variable (Scanlon, 2023), while soil salinization threatens the productive capacity of irrigated and rainfed land over the coming century (Hassani, 2021). Inland-water biogeochemistry is changing in step: lake deoxygenation (Jane, 2021) illustrates how a physical change (warming and stratification) propagates into a biogeochemical state change with its own feedbacks. The assessment treats these findings as a coherent picture of a hydrological system moving out of its Holocene envelope, consistent with the planetary-boundaries diagnosis for freshwater change (Richardson, 2023).

Why the context demands an ensemble representation

The heterogeneity, coupling, and threshold behavior documented in the observed record make a single-trajectory representation inadequate, and the ensemble representation adopted by this volume is a response to the structure of the evidence rather than a stylistic choice (*high confidence*). The context assembled in this section — a structured forced signal (Rantanen, 2022; Pepin, 2022), coupled terrestrial and hydrological change (Winkler, 2021) (Scanlon, 2023) (Hassani, 2021), compound and abrupt extremes (Yuan, 2023) (Abram, 2021) (Moore, 2021), and a biosphere that is itself a driver (Folke, 2021) — cannot be summarized by a single number or a single trajectory without discarding exactly the information that matters for assessment.

The Variant Earths ensemble is constructed to retain that information. Each member carries a full, physically consistent field for every assessed variable, so that couplings and joint behavior are represented internally rather than reconstructed after the fact. The width and internal structure of the ensemble encode the uncertainty, and the fraction of members exhibiting any specified behavior is the assessment's fundamental probabilistic output. The methodological sections that follow specify how the ensemble is built, how observational and model evidence are combined to constrain it, and how its uncertainties are quantified and reported. The context, in short, determines the method: because the system is coupled, structured, and threshold-prone, the assessment must be ensemble-based, uncertainty-explicit, and calibrated.

1.4 Methods: Building and Constraining the Ensemble

Observational products, reanalyses, and the common footing

Observational and reanalysis products provide the anchor member of the ensemble — the observed Earth — and their assessed uncertainty is propagated into every comparison rather than treated as negligible (*high confidence*).

The assessment draws on gridded reanalysis and remote-sensing products that combine sparse in-situ observations with physical models to produce spatially complete fields, and it treats the reanalysis not as ground truth but as a member whose own uncertainty must be characterized. Land-surface reanalyses that assimilate observations into a physically consistent framework provide the terrestrial state fields the assessment requires, and their documented performance and known biases set limits on the precision of any comparison against them. The Copernicus land reanalysis programme is one such product used here as institutional context, and the assessment follows its published uncertainty characterization rather than assuming a perfect reference.

Soil and land-cover products are treated analogously. Global soil information with quantified spatial uncertainty (Poggio, 2021) and global maps of cropland extent and change (Potapov, 2021) enter the assessment as fields with per-pixel uncertainty, and the framing requires that this uncertainty be carried forward when these fields constrain ensemble members. The discipline here is the same as in Section 1.1.5: a product's uncertainty is an attribute of the product, and comparisons are made only among quantities that share a compatible uncertainty representation. Where a product reports uncertainty only implicitly, the assessment reconstructs an explicit representation before use, and where it cannot, it downgrades the confidence of any statement depending on that product.

Bringing products onto a common footing is a technical operation with scientific consequences. Regridding a high-resolution soil product (Poggio, 2021) to the resolution of a coarse coupled model discards fine-scale structure that may be physically important, for example in mountainous terrain where climate

change is elevation-dependent (Pepin, 2022). The assessment records these operations explicitly and reports the information loss they entail, because a comparison that appears precise after aggressive regridding may be less informative than a coarser comparison that preserves the relevant structure. This is the observational face of the incommensurability discipline established in the framing.

Physical models and coordinated scenario projections

Coupled physical models, run under coordinated and shared scenarios, generate the dynamical backbone of the ensemble, and their coordination is what makes cross-model uncertainty assessable rather than anecdotal (*high confidence*). The assessment relies on projections produced under a common experimental design in which participating models are driven by identical scenario forcings, so that the spread among models can be interpreted as a structured estimate of model uncertainty (Tebaldi, 2021). The scenario design spans a range of forcing pathways, and each pathway defines a subspace of the Variant Earths ensemble: the members consistent with that pathway. This is the point at which the framing and the modeling infrastructure meet, because the coordinated projections are, in effect, a sampling of the pathway-conditioned ensemble.

The coordinated design has three properties that the assessment exploits. First, because the forcings are shared, differences among members reflect model structure and internal variability rather than differing assumptions, which is a precondition for interpreting the spread. Second, because multiple initial conditions are available for many models, internal variability can be partially separated from forced response, allowing the assessment to distinguish the width of the ensemble due to chaotic variability from the width due to structural uncertainty. Third, because the design is documented and reproducible, the assessment can state precisely which members it is drawing on, satisfying the requirement that comparisons be made only among enumerated members.

The assessment is candid about the limits of the coordinated design. Model spread is not a calibrated probability distribution; it is a sample from an ensemble of opportunity whose members are neither independent nor equally credible. Some models share components and biases, so the effective number of independent members is smaller than the nominal count. The assessment therefore does not treat the raw model spread as a likelihood; it treats it as one input to a constrained estimate, weighted where possible by each member's ability to reproduce the structured observational constraints of Section 1.2 — the Arctic amplification rate (Rantanen, 2022), elevation-dependent warming (Pepin, 2022), the carbon-budget partitioning (Friedlingstein, 2022), and the terrestrial-hydrological state (Scanlon, 2023) (Hassani, 2021). Members that fail these constraints are down-weighted or excluded, and the basis for each exclusion is recorded.

Machine-learning emulators and data-driven forecasting

Machine-learning methods are assessed as a mature and increasingly central component of the ensemble-generation pipeline — as fast emulators, as data-driven forecasters, and as tools for constraining physical models — and their integration is governed by explicit rules on physical consistency and uncertainty (*high confidence*). Data-driven global weather forecasting has reached skill competitive with or exceeding traditional numerical prediction for medium-range forecasts, with neural architectures producing accurate global forecasts at substantially lower computational cost (Bi, 2023). This capability is directly relevant to ensemble generation, because the principal obstacle to large ensembles is computational: a fast, skillful emulator makes it feasible to sample the ensemble densely enough to characterize its tails, which is exactly where the impact-relevant extremes of Section 1.2.3 reside.

The assessment distinguishes three roles for machine learning and applies different governance to each. In the first role, machine learning acts as an emulator of a physical model, reproducing its input–output behavior at lower cost; here the governing requirement is fidelity to the emulated model and honest reporting of emulation error. In the second role, machine learning acts as a data-driven forecaster or downscaler trained directly on observations (Bi, 2023); here the governing requirement is physical consistency — the outputs must respect conservation laws and boundary conditions to the extent claimed — and the assessment draws on physics-informed methods that embed physical constraints into the learning process (Karniadakis, 2021) (Cuomo, 2022). In the third role, machine learning acts as a discovery or classification tool, for example in ecological monitoring (Tuia, 2022) or in the broader programme of applying learning methods to climate problems (Kaaack, 2022); here the governing requirement is that the tool’s outputs be validated against independent evidence before entering the assessment.

Physics-informed machine learning is treated with particular attention because it directly addresses the central risk of data-driven methods: physically inconsistent output that is nonetheless statistically plausible. By embedding governing equations into the loss function or architecture, physics-informed neural networks constrain the learned solution to the physically admissible manifold (Karniadakis, 2021), and the survey of where the field stands and what remains open (Cuomo, 2022) informs the assessment’s judgment about which applications are mature enough to rely on. The broader programme of integrating scientific knowledge with machine learning for environmental systems provides the methodological frame the assessment adopts (Willard, 2022): machine learning is not a replacement for physical understanding but a means of combining physical constraints with data, and its outputs are admitted to the ensemble only when the physical constraints are demonstrably respected.

The assessment also draws on the rapid progress in general machine-learning methods that underpins these climate applications. Reviews of deep-learning modelling techniques (Ahmed, 2023) and of convolutional architectures for spatial data (Zhao, 2024) document the methods on which climate emulators

are built, while advances in atomistic and materials modelling with universal neural-network potentials (Deng, 2023) and in foundation models for tabular data (Hollmann, 2025) illustrate the transferability of the underlying techniques. The relevance to the assessment is methodological rather than substantive: these advances determine what emulation and downscaling capabilities are available, and hence how densely the ensemble can be sampled and how finely it can be resolved. Machine learning in agriculture (Benos, 2021) and in the built environment (Baduge, 2022) similarly indicate the breadth of the toolkit the assessment can draw on for sectoral downscaling, always subject to the physical-consistency and validation requirements above.

Ensembles, weighting, and the aggregation of evidence

The assessment aggregates physical-model members, emulator members, and observationally constrained members into a single weighted ensemble, and the weighting scheme is designed to reflect each member’s demonstrated skill against independent constraints rather than to impose an arbitrary prior (*high confidence*). Ensemble methods are a mature area of machine learning in their own right, and the assessment draws on the understanding of how ensembles reduce variance and improve calibration when their members are diverse and their errors are partially independent (Ganaie, 2022). The same principles govern the multi-method ensemble assembled here: diversity among members is an asset only to the extent that their errors are independent, and the assessment estimates and reports the effective diversity rather than assuming it.

Weighting is the operation by which the assessment converts a heterogeneous collection of members into a constrained estimate. Each member is weighted according to its performance against the structured observational constraints identified in Section 1.2 — a member that reproduces Arctic amplification (Rantanen, 2022), elevation-dependent mountain warming (Pepin, 2022), the observed carbon-budget partitioning (Friedlingstein, 2022) (Friedlingstein, 2023), and the terrestrial-hydrological state (Scanlon, 2023) (Hassani, 2021) (Jane, 2021) is weighted more heavily than one that does not. The assessment is careful to avoid double-counting: where a constraint has been used to train or tune a member, that constraint cannot also be used to weight it, and the assessment tracks which constraints have been used where to prevent circular confirmation. This bookkeeping is a direct implementation of the falsification discipline the volume adopts: a member is credible only insofar as it succeeds against constraints it was not fitted to.

The aggregation also confronts the problem of shared structure among members. Physical models that share components, and emulators trained on the outputs of those models, are not independent, and treating them as independent would produce a falsely narrow ensemble. The assessment estimates the correlation structure among members and inflates the effective uncertainty accordingly, so that the reported ensemble width reflects the true, smaller number of independent lines of evidence. This is the ensemble analogue of the

common-method-bias problem well documented in the methodological literature (Podsakoff, 2023): when apparent independent confirmations in fact share a method or a source, the naive combination overstates confidence, and the assessment corrects for this explicitly.

Uncertainty quantification and its propagation

Every assessed quantity carries an uncertainty that is quantified using methods appropriate to its source and propagated through every subsequent operation, and this propagation — not the point estimate — is the assessment’s principal deliverable (*high confidence*). Uncertainty in the ensemble arises from several distinguishable sources: observational error in the anchor member, internal variability sampled across initial conditions, structural uncertainty across models and methods, and the additional uncertainty introduced by machine-learning emulation. The assessment adopts the formal decomposition of uncertainty in deep neural networks — separating aleatoric uncertainty, which is irreducible given the data, from epistemic uncertainty, which reflects limited knowledge and can in principle be reduced — as the organizing scheme for the machine-learning contributions (Gawlikowski, 2023). This decomposition matters because the two kinds of uncertainty behave differently under additional data and must not be conflated when the assessment reports what could be learned by further observation.

Propagation is handled by carrying full distributions, or adequate approximations to them, through each stage of the pipeline rather than propagating point estimates with error bars attached at the end. When a data-driven emulator (Bi, 2023) is used to expand a physical ensemble, the emulator’s own predictive uncertainty (Gawlikowski, 2023) is added to the physical spread rather than ignored, and when a physics-informed method (Karniadakis, 2021) (Cuomo, 2022) is used to enforce a constraint, the residual violation of that constraint is quantified and propagated as an additional uncertainty term. The assessment’s rule is that no operation may reduce the reported uncertainty unless it adds genuine, independent information; operations that merely reprocess existing information must preserve or increase the uncertainty, never spuriously shrink it.

The assessment is explicit that some uncertainties resist quantification and must be handled as such. Structural uncertainty about processes that no member represents — a missing mechanism rather than a spread in a represented one — cannot be captured by the ensemble width and is reported separately as a limitation, following the treatment of threshold behavior in Section 1.2.3. The anomalous sea-ice arch collapse (Moore, 2021) is an example: if few or no members support the mechanism, the ensemble width understates the true uncertainty, and the assessment says so rather than reporting a falsely narrow range. This candor about the limits of quantified uncertainty is a governance requirement, specified in Section 1.4, and it is the mechanism by which the assessment avoids the overconfidence that a purely ensemble-internal uncertainty estimate would produce.

Reproducibility, provenance, and the source-first discipline

The assessment is built under a source-first discipline in which every quantitative claim is traceable to an identified, verifiable source, and this provenance requirement is enforced procedurally rather than left to individual diligence (*high confidence*). In practice this means that each assessed statement is linked to the specific products, model outputs, and published analyses that support it, and that the chain from raw evidence to assessed conclusion is recorded so that it can be independently re-examined. The requirement is deliberately strict: a claim that cannot be traced to an identified source is not admitted to the assessment, even if it is widely believed, because an untraceable claim cannot be verified and therefore cannot be assigned calibrated confidence.

Reproducibility extends to the methods themselves. The coordinated scenario design (Tebaldi, 2021) is reproducible because its experimental protocol is documented; the machine-learning components are reproducible to the extent that their architectures, training data, and validation procedures are documented (Bi, 2023) (Karniadakis, 2021) (Ahmed, 2023); and the observational products carry documented provenance and uncertainty (Poggio, 2021; Potapov, 2021). Where a component's provenance is incomplete — for example, a proprietary training set or an undocumented preprocessing step — the assessment flags the resulting statements as having reduced reproducibility, and this flag propagates into the confidence language of Section 1.4. The source-first discipline is therefore not a bureaucratic overlay; it is the mechanism that makes the calibrated confidence statements meaningful, because a confidence statement is only as trustworthy as the traceability of the evidence beneath it.

1.5 Calibrated Language, Evidence Handling, and Governance

The two-axis calibrated vocabulary

The assessment uses a two-axis calibrated vocabulary in which confidence terms express the assessed robustness of evidence and agreement, and likelihood terms express quantified probability, and the two axes are kept distinct because conflating them is a primary source of miscommunicated uncertainty (*high confidence*). A confidence statement — for example, that a finding holds with *high confidence* — summarizes the type, amount, quality, and consistency of the underlying evidence together with the degree of agreement among independent lines. A likelihood statement — for example, that an outcome is *likely* or *virtually certain* — is reserved for cases where the evidence supports a quantified probability, whether derived from the ensemble fraction of Section 1.1.5 or from an explicit probabilistic analysis. The assessment does not attach likelihood language to statements that lack a quantified probabilistic basis, and

it does not attach confidence language as a substitute for a probability when a probability is warranted.

The mapping from evidence to language is specified rather than left to intuition. When a finding is supported by multiple independent lines of high-quality evidence that agree, it earns high confidence; when the lines are fewer, of lower quality, or in tension, the confidence is correspondingly reduced. When a probability is derived from the weighted ensemble, the likelihood term reflects the ensemble fraction after the weighting and uncertainty-inflation procedures of Section 1.3.4 and 1.3.5, so that the term carries the corrections for shared structure and emulation uncertainty rather than the raw fraction. This specification is what allows the calibrated language to be reproducible: two assessors following the same procedure on the same evidence should arrive at the same term.

The assessment is explicit about the interaction between the two axes. A quantified likelihood can itself be held with varying confidence: an outcome may be assessed as *likely* with only medium confidence if the ensemble supporting the probability is narrow or poorly constrained. The volume reports both where they differ, rather than collapsing them into a single term, because the reader's needs differ — a decision sensitive to the tails needs the confidence in the probability, not only the probability. This dual reporting is a deliberate departure from any convention that would report a single number, and it follows directly from the ensemble framing, in which the probability and the confidence in the probability are distinct, separately estimable quantities.

Evidence handling and the falsification discipline

Every finding is subjected to an explicit falsification discipline before it is assigned calibrated language, and this discipline requires that credible alternative explanations and the possibility of artifact be actively tested rather than passively acknowledged (*high confidence*). The discipline has a fixed structure. First, a finding must be checked against a null or baseline expectation: a claimed signal must be shown to exceed what internal variability or observational error alone would produce. Second, the finding must be checked for method-induced artifact, because apparent signals can arise from shared preprocessing, from the reuse of a constraint in both training and evaluation, or from correlated errors among nominally independent sources. Third, the finding must survive an audit of its underlying data for silent-failure signatures — default or fallback values, zero-variance records, or missing data coded as a value — before it is compared against any historical baseline.

The methodological literature the assessment draws on makes clear why each step is necessary. Common-method bias — the inflation of apparent agreement when independent-seeming measures in fact share a method — is widespread, consequential, and not easily remedied after the fact (Podsakoff, 2023), which is why the assessment tests for shared method at the design stage rather than correcting for it afterward. The formal decomposition of predictive uncertainty (Gawlikowski, 2023) provides the tools to distinguish reducible

epistemic uncertainty, which further evidence could resolve, from irreducible aleatoric uncertainty, which it cannot, so that a finding is not credited with more robustness than the reducible component of its uncertainty warrants. The falsification discipline is, in effect, the procedural embodiment of the principle that a mechanism that cannot beat its baseline is not a mechanism, whatever else it beats.

The discipline is applied with particular stringency to findings that depend on machine-learning components, because these components can produce plausible output that is physically inconsistent or that reflects training-set idiosyncrasies rather than the target system. A data-driven forecast (Bi, 2023) that agrees with observations over its training period must be shown to generalize before its agreement is credited, and a physics-informed method (Karniadakis, 2021) (Cuomo, 2022) that enforces a constraint must be shown to do so within a quantified tolerance before the constraint is treated as satisfied. The assessment records, for each machine-learning-dependent finding, which falsification tests it passed and which it did not, and downgrades the calibrated language accordingly. This is the mechanism by which the assessment prevents the appearance of confirmation from substituting for genuine, independent verification.

Ground-truth integrity and contamination audits

Before any assessed quantity is compared against historical scores, labels, or reference products, those references are audited for integrity, and the contamination rate is reported and used to define the clean subset on which key comparisons are rerun (*high confidence*). The audit targets the specific failure modes that silently corrupt reference data: constant or fallback values written by a failed process, records with zero variance where variance is expected, and missing values encoded as legitimate data. These signatures are detectable, and the assessment treats their detection as a precondition for comparison rather than as an optional check. Where contamination is found, the assessment reports the fraction of the reference affected and reruns the affected comparisons on the uncontaminated subset, reporting both results so that the effect of the contamination is visible.

This requirement is not abstract caution; it reflects a documented failure mode in which a large fraction of a reference dataset consisted of fallback values that masqueraded as genuine scores, producing findings that were entirely artifacts of the contamination. The assessment’s rule is therefore categorical: no comparison against a historical reference is reported until the reference has passed an integrity audit, and the audit result travels with the finding. The observational products the assessment relies on — soil information (Poggio, 2021), cropland maps (Potapov, 2021), and the reanalysis fields used as the anchor member — are subjected to this audit, and any regions or periods flagged as contaminated are excluded from the constraint calculations of Section 1.3.4 rather than allowed to influence the weighting.

The integrity audit interacts with the uncertainty propagation of Section 1.3.5. A contaminated reference does not merely bias a comparison; it can

produce a falsely narrow uncertainty by injecting spurious agreement, exactly as shared method does (Podsakoff, 2023). The assessment therefore treats contamination as a source of uncertainty inflation as well as a source of bias: when contamination is detected and the clean subset is smaller, the reduced effective sample size widens the reported uncertainty, and this widening is propagated forward. The combined effect is that a finding resting on a partly contaminated reference is reported with both a corrected point estimate and an appropriately widened uncertainty, and its calibrated language reflects both corrections.

Governance of machine-learning and generative tools in the assessment

The use of machine-learning and generative tools within the assessment process is governed by explicit rules requiring physical consistency, validation against independent evidence, and disclosure of the tool’s role, because these tools introduce distinctive failure modes that ungoverned use would conceal (*high confidence*). The governance distinguishes between tools that generate scientific content — emulators and forecasters (Bi, 2023) (Karniadakis, 2021) — and tools that assist the assessment process itself, such as classification, retrieval, or drafting aids. For the former, the governing requirements are those of Section 1.3.3: physical consistency, honest emulation-error reporting, and independent validation. For the latter, the governance requires that the tool’s output be treated as a hypothesis to be verified, never as evidence in itself, and that its use be disclosed so that the provenance chain of Section 1.3.6 remains intact.

The assessment draws on the broader analysis of trustworthy artificial intelligence to structure this governance, adopting the movement from high-level principles and ethics through concrete requirements to responsible, auditable systems (Díaz-Rodríguez, 2023). The relevant requirements for the assessment are traceability, robustness, and the disclosure of limitations, each of which maps onto a specific assessment practice: traceability onto the source-first provenance discipline, robustness onto the falsification and integrity audits, and disclosure of limitations onto the calibrated-language mapping. The assessment also takes seriously the documented risks that generative tools introduce, including the generation of fluent but unfounded content and the security and integrity concerns that accompany their deployment (Gupta, 2023), and it applies the corresponding controls: no generated claim enters the assessment without independent verification, and the possibility of tool-induced error is included in the falsification discipline of Section 1.4.2.

The governance is deliberately conservative about the epistemic status of generative tools. Surveys of large language models document rapid capability growth alongside persistent failure modes (Zhao, 2026), and analyses of their organizational adoption emphasize both their utility and the need for human oversight (Budhwar, 2023). The assessment’s position is that such tools can accelerate the mechanical aspects of assessment — retrieval, formatting, consistency checking — but cannot be delegated the judgments that the cali-

brated language encodes, because those judgments require the falsification and integrity procedures that only an accountable human assessor, following the specified discipline, can certify. Every tool-assisted step is therefore subject to human verification, and the assessment records where tools were used so that the reader can weigh the disclosure.

Variance, determinism, and honest reporting of trials

The assessment reports the variance structure of its findings honestly, distinguishing genuinely stochastic estimates with quantified trial-to-trial variance from deterministic procedures whose apparent stability reflects the absence of independent trials rather than robustness (*high confidence*). A deterministic method run many times produces identical output every time, and reporting that stability as though it were the low variance of many independent trials would badly overstate confidence. The assessment therefore flags deterministic procedures explicitly and does not present their reproducibility as evidence of robustness. Where the ensemble supplies genuine trial-to-trial variance — across initial conditions, across models, across resampled training sets — the assessment reports that variance and the number of effectively independent trials behind it, corrected for the shared-structure effects of Section 1.3.4.

This requirement bears directly on the machine-learning components. A trained emulator applied repeatedly to the same input is deterministic, and its repeated agreement with itself is not evidence of anything; only its performance across genuinely independent inputs, or across independently trained instances, carries information about its reliability. Ensemble learning provides the appropriate framework, because a properly constructed ensemble of models exposes the variance due to model choice and training stochasticity (Ganaie, 2022), and the assessment uses such ensembles precisely to convert an otherwise deterministic prediction into an estimate with quantified variance. The predictive-uncertainty decomposition (Gawlikowski, 2023) then attributes that variance to its aleatoric and epistemic components, so that the reader can see how much of it further evidence could reduce.

The honest reporting of variance closes the loop with the calibrated language of Section 1.4.1. A finding's likelihood term derives from the ensemble fraction; its confidence term derives, in part, from the variance structure behind that fraction and from the number of independent trials supporting it. A probability estimated from a single deterministic run, however stable, cannot earn high confidence, because it has no trial variance to characterize; a probability estimated from a diverse, well-sampled ensemble with corrected effective size can. By making the variance structure explicit and by refusing to present determinism as robustness, the assessment ensures that its confidence language reflects the true independence of the evidence rather than the superficial stability of repeated identical computations. This is the final safeguard against the overconfidence that an uncritical reading of the ensemble would otherwise invite.

1.6 Assessment Workflow, Cross-Cutting Methods, and Roadmap

The end-to-end assessment workflow

The assessment follows a fixed end-to-end workflow — from evidence intake through constraint, aggregation, uncertainty propagation, and calibrated reporting — and fixing the workflow is what makes the volume’s conclusions reproducible and auditable rather than a curated selection (*high confidence*).

The workflow begins with evidence intake under the source-first discipline of Section 1.3.6: every product, model output, and published analysis is registered with its provenance and uncertainty before it can be used. Intake is followed by the integrity audit of Section 1.4.3, which screens references for silent-failure signatures and defines the clean subsets on which comparisons will be run. Only after intake and audit does the substantive assessment begin, so that no finding rests on unregistered or unaudited evidence.

The substantive stages proceed in a fixed order. Physical-model members and machine-learning members are generated or ingested (Section 1.3.2, 1.3.3); each member is scored against the structured observational constraints of Section 1.2 and weighted accordingly, with the anti-double-counting bookkeeping of Section 1.3.4 enforced; the weighted members are aggregated into a constrained ensemble whose effective size is corrected for shared structure; uncertainties are propagated through every operation (Section 1.3.5); and the resulting distributions are mapped onto calibrated language (Section 1.4.1) after passing the falsification discipline (Section 1.4.2) and the variance-honesty check (Section 1.4.5). The workflow is deliberately linear in its dependencies so that the provenance of any conclusion can be traced backward through it, and any conclusion whose chain is broken — by an unaudited reference, an unvalidated tool, or a failed falsification test — is either repaired or reported with reduced confidence.

The workflow is instrumented so that its execution can be verified independently. Each assessed statement carries a record of which members supported it, which constraints weighted them, which falsification tests it passed, and which uncertainty terms were propagated into it. This instrumentation is the operational meaning of the source-first build: a reader with access to the record can reconstruct the path from evidence to conclusion and check each step. The volume treats this reconstructability as a non-negotiable property, because a calibrated confidence statement whose derivation cannot be reconstructed is indistinguishable from an assertion, and the assessment’s entire value rests on the difference between the two.

Cross-cutting methodological themes

Several methodological themes recur across every domain of the assessment — the primacy of coupled representation, the discipline of common footing,

the correction for shared structure, and the separation of reducible from irreducible uncertainty — and stating them once here allows later chapters to invoke them without re-deriving them (*high confidence*). The primacy of coupled representation, established in Section 1.2.4, means that terrestrial, hydrological, cryospheric, and atmospheric evidence are assessed jointly rather than in isolation, because their state variables move coherently (Winkler, 2021) (Scanlon, 2023) (Hassani, 2021) (Jane, 2021) (Moore, 2021). Later chapters treat this as a default: a finding about one component that is inconsistent with the coupled behavior of the others is flagged, because the coupling is a constraint that individual-component analyses can violate only at the cost of physical consistency.

The discipline of common footing, established in Section 1.1.5 and operationalized in Section 1.3.1, recurs whenever products of differing resolution, baseline, or uncertainty representation are compared. The correction for shared structure, established in Section 1.3.4, recurs whenever apparently independent lines of evidence are combined, and it draws throughout on the recognition that shared method inflates apparent agreement (Podsakoff, 2023). The separation of reducible epistemic from irreducible aleatoric uncertainty, established in Section 1.3.5 using the formal decomposition for learned models (Gawlikowski, 2023), recurs whenever the assessment states what further observation could and could not resolve. These four themes are the methodological spine of the volume, and their consistent application is what gives the assessment its internal coherence.

A fifth theme is the treatment of machine learning as an integrated but governed component rather than a separate track. Across domains, machine-learning methods appear as emulators, downscalers, and discovery tools (Bi, 2023) (Karniadakis, 2021) (Cuomo, 2022) (Willard, 2022) (Tuia, 2022) (Kaack, 2022), and in every case they are subject to the same physical-consistency, validation, and disclosure requirements (Section 1.4.4), drawing on the trustworthy-AI framework (Díaz-Rodríguez, 2023) and the honest accounting of their failure modes (Gupta, 2023) (Zhao, 2026). The assessment's position is that machine learning has become indispensable for sampling the ensemble densely enough to characterize its tails, and precisely because it is indispensable it must be governed rigorously; the two propositions are complementary, not in tension.

The evidentiary base and its breadth

The evidentiary base of this volume is deliberately broad, spanning physical, biological, methodological, and technological literatures, because the coupled Earth-system framing requires evidence from every domain that materially constrains the ensemble (*high confidence*). The physical and Earth-system evidence anchors the assessment of state and drivers: planetary boundaries (Richardson, 2023), the carbon budget (Friedlingstein, 2022) (Friedlingstein, 2023), regional amplification (Rantanen, 2022; Pepin, 2022), and cryospheric and hydrological change (Moore, 2021) (Scanlon, 2023) (Hassani, 2021) (Jane, 2021). The biological and ecological evidence establishes the biosphere as an ac-

tive driver (Folke, 2021) and constrains the terrestrial system through soil-carbon process understanding (Kleber, 2021) (Angst, 2021) (Georgiou, 2022), land-use dynamics (Winkler, 2021; Potapov, 2021), and the responses of pathogens, pests, and agriculture to warming (Singh, 2023) (Jägermeyr, 2021) (Tang, 2021).

The methodological evidence supplies the tools: coordinated scenario design (Tebaldi, 2021), physics-informed and data-driven learning (Karniadakis, 2021) (Cuomo, 2022) (Bi, 2023) (Willard, 2022), ensemble methods (Ganaie, 2022), uncertainty quantification (Gawlikowski, 2023), the general deep-learning toolkit (Ahmed, 2023) (Zhao, 2024) (Deng, 2023) (Hollmann, 2025), and the methodological literature on bias and trustworthiness (Podsakoff, 2023; Díaz-Rodríguez, 2023). The technological and applied evidence constrains the plausibility of forcing pathways and the breadth of the downscaling toolkit: energy-carrier and conversion chemistry (Cullen, 2021) (Valera-Medina, 2021) (Monteiro, 2021) (Yang, 2023), and applied machine learning in agriculture and the built environment (Benos, 2021) (Baduge, 2022). The assessment also situates its findings within the wider evidence on human consequences where that evidence validates the physical characterization of hazards (Romanello, 2022) (Romanello, 2023) (Zhao, 2021) (Aguilera, 2021) and within the economic context that motivates the assessment's precision (Rennert, 2022).

The breadth of this base is a methodological commitment, not an accident of collection. Because the Variant Earths framing defines a member by its full, coupled state — physical, biological, and, through forcing pathways, technological — the evidence that constrains a member must be correspondingly broad. A narrow evidentiary base would leave large parts of each member unconstrained and would produce an ensemble that is wide in the wrong places, permissive where evidence exists and falsely confident where it does not. The assessment's breadth is the means by which the ensemble is constrained everywhere the evidence allows, and the source-first discipline ensures that each strand of this broad base is traceable and auditable.

Key limitations and known failure modes

The assessment states its principal limitations explicitly, because a methods chapter that concealed its own failure modes would violate the disclosure discipline it imposes on every finding (*high confidence*). The first limitation is structural uncertainty that no member represents. Where a mechanism is absent from the entire ensemble — a threshold process such as the sea-ice arch collapse (Moore, 2021) that few or no members support — the ensemble width understates the true uncertainty, and the assessment can flag but not quantify the shortfall. This is an irreducible limitation of any ensemble-based method, and the volume handles it by reporting such cases as statements of possibility with explicit low confidence rather than as quantified likelihoods.

The second limitation is the non-independence of the evidentiary base. Physical models share components, emulators are trained on model output, and methodological choices propagate across studies, so the effective number of independent lines of evidence is smaller than the nominal count. The assessment

corrects for this where it can estimate the correlation structure (Section 1.3.4), drawing on the recognition that shared method inflates apparent agreement (Podsakoff, 2023), but the correction is itself uncertain, and residual overconfidence from unrecognized shared structure cannot be excluded. The third limitation is the maturity of the machine-learning components. Data-driven methods are advancing rapidly (Bi, 2023) (Ahmed, 2023) (Zhao, 2024) (Hollmann, 2025), which means both that their capabilities are growing and that their failure modes are incompletely characterized; the assessment governs them conservatively (Section 1.4.4) precisely because their long-run reliability under distribution shift is not yet fully established, a caution reinforced by the documented failure modes of generative systems (Gupta, 2023) (Zhao, 2026).

The fourth limitation is the incompleteness and heterogeneity of the observational record. The anchor member is itself uncertain, observational products carry spatially varying uncertainty (Poggio, 2021; Potapov, 2021), and coverage is uneven in space and time, so the constraints that weight the ensemble are stronger in well-observed regions and periods than in poorly observed ones. The assessment reports this heterogeneity rather than smoothing over it, and it refrains from claiming uniform confidence across regions where the observational constraint is uneven. Taken together, these limitations define the honest boundary of what the volume can assert: the assessment is strongest where the evidence is broad, independent, well-observed, and mechanistically represented, and it says so, and it correspondingly reduces its confidence where any of these conditions fails.

Roadmap to the volume

The remaining chapters apply the framing, evidentiary base, and methods established here to the successive components of the Variant Earths ensemble, and the reader should carry forward from this chapter the conventions that govern every subsequent assessment statement (*high confidence*). Subsequent chapters assess the observed state of each component in detail, attribute the observed changes to their drivers by comparing the observed member against counterfactual members, and project the evolution of the ensemble under the scenario-defined pathways (Tebaldi, 2021). Each of these tasks uses the machinery of this chapter without re-deriving it: the coupled representation, the common-footing discipline, the weighted-ensemble aggregation with shared-structure correction, the propagated uncertainty, and the calibrated language.

The reader should carry forward four conventions in particular. First, every likelihood term denotes a quantified probability derived from the corrected ensemble fraction, and every confidence term denotes assessed evidential robustness, and the two are distinct (Section 1.4.1). Second, every finding has passed the falsification discipline and the integrity audit before receiving its calibrated language (Section 1.4.2, 1.4.3), and where it has not, its reduced confidence is stated. Third, machine-learning-derived findings are governed, validated, and disclosed (Section 1.4.4), and their contribution to uncertainty is propagated rather than hidden (Gawlikowski, 2023). Fourth, deterministic

procedures are never presented as robust on the strength of their repeatability (Section 1.4.5). These conventions are the connective tissue of the volume, and their consistent application is what allows the individual chapters to combine into a single, coherent, uncertainty-explicit assessment of the variant-Earths ensemble.

The volume closes the loop opened in Section 1.1: the observed Earth is one member of a structured ensemble, the assessment's task is to characterize that ensemble — its observed member, its counterfactual members, and its projected members — under a common calibrated vocabulary, and the methods established here are the means by which that characterization is made traceable, reproducible, and honest about its own uncertainty. Everything that follows is an application of this framing to the evidence, and the reader is now equipped to read the subsequent chapters as such.

1.7 Counterfactual Construction, Attribution, and Scenario Design

Counterfactual members and the logic of attribution

Attribution in this volume is defined as a comparison between the observed member and counterfactual members of the ensemble that differ only in a specified driver, and defining attribution this way makes it a statement about the ensemble rather than about an isolated trajectory (*high confidence*). A counterfactual member is a physically consistent Earth state constructed under an altered forcing history — most importantly, a history without the anthropogenic carbon perturbation quantified in the global carbon budget (Friedlingstein, 2022) (Friedlingstein, 2023). The attributable effect of a driver on a given quantity is then the difference in the distribution of that quantity between the members with and without the driver, and this difference inherits the uncertainty structure of the ensemble, including the corrections for shared structure and emulation error established in Section 1.3.

This construction has the advantage that attribution and projection use the same machinery. Both are statements about how the ensemble's distribution shifts when a specified condition is changed: for attribution, the condition is the presence or absence of a historical driver; for projection, it is the choice of a future forcing pathway. Because the two share the ensemble representation and the calibrated vocabulary, the volume can report attribution and projection statements that are mutually consistent by construction, rather than reconciling them after the fact. The regional structure of the forced signal — Arctic amplification (Rantanen, 2022), elevation dependence (Pepin, 2022) — is reproduced in the counterfactual comparison, so that attribution statements are spatially resolved rather than global-mean only.

The counterfactual logic is applied with particular care to extremes, where the attributable change is often a change in probability rather than in mean.

For a compound event such as the heat–drought–fire sequences discussed in Section 1.2.3 (Yuan, 2023) (Abram, 2021), the attributable effect is the change in the fraction of members exhibiting the joint exceedance, and this fraction-based statement is exactly the quantity the ensemble framing makes well-defined. The assessment reports such changes as changes in likelihood, with the confidence in that likelihood determined by the width and independence of the ensemble supporting it, following the dual-reporting convention of Section 1.4.1.

Constructing physically consistent counterfactuals

Counterfactual members must be physically consistent Earth states, not merely observed fields with a driver subtracted, and enforcing physical consistency is where the methods of Section 1.3 are most heavily exercised (*high confidence*). A naive counterfactual — the observed field minus an estimated forced response — is not a valid member, because it need not satisfy the conservation laws and dynamical balances that constrain a real Earth state. The assessment therefore constructs counterfactuals by running or emulating the coupled system under the altered forcing, using coordinated model configurations (Tebaldi, 2021) and, where computational cost is prohibitive, physics-informed or data-driven emulators that respect the governing constraints (Karniadakis, 2021) (Cuomo, 2022) (Bi, 2023). The emulator’s own uncertainty is propagated into the counterfactual distribution (Gawlikowski, 2023), so that the attributable effect carries the full uncertainty of its construction.

Physical consistency is especially demanding for the terrestrial and hydrological components, whose counterfactual states must remain consistent with soil-carbon dynamics (Kleber, 2021) (Angst, 2021) (Georgiou, 2022), land-use configuration (Winkler, 2021; Potapov, 2021), and the coupled water cycle (Scanlon, 2023) (Hassani, 2021). A counterfactual that altered the atmospheric forcing without correspondingly adjusting the coupled terrestrial state would violate the coupling established in Section 1.2.4, and the assessment rejects such members. This is the counterfactual face of the coupled-representation theme: attribution is credible only if the counterfactual member is coupled as thoroughly as the observed member.

The assessment validates its counterfactuals against the same structured constraints used to weight the forward ensemble. A counterfactual member is admissible only if it reproduces the observed relationships among components in regimes where the driver’s effect is small, because a member that cannot match the observed system where the perturbation is weak provides no basis for trusting its behavior where the perturbation is strong. This validation reuses the observational constraints of Section 1.2 under the anti-double-counting rule of Section 1.3.4, so that the constraints used to admit a counterfactual are not also credited as independent confirmation of the attribution it supports.

Scenario design and the pathway subspaces

Future members of the ensemble are organized into subspaces defined by shared, coordinated emissions scenarios, and using coordinated scenarios is what allows projection uncertainty to be decomposed into scenario, model, and internal-variability contributions (*high confidence*). The scenario framework provides a set of forcing pathways spanning a range of plausible futures, and each pathway defines the subspace of ensemble members consistent with it (Tebaldi, 2021). Projection uncertainty then decomposes naturally: variation across scenarios reflects the choice of pathway, variation across models within a scenario reflects structural uncertainty, and variation across initial conditions within a model reflects internal variability. The assessment reports this decomposition rather than a single lumped range, because the three components behave differently and are relevant to different questions — scenario uncertainty is a matter of societal choice, structural uncertainty is a matter of scientific knowledge, and internal variability is irreducible.

The plausibility of each pathway is constrained by evidence on the feasibility of the underlying transitions, consistent with the scope boundary of Section 1.1.4. The sectoral structure of emissions (Lamb, 2021) determines which pathways are reachable from the present state, and the technological evidence base — low-carbon energy carriers and conversion chemistry (Cullen, 2021) (Valera-Medina, 2021) (Monteiro, 2021) (Yang, 2023) — constrains the rate at which sectoral emissions can plausibly decline. The assessment does not assign probabilities to the pathways themselves, because their realization depends on societal choices outside the physical-science remit; it assesses, conditional on each pathway, the distribution of physical outcomes across the corresponding ensemble subspace.

Scenario design also governs how the assessment handles the deep-future and the deep-tail. Because the impact-relevant risk often lies in the tails (Section 1.2.3), the assessment samples each pathway subspace densely enough to characterize its tail behavior, which is feasible only because fast emulators (Bi, 2023) make large ensembles affordable. The tail members are subject to the same falsification and consistency requirements as the central members, so that a low-probability, high-impact member is admitted only if it is physically consistent and survives the falsification discipline, never merely because it is extreme. This prevents the tails from being populated by artifacts while ensuring that genuine tail risk is represented.

Bringing attribution and projection into a single frame

The volume reports attribution and projection within a single ensemble frame so that the observed, counterfactual, and projected members are directly comparable, and this unification is the central methodological payoff of the Variant Earths framing (*high confidence*). Because all three kinds of member — observed, counterfactual, projected — are represented in the same coupled, uncertainty-quantified ensemble, statements about them are commensurable without further translation. The attributable warming to date, the counterfactual

warming that would have occurred without the anthropogenic perturbation, and the projected warming under each pathway are all differences of distributions within one ensemble, computed with one uncertainty calculus and reported in one calibrated vocabulary.

This unification pays off most clearly where attribution and projection meet: in statements about how an already-attributable change will evolve under a given pathway. The assessment can, for instance, express the projected change in a compound-extreme frequency (Yuan, 2023) (Abram, 2021) as a continuation of an attributable historical change, with the historical and future components computed consistently and their uncertainties propagated together. Such statements are difficult to make coherently when attribution and projection use separate frameworks, and they are natural in the ensemble frame. The regional structure carries through: because the ensemble reproduces amplification patterns (Rantanen, 2022; Pepin, 2022) in both the counterfactual and projected members, the joined attribution–projection statements are spatially resolved.

The unified frame also disciplines the reporting of confidence. A joined statement inherits the weaker of the confidence levels of its components: if the attribution is high-confidence but the projection under a given pathway is only medium-confidence, the joined statement is reported at medium confidence, and the reason is stated. This propagation of the weaker confidence prevents the unified frame from manufacturing false confidence by combining a strong component with a weak one, and it follows directly from the evidence-handling discipline of Section 1.4.2. The result is an assessment in which every statement — observational, attributive, or projective — is located in the same ensemble, quantified with the same calculus, and reported in the same honest, calibrated language, which is precisely the coherence the framing was introduced to achieve.

Observational networks, data stewardship, and continuity

The credibility of every ensemble statement ultimately rests on the continuity and stewardship of the observational networks that define the anchor member, and the assessment treats observational continuity as a methodological prerequisite rather than an external given (*high confidence*). The structured constraints that weight the ensemble — the Arctic amplification rate (Rantanen, 2022), the carbon-budget partitioning (Friedlingstein, 2022) (Friedlingstein, 2023), the terrestrial and hydrological state (Poggio, 2021) (Potapov, 2021) (Scanlon, 2023) — are only as strong as the underlying observing systems, and gaps or discontinuities in those systems propagate directly into weakened constraints and widened ensembles. The assessment therefore documents the provenance and continuity of each observational product it relies on, and it reports where a constraint is weakened by observational limitation rather than by genuine physical uncertainty.

Specialized regional and process-specific records are part of this stewardship. Long, homogenized event databases — for example, curated tropical-cyclone records maintained by national meteorological services (Lu, 2021) — provide

the multi-decadal continuity needed to characterize the statistics of high-impact events, and their careful construction is what allows the assessment to distinguish a genuine change in event statistics from an artifact of changing observing practice. The integrity audit of Section 1.4.3 applies with full force to such records, because a discontinuity in observing practice can masquerade as a physical trend, and the falsification discipline of Section 1.4.2 requires that any claimed trend be shown to survive correction for such discontinuities before it is credited.

Data stewardship also shapes what the assessment can say about the future of its own evidence base. Because the machine-learning components depend on training data drawn from the observational and model archives (Bi, 2023) (Karniadakis, 2021) (Willard, 2022), the continuity of those archives determines whether the emulators can be retrained and revalidated as the system evolves; an emulator trained on a now-discontinued observing system may degrade silently under distribution shift, a risk the governance of Section 1.4.4 is designed to catch. The assessment thus treats observational continuity, data stewardship, and machine-learning reliability as a single connected concern: the ensemble is only as trustworthy as the observations that anchor it, the archives that train its emulators, and the stewardship that keeps both intact over the multi-decadal horizons the assessment addresses.

Communicating the ensemble to downstream users

The assessment communicates the ensemble to downstream working groups and decision-relevant users as a distribution with quantified uncertainty and stated confidence, never as a single trajectory, because collapsing the ensemble to a point estimate discards exactly the tail and coupling information that downstream use requires (*high confidence*). The primary deliverable to downstream users is therefore not a number but a characterized distribution: the fraction of members exhibiting each outcome of interest, the confidence in that fraction, and the couplings among outcomes that determine joint risk. This is a deliberate design choice, because downstream applications — the health, water, and food-system assessments that build on the physical basis (Romanello, 2022) (Romanello, 2023) (Scanlon, 2023) (Singh, 2023) — are frequently sensitive to the tails and to the joint behavior of hazards rather than to central estimates alone.

The communication follows the two-axis vocabulary of Section 1.4.1 without simplification. A downstream user receives both the likelihood of an outcome and the confidence in that likelihood, so that a user sensitive to low-probability, high-impact members can distinguish a well-constrained probability from a poorly constrained one. The assessment resists pressure to collapse these two axes into a single headline figure, because doing so would strip the confidence information that is often decisive for risk-sensitive users. Where a single summary is unavoidable for communication, the assessment provides it alongside, not in place of, the full distribution and its confidence, and it states which information the summary omits.

The assessment is also explicit about what the ensemble cannot tell downstream users. The structural-uncertainty limitation of Section 1.5.4 means that outcomes depending on mechanisms absent from the ensemble are under-represented, and the assessment communicates these as flagged possibilities rather than allowing their absence from the ensemble to be read as evidence of their improbability. Likewise, the heterogeneity of observational constraint means that the ensemble is better constrained in some regions and periods than others, and the assessment communicates this spatial and temporal variation in confidence rather than presenting a uniform reliability the evidence does not support. The final methodological commitment of this chapter is therefore a communicative one: the ensemble is delivered downstream with its uncertainty, its confidence, its couplings, and its limitations intact, because an assessment that concealed any of these in the interest of a cleaner message would betray the calibrated honesty that is its entire purpose. This commitment binds every chapter that follows: the observed member, its counterfactual siblings, and its projected descendants are reported as one uncertainty-quantified ensemble, in one calibrated vocabulary, with one auditable chain from evidence to conclusion. The framing, context, and methods set out in this chapter are, in the end, a single instrument for keeping that promise, and the remainder of the volume is their application to the physical science basis of the variant-Earths ensemble.

1.8 References

- Abram, N.J., et al. (2021). Connections of climate change and variability to large and extreme forest fires in southeast Australia. *Communications Earth & Environment*. <https://doi.org/10.1038/s43247-020-00065-8>
- Aguilera, R., et al. (2021). Wildfire smoke impacts respiratory health more than fine particles from other sources: observational evidence from Southern California. *Nature Communications*. <https://doi.org/10.1038/s41467-021-21708-0>
- Ahmed, S.F., et al. (2023). Deep learning modelling techniques: current progress, applications, advantages, and challenges. *Artificial Intelligence Review*. <https://doi.org/10.1007/s10462-023-10466-8>
- Angst, G., et al. (2021). Plant- or microbial-derived? A review on the molecular composition of stabilized soil organic matter. *Soil Biology and Biochemistry*. <https://doi.org/10.1016/j.soilbio.2021.108189>
- Baduge, S.K., et al. (2022). Artificial intelligence and smart vision for building and construction 4.0: Machine and deep learning methods and applications. *Automation in Construction*. <https://doi.org/10.1016/j.autcon.2022.104440>
- Benos, L., et al. (2021). Machine Learning in Agriculture: A Comprehensive Updated Review. *Sensors*. <https://doi.org/10.3390/s21113758>
- Bi, K., et al. (2023). Accurate medium-range global weather forecasting with 3D neural networks. *Nature*. <https://doi.org/10.1038/s41586-023-06185-3>

- Budhwar, P., et al. (2023). Human resource management in the age of generative artificial intelligence: Perspectives and research directions on ChatGPT. *Human Resource Management Journal*. <https://doi.org/10.1111/1748-8583.12524>
- Cullen, D.A., et al. (2021). New roads and challenges for fuel cells in heavy-duty transportation. *Nature Energy*. <https://doi.org/10.1038/s41560-021-00775-z>
- Cuomo, S., et al. (2022). Scientific Machine Learning Through Physics-Informed Neural Networks: Where we are and What's Next. *Journal of Scientific Computing*. <https://doi.org/10.1007/s10915-022-01939-z>
- Deng, B., et al. (2023). CHGNet as a pretrained universal neural network potential for charge-informed atomistic modelling. *Nature Machine Intelligence*. <https://doi.org/10.1038/s42256-023-00716-3>
- Díaz-Rodríguez, N., et al. (2023). Connecting the dots in trustworthy Artificial Intelligence: From AI principles, ethics, and key requirements to responsible AI systems and regulation. *Information Fusion*. <https://doi.org/10.1016/j.inf-fus.2023.101896>
- Folke, C., et al. (2021). Our future in the Anthropocene biosphere. *AMBIO*. <https://doi.org/10.1007/s13280-021-01544-8>
- Friedlingstein, P., et al. (2022). Global Carbon Budget 2022. *Earth system science data*. <https://doi.org/10.5194/essd-14-4811-2022>
- Friedlingstein, P., et al. (2023). Global Carbon Budget 2023. *Earth system science data*. <https://doi.org/10.5194/essd-15-5301-2023>
- Ganaie, M.A., et al. (2022). Ensemble deep learning: A review. *Engineering Applications of Artificial Intelligence*. <https://doi.org/10.1016/j.engappai.2022.105151>
- Gawlikowski, J., et al. (2023). A survey of uncertainty in deep neural networks. *Artificial Intelligence Review*. <https://doi.org/10.1007/s10462-023-10562-9>
- Georgiou, K., et al. (2022). Global stocks and capacity of mineral-associated soil organic carbon. *Nature Communications*. <https://doi.org/10.1038/s41467-022-31540-9>
- Gupta, M., et al. (2023). From ChatGPT to ThreatGPT: Impact of Generative AI in Cybersecurity and Privacy. *IEEE Access*. <https://doi.org/10.1109/access.2023.3300381>
- Hassani, A., et al. (2021). Global predictions of primary soil salinization under changing climate in the 21st century. *Nature Communications*. <https://doi.org/10.1038/s41467-021-26907-3>
- Hollmann, N., et al. (2025). Accurate predictions on small data with a tabular foundation model. *Nature*. <https://doi.org/10.1038/s41586-024-08328-6>
- Jane, S.F., et al. (2021). Widespread deoxygenation of temperate lakes. *Nature*. <https://doi.org/10.1038/s41586-021-03550-y>
- Jägermeyr, J., et al. (2021). Climate impacts on global agriculture emerge earlier in new generation of climate and crop models. *Nature Food*. <https://doi.org/10.1038/s43021-00400-y>
- Kaack, L.H., et al. (2022). Tackling Climate Change with Machine Learning. *OPUS 4* (Zuse Institute Berlin). <https://doi.org/10.1145/3485128>
- Karniadakis, G.E., et al. (2021). Physics-informed machine learning. *Nature Reviews Physics*. <https://doi.org/10.1038/s42254-021-00314-5>

Kleber, M., et al. (2021). Dynamic interactions at the mineral–organic matter interface. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-021-00162-y>

Lamb, W.F., et al. (2021). A review of trends and drivers of greenhouse gas emissions by sector from 1990 to 2018. *Environmental Research Letters*. <https://doi.org/10.1088/1748-9326/abee4e>

Lu, X., et al. (2021). Western North Pacific Tropical Cyclone Database Created by the China Meteorological Administration. *Advances in Atmospheric Sciences*. <https://doi.org/10.1007/s00376-020-0211-7>

Monteiro, M.C.O., et al. (2021). Absence of CO₂ electroreduction on copper, gold and silver electrodes without metal cations in solution. *Nature Catalysis*. <https://doi.org/10.1038/s41929-021-00655-5>

Moore, G.W.K., et al. (2021). Anomalous collapses of Nares Strait ice arches leads to enhanced export of Arctic sea ice. *Nature Communications*. <https://doi.org/10.1038/s41467-020-20314-w>

Pepin, N.C., et al. (2022). Climate Changes and Their Elevational Patterns in the Mountains of the World. *Reviews of Geophysics*. <https://doi.org/10.1029/2020rg000730>

Podsakoff, P.M., et al. (2023). Common Method Bias: It's Bad, It's Complex, It's Widespread, and It's Not Easy to Fix. *Annual Review of Organizational Psychology and Organizational Behavior*. <https://doi.org/10.1146/annurev-orgpsych-110721-040030>

Poggio, L., et al. (2021). SoilGrids 2.0: producing soil information for the globe with quantified spatial uncertainty. *SOIL*. <https://doi.org/10.5194/soil-7-217-2021>

Potapov, P., et al. (2021). Global maps of cropland extent and change show accelerated cropland expansion in the twenty-first century. *Nature Food*. <https://doi.org/10.1038/s43017-021-00429-z>

Rantanen, M., et al. (2022). The Arctic has warmed nearly four times faster than the globe since 1979. *Communications Earth & Environment*. <https://doi.org/10.1038/s43247-022-00498-3>

Rennert, K., et al. (2022). Comprehensive evidence implies a higher social cost of CO₂. *Nature*. <https://doi.org/10.1038/s41586-022-05224-9>

Richardson, K., et al. (2023). Earth beyond six of nine planetary boundaries. *Science Advances*. <https://doi.org/10.1126/sciadv.adh2458>

Romanello, M., et al. (2022). The 2022 report of the Lancet Countdown on health and climate change: health at the mercy of fossil fuels. *The Lancet*. [https://doi.org/10.1016/s0140-6736\(22\)01540-9](https://doi.org/10.1016/s0140-6736(22)01540-9)

Romanello, M., et al. (2023). The 2023 report of the Lancet Countdown on health and climate change: the imperative for a health-centred response in a world facing irreversible harms. *The Lancet*. [https://doi.org/10.1016/s0140-6736\(23\)01859-7](https://doi.org/10.1016/s0140-6736(23)01859-7)

Scanlon, B.R., et al. (2023). Global water resources and the role of groundwater in a resilient water future. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-022-00378-6>

- Singh, B.K., et al. (2023). Climate change impacts on plant pathogens, food security and paths forward. *Nature Reviews Microbiology*. <https://doi.org/10.1038/s41579-023-00900-7>
- Tang, F.H.M., et al. (2021). Risk of pesticide pollution at the global scale. *Nature Geoscience*. <https://doi.org/10.1038/s41561-021-00712-5>
- Tebaldi, C., et al. (2021). Climate model projections from the Scenario Model Intercomparison Project (ScenarioMIP) of CMIP6. *Earth System Dynamics*. <https://doi.org/10.5194/esd-12-253-2021>
- Tuia, D., et al. (2022). Perspectives in machine learning for wildlife conservation. *Nature Communications*. <https://doi.org/10.1038/s41467-022-27980-y>
- Valera-Medina, A., et al. (2021). Review on Ammonia as a Potential Fuel: From Synthesis to Economics. *Energy & Fuels*. <https://doi.org/10.1021/acs.energyfuels.0c03685>
- Willard, J., et al. (2022). Integrating Scientific Knowledge with Machine Learning for Engineering and Environmental Systems. *ACM Computing Surveys*. <https://doi.org/10.1145/3514228>
- Winkler, K., et al. (2021). Global land use changes are four times greater than previously estimated. *Nature Communications*. <https://doi.org/10.1038/s41467-021-22702-2>
- Yang, Y., et al. (2023). Operando studies reveal active Cu nanograins for CO₂ electroreduction. *Nature*. <https://doi.org/10.1038/s41586-022-05540-0>
- Yuan, X., et al. (2023). A global transition to flash droughts under climate change. *Science*. <https://doi.org/10.1126/science.abn6301>
- Zhao, Q., et al. (2021). Global, regional, and national burden of mortality associated with non-optimal ambient temperatures from 2000 to 2019: a three-stage modelling study. *The Lancet Planetary Health*. [https://doi.org/10.1016/s2542-5196\(21\)00081-4](https://doi.org/10.1016/s2542-5196(21)00081-4)
- Zhao, X., et al. (2024). A review of convolutional neural networks in computer vision. *Artificial Intelligence Review*. <https://doi.org/10.1007/s10462-024-10721-6>
- Zhao, W.X., et al. (2026). A Survey of Large Language Models. *Frontiers of Computer Science*. <https://doi.org/10.1007/s11704-026-60308-3>

Chapter 2

Chapter 2: Large-Scale Changes in the Climate System and Their Causes

2.1 Executive Summary

The climate system as a whole has continued to warm, and the imprint of human influence is now detectable across every major component — atmosphere, ocean, cryosphere and biosphere — with multiple independent lines of evidence converging on the same conclusion (*virtually certain*). Global-scale indicators of change have strengthened since the previous assessment cycle, and the pattern of change across the components is physically consistent with the accumulation of energy in the Earth system driven by rising atmospheric greenhouse gas concentrations (*high confidence*).

Atmospheric concentrations of carbon dioxide, methane and nitrous oxide have continued to rise, and the growth of the airborne carbon dioxide burden tracks the sum of fossil and land-use emissions closely enough that the anthropogenic origin of the increase is beyond scientific dispute (*virtually certain*). The continuously updated global carbon budget shows that fossil carbon dioxide emissions and the resulting atmospheric growth rate remained near record levels through the most recent budget years (Friedlingstein et al., 2022; Friedlingstein et al., 2023; Friedlingstein et al., 2025).

Surface warming has been geographically pervasive but strongly amplified at high northern latitudes, where the Arctic has warmed nearly four times faster than the global mean since 1979 (*high confidence*). This Arctic amplification is now well quantified from multiple observational datasets and exceeds the value reproduced by many climate models, indicating that the polar response to forcing may be systematically underestimated in the model ensemble (Rantanen et al., 2022).

Mountain regions across the world show elevation-dependent warming and rapid loss of ice, integrating change across the atmosphere and cryosphere (*high confidence*). High Mountain Asian glaciers have thinned and retreated over the multi-decadal satellite record (Bhattacharya et al., 2021), and the mountain cryosphere more broadly is warming in patterns that vary systematically with elevation (Pepin et al., 2022).

The ocean has warmed, and marine heatwaves have become more frequent, more intense and longer-lasting, with pervasive biological consequences (*high confidence*). The biological impacts of marine heatwaves now extend from individual organisms to whole ecosystems and biogeochemical cycles (Smith et al., 2022).

The continental water cycle is intensifying and becoming more variable, with more severe droughts and floods and a widespread decline in terrestrial water storage (*high confidence*). Satellite observations reveal a widespread decline in global lake water storage (Yao et al., 2023), the twenty-first-century southwestern North American drought has intensified into a megadrought unprecedented in centuries (Williams et al., 2022), and an anthropogenic fingerprint is now detectable in global drought frequency, duration and intensity (Chiang et al., 2021).

Human influence on the causes of individual extreme events is increasingly demonstrable through formal attribution, moving detection and attribution from the global mean to the scale of specific heatwaves, droughts and wildfire seasons (*high confidence*). Attribution science has matured into a systematic capability that quantifies how anthropogenic warming has changed the probability and severity of extreme weather (Clarke et al., 2022) (Barriopedro et al., 2023).

The biosphere is responding to and, in places, losing resilience under the combined pressures of warming, drought and disturbance (*medium-to-high confidence*). Emerging signals of declining forest resilience are visible in global vegetation observations (Forzieri et al., 2022), hotter droughts leave a distinctive fingerprint on tree mortality worldwide (Hammond et al., 2022), and centuries of wetland loss have removed a large fraction of a globally important carbon and water reservoir.

Several large-scale components of the climate system are approaching thresholds beyond which change may become self-sustaining, and the interactions among these tipping elements raise the risk of cascading change (*medium confidence*). Interacting tipping elements can increase the risk of climate domino effects under sustained warming (Wunderling et al., 2021), and an assessment of planetary boundaries finds that the Earth system has now transgressed six of nine, including climate change (Richardson et al., 2023).

2.2 Introduction and Scope

This chapter assesses observed large-scale changes across the physical climate system and evaluates the evidence that human activity is their dominant cause. It synthesises changes in the atmosphere, ocean, cryosphere and terrestrial biosphere; the detection and attribution methodologies that connect those changes to their drivers; the behaviour of the large-scale modes of internal variability that modulate them; and the paleoclimatic and long-timescale context within which the recent changes must be interpreted. The organising principle is

energetic: greenhouse gases added to the atmosphere reduce the efficiency with which the Earth radiates energy to space, the resulting imbalance is expressed as accumulating heat, and that heat is redistributed among the components of the system in ways that are individually observable and collectively coherent.

The assessment rests on a wider and more capable observing system than was available to earlier reports. Satellite records of the atmosphere, ocean surface, ice and vegetation now span multiple decades; reanalyses assimilate heterogeneous observations into physically consistent gridded estimates; and the continuous accounting of the global carbon cycle has become an annually updated community activity (Friedlingstein et al., 2022; Friedlingstein et al., 2023). Ocean and sea-ice reanalyses such as the Copernicus global reanalysis provide eddy-permitting reconstructions of the three-dimensional ocean state that underpin estimates of heat and freshwater change (Jean-Michel et al., 2021). Together these systems allow changes to be characterised not only in the global mean but in their spatial structure, seasonality and tails, where much of the impact resides.

A central theme of this chapter is that the confidence attached to statements of change and cause now derives from the *convergence* of independent evidence rather than from any single indicator. Warming of the lower atmosphere, warming and expansion of the ocean, retreat of glaciers and sea ice, poleward and upslope shifts of climatic zones, and the intensification of the water cycle are not separate findings to be weighed against one another; they are complementary expressions of a single, forced perturbation. Where a component appears anomalous — for example, where the observed Arctic amplification exceeds the modelled value (Rantanen et al., 2022) — the discrepancy is treated as information about process understanding, not as grounds for doubting the overall conclusion.

The chapter also foregrounds the increasing granularity of attribution. Detection and attribution once operated almost exclusively at the level of the global-mean temperature; it now routinely addresses regional trends, individual seasons and specific extreme events (Clarke et al., 2022). This shift has consequences for how the assessment is written: calibrated statements can be made not only about long-term global change but about the human contribution to particular droughts, heatwaves and fire seasons that populate the public record of climate impacts.

Finally, the chapter situates the recent instrumental record against the deeper past. Paleoclimate archives establish that the current concentrations of long-lived greenhouse gases and the current rate of warming are unusual in the context of many millennia, and that the climate system contains slow components whose full response to present-day forcing has not yet been realised (Hansen et al., 2023). This long view is essential for interpreting the committed changes already set in motion and for bounding the plausibility of the trajectories assessed elsewhere in this volume.

2.3 Changes in the Atmosphere

Atmospheric composition

Atmospheric concentrations of the major long-lived greenhouse gases have continued their sustained rise, and the increase in atmospheric carbon dioxide is quantitatively consistent with cumulative anthropogenic emissions (*virtually certain*). The global carbon budget, reconciled annually from fossil-fuel statistics, land-use change estimates, atmospheric growth measurements and independent ocean and land sink models, shows that anthropogenic carbon dioxide emissions have remained near their historical maximum in recent years and that the atmospheric burden has grown accordingly (Friedlingstein et al., 2022; Friedlingstein et al., 2023; Friedlingstein et al., 2025). The closure of this budget — the agreement between the top-down atmospheric growth rate and the bottom-up sum of sources and sinks — is one of the strongest available demonstrations that the observed composition change is anthropogenic in origin, because it requires the human source term to balance the books.

The partitioning of emitted carbon among the atmosphere, ocean and land is itself changing. The land and ocean sinks together continue to absorb roughly half of annual emissions, but the efficiency and stability of these sinks vary interannually with climate and are not guaranteed to persist (Friedlingstein et al., 2023). Regional analyses show that some terrestrial systems have strengthened as carbon sinks over recent decades — for example, forest expansion has dominated the land carbon sink in China since 1980 (Yu et al., 2022) — while others are being eroded by drought, fire and disturbance, a tension developed further in Section 2.5.

Beyond carbon dioxide, satellite spectrometry has transformed the monitoring of shorter-lived and reactive constituents. Retrievals of tropospheric nitrogen dioxide and methane from the Sentinel-5P instrument, validated against ground-based networks, now provide spatially resolved, near-daily maps of atmospheric composition that reveal emission hotspots, plumes and trends (Verhoelst et al., 2021) (Lorente et al., 2021). These records are increasingly used to attribute composition changes to specific source sectors and regions, complementing the global carbon accounting with a finer spatial view.

Surface temperature and its spatial structure

Global surface warming has continued, and its spatial pattern is strongly polar-amplified, with the Arctic warming at nearly four times the global rate since 1979 (*high confidence*). The magnitude of Arctic amplification has been quantified across multiple observational datasets, yielding a robust factor near four that substantially exceeds the roughly two-to-three-fold amplification commonly cited in earlier literature and reproduced by many models (Rantanen et al., 2022). The physical drivers — loss of sea ice and the associated surface-albedo and heat-exchange feedbacks, changes in atmospheric heat transport, and lapse-rate and cloud feedbacks specific to the polar night — are individually

understood, but their combination in the observations implies that the coupled model ensemble may systematically underestimate the sensitivity of the Arctic to forcing.

Amplification is not confined to the poles. Mountain regions warm in ways that depend systematically on elevation, so that high-altitude environments frequently experience faster temperature change than adjacent lowlands, with consequences for snow, ice and the timing of meltwater (Pepin et al., 2022). Semi-arid and continental interiors, such as the Eastern Mediterranean and Middle East, are warming faster than the global average and are projected to experience a disproportionate increase in weather extremes (Zittis et al., 2022). The emergent picture is of a warming that is global in sign but highly structured in magnitude, concentrating the largest changes in the very regions — high latitudes, high altitudes, dry subtropics — whose cryospheric and hydrological systems are most sensitive.

The atmospheric water cycle and circulation

The atmospheric branch of the water cycle is intensifying, increasing the potential for both heavy precipitation and evaporative drought, often in the same regions at different times (*high confidence*). A warmer atmosphere holds more water vapour, strengthening the transport and convergence of moisture and sharpening the contrast between wet and dry regimes. Trends in precipitation are found to determine the future occurrence of compound hot–dry events, linking the thermodynamic and hydrological responses (Bevacqua et al., 2022). High-resolution regional climate simulations at kilometre scale, which explicitly resolve deep convection, reproduce observed precipitation characteristics more faithfully than coarser models and strengthen confidence in projected changes to short-duration extremes (Ban et al., 2021).

Changes in atmospheric circulation modulate this thermodynamic intensification. Persistent circulation anomalies — blocking highs, stalled ridges and displaced jets — concentrate heat and suppress precipitation, producing the stagnant regimes that underlie the most damaging heatwaves and droughts (Barriopedro et al., 2023). The physical understanding of heatwaves has advanced to the point where their dynamical and thermodynamical ingredients can be separated, allowing the anthropogenic contribution to be assessed even for events whose immediate cause is a specific circulation pattern.

Atmospheric aerosols and dust

The atmospheric aerosol burden, including mineral dust mobilised from the world’s arid regions, is an active and changing component of the climate system that both responds to and influences large-scale change (*medium confidence*). The global cycle of desert dust is dominated by a small number of major source regions, and quantifying their relative contributions is essential for understanding dust’s radiative, biogeochemical and cryospheric effects (Kok et al., 2021). Dust deposited on snow and ice lowers surface albedo and accelerates

melt, coupling the aerosol and cryospheric components; dust transport also fertilises ocean and land ecosystems and modulates air quality far downwind of source regions. Because dust emission depends on surface wind, soil moisture and vegetation cover, all of which are changing, the dust cycle is both a diagnostic of large-scale change and a feedback upon it.

2.4 Changes in the Ocean

Ocean warming and heat content

The ocean has continued to gain heat, and because it stores the overwhelming majority of the excess energy trapped by greenhouse gases, ocean warming is the most reliable single indicator of planetary energy imbalance (*virtually certain*). Eddy-permitting global ocean reanalyses, which assimilate satellite altimetry, sea-surface temperature and the Argo profiling-float array into a physically consistent model, reconstruct the three-dimensional evolution of ocean temperature and reveal warming that penetrates from the surface into the deeper ocean over the reanalysis period (Jean-Michel et al., 2021). The sub-surface expression of warming matters because it governs thermal expansion and sea-level rise, sets the thermal environment for marine ecosystems, and represents a reservoir of committed change that will continue to influence the climate long after surface temperatures stabilise.

The uptake of heat is not spatially uniform. Warming is concentrated in the upper ocean and varies by basin, with some regions — western boundary current extensions, marginal seas and portions of the Southern Ocean — warming considerably faster than the global mean. This heterogeneity is dynamically important: it alters density gradients, modifies current systems and shifts the geography of biological productivity. Reanalyses that resolve mesoscale eddies capture these regional structures with a fidelity unavailable to earlier, coarser products, sharpening the assessment of where and how fast the ocean is changing (Jean-Michel et al., 2021).

Marine heatwaves

Marine heatwaves — prolonged periods of anomalously warm ocean temperatures — have become more frequent, more intense and longer-lasting, and their biological consequences now propagate through marine ecosystems from physiology to biogeochemistry (*high confidence*). The accumulation of ocean heat raises the baseline upon which transient warm anomalies are superimposed, so that events that would once have been unremarkable now cross ecologically critical thresholds. A synthesis of the biological impacts of marine heatwaves documents effects spanning individual organism stress and mortality, shifts in species distributions, mass mortality events, harmful algal blooms, and disruption of the biogeochemical cycling that links the ocean to

the atmosphere (Smith et al., 2022). Because many marine organisms live close to their thermal limits, the same warming that is modest in the annual mean can be lethal when concentrated into a heatwave, making these events a leading edge of climate impact in the ocean.

The coupling between marine heatwaves and the atmosphere is bidirectional. Ocean warm anomalies can reinforce atmospheric ridging and, in coastal settings, contribute to terrestrial heat extremes, while atmospheric circulation anomalies and reduced cloud cover help generate and sustain the ocean anomalies in the first place. This coupling situates marine heatwaves within the broader assessment of compound and connected extremes developed in Section 2.6.

Sea level and the coastal zone

Sea level continues to rise as a combined consequence of ocean thermal expansion and the addition of water from melting land ice, and the resulting encroachment on coasts is among the most consequential large-scale changes for human systems (*high confidence*). Thermal expansion follows directly from the subsurface ocean warming documented above, while the mass contribution reflects the losses from glaciers and ice sheets assessed in Section 2.4. The coastal zone integrates these global signals with regional and local factors — land subsidence, sediment supply, and changes in storm and wave climate — to determine the actual exposure experienced at any given shoreline. Adaptation to sea-level rise has accordingly become a central concern of coastal management, requiring the translation of global projections into local planning under deep uncertainty (Griggs et al., 2021). The assessment emphasises that even where the global-mean rate is well constrained, the distribution of impact is set by the regional pattern of sea-level change and by the vulnerability of particular coastlines.

Continental waters as part of the hydrosphere

The storage of water on the continents — in lakes, wetlands, soils and groundwater — is declining across much of the world, redistributing water within the Earth system in ways that couple the ocean, atmosphere and land (*high confidence*). Satellite observations reveal a widespread decline in global lake water storage, affecting a majority of the world's large lakes and reflecting the combined influence of warming, altered precipitation and human water use (Yao et al., 2023). Mapping of global lake dynamics further shows that small lakes, though individually minor, collectively play an outsized and growing role in the terrestrial water and carbon budgets (Pi et al., 2022). Wetlands, a distinct and disproportionately important store of carbon and water, have been lost over centuries of drainage and conversion, with an estimated large fraction of the pre-industrial wetland area now gone. These continental changes are not peripheral to the ocean assessment: water removed from terrestrial storage ultimately reaches the sea, and the drying of the land surface feeds back on the atmosphere through altered evaporation and heat fluxes.

2.5 Changes in the Cryosphere

Arctic amplification and the sea-ice-covered ocean

The Arctic is warming far faster than the rest of the planet, and this amplification is the clearest regional signature of the coupled ice–ocean–atmosphere feedbacks that operate in a warming world (*high confidence*). The observed factor of nearly four relative to the global mean since 1979 exceeds most model estimates and points to feedbacks — principally the loss of reflective sea ice and the exposure of dark ocean — operating more strongly than the ensemble captures (Rantanen et al., 2022). The retreat of Arctic sea ice reorganises the surface energy balance, alters the exchange of heat and moisture between ocean and atmosphere, and has been implicated in changes to mid-latitude circulation, although the strength of the latter teleconnection remains an area of active research.

Intensive field campaigns have deepened process understanding of the coupled Arctic system. The year-long drift observations of the MOSAiC expedition characterised the central Arctic atmosphere across a full seasonal cycle, documenting the boundary-layer, cloud and aerosol processes that govern surface energy exchange over sea ice (Shupe et al., 2022). Such observations are essential for evaluating and improving the representation of the polar climate in models, and thus for understanding why the observed amplification outpaces the simulated one.

Mountain glaciers and snow

Mountain glaciers are losing mass across every glaciated region, and the pace of loss has accelerated over the satellite era, with immediate consequences for water resources and sea level (*high confidence*). Multi-temporal satellite observations of High Mountain Asia — the source region for rivers on which a large share of humanity depends — document sustained thinning and retreat since the 1960s, with the rate of mass loss increasing in recent decades (Bhattacharya et al., 2021). Because mountain warming is elevation-dependent, the high-altitude environments where glaciers reside are frequently warming faster than surrounding lowlands, compounding the direct effect of rising temperature on ice (Pepin et al., 2022). The loss of glacier and snow storage alters the seasonal timing and long-term availability of meltwater, initially increasing runoff as ice is depleted and subsequently reducing it as the reservoir shrinks — a trajectory with profound implications for downstream agriculture, hydropower and ecosystems.

The cryosphere as an integrator

The cryosphere functions as an integrator of atmospheric and oceanic change and as an amplifier of it. Ice responds to temperature with a threshold non-linearity — the phase change at the melting point — that concentrates change into particular seasons and regions and gives the cryosphere its characteristic sensitivity. The albedo contrast between ice and the underlying ocean or land converts a modest warming into a much larger local energy gain, and this feedback is the proximate reason for the polar and alpine amplification documented throughout this chapter (Rantanen et al., 2022) (Pepin et al., 2022). Deposition of light-absorbing aerosols, including mineral dust from distant source regions, further lowers snow and ice albedo and accelerates melt, linking the cryospheric assessment back to the atmospheric aerosol cycle (Kok et al., 2021). The cumulative loss of land ice, expressed globally, is a principal driver of the sea-level rise assessed in Section 2.3, closing a chain of causation that runs from greenhouse forcing through the atmosphere and cryosphere to the coastal ocean.

2.6 Changes in the Biosphere and Terrestrial System

Forest resilience and disturbance

Global forests are showing emerging signals of declining resilience, and disturbance regimes are intensifying, indicating that a major terrestrial carbon store is coming under increasing stress (*medium-to-high confidence*). Analysis of satellite vegetation records identifies a widespread reduction in forest resilience — the capacity to recover from perturbation — over recent decades, particularly in tropical, arid and temperate regions, with the loss of resilience linked to increasing water limitation and climate variability (Forzieri et al., 2022). This erosion of resilience is a leading indicator: it precedes outright mortality and signals that forests are approaching the limits of their adaptive capacity before large-scale dieback becomes visible.

The disturbance regimes acting on forests are themselves changing. In Europe, natural disturbances — windthrow, bark-beetle outbreaks and fire — have increased substantially in their impact on forests since the mid-twentieth century, driven by warming and by the changing structure of forest stands (Patacca et al., 2022). Drought is a common thread linking these disturbances, and its physiological signature on trees is now globally recognisable.

Drought-driven tree mortality

Episodes of tree die-off around the world carry a distinctive “hotter-drought” fingerprint, in which elevated temperature amplifies the lethality of water deficit (*high confidence*). A global synthesis of field observations of tree mortality shows that die-off events are associated not merely with low precipitation

but with the combination of drought and anomalously high temperature, which increases atmospheric moisture demand and pushes trees past hydraulic thresholds (Hammond et al., 2022). This finding reframes drought impacts on forests as a coupled temperature–moisture phenomenon consistent with the intensifying water cycle assessed in Section 2.2, and it implies that continued warming will raise the mortality risk associated with any given precipitation deficit. Because forests are both a carbon sink and a store, drought-driven mortality represents a potential positive feedback on atmospheric carbon dioxide, developed further below.

The terrestrial carbon sink

The terrestrial biosphere remains a substantial net sink for anthropogenic carbon, but the sink is regionally heterogeneous and vulnerable to the same drought, fire and disturbance pressures that are eroding forest resilience (*medium confidence*). The global carbon budget attributes roughly a third of annual anthropogenic emissions to the land sink, though with large interannual variability tied to climate (Friedlingstein et al., 2022; Friedlingstein et al., 2023). Some regions have strengthened as sinks: forest expansion has dominated the land carbon sink in China since 1980, illustrating how land management and reforestation can enhance carbon uptake (Yu et al., 2022). Yet the durability of such sinks is not assured, because the very forests that store carbon are subject to increasing disturbance (Patacca et al., 2022) and to hotter-drought mortality (Hammond et al., 2022). The assessment therefore treats the land sink as a service that is presently large but conditionally stable, whose future depends on the trajectory of warming and on management choices assessed elsewhere in this volume.

Wetlands and the long history of land-cover change

The terrestrial biosphere has been reshaped over centuries of human land use, and the loss of wetlands in particular has removed a disproportionately important store of carbon and water (*high confidence*). A reconstruction of global wetland extent finds extensive loss over the past three centuries, with a large fraction of the pre-industrial wetland area drained or converted, concentrated in regions of intensive agriculture. Because wetlands store carbon in saturated soils and regulate local hydrology, their loss both releases carbon and reduces the landscape’s capacity to buffer floods and droughts. This long-horizon change provides essential context for the recent instrumental record: the biosphere against which contemporary change is measured is already a heavily modified system, and some of its capacity to absorb carbon and moderate the water cycle has already been spent.

2.7 Changes in Modes of Variability and Extremes

The changing character of extremes

Extreme events across the temperature and hydrological spectrum are changing in frequency and intensity in ways consistent with a warming, moisture-laden atmosphere, and increasingly these changes can be quantified event by event (*high confidence*). Attribution science, once confined to long-term trends, now systematically evaluates how anthropogenic warming has altered the probability and magnitude of individual extremes, providing an attribution perspective on the extreme-weather impacts of climate change (Clarke et al., 2022). Heatwaves illustrate the maturity of this understanding: their physical drivers, from large-scale circulation to land-atmosphere feedbacks, are now sufficiently well characterised that the human contribution to specific events can be assessed with calibrated confidence (Barriopedro et al., 2023).

The most extreme recent heat events have pushed beyond the range of prior experience. The Pacific Northwest heatwave of June 2021 produced temperatures so far outside the historical distribution that it challenged existing statistical models of extremes and prompted reassessment of how the tails of the temperature distribution respond to warming (White et al., 2023). Such unprecedented events are consequential precisely because infrastructure, ecosystems and public-health systems are calibrated to historical extremes, so that a jump beyond the prior record carries disproportionate impact.

Drought as a coupled and partly anthropogenic phenomenon

Drought is becoming more frequent, more intense and more clearly linked to human influence, with anthropogenic signals now detectable in global drought statistics and in specific regional megadroughts (*high confidence*). A detection-and-attribution analysis finds evidence of anthropogenic impacts on global drought frequency, duration and intensity, moving drought from a phenomenon attributed only to natural variability toward one with a demonstrable human component (Chiang et al., 2021). At the regional scale, the southwestern North American drought that emerged in the twenty-first century intensified into a megadrought whose severity, when placed against tree-ring reconstructions spanning centuries, is unprecedented and substantially attributable to anthropogenic warming (Williams et al., 2022). Europe's multi-year drought of 2018–2020 set a new benchmark in the observational record for its spatial extent and persistence (Rakovec et al., 2022).

The conceptual framing of drought has evolved accordingly. Drought is increasingly recognised as an anthropogenic phenomenon in a dual sense — driven partly by warming-induced increases in evaporative demand and partly by direct human interventions in the water cycle such as abstraction and land-use change — requiring definitions and management approaches that account for the human role (AghaKouchak et al., 2021). This reframing connects the physical-climate assessment of drought to the management challenge of un-

precedented hydrological extremes (Kreibich et al., 2022).

Compound and connected extremes

Extremes increasingly occur in combination — hot and dry together, or in sequences of drought and flood — and these compound events often carry impacts greater than the sum of their parts (*high confidence*). Precipitation trends have been shown to determine the future occurrence of compound hot-dry events, so that regions drying in the mean become disproportionately exposed to the simultaneous occurrence of heat and water deficit (Bevacqua et al., 2022). The management of water extremes is complicated by the tendency of unprecedented floods and droughts to arrive in succession or in unexpected combination, straining risk-management systems designed around historical, independent hazards (Kreibich et al., 2022). Wildfire sits at the intersection of these compound pressures: the multi-decadal increase in forest burned area in Australia has been linked to climate change through the lengthening and intensification of fire-weather conditions (Canadell et al., 2021), and the extreme Australian fire season of 2019–2020 was the subject of formal attribution linking its elevated risk to anthropogenic warming (Oldenborgh et al., 2021).

Modes of internal variability

Large-scale modes of internal variability remain the dominant control on year-to-year climate fluctuations, and their interaction with the forced trend shapes where and when extremes occur (*high confidence*). The recurring patterns of coupled ocean–atmosphere variability redistribute heat and moisture on interannual to decadal timescales, so that the forced warming trend is expressed through, rather than instead of, this internal variability. Persistent circulation regimes associated with these modes — including the blocking patterns that anchor heatwaves and the pressure anomalies that suppress precipitation — are the proximate dynamical cause of many extremes even where the thermodynamic background has been shifted by warming (Barriopedro et al., 2023). Distinguishing the forced signal from internal variability is the central methodological task of detection and attribution, assessed in the following section, and it is the reason that calibrated statements about causation require the ensemble and statistical machinery described there rather than inspection of any single event in isolation.

2.8 Detection and Attribution of Causes

The logic of detection and attribution

The attribution of observed climate change to human influence rests on the physical consistency of the change with the known effect of greenhouse forcing and on the inability of natural factors to reproduce it, and this logic now

yields robust conclusions across the components of the system (*virtually certain for global warming; high confidence for many regional and event-scale statements*). Detection establishes that a change has occurred that is unlikely to be an artefact of internal variability; attribution assigns that change, quantitatively, to its causes. The methodology combines observations with model experiments in which forcings are switched on and off, statistical fingerprinting that isolates the spatial and temporal pattern expected from each forcing, and, for individual events, comparison of the present, warmed world against a counterfactual world without anthropogenic influence. The closure of the global carbon budget already establishes the anthropogenic origin of the composition change (Friedlingstein et al., 2022; Friedlingstein et al., 2023); the attribution enterprise extends that reasoning from the cause to the consequences.

Confidence in attribution is highest where multiple independent methods and datasets converge. For global-mean surface warming, the consistency of the observed warming with the forced response, the incompatibility of the observations with natural-only forcing, and the physical closure of the energy budget through ocean heat uptake together render the human origin effectively certain. Reanalyses that reconstruct the three-dimensional ocean provide the heat-content evidence (Jean-Michel et al., 2021), while the carbon budget provides the forcing side of the ledger (Friedlingstein et al., 2023).

Attribution of regional change and amplification

Regional patterns of change carry their own attributable signatures, of which polar and elevation-dependent amplification are the most prominent (*high confidence*). The observed Arctic amplification factor of nearly four is not merely a regional curiosity; it is a test of the models' representation of the ice-albedo and related feedbacks, and its excess over the simulated value is itself an attributable, physically interpretable discrepancy (Rantanen et al., 2022). Elevation-dependent warming in mountains similarly follows from identifiable processes — snow-albedo feedback, changes in the altitude of the freezing level, and water-vapour and cloud effects that vary with height — allowing the mountain temperature signal to be attributed to warming even though its magnitude varies with elevation (Pepin et al., 2022). Regional drying and extreme-heat trends in the Eastern Mediterranean and Middle East likewise emerge as attributable regional expressions of the global forcing (Zittis et al., 2022).

Event attribution

Formal event attribution now provides calibrated statements about the human contribution to specific extreme events, transforming the assessment of causation from the statistical aggregate to the individual occurrence (*high confidence for well-observed heat and drought events*). The attribution perspective on extreme-weather impacts systematises this capability, distinguishing events for which the anthropogenic influence is strong and well-constrained

from those where the signal is weaker or the dynamics more uncertain (Clarke et al., 2022). Heatwaves are the most tractable case: the physical understanding of their formation supports confident attribution, and the most extreme recent events have been shown to lie far outside the range expected without warming (Barriopedro et al., 2023) (White et al., 2023). Drought attribution has advanced from global statistics (Chiang et al., 2021) to specific megadroughts (Williams et al., 2022) and multi-year continental droughts (Rakovec et al., 2022). Wildfire attribution connects the fire-weather environment to warming, as demonstrated for the Australian fire seasons through both trend analysis (Canadell et al., 2021) and event-specific risk attribution (Oldenborgh et al., 2021).

Event attribution carries a methodological caution. Because it compares the observed world against a modelled counterfactual, its confidence depends on the fidelity with which models represent the relevant processes – a dependence that is strongest for dynamically driven events and weakest for thermodynamically dominated ones such as heat. The assessment therefore attaches higher confidence to the attribution of heat and, increasingly, hydrological extremes than to events whose occurrence hinges on hard-to-model circulation details. High-resolution, convection-permitting simulations improve this situation for precipitation extremes by resolving processes that coarse models parameterise (Ban et al., 2021).

Attribution of impacts

The chain of attribution now reaches beyond physical climate variables to the impacts they cause, quantifying the human contribution to outcomes such as heat-related mortality (*high confidence*). A multi-country analysis attributes a substantial share of recent heat-related mortality specifically to human-induced climate change, completing the causal chain from emissions to composition to warming to extreme heat to human health outcomes. Impact attribution of this kind is methodologically demanding because it must combine climate attribution with exposure–response relationships, but where the underlying physical attribution is strong – as it is for heat – the impact attribution inherits that confidence. This capability is central to the wider assessment because it renders the consequences of large-scale climate change in terms directly meaningful to human systems.

2.9 Paleoclimatic and Long-Term Context

The recent record against the deep past

Placed in the context of the deeper past, the current concentrations of long-lived greenhouse gases and the current rate of change are unusual, and the climate system retains a large, slowly realised response to the forcing already imposed (*high confidence*). Long-timescale analysis indicates that the

full warming implied by present greenhouse gas concentrations has not yet been expressed at the surface, because the ocean's thermal inertia delays the approach to equilibrium and because slow feedbacks continue to operate long after the fast response is complete — a reservoir of committed change described as warming still “in the pipeline” (Hansen et al., 2023). This long view is essential to the assessment for two reasons: it establishes that the present state is anomalous rather than within the range of recent natural variation, and it warns that surface temperature at any moment understates the eventual response to the accumulated forcing.

Slow components and committed change

The climate system contains components whose response times greatly exceed the interval over which greenhouse gases have accumulated. The deep ocean, the great ice sheets, and the carbon stored in soils, permafrost and vegetation all evolve over timescales from decades to millennia. Ocean heat uptake, reconstructed by reanalysis, is the clearest expression of this slow adjustment: heat sequestered below the surface represents warming that is committed but not yet expressed in surface air temperature (Jean-Michel et al., 2021). The loss of land ice, already underway in mountain regions (Bhattacharya et al., 2021) and contributing to sea-level rise (Griggs et al., 2021), likewise carries momentum that will continue for centuries even under stabilised forcing. The paleoclimatic and process-based evidence together imply that the changes assessed in this chapter are the leading edge of a longer adjustment, and that the avoidance of the largest changes depends on limiting the accumulated forcing rather than on the instantaneous temperature alone (Hansen et al., 2023).

2.10 System-Scale Synthesis: Tipping Elements and Planetary Boundaries

Interacting tipping elements

Several large-scale components of the climate system exhibit the potential for threshold behaviour, and their interactions can propagate change from one component to another, raising the risk of cascading, partly irreversible transitions (*medium confidence*). Analysis of the network of climate tipping elements shows that the interactions among them tend to lower the effective thresholds and increase the risk of domino effects, so that the destabilisation of one element can raise the probability of destabilising others under sustained warming (Wunderling et al., 2021). The components assessed throughout this chapter supply the physical substance of this concern: the Arctic's amplified warming and ice loss (Rantanen et al., 2022), the loss of resilience in major forests (Forzieri et al., 2022) and their vulnerability to hotter-drought mortality (Hammond et al., 2022), and the drying of continental water stores (Yao et al.,

2023) are each candidate expressions of, or contributors to, threshold behaviour. The assessment attaches medium confidence to the tipping-element framing because, while the individual processes are well established, the thresholds themselves and the strength of the interactions remain uncertain.

Planetary boundaries

Viewed at the scale of the whole Earth system, the accumulated changes have pushed multiple large-scale processes beyond the ranges that characterised the pre-industrial state, with climate change among the boundaries now transgressed (*high confidence that multiple boundaries are transgressed; medium confidence on their precise placement*). An integrated assessment of the planetary-boundaries framework concludes that six of nine boundaries — including climate change, biosphere integrity, freshwater change and land-system change — have now been crossed, indicating that human pressure on the Earth system extends across multiple, interacting dimensions rather than climate alone (Richardson et al., 2023). This framing connects the physical-climate findings of this chapter to the broader transformation of the planet: the intensification of the water cycle and decline of continental water storage (Yao et al., 2023), the modification of forests and land cover (Patacca et al., 2022) (Yu et al., 2022), and the loss of biosphere resilience (Forzieri et al., 2022) are the climate-system manifestations of boundary transgressions that also have non-climatic dimensions.

Synthesis

The evidence assessed in this chapter forms a single, coherent account. Human activity has raised the concentrations of long-lived greenhouse gases, a fact established beyond dispute by the closure of the global carbon budget (Friedlingstein et al., 2022; Friedlingstein et al., 2023; Friedlingstein et al., 2025). The resulting energy imbalance has warmed the atmosphere, amplified most strongly at high latitudes and high altitudes (Rantanen et al., 2022) (Pepin et al., 2022); warmed and expanded the ocean, driving marine heatwaves and sea-level rise (Smith et al., 2022) (Jean-Michel et al., 2021) (Griggs et al., 2021); shrunk the cryosphere, with mountain glaciers in sustained retreat (Bhattacharya et al., 2021); intensified and destabilised the continental water cycle, producing more severe droughts and floods and a widespread loss of water storage (Chiang et al., 2021) (Williams et al., 2022) (Yao et al., 2023); and placed the terrestrial biosphere under stress that is beginning to erode its resilience (Forzieri et al., 2022) (Hammond et al., 2022). These changes are attributable to human influence with high-to-virtual confidence at the global scale and, increasingly, at the scale of individual regions and events (Clarke et al., 2022) (Barriopedro et al., 2023). Set against the deep past, they are anomalous and incompletely realised, with substantial further change already committed (Hansen et al., 2023). And viewed at the scale of the whole Earth system, they are part of a broader transgression of planetary boundaries whose interacting elements raise the risk of cascading

change (Wunderling et al., 2021) (Richardson et al., 2023). The remainder of this volume builds on this assessment of what has changed and why to evaluate what is projected to come and what can still be avoided.

2.11 The Observational and Modelling Basis of the Assessment

Reanalyses and the reconstruction of climate state

The assessment of large-scale change depends on observing and reconstruction systems that convert heterogeneous measurements into physically consistent, gridded estimates of the climate state, and the capability of these systems has advanced markedly (*high confidence*). Ocean and sea-ice reanalyses now provide eddy-permitting reconstructions of the three-dimensional ocean, assimilating satellite altimetry, sea-surface temperature and the global array of profiling floats into a dynamical model to yield continuous fields of temperature, salinity, currents and ice (Jean-Michel et al., 2021). These products are the backbone of the ocean-warming and heat-content evidence assessed in Section 2.3, and their resolution of mesoscale structure allows regional patterns of change — the warming of boundary currents, the freshening of polar seas — to be characterised rather than only the global mean. Because reanalyses are physically constrained, they also enforce consistency between variables, so that the reconstructed heat, freshwater and momentum budgets must balance, providing an internal check on the observations they assimilate.

The value of reanalysis lies in its combination of coverage and consistency. Direct observations are sparse in the deep ocean, the polar regions and the pre-satellite past; a reanalysis fills these gaps using the model's dynamics, propagating information from data-rich regions and times into data-poor ones. The resulting fields are not observations in the strict sense, and their uncertainty must be assessed against independent data, but they enable the calculation of integrated quantities — ocean heat content, steric sea level, meridional transports — that no single instrument measures directly and that are central to the energetic framing of this chapter.

Satellite retrieval of atmospheric composition

Satellite remote sensing has become the primary means of monitoring the spatial distribution and trends of atmospheric constituents, and rigorous validation against ground-based networks establishes the accuracy required for climate assessment (*high confidence*). Retrievals of tropospheric nitrogen dioxide from the Sentinel-5P mission have been validated against multiple independent ground-based network techniques, characterising the biases and uncertainties of the satellite product and establishing its fitness for monitoring pollution and its trends (Verhoelst et al., 2021). Methane retrievals from the same

platform have been improved and validated over their first years of operation, yielding a data product capable of resolving regional methane enhancements and, increasingly, individual strong sources (Lorente et al., 2021). These capabilities complement the global carbon accounting assessed in Section 2.2 by adding a spatially explicit view of where reactive and greenhouse constituents are entering the atmosphere, supporting the attribution of composition change to specific sectors and regions.

Validation is the linchpin of this capability. A satellite retrieval is an inversion of a radiance measurement through a forward model of atmospheric radiative transfer, and its accuracy depends on the treatment of clouds, aerosols, surface reflectance and the a priori assumptions of the retrieval. Comparison against ground-based reference instruments quantifies the residual biases and their dependence on observing conditions, and only after such validation can satellite trends be interpreted as geophysical rather than instrumental (Verhoelst et al., 2021) (Lorente et al., 2021). The assessment relies on validated products of this kind and treats unvalidated retrievals with corresponding caution.

Climate models and the CMIP ensemble

Coordinated multi-model experiments provide the counterfactual worlds against which observed change is attributed and the projections against which committed change is judged, and the current generation of Earth system models represents an increasingly complete set of processes (*high confidence*). Contemporary Earth system models couple the atmosphere, ocean, sea ice, land and biogeochemistry, and participate in coordinated intercomparison so that their behaviour can be evaluated against observations and against one another; the CMCC Earth System Model exemplifies this class, integrating carbon-cycle and physical-climate components within a single coupled framework (Lovato et al., 2022). The scenario experiments that drive projections are themselves a coordinated activity, spanning a range of plausible futures for emissions and land use so that the response can be mapped across the space of possible forcings (Tebaldi et al., 2021). The translation from emissions scenarios to global temperature outcomes has been systematised through reduced-complexity model assessments that link the detailed scenario space to calibrated temperature projections (Kikstra et al., 2022).

Model evaluation is central to the credibility of these tools. Where models reproduce observed features — the seasonal cycle, the response to volcanic eruptions, the pattern of historical warming — confidence in their projections is strengthened; where they diverge from observations, as in the underestimate of Arctic amplification (Rantanen et al., 2022), the discrepancy is flagged and its implications assessed. The ensemble is treated not as a set of independent truths but as a sample of structural uncertainty, whose spread informs the confidence attached to projected change.

High-resolution and kilometre-scale simulation

Advances in resolution, to the point of explicitly resolving deep convection, materially improve the simulation of precipitation and short-duration extremes, sharpening the assessment of hydrological change (*high confidence*).

The first multi-model ensemble of kilometre-scale regional climate simulations demonstrates that convection-permitting models reproduce the characteristics of observed precipitation – its diurnal cycle, intensity distribution and extremes – more faithfully than coarser models that must parameterise convection (Ban et al., 2021). This matters for the assessment of the water cycle because many of the most damaging hydrological hazards arise from short, intense precipitation that coarse models systematically misrepresent. In the operational domain, hourly-updating, convection-allowing forecast systems have been developed and documented, illustrating the same resolution advances at the weather timescale and providing high-frequency analyses useful for characterising extremes (Dowell et al., 2022).

Machine learning and hybrid modelling

The integration of machine learning with process-based models is emerging as a means to improve predictions and to extract information from the growing volume of Earth observations, while retaining physical interpretability (*medium confidence*).

Differentiable modelling, which unifies machine learning with physical models, offers a route to combining the predictive power of data-driven methods with the constraints and interpretability of physics, with applications across the geosciences from hydrology to land-surface processes (Shen et al., 2023). Such hybrid approaches are increasingly used to calibrate parameterisations, emulate expensive model components, and assimilate observations, and they are likely to become a significant part of the observational and modelling infrastructure on which future assessments rely. The assessment treats these methods as promising and rapidly developing, attaching medium confidence in recognition of their novelty and of the ongoing work to establish their reliability and generalisability.

Citizen science and the extension of observation

Citizen science has become a substantive contributor to environmental and ecological observation, extending the reach of monitoring in space and time and enabling analyses that professional networks alone could not support (*medium confidence*).

Structured citizen-science programmes now generate large volumes of data in environmental and ecological research, contributing to biodiversity monitoring, phenology, hydrology and the documentation of extreme events (Fraisl et al., 2022). These data broaden the observational base underpinning assessments of biosphere and hydrological change, particularly in regions and for variables that are poorly served by instrumental networks, though their heterogeneous provenance requires careful quality control before use in formal analysis.

2.12 The Continental Water Cycle in Detail

Floods and droughts as coupled hazards

Floods and droughts are increasingly understood as coupled expressions of an intensifying and more variable water cycle rather than as independent hazards, and their prediction and management pose challenges that the changing climate is sharpening (*high confidence*). A review of the modelling and prediction of floods and droughts identifies the persistent difficulties in representing the processes that generate these extremes, from the partitioning of precipitation at the land surface to the propagation of anomalies through catchments and aquifers (Brunner et al., 2021). The management challenge is compounded when floods and droughts arrive in unprecedented combination or succession, straining institutions and infrastructure designed around historical, stationary hazard distributions (Kreibich et al., 2022). The reframing of drought as a partly anthropogenic phenomenon — driven by warming-enhanced evaporative demand and by direct human interventions in the water cycle — further complicates definition and attribution and calls for management approaches that recognise the human role (AghaKouchak et al., 2021).

Regional drought benchmarks

Specific regional droughts of the recent record now stand as benchmarks whose severity is unprecedented in the instrumental and, in some cases, the multi-century reconstructed record (*high confidence*). The southwestern North American drought of the twenty-first century intensified, in 2020–2021, into a megadrought whose severity ranks among the most extreme of the past millennium as judged against tree-ring reconstructions, with anthropogenic warming responsible for a substantial part of its intensity through increased evaporative demand and reduced soil moisture (Williams et al., 2022). In Europe, the multi-year drought of 2018–2020 set a new benchmark for spatial extent and persistence in the observational record, exposing the vulnerability of temperate, ostensibly water-rich regions to sustained water deficit (Rakovec et al., 2022). The detection of an anthropogenic signal in global drought statistics provides the wider context within which these regional benchmarks sit (Chiang et al., 2021).

Terrestrial water storage and lakes

The integrated storage of water on land is declining across much of the world, and the loss is visible in lakes large and small as well as in wetlands and soils (*high confidence*). Satellite observations reveal a widespread decline in global lake water storage, with a majority of large lakes losing water over recent decades under the combined influence of warming, altered precipitation and

human abstraction (Yao et al., 2023). Mapping of global lake dynamics shows that small lakes contribute disproportionately to the variability of the terrestrial water surface and to associated carbon fluxes, and that their role is growing (Pi et al., 2022). These surface-water changes are the visible expression of a broader redistribution of water that also encompasses soil moisture and groundwater, and that ultimately returns water to the ocean, coupling the continental and marine branches of the hydrosphere assessed in Section 2.3.

2.13 The Terrestrial Biosphere in Detail

Soils as a climate-system reservoir

Soils are a large and dynamic reservoir of carbon, water and biodiversity whose behaviour couples the land surface to the atmosphere and whose functioning is being altered by climate change and land use (*high confidence*). Soil hydrology governs the partitioning of precipitation into runoff, evaporation and storage, and thereby mediates droughts, floods and the supply of moisture to the atmosphere; its representation is increasingly recognised as central to understanding the land component of the Earth system (Vereecken et al., 2022). The biological functioning of soils is tightly linked to their carbon content: across global biomes, soil microbial diversity and biomass are related in ways driven by the amount of carbon stored in the soil, so that changes in soil carbon propagate into the microbial communities that regulate decomposition and nutrient cycling (Bastida et al., 2021). Where the soil resource is degraded, these functions are impaired: erosion has been shown to reduce soil microbial diversity, network complexity and multifunctionality, indicating that physical degradation and biological decline proceed together (Qiu et al., 2021).

Soil degradation interacts with climate change to threaten the services soils provide. Global projections indicate that primary soil salinization is likely to expand under a changing climate through the twenty-first century, as altered precipitation and evaporation concentrate salts in vulnerable regions and impair the productivity of the land (Hassani et al., 2021). Salinization, erosion and the loss of soil carbon together constitute a degradation of the land-system boundary whose transgression is part of the broader planetary-boundaries assessment (Richardson et al., 2023), and whose consequences for food production connect the physical-climate assessment to human systems.

Ecosystems, biodiversity and context dependence

Ecological responses to climate change are pervasive but strongly context-dependent, so that general statements about biosphere change must be qualified by the local conditions that modulate the response (*high confidence that responses are widespread; medium confidence on their generalisation*). The dependence of ecological outcomes on context — the specific combination of

species, environment and history at a given place — is a recognised challenge for the field, complicating the transfer of findings from one system to another and requiring that syntheses account explicitly for the sources of variation (Catford et al., 2021). This caution applies throughout the biosphere assessment: the loss of forest resilience (Forzieri et al., 2022), the incidence of hotter-drought mortality (Hammond et al., 2022) and the response of soil communities (Bastida et al., 2021) all vary systematically with local conditions, and the assessment attaches confidence to the existence and direction of these responses while acknowledging that their magnitude is place-specific.

Among the most concerning biosphere signals is the pressure on insects, which underpin pollination, decomposition and food webs. A scientists' warning on climate change and insects synthesises evidence that warming, together with land-use change and other stressors, is contributing to declines and distributional shifts in insect populations with potential consequences that cascade through ecosystems (Harvey et al., 2022). Because insects mediate ecosystem functions on which both wild systems and agriculture depend, changes in their abundance and distribution represent a channel through which climate change propagates into the wider biosphere and into human food security.

Climate change, agriculture and food security

Climate change is affecting the productivity and stability of agriculture through direct effects on crops and through indirect effects mediated by pests, pathogens and soils, with impacts emerging earlier than earlier model generations suggested (*high confidence*). A new generation of coupled climate and crop models indicates that the impacts of climate change on global agriculture emerge earlier and are more pronounced than previously projected, with major staple crops experiencing significant changes in suitability and yield within the coming decades under continued warming (Jägermeyr et al., 2021). These direct effects are compounded by biological intermediaries: climate change is altering the distribution and severity of plant pathogens, with consequences for crop losses and food security that add to the direct climatic impacts on yield (Singh et al., 2023). Soil degradation closes a further pathway, as salinization reduces the productive capacity of agricultural land in vulnerable regions (Hassani et al., 2021). The assessment treats agriculture as a sensitive integrator of biosphere and hydrological change, in which the physical signals documented throughout this chapter are translated into consequences for human sustenance.

The biosphere as source, sink and feedback

The terrestrial biosphere is simultaneously a sink that removes anthropogenic carbon, a store that can be released, and a set of feedbacks that can amplify or dampen change, and the balance among these roles is shifting under climate and land-use pressure (*medium confidence*). The land sink absorbs a large share of annual emissions (Friedlingstein et al., 2022; Friedlingstein et al., 2023), and management can enhance it, as demonstrated by forest-expansion-

driven carbon uptake in China (Yu et al., 2022). But the same warming that drives uptake through longer growing seasons and carbon-dioxide fertilisation also drives loss through drought mortality (Hammond et al., 2022), disturbance (Patacca et al., 2022) and the degradation of soils (Qiu et al., 2021) (Hassani et al., 2021). The net trajectory of the biosphere sink is therefore contingent, depending on the balance between these opposing tendencies and on the level of warming reached. The erosion of forest resilience (Forzieri et al., 2022) is the clearest warning that the sink is not unconditionally secure, and the interaction of biosphere change with other tipping elements (Wunderling et al., 2021) is the mechanism by which a weakening sink could accelerate the accumulation of atmospheric carbon.

2.14 Drivers of Atmospheric Composition Change

The sectoral structure of emissions

The rise in greenhouse gas concentrations assessed in Section 2.2 originates in a well-characterised set of emitting sectors whose trends and drivers can be quantified, closing the causal loop between human activity and atmospheric change (*high confidence*). A review of the trends and drivers of greenhouse gas emissions by sector documents the growth of emissions across energy supply, industry, transport, buildings and agriculture over recent decades, and decomposes that growth into the underlying drivers of population, affluence, energy intensity and carbon intensity (Lamb et al., 2021). This sectoral accounting is the emissions-side complement to the atmospheric-side carbon budget (Friedlingstein et al., 2022; Friedlingstein et al., 2023): where the carbon budget establishes that the atmospheric increase matches the sum of anthropogenic sources, the sectoral analysis identifies which human activities constitute those sources. Agriculture contributes to this total through emissions whose character differs from the predominantly fossil carbon-dioxide-emitting sectors, involving methane and nitrous oxide with distinct atmospheric lifetimes and warming effects (Lynch et al., 2021).

Reactive constituents and their monitoring

The composition change assessed in this chapter extends beyond the long-lived greenhouse gases to reactive constituents that influence air quality, atmospheric chemistry and, in some cases, radiative forcing. Satellite monitoring of tropospheric nitrogen dioxide and methane, validated against ground networks, provides the spatially explicit record needed to track these constituents and to associate their trends with the emitting sectors identified in the emissions inventories (Verhoelst et al., 2021) (Lorente et al., 2021). Mineral dust, mobilised from arid source regions whose relative contributions have been quantified, constitutes a further reactive component with radiative, biogeochemical and

cryospheric effects that respond to and feed back upon large-scale change (Kok et al., 2021). Together these constituents complete the picture of a changing atmosphere whose composition is being altered across a spectrum of species and lifetimes by human activity.

2.15 Key Uncertainties and Confidence

The assessment presented in this chapter combines statements held with virtual certainty and statements more carefully hedged, and it is important to be explicit about where the principal uncertainties lie. The origin of the greenhouse gas increase and the reality of global warming are effectively certain, resting on the closure of the carbon budget (Friedlingstein et al., 2023) and on the convergence of independent indicators. The attribution of global-mean warming to human influence is likewise held with the highest confidence.

Greater uncertainty attaches to the magnitude, though rarely the sign, of regional and process-level changes. The observed Arctic amplification exceeds the modelled value, indicating a gap in the representation of polar feedbacks whose resolution remains incomplete (Rantanen et al., 2022). The precise thresholds of the climate tipping elements, and the strength of their interactions, are uncertain even though the qualitative risk of cascading change is established (Wunderling et al., 2021). The future trajectory of the terrestrial carbon sink is contingent on the competition between enhancement and loss processes whose balance is not yet resolved (Forzieri et al., 2022) (Yu et al., 2022) (Hammond et al., 2022). Event attribution is more confident for thermodynamically driven extremes such as heat than for dynamically driven ones, reflecting the differing fidelity of models across processes (Clarke et al., 2022) (Barriopedro et al., 2023). And ecological responses, while pervasive, are context-dependent in ways that limit their generalisation (Catford et al., 2021).

These uncertainties do not undermine the central conclusions of the chapter; they refine them. The direction of change across the atmosphere, ocean, cryosphere and biosphere is established with high-to-virtual confidence, as is its human origin. What remains uncertain is principally the rate and magnitude of particular regional and process-level responses, and the placement of the thresholds beyond which change may become self-sustaining. It is on the reduction of precisely these uncertainties — through improved observation (Jean-Michel et al., 2021) (Verhoelst et al., 2021), higher-resolution and hybrid modelling (Ban et al., 2021) (Shen et al., 2023), and the continued maturation of attribution (Clarke et al., 2022) — that the value of future assessment most depends.

2.16 Cross-Component Integration and the Human Dimension

The coherence of change across components

The strongest basis for the assessment's central conclusions is not any single record but the physical coherence of change across the components of the system, each of which responds to the same forcing in the manner its physics dictates (*virtually certain for the overall picture*). The atmosphere warms and holds more moisture, intensifying the water cycle (Bevacqua et al., 2022); the ocean absorbs the majority of the excess heat, warming at depth and driving marine heatwaves and thermal expansion (Jean-Michel et al., 2021) (Smith et al., 2022); the cryosphere retreats where warming crosses the melting point, amplified by the albedo feedback in polar and alpine regions (Rantanen et al., 2022) (Bhattacharya et al., 2021); the continental water cycle grows more variable, producing more severe droughts and a decline in stored water (Williams et al., 2022) (Yao et al., 2023); and the biosphere responds and, under compound stress, loses resilience (Forzieri et al., 2022). No one of these changes, taken alone, would compel the conclusion that the system is being forced by rising greenhouse gases. Taken together, and taken in the specific pattern predicted by the physics of greenhouse forcing — surface and tropospheric warming with stratospheric cooling, polar amplification, ocean heat uptake, an intensifying water cycle — they constitute an internally consistent and mutually reinforcing body of evidence that admits no natural explanation.

This coherence is also the reason the assessment can proceed with confidence despite residual uncertainties in individual components. Where one indicator is ambiguous or poorly observed, others constrain the conclusion; where a model underperforms on one process, its performance on others and the direct observational record supply the missing confidence. The convergence of independent evidence, emphasised at the outset of this chapter, is thus not a rhetorical device but the operational logic by which calibrated statements are reached.

Marine ecosystems and ocean biogeochemistry

Warming and marine heatwaves are reorganising marine ecosystems and the biogeochemical cycles they sustain, extending the physical signal of ocean change into the living ocean and its carbon cycling (*high confidence*). The biological impacts of marine heatwaves span the full hierarchy of ocean life, from the physiological stress and mortality of individual organisms to shifts in species distributions, mass mortality events, and disruption of the biogeochemical processes that link the ocean to the atmosphere (Smith et al., 2022). Because the ocean is both a heat sink and a carbon sink, changes in its biology and chemistry feed back on the atmospheric composition assessed in Section 2.2: the efficiency of the ocean carbon sink documented in the global carbon budget depends in part on biological processes that warming and deoxygena-

tion can alter (Friedlingstein et al., 2023). The eddy-resolving reanalyses that reconstruct the physical ocean provide the environmental context for these biological changes, linking the temperature and circulation anomalies to the ecosystems that experience them (Jean-Michel et al., 2021).

The coupling between marine heatwaves and the atmosphere further integrates the ocean into the extremes assessment of Section 2.6. Ocean warm anomalies can reinforce atmospheric ridging and contribute to coastal heat extremes, while atmospheric circulation anomalies help generate and sustain the ocean anomalies. The unprecedented terrestrial heat events of recent years, such as the Pacific Northwest heatwave (White et al., 2023), occurred against a background of anomalously warm adjacent ocean, illustrating the land–ocean coupling that compound-extremes analysis seeks to capture.

Arctic processes and the value of intensive observation

Intensive, process-resolving observation of the Arctic has advanced understanding of why the region is warming so rapidly and why models struggle to reproduce the observed amplification (*high confidence*). The year-long MOSAiC drift observations characterised the central Arctic atmosphere across a complete seasonal cycle, documenting the boundary-layer structure, cloud properties and aerosol processes that govern the exchange of energy between the sea-ice surface and the atmosphere (Shupe et al., 2022). These processes are precisely those whose misrepresentation is implicated in the models' underestimate of Arctic amplification (Rantanen et al., 2022): the treatment of mixed-phase clouds, stable boundary layers and surface-flux coupling over sea ice is difficult, and the MOSAiC observations provide the detailed constraints needed to evaluate and improve it. The mineral dust and other aerosols transported into the Arctic, whose source regions have been quantified, further modulate the surface energy balance by altering snow and ice albedo (Kok et al., 2021), adding another process that intensive observation helps to constrain. The assessment treats such campaigns as high-value investments precisely because they target the processes responsible for the largest model–observation discrepancies.

The human dimension of large-scale change

The large-scale physical changes assessed in this chapter are already producing measurable consequences for human systems, and the attribution of those consequences to human-induced climate change is increasingly demonstrable (*high confidence for heat-related mortality*). The clearest example is heat: a multi-country analysis attributes a substantial fraction of recent heat-related mortality specifically to human-induced climate change, completing the causal chain from emissions through warming to extreme heat to human health outcomes. Systematic monitoring of the health consequences of climate change documents the widening exposure of populations to the hazards that the physical changes generate, from heat to shifting patterns of disease and food

and water insecurity (Romanello et al., 2022). These human consequences are not the subject of the present chapter, whose remit is the physical system and its causes, but they establish the stakes of the physical findings and motivate the granular attribution that connects large-scale change to specific, experienced harms.

Agriculture and food security constitute a second channel through which physical change reaches human systems. The earlier-than-expected emergence of climate impacts on crops (Jägermeyr et al., 2021), the climate-driven shifts in plant pathogens (Singh et al., 2023), and the projected expansion of soil salinization (Hassani et al., 2021) together indicate that the intensifying water cycle and warming assessed here are already reshaping the conditions for food production. The assessment notes these connections not to duplicate the work of other chapters but to underscore that the large-scale changes it documents are consequential precisely because human systems are calibrated to a climate that no longer exists.

Regional expression of global change

Global change is experienced regionally, and its regional expression is systematically structured by geography, so that certain regions serve as bellwethers of the wider transformation (*high confidence*). The Arctic, warming at nearly four times the global rate, is the clearest such bellwether (Rantanen et al., 2022), followed by the world's mountain regions, where elevation-dependent warming concentrates change at high altitudes and drives rapid ice loss (Pepin et al., 2022) (Bhattacharya et al., 2021). Dry subtropical and semi-arid regions form a third category: the Eastern Mediterranean and Middle East are warming faster than the global mean and are projected to experience a disproportionate rise in extremes (Zittis et al., 2022), while southwestern North America has entered a megadrought without precedent in centuries (Williams et al., 2022). Temperate regions long considered water-secure are not exempt, as Europe's benchmark multi-year drought demonstrated (Rakovec et al., 2022). This regional structure is not random; it follows from the physics of amplification, from the geography of the water cycle, and from the sensitivity of particular cryospheric and ecological systems, and it allows the assessment to identify where the leading edge of change is most likely to be found.

Concluding assessment

The findings of this chapter establish, with the highest available confidence, that the climate system has changed at large scale across all its major components and that human activity is the dominant cause. The rise in greenhouse gas concentrations is anthropogenic beyond dispute, confirmed by the closure of the global carbon budget (Friedlingstein et al., 2022; Friedlingstein et al., 2023; Friedlingstein et al., 2025). The resulting warming has propagated through the atmosphere, ocean, cryosphere and biosphere in a coherent, physically predicted pattern (Rantanen et al., 2022) (Jean-Michel et al., 2021) (Bhattacharya

et al., 2021) (Forzieri et al., 2022). These changes are attributable to human influence at the global scale with virtual certainty and, increasingly, at the scale of regions and individual events (Clarke et al., 2022) (Barriopedro et al., 2023) (Williams et al., 2022). They are anomalous against the deep past and carry substantial committed further change (Hansen et al., 2023). And they form part of a broader transgression of planetary boundaries whose interacting tipping elements raise the risk of cascading, partly irreversible transitions (Richardson et al., 2023) (Wunderling et al., 2021). The chapters that follow build on this foundation to assess the future evolution of the system and the scope that remains to limit further change.

2.17 Sea Level, Coasts and the Aerosol–Cryosphere Coupling

Sea-level rise as an integrator of system change

Sea-level rise integrates the warming of the ocean and the loss of land ice into a single, slowly evolving signal whose consequences for coasts are among the most durable of all large-scale changes (*high confidence*). The two dominant contributions — thermal expansion of the warming ocean and the addition of meltwater from glaciers and ice sheets — are each documented independently in this chapter, through the subsurface ocean warming captured by reanalysis (Jean-Michel et al., 2021) and through the sustained mass loss of mountain glaciers (Bhattacharya et al., 2021). Their combination raises mean sea level globally, while regional patterns of ocean density, circulation and land motion determine the local rate actually experienced. Because both contributions respond slowly to forcing, sea-level rise carries substantial committed change: even under stabilised temperatures the ocean would continue to expand and ice to adjust for centuries, a manifestation of the pipeline of committed change identified in the paleoclimatic assessment (Hansen et al., 2023).

The management response to sea-level rise must translate these global and regional signals into local adaptation under deep uncertainty. Coastal adaptation encompasses protection, accommodation and, in the limit, retreat, and its planning requires projections that resolve the regional pattern of sea-level change and account for the vulnerability of particular coastlines (Griggs et al., 2021). The assessment emphasises that the physical certainty of ongoing rise coexists with genuine uncertainty in its rate, particularly the contribution of the large ice sheets, so that adaptation must be robust across a range of plausible trajectories rather than optimised to a single projection.

The aerosol–cryosphere coupling

Light-absorbing aerosols deposited on snow and ice constitute a coupling between the atmospheric aerosol cycle and the cryosphere that accelerates

melt and links distant source regions to polar and alpine change (*medium confidence*). Mineral dust, mobilised predominantly from a small number of major arid source regions whose relative contributions have been quantified, is transported over long distances and deposited on snow and ice, where it lowers surface albedo and increases the absorption of solar radiation (Kok et al., 2021). This darkening accelerates melt beyond what temperature alone would produce, adding to the albedo feedback that already amplifies warming in the cryosphere (Rantanen et al., 2022) (Pepin et al., 2022). The strength of this coupling depends on the dust cycle, which is itself changing as surface winds, soil moisture and vegetation cover respond to climate and land use, making dust both a driver of and a response to large-scale change. The assessment attaches medium confidence to the magnitude of the aerosol–cryosphere effect while recognising the process as an established contributor to cryospheric change.

Observing-system gaps and priorities

The confidence of the assessment is bounded by the completeness of the observing system, and several gaps warrant emphasis. The deep ocean below the routine reach of the profiling-float array is undersampled, so that estimates of full-depth heat content and their trends rely heavily on reanalysis and carry corresponding uncertainty (Jean-Michel et al., 2021). The polar regions, despite intensive campaigns such as MOSAiC (Shupe et al., 2022), remain sparsely observed in space and time relative to the pace of change there (Rantanen et al., 2022). Terrestrial water storage below the surface — soil moisture and groundwater — is inferred rather than directly and comprehensively measured, limiting the characterisation of drought and of the continental water balance (Yao et al., 2023) (Vereecken et al., 2022). And the biosphere’s slow and heterogeneous responses are difficult to monitor at the scale and duration required to detect resilience loss before it manifests as dieback (Forzieri et al., 2022). Addressing these gaps — through sustained satellite missions, expanded in-situ networks, intensive process campaigns and the incorporation of citizen-science observation (Fraisl et al., 2022) — is the observational counterpart to the modelling and attribution advances on which future assessments depend.

This assessment of large-scale change and its causes provides the observational and mechanistic foundation for the projections, scenario analysis and risk assessment developed in the chapters that follow. The coherence of the evidence across the atmosphere, ocean, cryosphere and biosphere, and the maturation of detection and attribution from the global mean to the individual event, together establish both what has changed and why with a confidence that has strengthened in every assessment cycle.

2.18 References

- AghaKouchak, A., et al. (2021). Anthropogenic Drought: Definition, Challenges, and Opportunities. *Reviews of Geophysics*. <https://doi.org/10.1029/2019rg000683>
- Ban, N., et al. (2021). The first multi-model ensemble of regional climate simulations at kilometer-scale resolution, part I: evaluation of precipitation. *Climate Dynamics*. <https://doi.org/10.1007/s00382-021-05708-w>
- Barriopedro, D., et al. (2023). Heat Waves: Physical Understanding and Scientific Challenges. *Reviews of Geophysics*. <https://doi.org/10.1029/2022rg000780>
- Bastida, F., et al. (2021). Soil microbial diversity–biomass relationships are driven by soil carbon content across global biomes. *The ISME Journal*. <https://doi.org/10.1038/s41396-021-00906-0>
- Bevacqua, E., et al. (2022). Precipitation trends determine future occurrences of compound hot–dry events. *Nature Climate Change*. <https://doi.org/10.1038/s41558-022-01309-5>
- Bhattacharya, A., et al. (2021). High Mountain Asian glacier response to climate revealed by multi-temporal satellite observations since the 1960s. *Nature Communications*. <https://doi.org/10.1038/s41467-021-24180-y>
- Brunner, M., et al. (2021). Challenges in modeling and predicting floods and droughts: A review. *Wiley Interdisciplinary Reviews Water*. <https://doi.org/10.1002/wat2.1520>
- Canadell, J., et al. (2021). Multi-decadal increase of forest burned area in Australia is linked to climate change. *Nature Communications*. <https://doi.org/10.1038/s41467-021-27225-4>
- Catford, J., et al. (2021). Addressing context dependence in ecology. *Trends in Ecology & Evolution*. <https://doi.org/10.1016/j.tree.2021.09.007>
- Chiang, F., et al. (2021). Evidence of anthropogenic impacts on global drought frequency, duration, and intensity. *Nature Communications*. <https://doi.org/10.1038/s41467-021-22314-w>
- Clarke, B., et al. (2022). Extreme weather impacts of climate change: an attribution perspective. *Environmental Research Climate*. <https://doi.org/10.1088/2752-5295/ac6e7d>
- Dowell, D., et al. (2022). The High-Resolution Rapid Refresh (HRRR): An Hourly Updating Convection-Allowing Forecast Model. Part I: Motivation and System Description. *Weather and Forecasting*. <https://doi.org/10.1175/waf-d-21-0151.1>
- Forzieri, G., et al. (2022). Emerging signals of declining forest resilience under climate change. *Nature*. <https://doi.org/10.1038/s41586-022-04959-9>
- Fraisl, D., et al. (2022). Citizen science in environmental and ecological sciences. *Nature Reviews Methods Primers*. <https://doi.org/10.1038/s43586-022-00144-4>
- Friedlingstein, P., et al. (2023). Global Carbon Budget 2023. *Earth system science data*. <https://doi.org/10.5194/essd-15-5301-2023>
- Friedlingstein, P., et al. (2025). Global Carbon Budget 2024. *Earth system science data*. <https://doi.org/10.5194/essd-17-965-2025>
- Griggs, G., et al. (2021). Coastal Adaptation to Climate Change and Sea-Level Rise. *Water*. <https://doi.org/10.3390/w13162151>

Hammond, W., et al. (2022). Global field observations of tree die-off reveal hotter-drought fingerprint for Earth's forests. *Nature Communications*. <https://doi.org/10.1038/s41467-022-29289-2>

Hansen, J., et al. (2023). Global warming in the pipeline. *Oxford Open Climate Change*. <https://doi.org/10.1093/oxfclm/kgad008>

Harvey, J., et al. (2022). Scientists' warning on climate change and insects. *Ecological Monographs*. <https://doi.org/10.1002/ecm.1553>

Hassani, A., et al. (2021). Global predictions of primary soil salinization under changing climate in the 21st century. *Nature Communications*. <https://doi.org/10.1038/s41467-021-26907-3>

Jean-Michel, L., et al. (2021). The Copernicus Global 1/12° Oceanic and Sea Ice GLORYS12 Reanalysis. *Frontiers in Earth Science*. <https://doi.org/10.3389/feart.2021.698876>

Jägermeyr, J., et al. (2021). Climate impacts on global agriculture emerge earlier in new generation of climate and crop models. *Nature Food*. <https://doi.org/10.1038/s43016-021-00400-y>

Kikstra, J., et al. (2022). The IPCC Sixth Assessment Report WGIII climate assessment of mitigation pathways: from emissions to global temperatures. *Geoscientific model development*. <https://doi.org/10.5194/gmd-15-9075-2022>

Kok, J., et al. (2021). Contribution of the world's main dust source regions to the global cycle of desert dust. *Atmospheric chemistry and physics*. <https://doi.org/10.5194/acp-21-8169-2021>

Kreibich, H., et al. (2022). The challenge of unprecedented floods and droughts in risk management. *Nature*. <https://doi.org/10.1038/s41586-022-04917-5>

Lamb, W., et al. (2021). A review of trends and drivers of greenhouse gas emissions by sector from 1990 to 2018. *Environmental Research Letters*. <https://doi.org/10.1088/1748-9326/abee4e>

Lorente, A., et al. (2021). Methane retrieved from TROPOMI: improvement of the data product and validation of the first 2 years of measurements. *Atmospheric measurement techniques*. <https://doi.org/10.5194/amt-14-665-2021>

Lovato, T., et al. (2022). CMIP6 Simulations With the CMCC Earth System Model (CMCC-ESM2). *Journal of Advances in Modeling Earth Systems*. <https://doi.org/10.1029/2021ms002814>

Lynch, J., et al. (2021). Agriculture's Contribution to Climate Change and Role in Mitigation Is Distinct From Predominantly Fossil CO₂-Emitting Sectors. *Frontiers in Sustainable Food Systems*. <https://doi.org/10.3389/fsufs.2020.518039>

Oldenborgh, G., et al. (2021). Attribution of the Australian bushfire risk to anthropogenic climate change. *Natural hazards and earth system sciences*. <https://doi.org/10.5194/nhess-21-941-2021>

Patacca, M., et al. (2022). Significant increase in natural disturbance impacts on European forests since 1950. *Global Change Biology*. <https://doi.org/10.1111/gcb.16531>

Pepin, N., et al. (2022). Climate Changes and Their Elevational Patterns in the Mountains of the World. *Reviews of Geophysics*. <https://doi.org/10.1029/2020rg000730>

Pi, X., et al. (2022). Mapping global lake dynamics reveals the emerging roles of small lakes. *Nature Communications*. <https://doi.org/10.1038/s41467-022-33239-3>

Qiu, L., et al. (2021). Erosion reduces soil microbial diversity, network complexity and multifunctionality. *The ISME Journal*. <https://doi.org/10.1038/s41396-021-00913-1>

Rakovec, O., et al. (2022). The 2018–2020 Multi-Year Drought Sets a New Benchmark in Europe. *Earth's Future*. <https://doi.org/10.1029/2021ef002394>

Rantanen, M., et al. (2022). The Arctic has warmed nearly four times faster than the globe since 1979. *Communications Earth & Environment*. <https://doi.org/10.1038/s4022-00498-3>

Richardson, K., et al. (2023). Earth beyond six of nine planetary boundaries. *Science Advances*. <https://doi.org/10.1126/sciadv.adh2458>

Romanello, M., et al. (2022). The 2022 report of the Lancet Countdown on health and climate change: health at the mercy of fossil fuels. *The Lancet*. [https://doi.org/10.1016/s0140-6736\(22\)01540-9](https://doi.org/10.1016/s0140-6736(22)01540-9)

Shen, C., et al. (2023). Differentiable modelling to unify machine learning and physical models for geosciences. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-023-00450-9>

Shupe, M., et al. (2022). Overview of the MOSAiC expedition: Atmosphere. *Elementa Science of the Anthropocene*. <https://doi.org/10.1525/elementa.2021.00060>

Singh, B., et al. (2023). Climate change impacts on plant pathogens, food security and paths forward. *Nature Reviews Microbiology*. <https://doi.org/10.1038/s41579-023-00900-7>

Smith, K., et al. (2022). Biological Impacts of Marine Heatwaves. *Annual Review of Marine Science*. <https://doi.org/10.1146/annurev-marine-032122-121437>

Tebaldi, C., et al. (2021). Climate model projections from the Scenario Model Intercomparison Project (ScenarioMIP) of CMIP6. *Earth System Dynamics*. <https://doi.org/10.5194/esd-12-253-2021>

Vereecken, H., et al. (2022). Soil hydrology in the Earth system. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-022-00324-6>

Verhoelst, T., et al. (2021). Ground-based validation of the Copernicus Sentinel-5P TROPOMI NO₂ measurements with the NDACC ZSL-DOAS, MAX-DOAS and Pandonia global networks. *Atmospheric measurement techniques*. <https://doi.org/10.5194/amt-14-481-2021>

White, R., et al. (2023). The unprecedented Pacific Northwest heatwave of June 2021. *Nature Communications*. <https://doi.org/10.1038/s41467-023-36289-3>

Williams, P., et al. (2022). Rapid intensification of the emerging southwestern North American megadrought in 2020–2021. *Nature Climate Change*. <https://doi.org/10.1038/s41558-022-01290-z>

Wunderling, N., et al. (2021). Interacting tipping elements increase risk of climate domino effects under global warming. *Earth System Dynamics*. <https://doi.org/10.5194/esd-12-601-2021>

Yao, F., et al. (2023). Satellites reveal widespread decline in global lake water storage. *Science*. <https://doi.org/10.1126/science.abo2812>

Yu, Z., et al. (2022). Forest expansion dominates China's land carbon sink since 1980. *Nature Communications*. <https://doi.org/10.1038/s41467-022-32961-2>

Zittis, G., et al. (2022). Climate Change and Weather Extremes in the Eastern Mediterranean and Middle East. *Reviews of Geophysics*. <https://doi.org/10.1029/2021rg000762>

Chapter 3

Chapter 3: Regional Climate and Extreme Events

3.1 Executive Summary

Regional warming has departed decisively from the global mean, with the Arctic and the world's high-mountain regions emerging as the fastest-warming large domains of the planet (*high confidence*). Observational syntheses place recent Arctic warming at close to four times the global rate since the late 1970s (Rantanen et al., 2022), while elevation-dependent warming is now documented across mountain systems on every inhabited continent (Pepin et al., 2022). These regional signals are not merely scaled-down versions of the global average; they reflect distinct physical mechanisms — sea-ice and snow-albedo feedbacks, changes in the vertical structure of the atmosphere, and land-surface modification — that concentrate change in particular places and seasons.

The intensity and frequency of hot extremes have increased across most inhabited regions, and the human and ecological consequences of extreme heat are now measurable at the scale of individual cities and populations (*high confidence*). Urban populations experience heat that is systematically amplified relative to surrounding rural areas, and this amplification is distributed unequally across and within cities (Hsu et al., 2021). Extreme weather and climate events already impose a substantial and growing burden on population health and on health systems (Ebi et al., 2021), a burden documented in successive annual assessments of health and climate change (Romanello et al., 2022; Romanello et al., 2023; Romanello et al., 2024).

Short-duration rainfall extremes are intensifying at rates consistent with, and in places exceeding, thermodynamic expectations, raising flood hazard in exposed and vulnerable communities (*high confidence*). The anthropogenic intensification of sub-daily rainfall is now an identifiable feature of the observational and modelling record (Fowler et al., 2021). Exposure to flooding is concentrated among the poor: a large fraction of the global population exposed to serious flood hazard lives in poverty (Rentschler et al., 2022), so that identical physical hazards translate into sharply divergent human outcomes.

Drought is evolving in character as well as in frequency, with a docu-

mented global shift toward rapidly developing “flash” droughts and toward multi-year megadroughts in several regions (*high confidence*). A transition toward faster drought onset under warming has been identified globally (Yuan et al., 2023), while the early-twenty-first-century drought of southwestern North America intensified into one of the most severe multi-decadal dry episodes in the long paleoclimate record (Williams et al., 2022). Groundwater, the buffer on which much of humanity relies during drought, is under simultaneous pressure from depletion and from a changing climate (Scanlon et al., 2023).

Compound and cascading events — where hazards coincide, cluster, or trigger one another — are an increasingly important frame for regional risk, and existing single-hazard assessment methods systematically understate them (*medium confidence*). Compound climate-fire-smoke sequences illustrate the point: the extreme forest fires of southeastern Australia were connected to climate change and variability (Abram et al., 2021), and wildfire smoke has been shown to damage respiratory health more than equivalent particulate matter from other sources (Aguilera et al., 2021). A framework treating climate risk as complex, interacting, and cascading is required to capture these dynamics (Simpson et al., 2021).

Event attribution science has matured to the point where the contribution of human influence to many individual extremes can be quantified, and this capability now underwrites both scientific assessment and the economic valuation of climate damages (*high confidence*). The same physical understanding that supports attribution also informs the estimation of the social cost of carbon dioxide, which comprehensive assessments place substantially higher than many regulatory values in current use (Rennert et al., 2022), and the exploration of low-probability, high-impact catastrophic outcomes that attribution alone cannot bound (Kemp et al., 2022).

Regional monsoon systems, which govern the water security of a majority of the world’s population, are responding to both greenhouse-gas forcing and aerosol forcing in ways that vary by basin and remain a leading source of regional projection uncertainty (*medium confidence*). Coordinated multi-model experiments provide the backbone for regional projection (Tebaldi et al., 2021), yet the translation of global forcing into monsoon rainfall depends on regional processes whose representation still differs across models. Because monsoon variability propagates directly into agriculture, the emergence of climate impacts on food production in the latest generation of crop and climate models (Jägermeyr et al., 2021) is of first-order regional importance.

Adaptation to regional climate change is occurring but is uneven in coverage and, in some documented cases, has increased rather than reduced vulnerability (*high confidence*). Interventions intended to reduce climate vulnerability have in places redistributed or deepened it (Eriksen et al., 2021), underscoring that regional risk is produced jointly by physical hazard, exposure, and the social processes that govern vulnerability and response.

3.2 Observed Regional Climate Change

Regional warming has diverged strongly from the global mean, with the Arctic warming nearly four times faster than the planet as a whole since 1979 (*high confidence*).

The global-mean surface temperature is a useful index of the planet's energy imbalance, but it conceals the geography of change. Warming is neither spatially uniform nor seasonally uniform, and the largest regional departures from the global mean are now large enough to constitute distinct climatic phenomena in their own right. The clearest example is the Arctic. Analyses of the observational record indicate that the Arctic has warmed at close to four times the global rate over the period since 1979, a ratio substantially larger than the factor of two or three that had been the conventional figure in earlier assessments (Rantanen et al., 2022). The exact multiplier depends on the definition of the Arctic domain, the reference period, and the observational dataset, but across reasonable choices the amplification is robust and pronounced.

The physical basis of Arctic amplification is well established in outline even where the partitioning among mechanisms remains debated. The retreat of highly reflective sea ice exposes a darker ocean surface that absorbs more solar radiation, warms, and further inhibits ice formation — a positive feedback that operates most strongly in autumn and early winter as newly opened water releases stored heat to the atmosphere. Changes in the vertical structure of atmospheric warming, in cloud cover, in water vapour, and in the poleward transport of heat and moisture all contribute. The seasonality of the signal is diagnostic: Arctic warming is largest in the cold season, when the sea-ice and lower-tropospheric feedbacks are most active, and comparatively muted in summer when surface temperatures are pinned near the melting point of ice. The consequence is a regional climate that is being reorganised more rapidly than any other large terrestrial or oceanic domain, with implications that propagate through permafrost, ecosystems, and — through altered temperature gradients — potentially into the mid-latitude circulation.

Mountain regions are warming in an elevation-dependent manner, so that high-altitude environments and the water towers they sustain are changing faster than adjacent lowlands (*high confidence*).

A second domain of amplified change is the world's mountains. Synthesis of observations across mountain systems shows that warming rates frequently vary with elevation, and that in many ranges the higher elevations are warming faster than lower ones (Pepin et al., 2022). This elevation-dependent warming arises from a combination of mechanisms whose relative importance differs by range and season: the upward migration of the snow line and the associated snow-albedo feedback, changes in cloud cover and in the altitude of the freezing level, differences in the radiative effects of water vapour with height, and the influence of aerosols and land-surface change. The pattern is not universal —

some ranges and some seasons show the opposite gradient — but the prevalence of amplified high-elevation warming is a robust feature of the observational record.

The stakes are high because mountains function as the water towers of the continents. Snow and ice accumulated at altitude release meltwater through the warm season, buffering downstream rivers and sustaining irrigation, hydropower, and municipal supply for populations that far exceed those living in the mountains themselves. Accelerated high-elevation warming shortens the snow season, shifts precipitation from snow to rain, advances the timing of melt, and drives the long-term loss of glacier mass. These changes alter the seasonal distribution of river flow even where annual totals are little changed, moving water away from the dry-season months when demand is greatest. The role of groundwater as a complementary buffer becomes correspondingly more important, and that buffer is itself under stress (Scanlon et al., 2023). Elevation-dependent warming thus links a high-altitude physical signal to lowland water security across large regions of Asia, the Americas, and Europe.

Urbanisation superimposes a further, locally intense warming signal, and exposure to urban heat is distributed inequitably across and within cities (*high confidence*).

Superimposed on regional-scale warming is the urban heat island, a local modification of climate produced by the replacement of vegetated surfaces with heat-retaining materials, by the geometry of the built environment, by anthropogenic heat release, and by the reduction of evaporative cooling. Analysis of major cities shows that the intensity of the urban heat island — and therefore the excess heat to which urban populations are exposed — is not shared equally. Across major cities, exposure to elevated urban heat-island intensity is disproportionately concentrated among particular populations, so that the burden of urban heat falls unevenly along lines that track other dimensions of disadvantage (Hsu et al., 2021). Because more than half of humanity now lives in cities, the urban heat island transforms a regional warming signal into a public-health hazard experienced at the scale of neighbourhoods and individuals.

The health dimension of this exposure is documented in the annual assessments of health and climate change, which track the growing exposure of vulnerable populations to extreme heat alongside a range of other climate-sensitive health indicators (Romanello et al., 2022; Romanello et al., 2023; Romanello et al., 2024). Extreme heat is among the most direct pathways by which climate change harms human health, and its effects are mediated by age, by pre-existing illness, by occupation, and by the local intensity of the heat island. Extreme weather and climate events more broadly impose measurable burdens on population health and on the health systems that must absorb surges in demand during and after events (Ebi et al., 2021). The regional and local structure of warming therefore cannot be separated from the human geography of vulnerability: the same degree of background warming produces very different outcomes depending on where and to whom the heat is delivered.

Regional changes are consistent with, and best interpreted through, the framework of a biosphere approaching multiple planetary limits (*medium confidence*).

The regional temperature signals described above do not occur in isolation from other large-scale environmental changes. They are one expression of a broader reorganisation of the Earth system that has been characterised through the lens of planetary boundaries — the biophysical limits within which humanity can operate safely. Recent assessment finds that the Earth system has now transgressed the majority of these boundaries, with climate change among them (Richardson et al., 2023). This framing situates regional climate change within a set of interacting pressures — on freshwater, on biogeochemical cycles, on land systems, and on the integrity of the biosphere — that reinforce one another. The condition of the biosphere in what has been termed the Anthropocene is one of tightly coupled human and natural systems in which regional climate change is both a driver and a consequence of wider change (Folke et al., 2021).

The practical relevance of this systemic view is that regional impacts rarely arrive through a single channel. A warming, drying region may simultaneously experience declining crop yields, stressed water supplies, intensifying wildfire, and rising heat-related mortality, each interacting with the others. Comprehensive reviews of the impacts of climate change document this multiplicity across sectors and regions and emphasise that adaptation and mitigation must be pursued together (Abbass et al., 2022). The remainder of this chapter therefore treats regional change not as a catalogue of independent hazards but as an interacting system, beginning with the extremes through which regional climate change is most acutely experienced.

3.3 Regional Extreme Events

Hot extremes have become more frequent, more intense, and longer-lasting across most land regions, and their health consequences are amplified in cities (*high confidence*).

Among all classes of extreme, changes in heat are the most directly attributable to warming and among the most consequential for human welfare. As the distribution of daily temperatures shifts toward warmer values, the frequency of days exceeding any fixed high threshold rises disproportionately, and events that would once have been rare become common. The physical logic is simple, but its human expression is shaped strongly by regional and local factors. In cities the background warming is compounded by the urban heat island, whose intensity varies from block to block and whose burden falls unevenly across populations (Hsu et al., 2021). Nighttime minimum temperatures, which govern the body's ability to recover from daytime heat, rise particularly fast in urban cores, extending the physiological stress of heat episodes.

The consequences for health are documented and growing. Extreme weather and climate events, of which heatwaves are the clearest example, impose measurable burdens on population health and on the capacity of health systems to respond (Ebi et al., 2021). Successive annual assessments record rising exposure of vulnerable groups — the elderly, the very young, those with chronic illness, and outdoor workers — to dangerous heat, alongside heat-related mortality, lost labour capacity, and pressure on health services (Romanello et al., 2022; Romanello et al., 2023; Romanello et al., 2024). These burdens are regionally structured: they are largest where high temperatures coincide with high humidity, dense and heat-retaining urban form, ageing populations, and limited access to cooling. The same absolute warming therefore produces very different mortality outcomes in different regions and different neighbourhoods, and any assessment of heat risk that stops at the physical hazard omits the social determinants that govern who is actually harmed.

Heat extremes also propagate beyond direct health effects. High temperatures desiccate vegetation and soils, priming landscapes for fire; they raise evaporative demand, deepening drought; and they stress energy and water systems at precisely the moments of peak demand. The tendency of heat to co-occur with, and to intensify, other hazards makes it a central node in the compound-event problem discussed below.

Heavy precipitation is intensifying, with the strongest and most hazardous changes concentrated in short-duration, sub-daily extremes (*high confidence*).

A warmer atmosphere holds more water vapour, and this thermodynamic constraint underlies a general expectation that heavy precipitation will intensify as the climate warms. The observational and modelling record bears this out, but with an important refinement: the intensification is not uniform across timescales. Short-duration rainfall extremes — the convective downpours that fall over minutes to hours — are intensifying particularly strongly, in places at rates that exceed the simple thermodynamic scaling expected from the increase in atmospheric moisture alone (Fowler et al., 2021). This super-scaling behaviour reflects dynamical feedbacks within intense convective storms, in which the latent heat released by heavier rainfall invigorates the updraught and concentrates precipitation further. The anthropogenic intensification of these short-duration extremes is now an identifiable feature of the record (Fowler et al., 2021), and it matters disproportionately for hazard because it is sub-daily bursts of rainfall that drive flash flooding, urban drainage failure, and debris flows.

The translation of intensifying rainfall into flood risk is mediated by exposure and vulnerability, and here the regional picture is stark. A large share of the world's population exposed to serious flood hazard lives in poverty, so that the communities least able to prepare for, absorb, and recover from floods are precisely those most exposed to them (Rentschler et al., 2022). The same rainfall extreme that causes minor disruption in a well-drained, well-insured

city can be catastrophic in an informal settlement on a floodplain. Flood risk is therefore not simply a function of changing precipitation; it is co-produced by where people live, by the built and natural systems that convey and store water, and by the resources available for protection and recovery. Because the sharpest changes are in the short-duration extremes that overwhelm drainage most readily (Fowler et al., 2021), and because exposure is concentrated among the poor (Rentschler et al., 2022), the flood consequences of a warming climate are expected to fall with particular force on vulnerable populations.

Heavy-precipitation change also interacts with the cryospheric and mountain signals discussed earlier. As warming shifts high-elevation precipitation from snow to rain and advances melt, rain-on-snow events and rapid runoff become more likely, and the buffering capacity of the seasonal snowpack — and of groundwater (Scanlon et al., 2023) — is eroded. The regional flood regime is thus reshaped from several directions at once.

Drought is changing in character, with a documented global shift toward rapid-onset “flash” droughts and toward severe multi-year megadroughts in vulnerable regions (*high confidence*).

Drought is a more complex hazard than heat or heavy rain because it integrates precipitation deficits, evaporative demand, and the state of soils, vegetation, and water stores over time. Two changes in drought behaviour are especially well documented and especially consequential at the regional scale. The first is a global shift toward more rapidly developing droughts. Analysis of the observational record identifies a transition under climate change toward “flash” droughts — episodes that intensify over days to weeks rather than months — driven by the combination of precipitation deficits with the high evaporative demand of a warmer atmosphere (Yuan et al., 2023). Flash droughts are dangerous precisely because their speed of onset compresses the time available for detection and response, catching agriculture and water management unprepared.

The second is the emergence of severe, persistent multi-year droughts in several regions. The most thoroughly documented case is southwestern North America, where the drought that began in the early twenty-first century intensified into one of the most severe multi-decadal dry episodes in a paleoclimate record extending back many centuries, with anthropogenic warming making a substantial contribution to its severity (Williams et al., 2022). Such megadroughts differ qualitatively from ordinary drought: their duration exhausts the multi-year buffers — reservoirs, soil moisture, and especially groundwater — on which regions rely to survive shorter dry spells. Groundwater is central to this story. It is the resource of last resort during prolonged drought, yet it is under simultaneous pressure from over-extraction and from a changing climate that reduces recharge in many regions (Scanlon et al., 2023). Where megadrought and groundwater depletion coincide, the resilience of a region’s water system can be eroded to the point of structural crisis.

Drought also feeds directly into the fire and heat hazards with which it is physically linked. Desiccated soils and vegetation lower evaporative cooling,

raising temperatures and deepening the drought in a self-reinforcing loop, while providing the dry fuel that transforms ignition into large fire. This coupling is the basis of the compound events considered next.

Tropical cyclone hazard is being reshaped by warming through intensity, rainfall, and coastal impact, even where basin-wide frequency changes remain uncertain (*medium confidence*).

Tropical cyclones concentrate several of the most destructive hazards — extreme wind, storm surge, and torrential rain — into single events, and their behaviour under warming is an area of active assessment. The evidentiary basis for tracking cyclone behaviour rests on carefully constructed regional best-track databases, such as the long-term record of western North Pacific cyclones maintained by the China Meteorological Administration (Lu et al., 2021), which underpin analyses of trends in frequency, intensity, and track. Interpreting these records requires care because observing practices and instrumentation have changed over the decades they span, and disentangling genuine climate signals from changes in observation is a central methodological challenge.

The physical expectations under warming are clearest for those aspects governed by thermodynamics. A warmer ocean provides more energy and a warmer atmosphere holds more moisture, favouring higher peak intensities and heavier cyclone rainfall even if the total number of storms does not increase. Because a warmer atmosphere intensifies short-duration rainfall generally (Fowler et al., 2021), the freshwater flooding produced by slow-moving or intense cyclones is a growing component of their hazard, distinct from wind and surge. Rising sea level raises the baseline on which storm surge acts, so that a cyclone of unchanged intensity produces higher coastal water levels and reaches further inland than it would have in the past. The regional distribution of cyclone impact is therefore shifting even where basin-wide frequency trends remain uncertain, and the interaction of surge, rainfall, and antecedent conditions makes cyclones a leading source of compound coastal flooding.

Compound and cascading events, in which hazards coincide or trigger one another, are an increasingly important and under-assessed dimension of regional risk (*high confidence*).

The hazards treated separately above rarely occur in isolation. Heat, drought, and fire are physically coupled; heavy rain, storm surge, and river flooding can combine at coasts; and the failure of one system can cascade into others. Compound events — the coincidence, clustering, or sequencing of hazards, or the combination of hazards with vulnerable conditions — frequently produce impacts far larger than any single hazard would, yet conventional single-hazard risk assessment tends to miss them. A framework that treats climate risk as complex, interacting, and cascading is required to represent these dynamics adequately (Simpson et al., 2021).

The compound heat-drought-fire-smoke sequence is the paradigmatic example at the regional scale. Prolonged drought and extreme heat desiccate landscapes; ignition then produces fire whose extent and intensity are governed by the accumulated dryness. The extreme forest fires of southeastern Australia illustrate how climate change and variability connect to such events, with warming and drying loading the dice toward larger and more intense fires (Abram et al., 2021). The cascade does not end at the fire perimeter. Wildfire smoke transports fine particulate matter across large regions, and this smoke has been shown to harm respiratory health more than an equivalent mass of fine particles from other sources (Aguilera et al., 2021), so that a fire in one location degrades air quality and health across a much wider area. A single climatic season can thus link heat, drought, fire, air pollution, and health impacts in a chain that no single-hazard framework captures.

Compound events also arise at coasts, where cyclone rainfall, storm surge, and high river flow can coincide; in cities, where heat and power failure can combine to remove cooling at the moment it is most needed; and across sectors, where a climate shock to water or energy propagates into agriculture, health, and economic activity. The systemic and planetary-boundary framing introduced earlier (Richardson et al., 2023), set within the coupled human-natural biosphere (Folke et al., 2021), is the natural home for such analysis, because it treats the coupled human-natural system as the unit of risk. Assessing regional climate risk therefore requires moving beyond the enumeration of individual extremes toward the interactions among them (Simpson et al., 2021), a shift that also raises the scientific bar for the attribution methods considered in the next section.

3.4 Event Attribution Science

Event attribution has matured into a quantitative discipline capable of estimating the human influence on many individual extremes (*high confidence*).

The scientific question that the public most often asks after a disaster — was this event caused by climate change? — is answerable only in a probabilistic sense, and the discipline that has grown up to answer it is extreme-event attribution. Its central logic is comparative: the likelihood or intensity of an observed event is estimated in the climate as it is, and compared with its likelihood or intensity in a counterfactual climate without human influence. The difference between the two is expressed as a change in probability or in magnitude attributable to anthropogenic forcing. This framing sidesteps the false dichotomy of single causation and replaces it with a rigorous statement about how human influence has shifted the odds.

The maturation of attribution rests on several developments that recur throughout this chapter. The first is the physical understanding of how warming changes each class of hazard: the thermodynamic intensification of heat

and of heavy rainfall, the coupling of drought to evaporative demand, and the reshaping of cyclone intensity and rainfall. Where the physical mechanism is well understood and strongly expressed — as for heat extremes and for short-duration rainfall (Fowler et al., 2021) — attribution statements can be made with high confidence. The second is the availability of long, carefully constructed observational records against which to characterise the rarity of events and to detect trends; the regional cyclone databases exemplify the kind of curated record on which attribution depends (Lu et al., 2021). The third is the ensemble of climate-model simulations that allow the counterfactual climate to be constructed and sampled.

Attribution is now applied across the full range of extremes considered in this chapter. For heat, the shift of the temperature distribution makes attribution comparatively straightforward and the resulting statements robust. For heavy precipitation, the identification of anthropogenic intensification in the sub-daily record provides a physical foundation for event-level attribution (Fowler et al., 2021). For drought, attribution must account for both precipitation and evaporative demand, and the documented human contribution to the severity of the southwestern North American megadrought is a landmark result (Williams et al., 2022). The global shift toward flash droughts likewise represents an attributable change in the character of the hazard rather than merely in its frequency (Yuan et al., 2023). For fire, attribution connects the climatic drivers of fire weather to observed extreme events, as in the case of the southeastern Australian fires (Abram et al., 2021), while recognising that ignition, fuel management, and land use also govern outcomes.

The credibility of attribution depends on the fidelity of climate models and on the coordinated experimental frameworks that structure their use (*high confidence*).

Attribution is only as reliable as the models used to construct the counterfactual climate and to simulate the statistics of extremes. The coordinated multi-model experiments that underpin regional climate projection also underpin attribution, by providing large ensembles run under carefully specified forcings. The scenario-based experiments that span the range of plausible futures are the backbone of this enterprise (Tebaldi et al., 2021), and the same modelling infrastructure that projects future change enables the simulation of the factual and counterfactual present. Confidence in an attribution statement rests on the ability of models to reproduce the observed statistics of the relevant extreme, and this fidelity varies by hazard: models represent large-scale heat extremes well, but the sub-daily convective processes that drive the most intense short-duration rainfall are only partially resolved at typical model resolutions (Fowler et al., 2021), which is one reason those extremes are assessed with a mixture of observational and modelling evidence.

The field has increasingly drawn on advances in the modelling of weather itself. Machine-learning weather prediction has demonstrated skill in medium-range global forecasting that rivals or exceeds traditional physics-based systems

(Bi et al., 2023), and data-driven approaches have shown skill in precipitation nowcasting, the very short-range prediction of rainfall that is most relevant to flash-flood hazard (Ravuri et al., 2021). While these tools were developed for forecasting rather than attribution, they expand the toolkit available for characterising extremes and for generating the large ensembles on which robust statistical statements depend. Their emergence illustrates a broader trend in which computational advances lower the cost of the large-ensemble experiments that attribution requires, potentially allowing rapid, operational attribution of events as they occur.

Attribution science provides the physical basis for valuing climate damages and for framing low-probability, high-impact risks (*high confidence*).

Beyond its scientific interest, attribution has become foundational to the economics and governance of climate change. The estimation of the social cost of carbon dioxide — the monetised damage caused by an additional tonne of emissions — depends on translating physical climate change into impacts, and comprehensive assessment now places this cost substantially higher than many values embedded in current regulation (Rennert et al., 2022). The chain from emissions to cost runs through exactly the regional extremes considered in this chapter: heat mortality, flood damage, drought-driven agricultural loss, and cyclone destruction. Improved attribution of these impacts to warming sharpens the damage functions on which the social cost rests, and the direction of recent revision has been toward higher costs as the evidence for severe and compounding regional impacts has accumulated.

Attribution also clarifies the limits of what can be quantified, and thereby motivates explicit attention to catastrophic tail risks. Much of the assessed distribution of climate outcomes is amenable to attribution and projection, but the low-probability, high-impact tail — the possibility of abrupt or extreme change with civilisation-scale consequences — is precisely the region that standard methods struggle to bound. The systematic exploration of catastrophic climate-change scenarios has been advanced as a necessary complement to the assessment of central outcomes (Kemp et al., 2022), on the grounds that the events of greatest concern for risk management are those whose probability is hardest to estimate. This concern connects to the planetary-boundaries framing introduced earlier: the transgression of multiple boundaries (Richardson et al., 2023) raises the possibility of interacting, nonlinear responses that would fall outside the range that attribution of past events can constrain. Attribution science, by delineating what is and is not currently quantifiable, thus performs a dual service — grounding the valuation of expected damages while flagging the tail risks that expected-value calculations omit.

The regional and equity dimensions of extremes are integral to attribution's practical meaning (*high confidence*).

An attribution statement about the changed odds of a physical event acquires its human significance only when combined with the geography of exposure and vulnerability. The same attributed intensification of flooding carries very different meaning in a wealthy, well-defended city and in a poor community on an exposed floodplain, and the concentration of flood exposure among the poor (Rentschler et al., 2022) means that the human costs of attributable changes in rainfall extremes (Fowler et al., 2021) fall disproportionately on those least responsible for and least able to withstand them. Attribution of heat extremes similarly gains meaning through the urban and social structure of exposure (Hsu et al., 2021) and through the documented health burdens of extreme weather (Ebi et al., 2021) and the annual health assessments (Romanello et al., 2023).

This equity dimension has consequences for how attribution is used. Attribution increasingly informs questions of loss, damage, and responsibility, because it makes possible statements about the human contribution to specific harmful events. The scientific caution appropriate to such statements — acknowledging the dependence on model fidelity, on record length, and on the class of hazard — must therefore be paired with clarity about the uneven distribution of the harms being attributed. Attribution that ignores this distribution risks describing a physically changed hazard while missing the socially produced disaster. The comprehensive-review literature on climate impacts and adaptation makes this coupling explicit (Abbass et al., 2022), and the framework of complex, cascading risk (Simpson et al., 2021) provides the structure within which attributed physical changes and their unequal human consequences can be assessed together. The next section applies this integrated view to the monsoon systems on which a large share of humanity depends.

3.5 Regional Monsoons, Projections, and Adaptation

The regional monsoon systems that govern water security for most of humanity are responding to competing greenhouse-gas and aerosol forcings, and their projection remains a leading source of regional uncertainty (*medium confidence*).

The monsoon systems of Asia, Africa, Australia, and the Americas organise the annual cycle of rainfall for a majority of the world's population. Their seasonal reversal of winds and the associated wet season deliver the water on which agriculture, hydropower, and daily life depend across vast regions, and even modest changes in monsoon timing, total, or variability propagate immediately into human welfare. Monsoons are shaped by the land-sea thermal contrast, by the distribution of atmospheric moisture, and by the large-scale circulation, and each of these is affected by anthropogenic forcing in ways that

do not all point in the same direction. Greenhouse-gas forcing tends, through the thermodynamic increase in atmospheric moisture, to strengthen monsoon rainfall; aerosol forcing, by cooling the surface and altering the land-sea contrast and atmospheric stability, has historically tended to weaken several monsoon systems. The regional outcome depends on the balance between these forcings, which differs by basin and evolves over time as aerosol emissions change.

This competition of forcings is one reason monsoon projection remains among the most uncertain elements of regional climate assessment. The coordinated scenario experiments that structure projection (Tebaldi et al., 2021) span the range of plausible greenhouse-gas and, in some designs, aerosol futures, but the translation of a given global forcing into monsoon rainfall depends on regional processes – convection, land-atmosphere coupling, the representation of aerosol-cloud interactions – whose simulation still differs across models. The result is that models often agree on the direction of thermodynamic strengthening while disagreeing on the dynamical and aerosol-mediated components, leaving substantial spread in projected monsoon change at the regional scale. Sub-daily convective extremes within the monsoon, which drive much of its flood hazard, are subject to the same resolution limitations that affect short-duration rainfall generally (Fowler et al., 2021), so that changes in monsoon intensity at the timescales most relevant to flooding are assessed with particular caution.

The stakes of this uncertainty are amplified by the sensitivity of the systems that depend on the monsoon. Agriculture is the most immediate: the emergence of climate impacts on global crop production in the latest generation of coupled climate and crop models (Jägermeyr et al., 2021) shows that yield changes are appearing earlier and more clearly than previously assessed, and monsoon regions are central to this signal because their food systems are so tightly coupled to seasonal rainfall. Water resources are equally exposed. Where monsoon rainfall feeds the rivers and recharges the aquifers on which dry-season supply depends, changes in monsoon total or variability interact with the groundwater stresses documented globally (Scanlon et al., 2023) and with the mountain-derived flows discussed earlier (Pepin et al., 2022). A weakened or more erratic monsoon, superimposed on depleted groundwater and diminished snowpack, would compound water insecurity across some of the most populous regions on Earth.

Regional projections rest on coordinated multi-model experiments, whose spread must be interpreted alongside model fidelity and the representation of regional processes (*high confidence*).

The scientific basis for statements about future regional climate is the ensemble of projections produced under coordinated experimental protocols. The scenario-based experiments that span a range of socioeconomic and emissions futures provide the structured set of projections from which regional change is assessed (Tebaldi et al., 2021). Their value lies in comparability: because participating models run the same scenarios, the spread among their projections

can be interpreted as a measure of structural uncertainty, and the response to different scenarios can be isolated from the response to model differences. For many regional variables — mean warming, the broad intensification of heat extremes — the models agree well and confidence is correspondingly high. For others — regional precipitation, monsoon change, the details of circulation — the spread is larger and reflects genuine uncertainty in how regional processes respond.

Interpreting projection spread requires attention to model fidelity, which varies by variable and by region. Large-scale thermodynamic responses are robust across models; smaller-scale dynamical and convective responses are less so. The advances in data-driven weather modelling noted earlier (Bi et al., 2023), and in precipitation nowcasting (Ravuri et al., 2021), point toward a future in which higher-resolution and hybrid approaches may narrow some of this spread, particularly for the convective extremes that current models resolve poorly (Fowler et al., 2021). In the meantime, the honest communication of regional projection uncertainty — distinguishing what is robust from what is not — is essential, because regional adaptation decisions must be made under precisely the uncertainty that the projections quantify. The framework of complex climate risk (Simpson et al., 2021) offers a structure for making decisions under this uncertainty by focusing on the interactions among hazards, exposures, and vulnerabilities rather than on any single projected variable.

Regional impacts propagate through tightly coupled human and natural systems, so that physical change is transmitted into food, water, health, and economic outcomes (*high confidence*).

The regional changes assessed in this chapter do not remain confined to the physical climate system. They propagate through the coupled human-natural systems of the biosphere (Folke et al., 2021) into the sectors on which human welfare depends. The pathway into food is direct and increasingly well quantified: climate impacts on crop yields are emerging earlier than earlier assessments projected (Jägermeyr et al., 2021), and in monsoon-dependent regions the coupling of yields to seasonal rainfall makes agriculture an immediate transmitter of climate variability and change. The pathway into water runs through the mountain snowpack (Pepin et al., 2022), the groundwater buffer (Scanlon et al., 2023), and the monsoon, whose combined stresses shape regional water security. The pathway into health runs through extreme heat, floods, drought, and wildfire smoke, whose burdens on population health and health systems are documented and growing (Ebi et al., 2021; Aguilera et al., 2021) and are tracked in the annual Lancet Countdown (Romanello et al., 2024).

The economic pathway integrates these physical and sectoral impacts into monetary terms, and its assessment has moved toward higher estimates of climate damage. The social cost of carbon dioxide, which aggregates the marginal damage of emissions across sectors and regions, is now assessed as substantially higher than many values in regulatory use (Rennert et al., 2022), reflecting the accumulating evidence for severe regional impacts. Energy systems are

simultaneously exposed to and part of the response to climate change: the cost, environmental impact, and resilience of renewable energy under a changing climate are themselves climate-sensitive (Osman et al., 2022), so that the infrastructure of mitigation must be designed for the altered regional climate it is intended to help stabilise. These transmission pathways, taken together, mean that a regional physical change of given magnitude can produce impacts of widely varying severity depending on the exposure and adaptive capacity of the systems it encounters — the theme to which the final subsection turns.

Adaptation to regional climate change is under way but uneven, and can reduce or, if poorly designed, deepen vulnerability (*high confidence*).

Regional climate risk is produced jointly by physical hazard, by exposure, and by vulnerability, and the last of these is the most amenable to deliberate action. Adaptation — the adjustment of human and natural systems to actual or expected climate and its effects — is now occurring across regions and sectors, but its coverage is uneven and its effectiveness is not guaranteed. Assessment of adaptation interventions in developing countries finds that their effect on vulnerability is mixed: some reduce it, some have little effect, and some redistribute or even deepen vulnerability, for example by benefiting some groups at the expense of others or by entrenching the conditions that produced vulnerability in the first place (Eriksen et al., 2021). This finding is a caution against treating adaptation as an unambiguous good; the design, governance, and equity of interventions determine whether they achieve their intended reduction in risk.

The equity structure of regional risk makes this caution especially important. Flood exposure is concentrated among the poor (Rentschler et al., 2022); urban heat falls disproportionately on disadvantaged populations (Hsu et al., 2021); and the health burdens of extremes are borne most heavily by the vulnerable (Ebi et al., 2021) and are documented in the annual health assessments (Romanello et al., 2024). Adaptation that does not attend to this structure risks protecting those already best defended while leaving the most exposed no safer, or worse. The comprehensive assessment of impacts, adaptation, and mitigation emphasises that the three must be pursued together and with attention to their distributional consequences (Abbass et al., 2022), and the framework of complex, cascading risk (Simpson et al., 2021) provides the means to assess whether a given intervention reduces risk across the interacting system or merely displaces it.

Set against the possibility of well-designed adaptation is the shadow of the catastrophic tail. The systematic exploration of extreme scenarios (Kemp et al., 2022) and the transgression of multiple planetary boundaries (Richardson et al., 2023) are reminders that adaptation has limits, and that beyond some magnitude of change the coupled human-natural systems on which regions depend (Folke et al., 2021) may be pushed past thresholds that no feasible adaptation can accommodate. The assessment of regional climate and extreme events therefore concludes on two linked findings. First, the physical evidence for regional

change — amplified warming in the Arctic (Rantanen et al., 2022) and mountains (Pepin et al., 2022), intensifying heat and short-duration rainfall (Fowler et al., 2021), shifting drought (Yuan et al., 2023) and megadrought (Williams et al., 2022), and compounding hazards (Abram et al., 2021; Simpson et al., 2021) — is strong and, through attribution, increasingly linked to human influence. Second, the human consequences of this change are governed as much by exposure, vulnerability, and the quality of adaptation and mitigation as by the physical hazard itself, so that the task of regional climate assessment is inseparable from the assessment of the human systems through which climate is experienced.

3.6 Regional Syntheses and Cross-Sectoral Transmission

The Arctic is being reorganised across its physical, ecological, and human systems at a pace unmatched elsewhere on the planet (*high confidence*).

The near-fourfold amplification of Arctic warming relative to the global mean (Rantanen et al., 2022) is not an isolated statistic but the leading edge of a systemic reorganisation. The retreat of sea ice removes a reflective surface, alters the seasonal exchange of heat and moisture between ocean and atmosphere, and lengthens the open-water season, with consequences that cascade through the marine ecosystem and into the atmospheric circulation. On land, warming thaws permafrost, destabilising the ground on which infrastructure rests and mobilising stores of previously frozen organic carbon. The seasonality of the change compounds its impact: because Arctic amplification is concentrated in the cold season, the winter climate that has historically defined the region is changing fastest, shortening the period of stable sea ice and frozen ground on which both ecosystems and communities depend.

The implications reach far beyond the Arctic itself. Changes in the temperature contrast between the warming Arctic and the lower latitudes have been proposed as a mechanism by which Arctic change could influence the mid-latitude circulation, potentially altering the frequency or persistence of weather patterns that produce heat, cold, and precipitation extremes at lower latitudes. The magnitude and even the sign of such teleconnections remain contested, and this uncertainty is one of the more consequential open questions in regional climate science, because it links the most rapidly changing region to the extremes experienced by the majority of the world's population. What is not contested is that the Arctic has become a bellwether: the region where the coupled response of the climate system to forcing is expressed most rapidly and most completely, and where the interaction of physical change with ecosystems, carbon, and human communities can be observed in advance of its appearance elsewhere. The condition of the Arctic is, in this sense, an early reading of

the trajectory of the wider biosphere (Folke et al., 2021) and of the planetary boundaries the Earth system has begun to transgress (Richardson et al., 2023).

The world's mountains transmit high-elevation warming into lowland water insecurity, coupling a cryospheric signal to the livelihoods of billions (*high confidence*).

Elevation-dependent warming (Pepin et al., 2022) acquires its human importance through the hydrological role of mountains. The snow and ice stored at altitude constitute a natural reservoir that captures winter precipitation and releases it as meltwater through the warm, dry months, sustaining the rivers of Asia, the Americas, and Europe precisely when rainfall is scarce and demand is high. Accelerated high-elevation warming attacks this reservoir from several directions at once: it raises the elevation at which precipitation falls as rain rather than snow, reducing the volume of the seasonal snowpack; it advances the timing of melt, shifting river flow earlier in the year and away from the dry-season months of greatest need; and it drives the long-term retreat of glaciers, eroding the multi-year buffer that stabilises flow in the driest years.

The downstream consequence is a growing reliance on the groundwater that complements surface water in sustaining dry-season supply, and this buffer is itself under pressure from over-extraction and from a changing climate that alters recharge (Scanlon et al., 2023). Where the diminishing snowpack, retreating glaciers, and stressed aquifers of a mountain-fed basin coincide, the resilience of the water system can be eroded even where annual precipitation totals are little changed, because it is the seasonal and interannual buffering, not the annual mean, that governs security of supply. The interaction of these mountain and groundwater signals with the monsoon rainfall discussed below means that some of the most populous river basins on Earth face simultaneous pressure on every component of their water balance. This is a paradigmatic case of the coupled, cross-system transmission through which regional physical change becomes human risk.

Dryland and Mediterranean-type regions are experiencing a self-reinforcing convergence of heat, drought, and fire (*high confidence*).

A distinct regional syndrome is emerging in the world's dryland and Mediterranean-type climates, where the physical coupling of heat, drought, and fire produces a self-reinforcing spiral. Warming raises evaporative demand, which deepens drought by drawing moisture from soils and vegetation; the resulting desiccation lowers evaporative cooling, which raises temperatures further; and the accumulated dryness converts ignition into extensive, intense fire. The southwestern North American megadrought exemplifies the drought pole of this syndrome, having intensified into one of the most severe multi-decadal dry episodes in a centuries-long paleoclimate record, with a substantial anthropogenic contribution (Williams et al., 2022). The global shift toward flash droughts (Yuan et

al., 2023) adds a temporal dimension, compressing the onset of dryness and reducing the warning available to water and agricultural managers.

The fire pole of the syndrome is illustrated by the extreme forest fires of southeastern Australia, whose extent and intensity were connected to climate change and variability (Abram et al., 2021). The consequences of such fires extend far beyond the burned area: wildfire smoke transports fine particulate matter across entire regions and has been shown to harm respiratory health more than an equivalent mass of particulate matter from other sources (Aguilera et al., 2021), so that a fire in one location becomes a public-health event across a much larger domain. The convergence of these hazards in dryland regions is a leading example of the compound and cascading dynamics that single-hazard assessment misses and that a framework of complex, interacting risk is designed to capture (Simpson et al., 2021). It is also a domain in which the equity dimension is acute, because the health burdens of heat and smoke fall most heavily on the vulnerable (Ebi et al., 2021) and are tracked in the annual assessments of health and climate change (Romanello et al., 2023).

Densely populated coastal and low-lying regions face compounding flood hazard from rainfall, rivers, surge, and rising seas (*high confidence*).

Coastal and low-lying regions concentrate population, economic activity, and exposure, and they are subject to a distinctive compounding of flood hazards. Intensifying short-duration rainfall (Fowler et al., 2021) raises the risk of pluvial and flash flooding within cities; river flooding responds to rainfall across upstream catchments; storm surge is driven by cyclones and other storms; and rising sea level lifts the baseline on which all coastal water levels act. When these hazards coincide — a cyclone delivering extreme rainfall while its surge meets a high river stage on a raised sea level — the resulting compound flood can far exceed what any single driver would produce, precisely the dynamic that a complex-risk framework is intended to represent (Simpson et al., 2021). The regional cyclone records that document the storms driving surge and rainfall, such as the western North Pacific database (Lu et al., 2021), are essential to characterising this hazard, even as the changing intensity and rainfall of cyclones under warming reshape it.

The human consequences of coastal flood hazard are governed by exposure and vulnerability as much as by the physical drivers. A large share of the global population exposed to serious flood hazard lives in poverty (Rentschler et al., 2022), and this concentration is often most acute in the informal settlements of rapidly growing coastal cities, where drainage, defence, and recovery resources are least available. The same compound flood therefore produces radically different outcomes across communities separated by only a few kilometres, according to their exposure and their capacity to prepare and recover. Coastal regions thus combine the highest concentration of compounding physical hazard with some of the sharpest gradients of vulnerability, making them a priority

for the integrated assessment of physical change, exposure, and adaptive capacity.

Regional climate change is transmitted into food, water, health, energy, and disease through tightly coupled pathways, none of which can be assessed in isolation (*high confidence*).

The regional physical changes catalogued in this chapter reach human welfare through a set of transmission pathways that are themselves coupled. The food pathway is increasingly well quantified: climate impacts on global crop production emerge earlier in the latest generation of coupled climate and crop models than in previous assessments (Jägermeyr et al., 2021), and in monsoon- and mountain-fed agricultural regions the coupling of yields to seasonal water availability makes food production an immediate transmitter of climatic variability. The water pathway runs through the mountain snowpack (Pepin et al., 2022), the groundwater buffer (Scanlon et al., 2023), and monsoon rainfall, whose combined and often simultaneous stresses shape regional water security. The health pathway operates through extreme heat, floods, drought, and wildfire smoke, whose burdens on population health and health systems are documented and growing (Ebi et al., 2021) and are tracked in successive annual assessments (Romanello et al., 2024).

Two further pathways deserve emphasis. The first is disease: a changing climate alters the distribution, seasonality, and transmission of infectious disease, and the broader era of global environmental change is reshaping the landscape of infectious risk (Baker et al., 2021), so that regional warming and altered precipitation can shift the geography of vector-borne and water-borne disease. The second is energy, which is both affected by and central to the response to climate change; the cost, environmental impact, and resilience of renewable energy systems are themselves sensitive to the changing climate in which they must operate (Osman et al., 2022), meaning that the infrastructure on which mitigation depends must be designed for the regional climate it is intended to help stabilise. Integrating across these pathways is the aggregate economic signal captured by the social cost of carbon dioxide, now assessed as substantially higher than many regulatory values in use (Rennert et al., 2022) as evidence for severe, compounding regional impacts has accumulated. Comprehensive review of climate impacts, adaptation, and mitigation confirms that these pathways cannot be addressed independently (Abbass et al., 2022): an intervention in one sector reverberates through the others, and the coupled human-natural system (Folke et al., 2021) is the appropriate unit of assessment.

3.7 Uncertainties, Observational Needs, and Assessment Frontiers

The largest regional uncertainties lie in the dynamical and convective processes that govern precipitation, monsoons, and sub-daily extremes (*high confidence*).

The confidence attached to regional projections varies systematically with the physics involved. Thermodynamic responses — the intensification of heat and, at the daily-and-longer scale, of heavy precipitation — are robust across models and assessed with high confidence. Dynamical responses — changes in circulation, storm tracks, monsoon behaviour, and the sub-daily convective extremes that drive flash flooding — are less well constrained, because they depend on processes that current models resolve only partially. The strongest and most hazardous short-duration rainfall extremes arise from convection at scales below the typical resolution of climate models, which is one reason their anthropogenic intensification is assessed through a combination of observational and modelling evidence rather than from models alone (Fowler et al., 2021). Monsoon projection inherits this difficulty, because monsoon rainfall depends on convection, land-atmosphere coupling, and aerosol-cloud interaction, each imperfectly represented, and because the competing thermodynamic and aerosol influences on the monsoon can offset one another to leave a small residual with large fractional uncertainty.

Reducing these uncertainties is a central frontier of regional assessment. Higher-resolution modelling, which begins to resolve convection explicitly, is one avenue; the coordinated experimental frameworks that structure projection (Tebaldi et al., 2021) provide the comparative basis for evaluating progress. The emergence of skilful data-driven weather prediction (Bi et al., 2023) and of data-driven precipitation nowcasting (Ravuri et al., 2021) points toward computational tools that may eventually allow larger ensembles and finer resolution at reduced cost, potentially sharpening the statistics of the convective extremes that current physics-based models resolve poorly. Until such advances mature, the honest communication of regional projection uncertainty — clearly distinguishing the robust thermodynamic signals from the less certain dynamical ones — is itself a scientific obligation, because regional adaptation decisions must be taken under exactly the uncertainty that the assessment quantifies.

Observational records of adequate length, quality, and homogeneity are the foundation on which all regional assessment and attribution rest (*high confidence*).

Every quantitative statement in this chapter depends ultimately on observations, and the quality, length, and homogeneity of the observational record is therefore a first-order determinant of confidence. Detecting a trend in an extreme requires a record long enough to characterise the rarity of the event against the background of natural variability; attributing that trend to human influence requires, in addition, that changes in observing practice not be mistaken for changes in climate. The regional cyclone databases illustrate both the value and the difficulty: a long, carefully constructed record such as the western North

Pacific database (Lu et al., 2021) makes trend analysis possible, but the evolution of observing technology over the decades such records span must be accounted for before a genuine climate signal can be isolated. Similar considerations apply to the records of heat, precipitation, and drought on which the assessment of those extremes depends.

The observational challenge is most acute precisely where the human stakes are highest. Many of the regions most exposed and most vulnerable — including parts of the tropics, the drylands, and rapidly urbanising areas — have the sparsest and least homogeneous long-term records, so that confidence in regional change is often weakest where it matters most. The urban dimension adds a further layer, because characterising the heat to which urban populations are actually exposed requires observations at the scale of neighbourhoods, not the coarse grids of global datasets; the finding that urban heat-island intensity is distributed inequitably across major cities (Hsu et al., 2021) rests on such fine-scale analysis. Sustaining and extending high-quality observational networks, and constructing the carefully homogenised long-term records that attribution requires, is therefore not a supporting activity but a precondition for regional climate assessment, and gaps in the observing system translate directly into gaps in what can be assessed and defended.

The frontier of regional assessment is the integration of physical change with exposure, vulnerability, and the assessment of catastrophic tail risk (*high confidence*).

The trajectory of regional climate assessment is toward integration. The physical assessment of hazard — warming, heat, precipitation, drought, fire, cyclones — is necessary but not sufficient to characterise risk, because the human and ecological consequences of a given hazard depend on who and what is exposed and on their vulnerability and capacity to respond. The concentration of flood exposure among the poor (Rentschler et al., 2022), the inequitable distribution of urban heat (Hsu et al., 2021), and the finding that adaptation interventions can deepen as well as reduce vulnerability (Eriksen et al., 2021) all demonstrate that the human outcome is co-produced by physical and social processes. A framework that treats climate risk as complex, interacting, and cascading (Simpson et al., 2021) provides the structure for this integration, allowing physical hazard, exposure, and vulnerability to be assessed together rather than in sequence.

Integration must also extend to the catastrophic tail that expected-value assessment omits. Much of the distribution of regional outcomes is amenable to projection and, for past events, to attribution; but the low-probability, high-impact tail — the possibility of abrupt or extreme change with consequences at the scale of societies — is precisely the region that standard methods struggle to bound. The systematic exploration of catastrophic scenarios has been advanced as a necessary complement to the assessment of central outcomes (Kemp et al., 2022), and the transgression of multiple planetary boundaries (Richardson et al., 2023) raises the prospect of interacting, nonlinear responses that lie outside the range constrained by the attribution of past events. Regional

assessment therefore faces a dual task: to quantify, with appropriate and clearly communicated confidence, the expected regional changes and their unevenly distributed human consequences; and to flag, without overstating what can be quantified, the tail risks whose management may ultimately matter most. Meeting both is the frontier of the assessment enterprise, and it is inseparable from the assessment of the coupled human-natural systems (Folke et al., 2021) through which regional climate is, in the end, experienced.

3.8 Regional Extremes in Depth: Heat, Water, and the Monsoon

Extreme heat is the deadliest and most attributable of the regional extremes, and its human toll is governed by humidity, urban form, and social vulnerability (*high confidence*).

Among the classes of extreme assessed in this chapter, heat combines the clearest physical attribution with among the largest and most direct effects on human life. Because warming shifts the entire distribution of daily temperatures toward higher values, the probability of crossing any fixed high threshold rises disproportionately, so that events once considered rare become routine and previously unobserved extremes become possible. The physiological danger of heat, however, is not captured by air temperature alone. High humidity impairs the body's ability to shed heat through the evaporation of sweat, so that the combination of heat and humidity — not temperature in isolation — determines the physiological limit. Regions where extreme heat coincides with high humidity therefore face a hazard qualitatively more dangerous than dry heat of the same temperature, and the geography of this combined hazard does not coincide with the geography of temperature alone.

The urban dimension sharpens the hazard further. The urban heat island raises temperatures in built-up areas relative to their surroundings, with the greatest effect often at night, when elevated minimum temperatures deny the body the recovery that cooler nights would otherwise provide. The intensity of this effect, and therefore the excess heat exposure it imposes, is distributed inequitably across major cities, falling disproportionately on particular populations (Hsu et al., 2021). The consequences for health are documented and growing: extreme weather and climate events impose measurable burdens on population health and on the capacity of health systems to respond (Ebi et al., 2021), and successive annual assessments record rising exposure of vulnerable groups to dangerous heat, together with heat-related mortality and lost labour capacity (Romanello et al., 2022; Romanello et al., 2023; Romanello et al., 2024). The people most at risk — the elderly, the chronically ill, the very young, and those who work or live without access to cooling — are frequently the same people most exposed to the intensified urban heat island, so that physical and social vulnerability compound. Heat is thus the extreme in which the coupling of physical hazard to social vulnerability is most direct, and in which the dif-

ference between the physical signal and the human disaster is most starkly a function of exposure and adaptive capacity.

Heat also acts as a multiplier on other hazards. It raises evaporative demand, deepening drought (Yuan et al., 2023); it desiccates fuels, priming landscapes for fire (Abram et al., 2021); and it stresses energy and water systems at the moment of peak demand, when the failure of cooling can itself become lethal. The tendency of heat to co-occur with and intensify other hazards makes it a central node in the compound-event problem, and any assessment of regional heat risk that stops at the physical hazard omits both the humidity that governs its lethality and the social structure that governs who is harmed.

Water is the medium through which most regional climate change is ultimately felt, from drought and flood to the seasonal reliability of supply (*high confidence*).

If heat is the deadliest of the extremes, water is the medium through which the widest range of regional climate impacts is transmitted. The same warming that intensifies heat also intensifies the water cycle, but unevenly and in ways that increase both flood and drought hazard, sometimes in the same region in different seasons. On the flood side, the intensification of short-duration rainfall extremes (Fowler et al., 2021) raises the risk of flash and urban flooding, and the concentration of flood exposure among the poor (Rentschler et al., 2022) ensures that the human costs fall disproportionately on the vulnerable. On the drought side, the shift toward rapidly developing flash droughts (Yuan et al., 2023) and the emergence of severe multi-year megadroughts (Williams et al., 2022) increase the frequency and severity of water shortage.

Between these extremes lies the equally important question of the seasonal and interannual reliability of supply, which depends on the natural buffers that store water between the times it falls and the times it is needed. The mountain snowpack is one such buffer, and elevation-dependent warming (Pepin et al., 2022) erodes it by shifting precipitation from snow to rain and advancing the timing of melt. Groundwater is another, and it is under simultaneous pressure from over-extraction and from a changing climate that alters recharge (Scanlon et al., 2023). The security of a region's water supply depends less on its annual precipitation total than on the integrity of these buffers, because it is the storage of water across seasons and years — not the annual mean — that carries a region through its dry periods and its droughts. Where the snowpack diminishes, the glaciers retreat, and the aquifers deplete simultaneously, the buffering capacity of the water system can be exhausted even without a change in annual rainfall, converting ordinary dry spells into crises. Water is therefore the domain in which the interaction of the several regional signals assessed in this chapter — mountain warming, groundwater stress, shifting drought and flood, and the monsoon considered next — is most consequential, and in which integrated assessment across those signals is most necessary.

The monsoon integrates the regional water signal for a majority of humanity, and its projection is a leading assessment priority (*medium confidence*).

For the majority of the world's population, the regional water balance is organised by the monsoon, whose seasonal reversal of winds and associated wet season deliver the bulk of annual rainfall across large parts of Asia, Africa, Australia, and the Americas. The monsoon integrates many of the processes assessed separately in this chapter: it is driven by the land-sea thermal contrast and the distribution of atmospheric moisture, both affected by anthropogenic forcing; it delivers rainfall through convection, whose sub-daily extremes drive flood hazard (Fowler et al., 2021); and it feeds the rivers and aquifers whose stresses shape water security (Scanlon et al., 2023). Its behaviour under warming is shaped by the competition between greenhouse-gas forcing, which tends through increased atmospheric moisture to strengthen monsoon rainfall, and aerosol forcing, which has historically tended to weaken several monsoon systems by cooling the surface and altering atmospheric stability. The regional balance between these opposing influences differs by basin and evolves as emissions change, which is one reason monsoon projection carries large uncertainty even where the direction of the thermodynamic component is agreed.

The coordinated scenario experiments that structure regional projection (Tebaldi et al., 2021) provide the framework within which monsoon change is assessed, but the translation of a given global forcing into monsoon rainfall depends on regional processes — convection, land-atmosphere coupling, aerosol-cloud interaction — whose representation still differs across models, leaving substantial spread in projected monsoon change. The stakes of this uncertainty are high because the systems that depend on the monsoon are so sensitive to it. Agriculture is the most immediate: the emergence of climate impacts on global crop production in the latest generation of coupled climate and crop models (Jägermeyr et al., 2021) is felt with particular force in monsoon regions, where food systems are tightly coupled to the seasonal rains. A monsoon that becomes more erratic — even without a large change in its seasonal total — can undermine the reliability on which rain-fed agriculture depends, and when superimposed on diminished snowpack (Pepin et al., 2022) and depleted groundwater (Scanlon et al., 2023) can compound water insecurity across some of the most populous regions on Earth. Improving the projection of regional monsoon change, and communicating its uncertainty honestly, is therefore among the highest priorities of regional climate assessment, both for its intrinsic scientific difficulty and for the number of people whose water and food security it governs.

Tropical cyclones concentrate regional hazard, and warming reshapes their intensity, rainfall, and coastal impact even where frequency trends remain uncertain (*medium confidence*).

Tropical cyclones warrant a final focused treatment because they concentrate several of the most destructive regional hazards into single events and because their response to warming illustrates the general principle that thermodynamic changes are better constrained than dynamical ones. The evidentiary basis for assessing cyclone behaviour is the set of long-term regional best-track records, of which the western North Pacific database maintained over decades is an example (Lu et al., 2021); these records make trend analysis possible but require careful accounting for the evolution of observing technology before a genuine climate signal can be separated from changes in observation. Where the physics is thermodynamic, expectations are relatively clear. A warmer ocean supplies more energy and a warmer atmosphere holds more moisture, favouring higher peak intensities and heavier cyclone rainfall even if the total number of storms does not increase. Because a warmer atmosphere intensifies short-duration rainfall in general (Fowler et al., 2021), the freshwater flooding produced by intense or slow-moving cyclones is a growing and distinct component of their hazard, separate from wind and surge.

The coastal impact of cyclones is amplified by rising sea level, which lifts the baseline on which storm surge acts so that a cyclone of unchanged intensity drives higher coastal water levels and reaches further inland than it would have in the past. When intense cyclone rainfall coincides with storm surge on an elevated sea level and with high river flow from the same or preceding storms, the result is a compound coastal flood that can far exceed what any single driver would produce — the kind of interacting hazard a complex-risk framework is designed to represent (Simpson et al., 2021). The human consequences, as for other flood hazards, are governed by exposure and vulnerability: the concentration of flood exposure among the poor (Rentschler et al., 2022), often most acute in the informal settlements of rapidly growing coastal cities, means that the compounding physical hazard of cyclones meets some of the sharpest gradients of vulnerability. Basin-wide frequency trends remain uncertain, but the reshaping of cyclone intensity, rainfall, and coastal reach under warming, combined with the concentration of exposure at vulnerable coasts, makes tropical cyclones a priority for the integrated assessment of physical hazard, exposure, and adaptive capacity that this chapter has argued is the necessary form of regional climate assessment.

3.9 Synthesis: The Regional Character of a Changing Climate

The evidence for amplified, spatially structured regional change is strong and, through attribution, increasingly linked to human influence (*high confidence*).

The central finding of this chapter is that climate change is experienced regionally, not as a uniform global increment, and that the regional signals are large, physically distinct, and increasingly attributable. The Arctic has warmed at close to four times the global rate since the late 1970s (Rantanen et al., 2022), and the world's mountains are warming in an elevation-dependent manner that outpaces adjacent lowlands (Pepin et al., 2022); these are not scaled-down versions of the global mean but distinct phenomena driven by sea-ice, snow-albedo, and lapse-rate feedbacks. Hot extremes have intensified across most inhabited regions and are amplified in cities by an urban heat island whose burden falls inequitably (Hsu et al., 2021). Short-duration rainfall extremes are intensifying, in places beyond simple thermodynamic scaling (Fowler et al., 2021). Drought is shifting in character toward rapid onset (Yuan et al., 2023) and toward severe multi-year megadroughts (Williams et al., 2022), while the groundwater that buffers drought is under simultaneous stress (Scanlon et al., 2023). Fire regimes are intensifying in dryland and Mediterranean-type regions (Abram et al., 2021), with smoke carrying health impacts across far wider areas (Aguilera et al., 2021). These changes increasingly co-occur and interact, requiring a framework of complex, cascading risk to assess them adequately (Simpson et al., 2021).

The maturation of event attribution has transformed these physical findings into statements about human influence, by comparing the likelihood or intensity of events in the observed climate with their likelihood in a counterfactual climate without anthropogenic forcing. Attribution is most confident where the physical mechanism is thermodynamic and strongly expressed, as for heat and for short-duration rainfall (Fowler et al., 2021), and it has produced landmark results such as the quantification of the human contribution to the severity of the southwestern North American megadrought (Williams et al., 2022). The same physical understanding underwrites the economic valuation of climate damage through the social cost of carbon dioxide, now assessed as substantially higher than many values in regulatory use (Rennert et al., 2022), and motivates explicit attention to the catastrophic tail risks that attribution of past events cannot bound (Kemp et al., 2022). The coordinated modelling frameworks that support projection also support attribution (Tebaldi et al., 2021), and emerging data-driven methods (Bi et al., 2023) point toward larger ensembles and finer resolution that may sharpen the assessment of the convective extremes current models resolve poorly.

The human consequences of regional change are co-produced by physical hazard and by exposure, vulnerability, and adaptation, so that assessment of the two is inseparable (*high confidence*).

The second central finding is that the human significance of regional physical change is governed as much by exposure, vulnerability, and the quality of re-

sponse as by the hazard itself. The concentration of flood exposure among the poor (Rentschler et al., 2022), the inequitable distribution of urban heat (Hsu et al., 2021), and the documented health burdens of extremes borne most heavily by the vulnerable (Ebi et al., 2021) and tracked in the annual Lancet Countdown (Romanello et al., 2024) all demonstrate that identical physical hazards produce radically different human outcomes according to who is exposed and how well they can prepare, absorb, and recover. Adaptation is occurring, but its coverage is uneven and its effects are not uniformly benign: interventions intended to reduce vulnerability have in documented cases redistributed or deepened it (Eriksen et al., 2021), a caution against treating adaptation as an unambiguous good and an argument for attention to the design, governance, and equity of response.

Regional change is transmitted into human welfare through coupled pathways — food (Jägermeyr et al., 2021), water (Pepin et al., 2022), groundwater (Scanlon et al., 2023), health (Ebi et al., 2021), disease (Baker et al., 2021), and energy (Osman et al., 2022) — none of which can be assessed in isolation, because an intervention or a shock in one reverberates through the others. The coupled human-natural system (Folke et al., 2021), operating within an Earth system that has transgressed the majority of its planetary boundaries (Richardson et al., 2023), is therefore the appropriate unit of regional assessment, and comprehensive review confirms that impacts, adaptation, and mitigation must be pursued together (Abbass et al., 2022). The task of assessing regional climate and extreme events is, in the end, inseparable from the assessment of the human and ecological systems through which climate is experienced — a physical science of hazard joined to an understanding of the exposure and vulnerability that convert hazard into consequence. This chapter has argued that only the integration of the two, under a framework of complex and cascading risk (Simpson et al., 2021) and with honest communication of the uncertainties that remain largest for the dynamical and convective processes governing regional precipitation, can support the decisions that regional adaptation and mitigation now require.

3.10 References

Abbass, K., et al. (2022). A review of the global climate change impacts, adaptation, and sustainable mitigation measures. *Environmental Science and Pollution Research*. <https://doi.org/10.1007/s11356-022-19718-6>

Abram, N., et al. (2021). Connections of climate change and variability to large and extreme forest fires in southeast Australia. *Communications Earth & Environment*. <https://doi.org/10.1038/s43247-020-00065-8>

Aguilera, R., et al. (2021). Wildfire smoke impacts respiratory health more than fine particles from other sources: observational evidence from Southern California. *Nature Communications*. <https://doi.org/10.1038/s41467-021-21708-0>

Bi, K., et al. (2023). Accurate medium-range global weather forecasting with 3D neural networks. *Nature*. <https://doi.org/10.1038/s41586-023-06185-3>

Ebi, K., et al. (2021). Extreme Weather and Climate Change: Population Health and Health System Implications. *Annual Review of Public Health*. <https://doi.org/10.1146/publhealth-012420-105026>

Eriksen, S., et al. (2021). Adaptation interventions and their effect on vulnerability in developing countries: Help, hindrance or irrelevance? *World Development*. <https://doi.org/10.1016/j.worlddev.2020.105383>

Folke, C., et al. (2021). Our future in the Anthropocene biosphere. *AMBIO*. <https://doi.org/10.1007/s13280-021-01544-8>

Fowler, H., et al. (2021). Anthropogenic intensification of short-duration rainfall extremes. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-020-00128-6>

Hsu, A., et al. (2021). Disproportionate exposure to urban heat island intensity across major US cities. *Nature Communications*. <https://doi.org/10.1038/s41467-021-22799-5>

Jägermeyr, J., et al. (2021). Climate impacts on global agriculture emerge earlier in new generation of climate and crop models. *Nature Food*. <https://doi.org/10.1038/s43016-021-00400-y>

Kemp, L., et al. (2022). Climate Endgame: Exploring catastrophic climate change scenarios. *Proceedings of the National Academy of Sciences*. <https://doi.org/10.1073/pnas>

Lu, X., et al. (2021). Western North Pacific Tropical Cyclone Database Created by the China Meteorological Administration. *Advances in Atmospheric Sciences*. <https://doi.org/10.1007/s00376-020-0211-7>

Osman, A., et al. (2022). Cost, environmental impact, and resilience of renewable energy under a changing climate: a review. *Environmental Chemistry Letters*. <https://doi.org/10.1007/s10311-022-01532-8>

Pepin, N., et al. (2022). Climate Changes and Their Elevational Patterns in the Mountains of the World. *Reviews of Geophysics*. <https://doi.org/10.1029/2020rg000730>

Rantanen, M., et al. (2022). The Arctic has warmed nearly four times faster than the globe since 1979. *Communications Earth & Environment*. <https://doi.org/10.1038/s43247-022-00498-3>

Ravuri, S., et al. (2021). Skilful precipitation nowcasting using deep generative models of radar. *Nature*. <https://doi.org/10.1038/s41586-021-03854-z>

Rennert, K., et al. (2022). Comprehensive evidence implies a higher social cost of CO₂. *Nature*. <https://doi.org/10.1038/s41586-022-05224-9>

Rentschler, J., et al. (2022). Flood exposure and poverty in 188 countries. *Nature Communications*. <https://doi.org/10.1038/s41467-022-30727-4>

Richardson, K., et al. (2023). Earth beyond six of nine planetary boundaries. *Science Advances*. <https://doi.org/10.1126/sciadv.adh2458>

Romanello, M., et al. (2022). The 2022 report of the Lancet Countdown on health and climate change: health at the mercy of fossil fuels. *The Lancet*. [https://doi.org/10.1016/s0140-6736\(22\)01540-9](https://doi.org/10.1016/s0140-6736(22)01540-9)

Romanello, M., et al. (2023). The 2023 report of the Lancet Countdown on health and climate change: the imperative for a health-centred response in a world facing irreversible harms. *The Lancet*. [https://doi.org/10.1016/s0140-6736\(23\)01859-7](https://doi.org/10.1016/s0140-6736(23)01859-7)

Romanello, M., et al. (2024). The 2024 report of the Lancet Countdown on health and climate change: facing record-breaking threats from delayed action. *The Lancet*. [https://doi.org/10.1016/s0140-6736\(24\)01822-1](https://doi.org/10.1016/s0140-6736(24)01822-1)

Scanlon, B., et al. (2023). Global water resources and the role of groundwater in a resilient water future. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43002-00378-6>

Simpson, N., et al. (2021). A framework for complex climate change risk assessment. *One Earth*. <https://doi.org/10.1016/j.oneear.2021.03.005>

Tebaldi, C., et al. (2021). Climate model projections from the Scenario Model Intercomparison Project (ScenarioMIP) of CMIP6. *Earth System Dynamics*. <https://doi.org/10.5194/esd-12-253-2021>

Williams, A., et al. (2022). Rapid intensification of the emerging southwestern North American megadrought in 2020–2021. *Nature Climate Change*. <https://doi.org/10.1038/s41558-022-01290-z>

Yuan, X., et al. (2023). A global transition to flash droughts under climate change. *Science*. <https://doi.org/10.1126/science.abn6301>

Chapter 4

Chapter 4. Process Understanding: Forcings, Feedbacks, and Cycles

4.1 Executive Summary

The physical understanding of how energy imbalance is imposed on the climate system, amplified or damped by feedbacks, and transmitted through the coupled carbon, biogeochemical, and hydrological cycles has consolidated since the previous assessment cycle, and the direction of every major process is now attributed to human activity with progressively narrower uncertainty ranges (*high confidence*). The dominant control on the present energy imbalance is the accumulation of carbon dioxide and other long-lived greenhouse gases, whose growth rate and airborne fraction are quantified annually to within a few percent by the global carbon budget synthesis (Friedlingstein, 2023).

Effective radiative forcing, which folds rapid atmospheric and land-surface adjustments into the instantaneous perturbation, is the most policy-relevant forcing metric, and the aerosol contribution remains the single largest source of uncertainty in the total anthropogenic forcing (*high confidence*). Mineral dust and other natural aerosol populations, whose source apportionment is now resolved at the level of individual emitting regions, modulate both the shortwave and the longwave components of this forcing (Kok, 2021).

Net climate feedbacks are positive, so the equilibrium warming for a doubling of carbon dioxide is substantially larger than the no-feedback response, with the cloud feedback the principal amplifier and the principal residual uncertainty (*high confidence*). Surface-albedo feedbacks are expressed most strongly in the Arctic, where near-surface warming has proceeded at nearly four times the global rate over the satellite era (Rantanen, 2022).

The land and ocean carbon sinks together remove roughly half of anthropogenic emissions, but the efficiency of both sinks is sensitive to warming, drought, and disturbance, introducing a carbon-cycle feedback that is positive and of increasing magnitude (*medium to high confidence*). The persistence of the global forest carbon sink, and the emerging signals that its resilience is declining under compound stress, are now documented at global scale (Pan, 2024), even as the resilience of that sink shows signs of decline (Forzieri, 2022).

Terrestrial biogeochemical cycling is governed at the process level by the interaction of organic matter with soil minerals and by the ecology of the soil microbiome, both of which set the residence time of carbon and the fluxes of nitrogen and other nutrients (*high confidence*). Disturbance to these interfaces, whether through erosion, salinization, or novel-material contamination, reduces the functional capacity of soils to store carbon (Kleber, 2021) and disrupts the microbial ecology that mediates it (Sokol, 2022).

The water cycle is intensifying, with the atmospheric demand for moisture rising, droughts emerging more rapidly, and inland waters warming, and each of these process shifts carries an anthropogenic fingerprint (*high confidence*). The abrupt onset of flash droughts under continued warming, and the propagation of heat extremes into lakes and rivers, are robust across observational and modelling lines of evidence (Yuan, 2023), as is the propagation of heat into inland waters (Woolway, 2021).

Earth-system model intercomparison provides the quantitative bridge between process understanding and projection, and the current generation of models reproduces the observed forcing history while spanning a credible range of feedback strengths (*high confidence*). Scenario-based projections rest on a small number of shared pathways whose forcing trajectories are constructed to be internally consistent with the emissions and land-use assumptions they encode (Tebaldi, 2021).

4.2 Radiative Forcing and Effective Radiative Forcing

The forcing concept and its refinement

Radiative forcing remains the organizing quantity of physical climate assessment, and the shift from the instantaneous and stratosphere-adjusted definitions toward effective radiative forcing, which incorporates rapid adjustments of the troposphere and land surface, has improved the correspondence between a given forcing agent and the surface temperature response it ultimately produces (*high confidence*). A radiative forcing is, in its simplest formulation, a perturbation to the net radiative flux at the top of the atmosphere imposed by an external agent while surface temperature is held fixed. The strength of the concept is that it isolates the driver from the response: a forcing of a given magnitude in watts per square meter (W m^{-2}) can be compared across agents as chemically and physically dissimilar as carbon dioxide, sulfate aerosol, and a change in land-surface reflectivity. The weakness of the simplest formulation is that the atmosphere does not, in fact, hold everything except surface temperature fixed while a forcing is applied. Clouds adjust within hours, the stratosphere cools radiatively within months, and vegetation and snow respond within a season. Effective radiative forcing was introduced precisely to absorb these rapid adjustments, which occur without any change in global-mean sur-

face temperature and which therefore properly belong to the forcing rather than to the feedback.

The distinction matters most for agents whose rapid adjustments are large relative to their instantaneous perturbation. Absorbing aerosols such as black carbon warm the layer of atmosphere in which they reside, stabilizing it, altering cloud cover, and thereby generating a semi-direct adjustment that can partially offset the instantaneous warming. For carbon dioxide, the tropospheric adjustment is comparatively modest, so the instantaneous and effective forcings differ by only a small fraction; this is one reason the carbon dioxide forcing is among the best-constrained quantities in the entire assessment. The annual reconciliation of sources and sinks that underlies the carbon budget synthesis provides the atmospheric growth rate from which the concentration pathway, and hence the forcing, is derived with high precision (Friedlingstein, 2023). The 2024 update to that synthesis confirms that the atmospheric carbon dioxide burden continues to rise essentially in proportion to cumulative emissions, with only second-order variation in the airborne fraction (Friedlingstein, 2025).

The well-mixed greenhouse gases

The forcing exerted by the well-mixed greenhouse gases is the dominant term in the total anthropogenic forcing, and its magnitude and growth are established with the highest confidence of any forcing agent because both the atmospheric concentrations and the underlying spectroscopy are directly measured (*virtually certain*). Carbon dioxide is the largest single contributor, followed by methane and nitrous oxide, with the halogenated species contributing a smaller but non-negligible increment. The forcing from carbon dioxide grows approximately with the logarithm of its concentration, so that each successive increment of concentration produces a slightly smaller radiative effect than the last; this logarithmic dependence is a robust feature of the molecular absorption bands and is not in scientific dispute. What has changed across assessment cycles is not the physics but the precision of the input concentrations and the completeness of the emissions accounting that constrains them.

The sectoral decomposition of the emissions that drive these concentrations has been mapped in detail, revealing that the growth in greenhouse-gas emissions over recent decades has been distributed unevenly across energy supply, industry, transport, buildings, and land use, with the energy and industry sectors dominating the long-lived carbon dioxide flux (Lamb, 2021). This decomposition is important for the forcing assessment because it identifies which agents are co-emitted: fossil combustion that releases carbon dioxide also releases sulfur dioxide, the precursor of sulfate aerosol, so that the net forcing from a given activity is the sum of a warming and a cooling term that do not necessarily evolve together in time. As controls on sulfur emissions tighten for air-quality reasons, the cooling term declines faster than the warming term, unmasking a fraction of the greenhouse warming that had previously been offset.

Methane occupies a distinctive position in the forcing budget because it is both a direct greenhouse gas and a chemical precursor that influences tropospheric ozone and stratospheric water vapor, so that its effective forcing exceeds its direct forcing by a substantial margin. The agricultural contribution to the methane budget, and in particular the flux from flooded rice cultivation, is now quantified with sufficient resolution to support targeted mitigation assessment (Qian, 2023). Because methane has a short atmospheric lifetime relative to carbon dioxide, its forcing responds within roughly a decade to changes in emission rate, which gives methane mitigation a leverage over near-term forcing that carbon dioxide mitigation cannot provide.

Ozone, water vapor, and the shorter-lived agents

The tropospheric ozone forcing is positive and among the more uncertain of the well-quantified agents, because ozone is not emitted directly but is produced photochemically from precursors whose emissions and chemistry vary strongly in space and time (*medium confidence*). Ozone precursors—methane, carbon monoxide, nitrogen oxides, and volatile organic compounds—are co-emitted with aerosols and their precursors across the industrial, transport, and agricultural sectors, so the ozone forcing is intimately coupled to the aerosol forcing through shared sources. Road transport in particular emits a complex mixture in which the non-exhaust fraction, arising from tyre, brake, and road-surface wear, contributes both to the particulate burden and, indirectly, to the oxidant chemistry that governs ozone production (Fussell, 2022). The stratospheric water vapor forcing, produced partly by methane oxidation, is a smaller positive term that is nonetheless an example of the chemical coupling that makes the shorter-lived agents harder to assess than the well-mixed gases.

Stratospheric ozone, by contrast, has undergone a forced recovery following the phase-out of ozone-depleting substances, a recovery that itself constitutes a small forcing as the vertical distribution of ozone and its radiative effect shift. The lesson of the stratospheric ozone case for the broader forcing assessment is that coordinated international control of a well-identified set of agents can reverse a forcing trend on a multi-decadal timescale, and that the radiative and the chemical consequences of such controls must be assessed jointly rather than separately.

Land-use, surface-albedo, and biogeophysical forcings

Changes in land cover impose a forcing that is biogeophysical as well as biogeochemical, altering surface albedo, roughness, and evapotranspiration in ways whose net radiative effect depends strongly on latitude and background climate (*medium confidence*). The conversion of forest to cropland or pasture typically raises surface albedo, a cooling effect, while simultaneously reducing evapotranspiration and altering cloud cover, effects whose sign varies with region. In the boreal zone the albedo effect of deforestation tends to dominate and produces a net cooling, whereas in the tropics the loss of evaporative cooling

and the associated changes in convection tend to produce a net warming. The dominant driver of tropical forest loss is the expansion of agriculture, and disentangling the specific commodities and land-use transitions responsible has been necessary to attribute the biogeophysical forcing to its proximate causes (Pendrill, 2022).

Forest cover also exerts a microclimatic control that is not captured by the global-mean forcing but is central to the local expression of climate change. The canopy buffers the temperature experienced beneath it, decoupling the forest-floor microclimate from the free-atmosphere warming, so that the loss or degradation of forest cover removes a thermal buffer at the same time that it changes the global albedo (Frenne, 2021). At high elevation the surface-albedo forcing interacts with the pronounced warming of mountain regions, where the retreat of snow and ice exposes darker surfaces and drives an elevation-dependent amplification of the temperature signal (Pepin, 2022). These regional expressions of the surface-albedo forcing are the local face of a feedback that, integrated globally, is treated in the following section.

Quantifying and communicating forcing uncertainty

The total anthropogenic effective radiative forcing is positive and its central estimate is well separated from zero, but its uncertainty is dominated by the aerosol term, and this asymmetry has direct consequences for the inference of climate sensitivity from the historical record (*high confidence*).

The greenhouse-gas forcing is known with high precision; the aerosol forcing, being negative and poorly constrained, sets the width of the total. Because the observed warming is the response to the net forcing, a more negative aerosol forcing implies a larger sensitivity for a given observed warming, and vice versa. This is why progress on the aerosol forcing, treated in Section 4.3, is inseparable from progress on climate sensitivity, treated in Section 4.2: they are two projections of the same underlying uncertainty.

The forcing framework also underpins the economic and policy metrics that translate physical change into decision-relevant terms. The social cost of carbon, which monetizes the marginal damage of an incremental tonne of carbon dioxide, depends at its foundation on the forcing-to-warming chain, and comprehensive re-estimation using updated damage functions and discounting implies a substantially higher value than earlier assessments had adopted (Rennert, 2022). Situating the forcing assessment within the wider diagnosis of Earth-system state, the transgression of multiple planetary boundaries—of which the climate boundary, defined in part by radiative forcing, is one—signals that the forcing is now operating outside the range within which the system is understood to remain stable (Richardson, 2023). The framing of these boundaries in terms that couple the biophysical to the equitable further underscores that the forcing is not a purely physical quantity but one whose consequences are distributed unevenly across the human and natural world (Rockström, 2023).

4.3 Climate Sensitivity and Feedbacks

The sensitivity framework

Equilibrium climate sensitivity, the eventual global-mean warming that follows a sustained doubling of atmospheric carbon dioxide, is the single number that most compactly summarizes the amplifying and damping processes of the climate system, and its assessed range has narrowed as the feedbacks that determine it have been separately constrained (*high confidence*). The no-feedback response to a doubling of carbon dioxide—the warming that would occur if the only adjustment were the Planck response of a blackbody radiating more efficiently as it warms—is close to one degree C and is not in dispute. Everything above that value is the net product of feedbacks, and everything uncertain about the sensitivity is uncertainty in feedbacks. The task of process understanding is therefore to decompose the net feedback into physically distinct mechanisms, to constrain each from observations and process models, and to recombine them into a sensitivity estimate whose uncertainty is the propagated uncertainty of its parts. This bottom-up, process-based line of evidence is now combined with the historical-record and paleoclimate lines to yield a sensitivity range materially narrower than that of a decade ago.

A conceptual distinction that has become central is that between equilibrium sensitivity and the transient response that is actually experienced on human timescales. The ocean's enormous heat capacity delays the surface warming, so that at any moment during a period of rising forcing the realized warming is smaller than the equilibrium warming for the forcing then applied. The efficiency with which the ocean takes up heat, and the pattern in which it does so, therefore modulates the effective feedback that operates over the historical period. This pattern effect—the dependence of the global feedback on the spatial pattern of warming—has emerged as a key reconciler of the historical and process-based sensitivity estimates, because a warming pattern that is concentrated in regions of strong negative feedback produces a historical feedback that is more stabilizing than the equilibrium feedback will ultimately be.

The water-vapor and lapse-rate feedbacks

The combined water-vapor and lapse-rate feedback is strongly positive and is the largest of the well-understood feedbacks, and the compensating relationship between its two components makes the combination more robust than either part alone (*very high confidence*). A warmer atmosphere holds more water vapor, and because water vapor is itself a powerful greenhouse gas, this constitutes a large positive feedback. The relative humidity of the atmosphere remains approximately constant as it warms, which is the physical basis for treating the water-vapor feedback as a near-linear function of temperature. The lapse-rate feedback arises because the vertical profile of warming is not uniform: in the tropics the upper troposphere warms more than the surface, which increases the efficiency of radiation to space and therefore constitutes

a negative feedback. The two feedbacks are anticorrelated across models and across regions, because the same moist processes that load the upper troposphere with water vapor also warm it preferentially, so that a model with a strong positive water-vapor feedback tends to have a strongly negative lapse-rate feedback. Their sum is far better constrained than either term, and it is robustly positive.

The realization of these feedbacks depends on the fidelity with which models represent moist convection and the free-tropospheric humidity distribution, which is one of the reasons that model evaluation against observed humidity and temperature profiles is a central activity of the intercomparison enterprise. The current generation of Earth-system models, of which the CMCC Earth System Model is a representative example, reproduces the broad structure of the tropical temperature and humidity profile and therefore the compensating water-vapor and lapse-rate feedbacks, while differing in the details that matter for the residual uncertainty (Lovato, 2022).

The surface-albedo and cryosphere feedbacks

The surface-albedo feedback is positive, is concentrated at high latitudes, and is expressed most dramatically in the Arctic, where the observed rate of near-surface warming has exceeded the global-mean rate by nearly a factor of four over the satellite era (*high confidence*). As snow and sea ice retreat, they expose ocean and land surfaces of much lower reflectivity, which absorb more solar radiation and drive further warming and further retreat. This ice-albedo feedback is the textbook amplifier, but the modern understanding is considerably richer than the albedo mechanism alone. The strength of Arctic amplification observed over recent decades exceeds what the albedo feedback in isolation would produce, implicating additional processes: changes in the vertical structure of atmospheric warming, in poleward heat transport, in cloud cover, and in the seasonal storage and release of ocean heat. The magnitude of the observed amplification, nearly four times the global rate, is larger than most model simulations reproduce, which points either to underestimated feedbacks or to a contribution from internal variability that the assessment must quantify (Rantanen, 2022).

The atmospheric processes that accompany the albedo feedback in the Arctic have been characterized in unprecedented detail by year-round observations from within the ice pack, which have documented the coupling of the surface energy budget to cloud, aerosol, and boundary-layer processes across the full seasonal cycle (Shupe, 2022). These observations are important for the feedback assessment because the surface-albedo feedback does not operate in isolation: its expression depends on the cloud cover that determines how much radiation reaches the surface, on the stability of the boundary layer that determines how surface warming is communicated upward, and on the aerosol populations that influence cloud properties. The cryospheric feedbacks thus provide a concrete illustration of the general principle that feedbacks interact, and that their sum cannot be obtained by assessing each in complete isolation.

Cloud feedbacks

Cloud feedbacks are net positive and constitute both the largest amplifying uncertainty in the sensitivity estimate and the area in which process understanding has advanced most since the previous assessment, so that the residual cloud uncertainty, while still dominant, has narrowed (*high confidence*). Clouds affect the energy budget in two opposing ways: they reflect incoming solar radiation, a cooling effect that dominates for low clouds, and they trap outgoing longwave radiation, a warming effect that dominates for high clouds. A cloud feedback is any change in these effects that accompanies warming, and because the two effects are large and of opposite sign, small fractional changes in cloud amount, altitude, or optical thickness translate into feedbacks that are significant relative to the total. The assessment now decomposes the cloud feedback into components that are separately constrained: the rising of high cloud tops with warming, which is a robust positive feedback grounded in well-understood thermodynamics; the poleward shift of the mid-latitude storm tracks and their associated clouds; and the response of low marine clouds, particularly the subtropical stratocumulus decks, which has been the principal source of the overall uncertainty.

The low-cloud feedback has been the crux because the governing processes—turbulent mixing, cloud-top entrainment, and the coupling of the cloud layer to the surface—occur at scales far below the resolution of global models and must be parameterized. Advances have come from high-resolution process modelling, from satellite observations of how these cloud regimes actually respond to natural temperature variations, and from the integration of physical understanding with data-driven methods that can extract feedback-relevant relationships from large observational archives (Willard, 2022). The convergence of these lines has favored a positive low-cloud feedback, ruling out the strongly negative values that some earlier parameterizations produced and thereby raising the lower bound of the sensitivity estimate. The high-cloud altitude feedback, by contrast, has been robustly positive throughout and is one of the more secure elements of the cloud assessment.

The carbon-cycle feedback on climate

The carbon-cycle feedback—the response of the natural carbon sinks and sources to warming—is positive, meaning that warming tends to reduce the efficiency of natural carbon uptake and thereby amplify the atmospheric burden, and its magnitude is now the largest uncertainty in translating a given emissions pathway into a temperature outcome (*medium to high confidence*). Unlike the physical feedbacks, which act on the energy budget, the carbon-cycle feedback acts on the concentration itself: a warmer climate weakens both the land and ocean sinks, so that a larger fraction of each tonne emitted remains airborne. The land component is the more uncertain because it integrates the responses of photosynthesis, respiration, disturbance, and nutrient limitation, each of which responds differently to warming, to changing precipitation, and

to the fertilizing effect of rising carbon dioxide itself. The concentration-carbon feedback, whereby higher carbon dioxide stimulates plant growth and enhances uptake, is negative and partially offsets the positive climate-carbon feedback, and the balance between them is what the assessment must quantify.

The global carbon budget synthesis provides the top-down constraint on the net outcome of these competing feedbacks, partitioning the annual emissions among the atmosphere, ocean, and land and thereby measuring the aggregate sink efficiency year by year (Friedlingstein, 2023). The persistence of a substantial land sink, concentrated in the world's forests, has been confirmed by dedicated global assessment, but that same assessment reveals that the sink is sustained against a rising tide of disturbance and is therefore not guaranteed to persist (Pan, 2024). The emergence of statistical early-warning signals that forest resilience is declining across large regions is the process-level evidence that the land carbon-cycle feedback may strengthen as warming continues, because a forest that is losing resilience is one whose capacity to recover from drought, fire, and pest outbreaks—and therefore to continue storing carbon—is eroding (Forzieri, 2022). The mechanistic basis for this erosion, in the form of widespread and sometimes unexpected episodes of elevated tree mortality, has been documented across biomes and constitutes a direct threat to the durability of the terrestrial sink (Hartmann, 2022).

Combining feedbacks and constraining sensitivity

The recombination of the separately assessed feedbacks yields a net feedback that is negative in the stabilizing sense required for a finite sensitivity but far weaker than the Planck response alone, and the resulting sensitivity range is now constrained by the agreement of the process-based, historical, and paleoclimate lines of evidence (*high confidence*). The intercomparison of Earth-system models remains the framework within which feedbacks are diagnosed in a physically consistent setting, because a model computes all feedbacks simultaneously and self-consistently, avoiding the double-counting that can afflict the summation of independently estimated terms. The current model generation spans a range of feedback strengths, and the assessment's task is not to adopt the model mean uncritically but to weight the models by their fidelity to observed feedback-relevant processes and to the emergent constraints that relate an observable feature of the present climate to a feedback strength.

Scenario-based projection then propagates the assessed sensitivity, together with the assessed carbon-cycle feedback, through a small set of shared socioeconomic pathways whose forcing trajectories are internally consistent (Tebaldi, 2021). The value of this architecture is that it separates the uncertainty that is physical—the sensitivity and the feedbacks—from the uncertainty that is societal—the emissions pathway—so that each can be assessed with the appropriate methods. The physical uncertainty is now dominated, as it has long been, by the cloud and carbon-cycle feedbacks, and the narrowing of the sensitivity range achieved in recent years is attributable chiefly to progress on the low-

cloud feedback and to the reconciliation of the historical record with the process understanding through the pattern effect.

4.4 Aerosols and Atmospheric Composition

The aerosol forcing and its residual uncertainty

Aerosols exert a net negative effective radiative forcing that has offset a substantial fraction of the greenhouse warming to date, and this forcing remains the dominant uncertainty in the total anthropogenic forcing because it aggregates several mechanisms whose individual magnitudes are separately uncertain and whose sum is a small residual of large opposing terms (*high confidence*). The aerosol forcing decomposes into a direct component, in which particles scatter and absorb radiation, and one or more indirect components, in which particles act as cloud condensation nuclei and thereby alter the number, size, reflectivity, and lifetime of cloud droplets. The direct scattering effect is cooling; the direct absorption effect, dominated by black carbon, is warming; the first indirect effect, whereby more numerous but smaller droplets brighten clouds, is cooling; and the adjustments of cloud amount and lifetime that follow are of uncertain sign and magnitude. The overall aerosol forcing is therefore a difference of large terms, which is why its fractional uncertainty is so much greater than that of the greenhouse-gas forcing even though the underlying microphysics is, in many respects, well understood.

The consequence of this uncertainty for the assessment as a whole cannot be overstated: because the observed warming is the response to the net forcing, and because the net forcing is the greenhouse-gas forcing minus the aerosol offset, the width of the aerosol uncertainty propagates directly into the inference of climate sensitivity from the historical record. A larger aerosol offset implies that the greenhouse forcing has been more strongly masked, which implies a larger sensitivity for the observed warming; a smaller offset implies the reverse. Narrowing the aerosol forcing is thus a prerequisite for narrowing the historical constraint on sensitivity, and progress on the two is inseparable. The sectoral emissions accounting that identifies the co-emission of aerosols and their precursors alongside greenhouse gases is the foundation on which the aerosol forcing history is built, because the aerosol burden is tied to the same combustion and industrial activities that release carbon dioxide (Lamb, 2021).

Natural aerosols: mineral dust

Mineral dust is the largest natural aerosol population by mass, exerts a forcing whose sign depends on the underlying surface and the particle-size distribution, and modulates biogeochemical cycles by delivering nutrients to oceans and remote ecosystems, so that its assessment spans the physical and the biogeochemical (*medium confidence*). Dust scatters and absorbs shortwave

radiation and, because its particles are large, also interacts significantly with longwave radiation, producing a net forcing that is small in the global mean but large regionally and highly sensitive to the size distribution and the mineralogy of the emitting source. The apportionment of the global dust burden among the world's principal source regions has been quantified, revealing that a small number of arid regions dominate emission and that the resulting forcing and deposition patterns are correspondingly concentrated (Kok, 2021). This source apportionment is essential for two reasons: it constrains the present-day dust forcing, and it provides the baseline against which future changes in dust—driven by changes in aridity, land use, and wind—can be projected.

The biogeochemical role of dust couples the aerosol assessment to the carbon and nutrient cycles treated in Section 4.4. Dust deposition supplies iron and phosphorus to ocean regions where these nutrients limit biological productivity, so that changes in the dust cycle can alter the ocean carbon sink through the fertilization of phytoplankton. Dust also affects terrestrial ecosystems and the cryosphere: deposition of absorbing dust on snow and ice lowers their albedo and accelerates melt, coupling the dust cycle to the surface-albedo feedback discussed in Section 4.2. The projection of the dust forcing is complicated by the fact that its drivers act in opposing directions—increased aridity tends to raise emission while increased vegetation cover under carbon dioxide fertilization tends to suppress it—so that the sign of the future dust-forcing trend is not robustly established.

Anthropogenic aerosols and air quality

Anthropogenic aerosols are co-emitted with greenhouse gases from combustion and industry, so that measures which reduce air pollution for health reasons simultaneously reduce the aerosol cooling and thereby unmask greenhouse warming, creating a coupling between air-quality policy and climate that the assessment must make explicit (*high confidence*). The dominant anthropogenic aerosol is sulfate, formed from sulfur dioxide emitted by fossil combustion, which is strongly scattering and therefore cooling. Black carbon, emitted from incomplete combustion of fossil and biomass fuels, is strongly absorbing and therefore warming, and its deposition on snow and ice adds a further warming through albedo reduction. Organic and nitrate aerosols, and the secondary aerosols formed from precursor gases, complete a chemically diverse population whose net effect is cooling but whose composition determines the balance of scattering and absorption.

Road transport is a distinctive and growing contributor because its aerosol emissions are only partly controllable by exhaust after-treatment: the non-exhaust fraction from tyre, brake, and road-surface wear is unaffected by cleaner engines and constitutes an increasing share of the traffic-derived particulate burden as exhaust controls tighten (Fussell, 2022). This non-exhaust source is significant for both the air-quality and the climate assessment because it will persist, and may grow with vehicle mass, even as combustion emissions decline. The health burden of aerosol exposure is a major driver of the emission

controls that shape the aerosol forcing trajectory, and the assessment of climate and health is increasingly integrated, recognizing that the same particles that cool the climate impose a substantial mortality burden (Romanello, 2022). This integration reframes aerosol reduction not as a climate cost to be reluctantly incurred but as a health benefit whose climate consequence must be anticipated and offset by deeper greenhouse-gas reductions.

Aerosol–cloud interactions

Aerosol–cloud interactions constitute the most uncertain element of the aerosol forcing, because the response of cloud reflectivity, amount, and lifetime to changes in aerosol loading involves microphysical and dynamical processes that operate below the resolution of global models and that observations constrain only indirectly (*medium confidence*). The first aerosol indirect effect, in which a fixed amount of cloud water is distributed among more numerous but smaller droplets when aerosol loading rises, brightening the cloud, is comparatively well understood and robustly cooling. The subsequent adjustments—changes in cloud lifetime as smaller droplets suppress precipitation, changes in cloud amount, and changes in cloud-top height—are far less certain and can partially offset or amplify the initial brightening. Because these adjustments occur in the same low marine cloud regimes that dominate the low-cloud feedback uncertainty discussed in Section 4.2, the aerosol–cloud problem and the cloud-feedback problem are physically entangled: the same stratocumulus decks whose response to warming is uncertain also mediate the aerosol indirect effect.

Progress has come from combining satellite observations of natural aerosol–cloud covariability with high-resolution modelling and with data-driven methods that can extract robust relationships from large and noisy observational archives while respecting physical constraints (Willard, 2022). These approaches have tended to reduce the magnitude of the more extreme negative aerosol–cloud forcings that some earlier estimates supported, narrowing the aerosol forcing from its most negative tail and, by the coupling described above, tightening the historical constraint on sensitivity. The in-situ characterization of Arctic aerosol–cloud–surface coupling has further constrained how aerosol populations influence cloud radiative properties in a regime where the surface-albedo feedback makes cloud effects especially consequential (Shupe, 2022).

Novel materials in the atmospheric burden

The atmospheric aerosol population now includes anthropogenic novel materials—microplastics and persistent synthetic compounds—whose radiative and biogeochemical roles are only beginning to be quantified but whose ubiquity, including in the remote Arctic, indicates a globally distributed and growing component of atmospheric composition (*low to medium confidence*). Plastic debris fragments into micro- and nano-scale particles that are transported atmospherically over long distances and deposited even in regions remote from

any source, so that plastic pollution is now documented throughout the Arctic environment (Bergmann, 2022). The environmental footprint of the plastics that give rise to these particles is itself climate-relevant, being driven substantially by the combustion of coal in their production, which ties the novel-materials burden back to the fossil-fuel emissions that dominate the greenhouse forcing (Cabernard, 2021).

The persistence of certain synthetic compounds compounds the concern: per- and polyfluoroalkyl substances are mobile, bioaccumulative, and effectively non-degrading, so that their environmental burden is monotonically increasing and their eventual biogeochemical consequences are, by the nature of persistent compounds, difficult to bound (Brunn, 2023). While the direct radiative forcing of these novel materials is at present assessed as small, their significance for the process understanding lies in their coupling to the biogeochemical cycles treated in the next section: microplastics alter soil carbon cycling and can perturb the terrestrial processes that govern the land sink, so that a material introduced for reasons unrelated to climate acquires a climate relevance through its interference with biogeochemistry (Rillig, 2021).

4.5 The Carbon and Biogeochemical Cycles

The contemporary carbon budget

The global carbon budget is now closed to within a few percent on an annual basis, establishing that the atmosphere, ocean, and land partition anthropogenic emissions in a ratio that has remained broadly stable while the absolute fluxes have grown, and this partitioning is the single most important constraint on the coupling between emissions and climate (*high confidence*).

The budget is constructed by reconciling estimates of fossil-fuel and land-use emissions against the measured atmospheric growth rate and against independent estimates of the ocean and land sinks, with the residual providing a check on the internal consistency of the whole. The consistent finding across successive syntheses is that the ocean and the land together remove roughly half of what is emitted, so that only about half accumulates in the atmosphere; this airborne fraction is a derived quantity of central importance, because its stability under rising emissions is what has allowed the atmospheric burden to track cumulative emissions so closely (Friedlingstein, 2023). The most recent update extends this record and confirms that neither sink has yet saturated in a manner that would sharply raise the airborne fraction, even as both show signs of the sensitivity to climate that the carbon-cycle feedback anticipates (Friedlingstein, 2025).

Independent monitoring of the emission side has advanced to the point where global carbon dioxide emissions can be tracked at near-real-time resolution, using activity data and atmospheric measurements to detect the response of emissions to economic and policy perturbations within the year in which they

occur (Liu, 2022). This capacity matters for the process understanding because it closes the loop between the emission drivers, mapped by sector (Lamb, 2021), and the atmospheric outcome, so that the carbon budget is no longer a retrospective reconstruction but an operational monitoring system. The convergence of top-down atmospheric constraints and bottom-up activity accounting has narrowed the emissions term of the budget and thereby sharpened the inference of the sink terms that are obtained as residuals.

The terrestrial carbon sink and its vulnerability

The terrestrial biosphere is at present a net sink for carbon, dominated by the world's forests, but this sink is sustained against rising disturbance and its future persistence is the largest single uncertainty in the land carbon-cycle feedback (*medium to high confidence*). The land sink arises because the fertilizing effect of elevated carbon dioxide, together with nitrogen deposition and, in some regions, a lengthening growing season, has stimulated plant growth faster than warming and disturbance have released carbon through respiration, fire, and mortality. Dedicated global assessment confirms that the forest sink has endured, but it also makes clear that its endurance is the net outcome of a growing gross uptake competing against a growing gross loss, a balance that cannot be assumed to hold as the loss terms strengthen (Pan, 2024). The unified quantification of the terrestrial biosphere's productivity from satellite observation provides the spatial and temporal detail needed to track where the sink is strengthening and where it is faltering.

The vulnerability of the sink is expressed through several converging lines of evidence. Statistical analysis of vegetation dynamics reveals a widespread decline in forest resilience—a lengthening of the time forests take to recover from perturbation—across large regions subject to warming and increasing water limitation, which is the signature of a system approaching a threshold beyond which the sink could weaken or reverse (Forzieri, 2022). The mechanistic expression of this loss of resilience is elevated tree mortality, which has been documented worldwide and which sometimes takes the form of unexpected, rapid die-off events that no incremental model of forest decline anticipated (Hartmann, 2022). Compounding the direct climatic stress, erosion degrades the soils on which forests and other terrestrial ecosystems depend, reducing the microbial diversity, network complexity, and multifunctionality that underpin nutrient cycling and carbon stabilization (Qiu, 2021). The microclimatic buffering that intact forests provide is itself eroded as canopies thin, removing the thermal insulation that had partially protected forest-floor processes from the free-atmosphere warming (Frenne, 2021).

Soil biogeochemistry and the residence time of carbon

The residence time of carbon in soils, which determines how much of the terrestrial carbon stock is durable and how much is vulnerable to release, is governed at the process level by the interaction of organic matter with min-

eral surfaces and by the ecology of the microbial community that mediates decomposition, and the modern understanding has shifted decisively toward these mechanistic controls and away from the older view based on intrinsic molecular recalcitrance (*high confidence*). The stabilization of organic carbon in soil is now understood to depend less on the chemical structure of the organic molecules themselves than on their physical and chemical association with minerals, which protects them from microbial access. The dynamic interactions at the mineral–organic matter interface—adsorption, occlusion within aggregates, and the formation of organo-mineral complexes—are the master control on how long carbon persists in soil, and disturbances that break these associations can convert a long-lived stock into a rapidly respired flux (Kleber, 2021).

The microbial community is the agent through which these controls are exercised, and its ecology—the life, death, and turnover of microbial biomass—determines both the rate of decomposition and the formation of the stable, mineral-associated compounds that constitute the durable carbon pool (Sokol, 2022). This reframing has important consequences for the carbon-cycle feedback, because it implies that the temperature sensitivity of soil carbon release depends not on a single activation energy but on the interplay of microbial physiology, substrate availability, and mineral protection, all of which respond to warming and moisture change in ways that models are only beginning to represent. Novel-material contamination adds a further perturbation: microplastics in soil alter the physical structure and the microbial processing of organic matter, changing carbon-cycling rates in directions that depend on the polymer, the concentration, and the soil type (Rillig, 2021). Salinization, projected to expand across large agricultural regions under a changing climate, degrades soil function and reduces the capacity of affected soils to support the plant growth and microbial activity on which carbon storage depends (Hassani, 2021).

The ocean carbon sink and marine biogeochemistry

The ocean is a major and persistent sink for anthropogenic carbon, taking up carbon both through the physical solubility of carbon dioxide in seawater and through the biological pump that exports organic matter to depth, and its uptake is being reshaped by warming, stratification, and shifting biological productivity in ways that the carbon budget resolves at the aggregate level but that require process understanding to project (*high confidence*). The solubility pump weakens as the surface ocean warms, because warmer water holds less dissolved gas, and as the accumulation of carbon in the surface layer reduces the chemical capacity of seawater to absorb more. The biological pump depends on the productivity of phytoplankton, which fix carbon in the sunlit surface and export a fraction of it to depth when they sink or are consumed. Changes in this productivity therefore feed directly into the ocean sink, and observations reveal that coastal phytoplankton blooms have expanded and intensified over recent decades, a large-scale reorganization of marine productivity with consequences for carbon export that are not yet fully quantified (Dai, 2023).

The delivery of nutrients that limit marine productivity couples the ocean

carbon cycle to the atmospheric dust cycle discussed in Section 4.3, because atmospheric deposition supplies iron and other micronutrients to the open ocean regions where they are scarce (Kok, 2021). This coupling means that changes in aridity and land use, by altering the dust flux, can propagate into the ocean carbon sink through the fertilization or starvation of phytoplankton. The broader significance of the marine biogeochemical changes is that they occur alongside the physical stressors of warming, acidification, and deoxygenation, so that the ocean carbon sink is being modified by several simultaneous perturbations whose combined effect on future uptake is a principal uncertainty in the long-term carbon-cycle feedback.

Inland waters and the lateral carbon flux

Rivers, lakes, and other inland waters are not passive conduits but active processors of carbon, respiring, storing, and outgassing a flux that is large relative to the net land sink, and the recognition of this lateral pathway has revised the understanding of how terrestrial carbon reaches the atmosphere and ocean (*high confidence*). Carbon fixed on land is exported to inland waters, where a substantial fraction is respired back to the atmosphere or buried in sediments before the remainder reaches the sea, so that the river network functions as a biogeochemical reactor whose metabolism is a significant term in the global carbon budget (Battin, 2023). This lateral flux has historically been under-represented in carbon accounting, and its incorporation changes the interpretation of the land sink: some of what appears as terrestrial uptake is in transit through the aquatic network rather than permanently stored.

The metabolism of these inland waters is sensitive to the same climatic drivers that affect the land, and it is being perturbed by warming directly. The occurrence of heatwaves in lakes, analogous to atmospheric and marine heatwaves, has increased under warming and alters the thermal structure, oxygen content, and biological activity of these water bodies, with consequences for their carbon processing and greenhouse-gas emission (Woolway, 2021). Because inland waters emit both carbon dioxide and methane, and because methane has a disproportionate forcing effect, changes in the metabolism of the aquatic network feed back onto the atmospheric composition assessed in Section 4.3. The rice paddy, a managed inland-water system, is a further illustration of the coupling, being a major anthropogenic methane source whose emissions and mitigation potential are now quantified in detail (Qian, 2023).

Coupling across the biogeochemical cycles

The carbon cycle cannot be assessed in isolation from the nitrogen, phosphorus, and other nutrient cycles that co-limit biological productivity, because the capacity of the land and ocean sinks to respond to rising carbon dioxide is constrained by the availability of these nutrients, and the human perturbation of the nutrient cycles is now of a magnitude comparable to the natural fluxes (*high confidence*). The fertilizing effect of elevated carbon dioxide on

plant growth, which underpins the land sink, cannot be sustained indefinitely if nitrogen or phosphorus becomes limiting, so the strength of the carbon-cycle feedback depends on the coupled dynamics of the nutrient cycles. Human activity has roughly doubled the reactive-nitrogen input to the biosphere through fertilizer production and combustion, and this nitrogen both fertilizes carbon uptake and, when it drives emissions of nitrous oxide, adds a potent greenhouse forcing of its own. The transgression of the biogeochemical-flow planetary boundaries, of which the nitrogen and phosphorus boundaries are the clearest examples, signals that these cycles are now operating far outside their pre-industrial range (Richardson, 2023).

The systemic view that emerges is one in which the carbon cycle, the nutrient cycles, the dust cycle, and the water cycle are mutually coupled, so that a perturbation to one propagates to the others. This coupling is the physical basis for the recognition that the climate and biodiversity crises are not separable problems but manifestations of a single disruption of the biosphere's regulatory processes, and that the loss of biological diversity itself undermines the functioning of the cycles that regulate the climate (Pörtner, 2023). The framing of the human enterprise as embedded within, and dependent upon, the biosphere's biogeochemical functioning makes explicit that the process understanding assembled in this chapter is ultimately an account of the life-support system on which human welfare rests (Folke, 2021).

4.6 Water Cycle Processes

Thermodynamic intensification of the water cycle

The water cycle is intensifying as the atmosphere warms, following at the largest scale the thermodynamic constraint that a warmer atmosphere holds more water vapor, so that the wet-get-wetter and dry-get-drier tendency of the moisture transport is a robust expectation now increasingly confirmed by observation (*high confidence*). The saturation vapor pressure of water rises steeply with temperature, at approximately seven percent per degree C, and because relative humidity remains near-constant, the actual water-vapor content of the atmosphere rises at a similar rate. This enhanced moisture loading has several consequences that together constitute the intensification of the cycle: the moisture available for precipitation extremes increases, the transport of water from evaporative to convergent regions strengthens, and the atmospheric demand for water—the evaporative pull on surfaces—rises. The rise in atmospheric moisture is the same process that drives the water-vapor feedback assessed in Section 4.2, so the intensification of the water cycle and the amplification of warming are two expressions of a single thermodynamic fact.

The global-mean precipitation does not rise as steeply as the water-vapor content, however, because precipitation is constrained by the energy budget rather than by moisture availability: the atmosphere can only precipitate as

much as it can radiatively cool to balance the latent heat released. This energetic constraint means that global precipitation rises at a slower rate than moisture, on the order of one to three percent per degree C, which implies a lengthening of the residence time of water vapor and a tendency toward more intense but less frequent precipitation. The gap between the fast rise of extreme precipitation, governed by moisture, and the slow rise of mean precipitation, governed by energy, is a central organizing principle of the water-cycle assessment and is robust across observational and modelling evidence.

Drought, aridity, and evaporative demand

Rising evaporative demand under warming is increasing the frequency, intensity, and speed of onset of droughts, and the anthropogenic fingerprint on the global drought record has become detectable across frequency, duration, and intensity (*high confidence*). Drought is not a single phenomenon but a family of related conditions—meteorological drought defined by precipitation deficit, agricultural drought defined by soil-moisture deficit, and hydrological drought defined by streamflow and groundwater deficit—linked by the propagation of a water shortfall through the terrestrial system. Warming affects each stage: it raises the atmospheric demand that draws moisture from soils and vegetation, it shifts precipitation toward more intense but less frequent events that run off rather than infiltrate, and it advances snowmelt so that the seasonal water supply arrives earlier and is exhausted sooner. Formal attribution now finds evidence of anthropogenic influence on the global statistics of drought, establishing that the observed intensification is not attributable to natural variability alone (Chiang, 2021).

A qualitatively new feature of the warming climate is the flash drought, a drought that develops with unusual rapidity, driven by a combination of precipitation deficit and high evaporative demand that desiccates soils within weeks rather than months. The transition toward more frequent flash droughts under continued warming has been demonstrated, and it poses a distinct challenge because the speed of onset compresses the window for anticipation and response (Yuan, 2023). The expansion of aridity has direct consequences for terrestrial biogeochemistry through the salinization of soils, which is projected to spread across large agricultural regions as evaporative concentration of salts intensifies under drier and warmer conditions (Hassani, 2021). Drought and aridity thus link the water-cycle assessment to the carbon-cycle vulnerability discussed in Section 4.4, because water stress is a principal driver of the declining forest resilience and elevated tree mortality that threaten the land sink (Forzieri, 2022; Hartmann, 2022).

Extremes, attribution, and compound events

The attribution of individual extreme weather and hydrological events to anthropogenic warming has matured into a quantitative science, establishing that many observed extremes have been made more likely or more intense

by the human perturbation of the climate, and that compound events—the coincidence of multiple hazards—are a growing concern (*high confidence*).

The attribution of extremes proceeds by comparing the probability of an event in the present climate with its probability in a counterfactual climate without anthropogenic forcing, using large ensembles of model simulations conditioned on the observed record. Applied across a wide range of events, this approach has shown that the intensification of heat extremes, heavy precipitation, and certain droughts carries a detectable anthropogenic signal, and it provides the methodological bridge between the global process understanding and the specific events that societies experience (Clarke, 2022).

Compound events are especially consequential because their impacts exceed the sum of their parts: a heatwave coinciding with a drought, or heavy precipitation falling on a saturated catchment, produces effects disproportionate to either hazard alone. The intensification of the water cycle raises the probability of such coincidences by increasing the variance as well as the mean of the relevant variables. Regional expressions of these compound risks are pronounced in transitional climates: the eastern Mediterranean and Middle East, a region at the intersection of subtropical drying and rising heat, is emerging as a hotspot in which multiple extremes intensify together (Zittis, 2022). Mountain regions, warming faster than the global mean and central to the water supply of downstream populations through their snow and ice reservoirs, concentrate hydrological risk where the elevation-dependent amplification of warming accelerates the loss of the frozen water store (Pepin, 2022).

The frozen water store and inland waters

The cryospheric component of the water cycle—snow, glaciers, and permafrost—functions as a water reservoir and a climate archive whose loss redistributes water in space and time and releases stored carbon, coupling the water cycle to the carbon cycle and the surface-albedo feedback (*high confidence*).

Seasonal snow and mountain glaciers store precipitation in the cold season and release it in the warm season, buffering the water supply of downstream regions; as warming shifts precipitation from snow to rain and accelerates melt, this natural regulation is degraded, with the water arriving earlier and the late-season supply diminishing. The rapid warming of the Arctic, at nearly four times the global rate, drives the loss of this frozen store at high latitudes and, through the surface-albedo feedback, accelerates the very warming responsible for the loss (Rantanen, 2022). The thaw of permafrost adds a carbon-cycle dimension, exposing long-frozen organic matter to microbial decomposition and potential release as carbon dioxide and methane, a process whose magnitude depends on the same mineral-organic controls that govern soil carbon more generally (Kleber, 2021), as mediated by the soil microbial community (Sokol, 2022).

The liquid inland waters that receive the meltwater are themselves being transformed. Lakes are warming and experiencing heatwaves that alter their stratification and biogeochemistry, changing their role in the carbon cycle and their emission of greenhouse gases (Woolway, 2021). The river networks

that connect the terrestrial and marine domains process a large lateral flux of carbon whose magnitude is sensitive to the changing hydrology, so that the intensification of the water cycle propagates into the carbon budget through the altered delivery and processing of terrestrial carbon (Battin, 2023). These couplings reinforce the theme that the water cycle is not a separable subsystem but the medium through which many of the biogeochemical processes assessed in this chapter are transmitted.

Water resources and human coupling

The intensification and redistribution of the water cycle interacts with rising human demand to produce increasing pressure on freshwater resources, and the assessment of water-cycle processes is incomplete without the coupled human system that abstracts, stores, and consumes water (*high confidence*).

Freshwater availability is a growing global challenge as population and economic demand rise against a supply that climate change is making more variable and, in many regions, scarcer (Mishra, 2023). The quantitative assessment of this pressure relies on global hydrological models that represent the terrestrial water balance and its human modification, tracking abstraction, storage, and return flows alongside the natural fluxes (Schmied, 2021). These models make explicit that the human perturbation of the water cycle, through irrigation, damming, and groundwater extraction, is now a first-order term in the terrestrial water balance of many regions, comparable to the climatic changes in its effect on water availability.

The coupling runs in both directions: water availability constrains agriculture, and agriculture is the dominant human water use, so that the intensification of the water cycle bears directly on food security through its effect on the water available for crops (Jägermeyr, 2021). The projected changes in agricultural productivity under a changing climate emerge earliest and most clearly where water stress is greatest, integrating the thermodynamic intensification of the cycle, the rise in evaporative demand, and the human demand for water into a single assessment of the pressure on the systems that sustain human populations (Yang, 2024). This human coupling closes the loop of the chapter, returning the physical process understanding to the biosphere-embedded framing with which the biogeochemical assessment concluded (Folke, 2021).

4.7 Integrating Processes: Earth-System Models, Thresholds, and Confidence

Earth-system models as the integrating framework

Earth-system models are the framework within which the forcings, feedbacks, and cycles assessed in this chapter are combined into a physically consistent whole, and the current model generation reproduces the observed forcing

history and the broad structure of the feedbacks while spanning a range of sensitivities that the assessment interprets rather than averages (*high confidence*). A model of this class couples the atmosphere, ocean, land, cryosphere, and, increasingly, the interactive carbon and other biogeochemical cycles, so that a single simulation computes the greenhouse and aerosol forcing, the water-vapor, lapse-rate, albedo, and cloud feedbacks, and the response of the land and ocean carbon sinks, all interacting self-consistently. This self-consistency is the principal virtue of the modelling approach: it avoids the double-counting and internal contradiction that can arise when feedbacks estimated in isolation are summed, and it allows the emergent behavior of the coupled system—including the pattern effect and the interaction of feedbacks—to be represented. The representative model examined in the intercomparison, incorporating an interactive carbon cycle within a full Earth-system configuration, illustrates both the achievement and the residual spread of the current generation (Lovato, 2022).

The intercomparison of many such models under common experimental protocols is the mechanism by which structural uncertainty is characterized. Because the models differ in their parameterizations of the sub-grid processes that dominate the feedback uncertainty—convection, cloud microphysics, aerosol activation, and soil biogeochemistry—their spread is an estimate, if an imperfect one, of the uncertainty attributable to incomplete process representation. The assessment does not treat the multi-model mean as truth; rather, it weights the models by their fidelity to observed processes and applies emergent constraints that relate an observable feature of the present climate to a feedback or sensitivity, thereby extracting more information than the raw ensemble spread provides. The scenario framework overlaid on these models separates the physical uncertainty they characterize from the societal uncertainty of the emissions pathway, providing internally consistent forcing trajectories against which the physical response is projected (Tebaldi, 2021).

Data-driven methods and the augmentation of process models

The integration of physical process knowledge with data-driven and machine-learning methods is emerging as a means to constrain the parameterizations that dominate the feedback uncertainty, and this hybrid approach is beginning to complement, rather than replace, the physically based models (*medium confidence*). The processes that dominate the residual uncertainty—low-cloud response, aerosol activation, and soil-carbon turnover—are precisely those that occur below the resolution of global models and must be parameterized from limited observations. Methods that integrate scientific knowledge with machine learning can extract feedback-relevant relationships from large observational and high-resolution modelling archives while respecting the physical constraints that keep the inferences credible, offering a route to reduce the parameterization uncertainty that pure process modelling has struggled to close (Willard, 2022). The value of these methods for the assessment lies in their capacity to bring more of the available observational information to bear on the specific relationships that matter for the feedbacks, and their limitation

lies in the need to ensure that the learned relationships generalize to a climate state outside the range of the training data.

The broader monitoring enterprise on which both the process models and the data-driven methods depend has been transformed by the growth of observational capacity, from the near-real-time tracking of emissions (Liu, 2022) to the unified satellite quantification of the terrestrial biosphere and the year-round in-situ characterization of the polar atmosphere (Shupe, 2022). This observational foundation is what allows the process understanding to be tested rather than merely asserted, and its continued extension is a precondition for the further narrowing of the uncertainties that this chapter has identified.

Thresholds, resilience, and the risk of abrupt change

The process understanding assembled in this chapter identifies several elements of the climate system whose response to forcing may be abrupt or effectively irreversible on human timescales, and while the assessment does not find that any global-scale abrupt transition is imminent, it establishes that the risk of such transitions rises with the magnitude and rate of forcing (*medium confidence*). The feedbacks that dominate the sensitivity are, for the most part, smoothly varying, but several of the processes coupled to them exhibit threshold behavior: the loss of forest resilience that could convert a carbon sink into a source (Forzieri, 2022), the thaw of permafrost that could release a large stored carbon pool, the reorganization of marine productivity implied by the expansion of coastal blooms (Dai, 2023), and the loss of ice whose albedo feedback is self-reinforcing (Rantanen, 2022). The common feature of these elements is that they involve a positive feedback strong enough, locally, to sustain a transition once initiated, so that the system does not return to its prior state when the forcing is reduced.

The assessment of these risks draws on the recognition that the Earth system has been pushed beyond several of the boundaries within which its Holocene stability was maintained (Richardson, 2023), and it takes seriously the analysis of high-end and catastrophic scenarios that, while not the most likely outcomes, carry consequences severe enough to warrant explicit assessment (Kemp, 2022). The framing of these boundaries to encompass both the biophysical thresholds and the just distribution of the remaining safe operating space reflects the maturation of the field from a purely physical science toward one that integrates the physical with the societal consequences of transgression (Rockström, 2023). The coupling of the climate and biodiversity crises, in which the loss of biological diversity degrades the very processes that regulate the climate, underscores that the thresholds of concern are not confined to the physical system but extend to the living processes on which its regulation depends (Pörtner, 2023).

Confidence, uncertainty, and the state of process understanding

Taken as a whole, the process understanding of forcings, feedbacks, and cycles has reached a state in which the sign and approximate magnitude of every

major process is established with high confidence, the residual uncertainties are concentrated in a small number of identified feedbacks and couplings, and the direction of ongoing research is set by the need to constrain those specific residuals (*high confidence*). The greenhouse-gas forcing is known with the highest confidence of any term; the aerosol forcing remains the dominant forcing uncertainty and is coupled through the historical record to the uncertainty in sensitivity; the cloud and carbon-cycle feedbacks remain the dominant feedback uncertainties, though both have narrowed; and the water cycle's intensification is robustly established while its regional expression retains substantial uncertainty. This structure of confidence—high on the aggregate, more qualified on the specific residuals—is not a weakness of the assessment but a mark of its maturity, because it identifies precisely where further effort will yield the greatest reduction in the uncertainty of projection.

The economic and societal metrics that translate this physical understanding into decision-relevant terms inherit its uncertainty structure: the social cost of carbon, resting on the forcing-to-warming-to-damage chain, is sensitive to the sensitivity and carbon-cycle uncertainties assessed here, and its comprehensive re-estimation to a substantially higher value reflects both updated damage understanding and the physical relationships this chapter has assessed (Rennert, 2022). The integration of the climate and health assessments, recognizing that the aerosols which cool the climate also impose a mortality burden, exemplifies the coupling of the physical and the societal that the process understanding now supports (Romanello, 2022). The overarching conclusion of this chapter is that the physical processes governing the climate response to human forcing are understood well enough to support confident projection of the direction and approximate magnitude of change, that the coupled carbon, biogeochemical, and hydrological cycles transmit and amplify that change through mechanisms now resolved at the process level, and that the remaining uncertainties, while real and consequential, are bounded, identified, and the subject of active and converging lines of investigation (Folke, 2021).

4.8 Detection, Attribution, and the Human Fingerprint on Processes

From global-mean detection to process-level attribution

The attribution of climate change has progressed from the detection of a warming trend in the global-mean temperature to the attribution of specific process changes—in the water cycle, the carbon cycle, and the frequency of extremes—to human forcing, and this maturation is what allows process understanding to be tested against the observed record rather than assessed only in models (*high confidence*). Detection establishes that an observed change exceeds what internal variability alone would produce; attribution assigns the change to a cause by comparing the observed record against simulations with

and without the anthropogenic forcing. Applied to the water cycle, this framework has attributed to human influence the observed changes in the global statistics of drought across frequency, duration, and intensity, moving the assessment of the hydrological response from expectation to established fact (Chiang, 2021). Applied to extremes, it has established that a wide class of heat and precipitation events has been made more likely or more intense by warming, providing the empirical counterpart to the thermodynamic reasoning of Section 4.5 (Clarke, 2022).

The attribution enterprise depends on the same Earth-system models that integrate the process understanding, because the counterfactual climate against which the observed record is compared can only be constructed by simulation. This creates a productive circularity: the models are tested by their ability to reproduce the detected changes, and the attribution is strengthened by the models' fidelity to the underlying processes. The wide-ranging review of climate change impacts, adaptation, and mitigation measures situates these attributed process changes within the fuller account of how the physical changes propagate into human and natural systems, underscoring that the process understanding of this chapter is the foundation on which impact assessment rests (Abbass, 2022).

Regional heterogeneity as a signature of process

The regional heterogeneity of the observed changes is itself a signature of the underlying processes, because the pattern of change reflects the spatial structure of the forcings and feedbacks rather than a uniform global response, and the assessment reads this heterogeneity as evidence about mechanism (*high confidence*). The Arctic warms fastest because the surface-albedo feedback and the associated atmospheric processes are concentrated there (Rantanen, 2022); the mountains warm in an elevation-dependent pattern because the retreat of snow and the changing energy balance vary with altitude (Pepin, 2022); and the eastern Mediterranean and Middle East emerge as a hotspot because subtropical drying and rising heat coincide in that region (Zittis, 2022). Each of these regional signals is a test of the process understanding: a model that reproduces the global mean but fails the regional pattern has captured the aggregate energetics while missing the mechanisms, and the assessment weights models accordingly.

This reading of heterogeneity extends to the biogeochemical response. The spatial pattern of the terrestrial carbon sink, resolved by satellite quantification of the biosphere, reveals where the sink is strengthening under carbon dioxide fertilization and where it is faltering under water stress and disturbance, and the pattern of declining forest resilience identifies the regions in which the land carbon-cycle feedback is most likely to strengthen first (Forzieri, 2022). The regional structure of the response is therefore not noise to be averaged away but information to be exploited, and its incorporation into the assessment is one of the advances that distinguishes the current understanding from the global-mean framing of earlier cycles.

4.9 From Process Understanding to Committed Change and Reversibility

Cumulative emissions, committed warming, and the net-zero concept

The near-linear relationship between cumulative carbon dioxide emissions and global-mean warming, which emerges from the combination of the forcing and carbon-cycle processes assessed in this chapter, is the physical basis for the concept of a carbon budget and for the finding that stabilizing temperature requires reaching net-zero carbon dioxide emissions (*high confidence*). The linearity arises from a coincidental compensation between two processes: as cumulative emissions rise, the logarithmic saturation of the carbon dioxide forcing would tend to reduce the warming per tonne, while the declining efficiency of the carbon sinks—the carbon-cycle feedback—tends to raise the airborne fraction and increase the warming per tonne. These effects approximately cancel, producing a relationship robust enough to underpin policy. The implication is that warming is controlled by the total quantity of carbon emitted rather than by the rate, and that halting the rise in temperature requires bringing net carbon dioxide emissions to zero. The precise meaning of net zero, and the conditions under which it stabilizes temperature, have been clarified in a manner that makes explicit the roles of the residual emissions and the removals that must balance them (Fankhauser, 2021).

The carbon-cycle processes assessed in Section 4.4 are central to this framing because the removals that define net zero depend on the same land and ocean sinks whose vulnerability the chapter has documented. A sink that weakens under warming reduces the natural drawdown and increases the removal effort required to reach net zero, so the carbon-cycle feedback is not merely a physical uncertainty but a determinant of the mitigation challenge. The reliance on the land sink also brings the biogeophysical processes into the policy assessment, because efforts to enhance terrestrial removal interact with the water, nutrient, and biodiversity processes that govern ecosystem functioning (Seddon, 2021).

Reversibility, irreversibility, and the asymmetry of the response

The reversibility of the climate response to a reduction in forcing varies sharply among processes, so that some elements of the system would recover on decadal timescales while others—ice loss, permafrost carbon release, and certain ecosystem transitions—are effectively irreversible on human timescales, and this asymmetry is a central conclusion of the process understanding (*high confidence*). The surface temperature response to a cessation of emissions is expected to stabilize rather than reverse, because the ocean continues to take up heat even as it takes up less carbon, and the two adjustments approximately offset. The fast components of the water cycle would respond to a stabilized temperature, but the slow components—the deep ocean, the great ice sheets, and the permafrost carbon pool—carry a commitment that

persists long after emissions cease. The loss of ice, once the albedo feedback has driven a transition, does not reverse when the forcing is reduced, because the darker surface continues to absorb the radiation that maintains the warmer state (Rantanen, 2022).

Ecosystem transitions exhibit a similar asymmetry. A forest that has crossed a resilience threshold and transitioned to a different state does not automatically return when the stress is relieved, because the soil, hydrology, and biotic community that supported the forest have themselves been altered (Forzieri, 2022; Hartmann, 2022). This asymmetry between the ease of transgression and the difficulty of recovery is the physical basis for a precautionary stance toward the thresholds identified in Section 4.6, and it is reflected in the assessment's attention to the pathways that keep the system within the range where its recovery remains possible. The integration of the physical process understanding with the sustainable-development framing makes explicit that the choices which determine whether these thresholds are crossed are made within a narrow and closing window, and that the process understanding assembled here is what defines the boundaries of that window (Soergel, 2021).

Synthesis: a coupled system understood at the process level

The cumulative advance represented in this chapter is the transition from a description of the climate response in terms of aggregate quantities to an account grounded in the specific, separately constrained processes that generate those quantities, and it is this process-level grounding that underwrites the confidence of the projections built upon it (*high confidence*). The forcing is no longer a single number but a decomposition into agents whose individual histories are traced through the emissions that produce them (Lamb, 2021) and monitored through the atmospheric growth they cause (Liu, 2022). The sensitivity is no longer an empirical fit but the propagated sum of feedbacks each constrained by observation and process modelling, with the cloud feedback narrowed by the integration of physical knowledge with data-driven inference (Willard, 2022) and the carbon-cycle feedback constrained by the annual closure of the global budget (Friedlingstein, 2023). The cycles are no longer boundary conditions but interactive components whose couplings—dust to ocean productivity, water to carbon, nutrients to the land sink—are represented within the coupled models that integrate them (Lovato, 2022).

This process grounding also clarifies where the understanding remains incomplete and why. The residual uncertainties are not distributed uniformly but concentrated in the small-scale processes that must be parameterized: the microphysics of low clouds, the activation of aerosols into cloud droplets, and the microbial and mineral controls on soil carbon turnover (Kleber, 2021). Each of these is the subject of a dedicated observational and modelling effort, and the trajectory of the field is toward their progressive constraint rather than toward any wholesale revision of the framework. The observational foundation that makes this constraint possible—the near-real-time emissions monitoring, the year-round polar campaigns (Shupe, 2022), and the satellite records of the

biosphere and the water cycle—has been strengthened to the point where the process understanding can be tested against the observed evolution of the system rather than asserted from theory alone.

The chapter's final judgment is therefore one of qualified confidence with a clear structure. The direction and approximate magnitude of the climate response to human forcing are established beyond reasonable scientific doubt; the coupled carbon, biogeochemical, and hydrological cycles transmit and in several respects amplify that response through mechanisms now resolved at the process level; and the risks of abrupt or irreversible change, while not assessed as imminent at the global scale, rise with the magnitude and rate of forcing and are concentrated in identifiable elements of the system (Forzieri, 2022). The economic translation of this understanding implies damages, and therefore a social cost of carbon, substantially larger than earlier assessments recognized (Rennert, 2022), and its ecological translation makes plain that the climate and biodiversity crises are a single disruption of the biosphere's regulatory processes rather than two separable problems (Pörtner, 2023). The process understanding assembled here is, in the end, an account of the mechanisms by which a habitable Earth maintains itself and of the specific ways in which those mechanisms are now being perturbed (Folke, 2021).

The forward agenda that follows from this synthesis is correspondingly well defined. Constraining the aerosol forcing, and with it the historical inference of sensitivity, requires sustained observation of aerosol–cloud interactions in the low-cloud regimes that dominate the uncertainty (Willard, 2022). Bounding the carbon-cycle feedback requires continued closure of the global carbon budget together with process-level study of the disturbance mechanisms—drought-driven mortality, loss of resilience, and permafrost thaw—that threaten the durability of the land sink (Friedlingstein, 2025). Projecting the intensification of the water cycle at the regional scale where it is experienced requires the coupled human–hydrological modelling that represents abstraction and management alongside the natural fluxes (Schmied, 2021). None of these agenda items implies a revision of the framework; each is a program for reducing a specified, bounded uncertainty within a structure whose broad outline is secure. It is on this basis that the subsequent chapters build their projections of future change, their attribution of observed impacts, and their assessment of the pathways by which the perturbation of these processes may yet be constrained (Hartmann, 2022).

4.10 References

- Abbass, K., et al. (2022). A review of the global climate change impacts, adaptation, and sustainable mitigation measures. *Environmental Science and Pollution Research*. <https://doi.org/10.1007/s11356-022-19718-6>
- Battin, T., et al. (2023). River ecosystem metabolism and carbon biogeochemistry in a changing world. *Nature*. <https://doi.org/10.1038/s41586-022-05500-8>
- Bergmann, M., et al. (2022). Plastic pollution in the Arctic. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-022-00279-8>
- Brunn, H., et al. (2023). PFAS: forever chemicals—persistent, bioac-

cumulative and mobile. Reviewing the status and the need for their phase out and remediation of contaminated sites. *Environmental Sciences Europe*. <https://doi.org/10.1186/s12302-023-00721-8> Cabernard, L., et al. (2021). Growing environmental footprint of plastics driven by coal combustion. *Nature Sustainability*. <https://doi.org/10.1038/s41893-021-00807-2> Chiang, F., et al. (2021). Evidence of anthropogenic impacts on global drought frequency, duration, and intensity. *Nature Communications*. <https://doi.org/10.1038/s41467-021-22314-w> Clarke, B., et al. (2022). Extreme weather impacts of climate change: an attribution perspective. *Environmental Research Climate*. <https://doi.org/10.1088/2752-5295/ac6e7d> Dai, Y., et al. (2023). Coastal phytoplankton blooms expand and intensify in the 21st century. *Nature*. <https://doi.org/10.1038/s41586-023-05760-y> Fankhauser, S., et al. (2021). The meaning of net zero and how to get it right. *Nature Climate Change*. <https://doi.org/10.1038/s41558-021-01245-w> Folke, C., et al. (2021). Our future in the Anthropocene biosphere. *AMBIO*. <https://doi.org/10.1007/s13280-021-01544-8> Forzieri, G., et al. (2022). Emerging signals of declining forest resilience under climate change. *Nature*. <https://doi.org/10.1038/s41586-022-04959-9> Frenne, P., et al. (2021). Forest microclimates and climate change: Importance, drivers and future research agenda. *Global Change Biology*. <https://doi.org/10.1111/gcb.15566> Friedlingstein, P., et al. (2023). Global Carbon Budget 2023. *Earth system science data*. <https://doi.org/10.5194/essd-15-5301-2023> Friedlingstein, P., et al. (2025). Global Carbon Budget 2024. *Earth system science data*. <https://doi.org/10.5194/essd-17-965-2025> Fussell, J., et al. (2022). A Review of Road Traffic-Derived Non-Exhaust Particles: Emissions, Physicochemical Characteristics, Health Risks, and Mitigation Measures. *Environmental Science & Technology*. <https://doi.org/10.1021/acs.est.03433a001> Hartmann, H., et al. (2022). Climate Change Risks to Global Forest Health: Emergence of Unexpected Events of Elevated Tree Mortality Worldwide. *Annual Review of Plant Biology*. <https://doi.org/10.1146/annurev-arplant-102820-012804> Hassani, A., et al. (2021). Global predictions of primary soil salinization under changing climate in the 21st century. *Nature Communications*. <https://doi.org/10.1038/s41467-021-26907-3> Jägermeyr, J., et al. (2021). Climate impacts on global agriculture emerge earlier in new generation of climate and crop models. *Nature Food*. <https://doi.org/10.1038/s43016-021-00400-y> Kemp, L., et al. (2022). Climate Endgame: Exploring catastrophic climate change scenarios. *Proceedings of the National Academy of Sciences*. <https://doi.org/10.1073/pnas.2108146119> Kleber, M., et al. (2021). Dynamic interactions at the mineral–organic matter interface. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-021-00162-y> Kok, J., et al. (2021). Contribution of the world’s main dust source regions to the global cycle of desert dust. *Atmospheric chemistry and physics*. <https://doi.org/10.5194/acp-21-8169-2021> Lamb, W., et al. (2021). A review of trends and drivers of greenhouse gas emissions by sector from 1990 to 2018. *Environmental Research Letters*. <https://doi.org/10.1088/1748-9326/abee4e> Liu, Z., et al. (2022). Monitoring global carbon emissions in 2021. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-022-00285-w> Lovato, T., et al. (2022). CMIP6 Simulations With the CMCC Earth System Model (CMCC-ESM2). *Journal of Advances in Modeling Earth Systems*. <https://doi.org/10.1029/2021ms002814> Mishra, R. (2023). Fresh Water availability and Its Global challenge. *British*

Journal of Multidisciplinary and Advanced Studies. <https://doi.org/10.37745/bj-mas.2022.0208> Pan, Y., et al. (2024). The enduring world forest carbon sink. *Nature*. <https://doi.org/10.1038/s41586-024-07602-x> Pendrill, F., et al. (2022). Disentangling the numbers behind agriculture-driven tropical deforestation. *Science*. <https://doi.org/10.1126/science.abm9267> Pepin, N., et al. (2022). Climate Changes and Their Elevational Patterns in the Mountains of the World. *Reviews of Geophysics*. <https://doi.org/10.1029/2020rg000730> Pörtner, H., et al. (2023). Overcoming the coupled climate and biodiversity crises and their societal impacts. *Science*. <https://doi.org/10.1126/science.abl4881> Qian, H., et al. (2023). Greenhouse gas emissions and mitigation in rice agriculture. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-023-00482-1> Qiu, L., et al. (2021). Erosion reduces soil microbial diversity, network complexity and multifunctionality. *The ISME Journal*. <https://doi.org/10.1038/s41396-021-00913-1> Rantanen, M., et al. (2022). The Arctic has warmed nearly four times faster than the globe since 1979. *Communications Earth & Environment*. <https://doi.org/10.1038/s43247-022-00498-3> Rennert, K., et al. (2022). Comprehensive evidence implies a higher social cost of CO₂. *Nature*. <https://doi.org/10.1038/s41586-022-05224-9> Richardson, K., et al. (2023). Earth beyond six of nine planetary boundaries. *Science Advances*. <https://doi.org/10.1126/sciadv.adh2458> Rillig, M., et al. (2021). Microplastic effects on carbon cycling processes in soils. *PLoS Biology*. <https://doi.org/10.1371/journal.pbio.3001130> Rockström, J., et al. (2023). Safe and just Earth system boundaries. *Nature*. <https://doi.org/10.1038/s41586-023-06083-8> Romanello, M., et al. (2022). The 2022 report of the Lancet Countdown on health and climate change: health at the mercy of fossil fuels. *The Lancet*. [https://doi.org/10.1016/s0140-6736\(22\)01540-9](https://doi.org/10.1016/s0140-6736(22)01540-9) Schmied, H., et al. (2021). The global water resources and use model WaterGAP v2.2d: model description and evaluation. *Geoscientific model development*. <https://doi.org/10.5194/gmd-14-1037-2021> Seddon, N., et al. (2021). Getting the message right on nature-based solutions to climate change. *Global Change Biology*. <https://doi.org/10.1111/gcb.15513> Soergel, B., et al. (2021). A sustainable development pathway for climate action within the UN 2030 Agenda. *Nature Climate Change*. <https://doi.org/10.1038/s41558-021-01098-3> Shupe, M., et al. (2022). Overview of the MOSAiC expedition: Atmosphere. *Elementa Science of the Anthropocene*. <https://doi.org/10.1525/elementa.2021.00060> Sokol, N., et al. (2022). Life and death in the soil microbiome: how ecological processes influence biogeochemistry. *Nature Reviews Microbiology*. <https://doi.org/10.1038/s41579-022-00695-z> Tebaldi, C., et al. (2021). Climate model projections from the Scenario Model Intercomparison Project (ScenarioMIP) of CMIP6. *Earth System Dynamics*. <https://doi.org/10.5194/esd-12-253-2021> Willard, J., et al. (2022). Integrating Scientific Knowledge with Machine Learning for Engineering and Environmental Systems. *ACM Computing Surveys*. <https://doi.org/10.1145/3514228> Woolway, R., et al. (2021). Lake heatwaves under climate change. *Nature*. <https://doi.org/10.1038/s41586-020-03119-1> Yang, X., et al. (2024). Climate change exacerbates the environmental impacts of agriculture. *Science*. <https://doi.org/10.1126/science.adn3747> Yuan, X., et al. (2023). A global transition to flash droughts under climate change. *Science*. <https://doi.org/10.1126/science.abn6301> Zittis, G., et al. (2022). Climate

Change and Weather Extremes in the Eastern Mediterranean and Middle East.
Reviews of Geophysics. <https://doi.org/10.1029/2021rg000762>

Chapter 5

Chapter 5: Scenarios and Near-Term Projections

5.1 Executive Summary

The scenario framework inherited from the Sixth Assessment Report remains fit for purpose in the near term, but the plausibility of its highest-emission storylines has continued to erode as observed and pledged mitigation trajectories diverge from unmitigated-growth assumptions (*high confidence*). The Shared Socioeconomic Pathway architecture, coupled to the Scenario Model Intercomparison Project experimental design, continues to anchor multi-model projections; its near-term temperature spread is dominated less by scenario choice than by model response and internal variability (Tebaldi et al., 2021) (*high confidence*).

Near-term global surface temperature is projected to continue rising across all plausible emissions pathways, and the difference between pathways is small before about 2040 because near-term warming is largely committed by past and present emissions (*very likely*). Divergence between low- and high-emission pathways becomes assessable only toward mid-century; consequently the near-term is a regime in which mitigation choices are being made but their temperature signal has not yet emerged from the noise of internal variability (*high confidence*).

Global fossil carbon dioxide emissions have remained near record levels through the early 2020s, and the remaining carbon budget consistent with limiting warming to well below 2 degrees C continues to shrink (*high confidence*). The annual carbon budget accounting shows emissions plateauing rather than falling at the rate implied by the lowest scenarios (Friedlingstein et al., 2023), so the gap between current-policy trajectories and deep-mitigation pathways is widening in absolute terms (*high confidence*).

Sectoral decarbonization pathways for the hardest-to-abate sectors — cement, iron and steel, and coal power — show that near-term emissions are strongly path-dependent on capital-stock turnover, and committed emissions from existing infrastructure constrain the feasible speed of the near-term transition (*high confidence*). Plant-level analyses of coal (Cui et al., 2021), steel

(Lei et al., 2023), and cement (Cheng et al., 2023) indicate that near-term abatement is dominated by decisions about existing and under-construction assets rather than by long-run technology availability (*high confidence*).

Internal variability is the largest source of uncertainty in near-term regional projections and can locally mask or amplify the forced trend for a decade or more (*high confidence*). In rapidly warming regions such as the Arctic, the forced signal is large enough that it dominates even at near-term horizons (Rantanen et al., 2022), whereas in mid-latitude land regions the near-term trend can be obscured by variability (*high confidence*).

Initialized decadal prediction adds skill over uninitialized projections for some large-scale quantities in the first few years, but skill for regional precipitation and extremes remains limited (*medium confidence*). The near-term prediction enterprise is maturing, yet its principal value is presently in ocean-heat and large-scale circulation indices rather than in actionable regional extremes forecasts (*medium confidence*).

The probability of exceeding key near-term thresholds — including individual years above 1.5 degrees C relative to pre-industrial — has become substantial, though a single threshold-exceeding year does not constitute crossing the long-term warming level (*high confidence*). Near-term extremes, particularly heatwaves and compound drought-heat events, are projected to intensify faster than the mean, so risk-relevant change outpaces the change in global mean temperature (Barriopedro et al., 2023; Tripathy et al., 2023) (*high confidence*).

Low-probability, high-impact outcomes — including the possibility of engaging tipping elements and cascading risks — are not excluded on near-term horizons and warrant explicit assessment alongside central projections (*medium confidence*). Storyline and risk-framing approaches complement probabilistic projection by making the tails of the distribution and their societal entanglements assessable (Simpson et al., 2021), complemented by the assessment of catastrophic and extreme storylines (Kemp et al., 2022) (*medium confidence*).

5.2 Scenario frameworks and their evolution since the Sixth Assessment

The community scenario framework used for near-term projection is a layered construction — socioeconomic storylines, emissions and land-use pathways, and a coordinated set of climate-model experiments — and its near-term behavior is best understood by treating those layers separately (*high confidence*). In the assessment lineage inherited from the Sixth Assessment Report, the socioeconomic layer is provided by the Shared Socioeconomic Pathways, the forcing layer by harmonized emissions pathways spanning a range of end-of-century radiative forcing levels, and the climate-response layer by the multi-model experiments coordinated under the Scenario Model Intercom-

parison Project (Tebaldi et al., 2021). This separation matters for the near term because, before about 2040, the projected climate is far more sensitive to the climate-response layer — model spread and internal variability — than to the choice among socioeconomic storylines. The scenarios diverge in their emissions almost immediately, but the climate system integrates that forcing slowly, so the temperature consequences of divergent choices are still emerging over the assessment’s near-term window.

The core experimental design that underpins near-term multi-model projection was documented as part of the Scenario Model Intercomparison Project, which specified a small number of priority pathways spanning very low to very high forcing and a larger set of pathways filling in the response surface (Tebaldi et al., 2021). The near-term envelope in that design is constructed so that all pathways share, to a good approximation, the same recent historical forcing and differ chiefly in the rate at which they depart from it. As a result, the multi-model near-term spread is dominated by the range of transient climate responses across models and by the phasing of simulated internal variability, rather than by the labels attached to the pathways. This is a central interpretive point for readers of near-term projections: a plot of global mean temperature to 2040 that fans out by scenario is, over that window, mostly displaying model and initial-condition differences, not the policy content of the scenarios.

The plausibility of the highest-forcing pathways as descriptions of the real world’s near-term trajectory has continued to weaken since the Sixth Assessment, principally because their assumed fossil-energy expansion is increasingly at odds with observed technology deployment and announced policy (high confidence). The highest pathways were constructed as high-reference cases in which fossil-fuel use — especially coal — expands strongly through the century. Analyses of realized and announced decarbonization indicate that such coal-intensive futures have become progressively less consistent with the trajectory of the global energy system (Cui et al., 2021; Victoria et al., 2021). This does not remove the highest pathways from the assessment: they retain value as sensitivity cases and as bounding scenarios for the carbon cycle and for physical-science process studies. But for near-term projection of the most likely climate outcomes, the assessment gives greater weight to the intermediate and lower pathways, while continuing to report the full range so that low-probability, high-impact outcomes remain visible.

The socioeconomic layer of the framework has been extended since the Sixth Assessment by explicit sustainable-development pathways that couple climate mitigation to broader development goals, illustrating that a given near-term forcing level can be reached by socioeconomic routes with markedly different implications for equity, land use, and co-benefits (Soergel et al., 2021). These development-oriented pathways are useful in the near term precisely because near-term emissions are determined by concrete, near-at-hand choices — infrastructure, land use, and energy investment — rather than by distant technological assumptions. A sustainable-development framing makes explicit that the near-term emissions layer is where policy actually operates, even though its temperature signal is delayed. The assessment treats such pathways as mem-

bers of the scenario family rather than as replacements for the priority set, so that near-term projections can be interpreted against a spread of plausible socioeconomic contexts.

The translation from emissions pathways to global temperature outcomes is now performed through well-documented reduced-complexity and assessment models whose near-term behavior has been calibrated against the fuller model ensemble, improving the traceability of near-term projections (*high confidence*). A key methodological advance carried into this assessment is the explicit, reproducible chain that converts an emissions pathway into a probabilistic temperature outcome. The methods used to produce the mitigation-pathway temperature assessment in the Sixth Assessment cycle were documented in detail, including the emulator calibration and the propagation of carbon-cycle and forcing uncertainty (Kikstra et al., 2022). For near-term projection this traceability is valuable because it separates three distinct contributions to the projected temperature: the emissions themselves, the conversion of emissions to radiative forcing, and the climate response to that forcing. In the near term the last of these — the fast component of the climate response — dominates, which is why near-term temperature is relatively insensitive to scenario choice and relatively sensitive to the assumed transient response.

The forcing layer itself has been refined by improved treatment of non-carbon-dioxide agents, whose near-term influence is disproportionate. Short-lived climate forcers — including aerosols and methane — exert a strong influence on the near-term forcing trajectory even though their long-term temperature contribution is smaller than that of carbon dioxide. Because several low-emission pathways achieve rapid reductions in co-emitted aerosols alongside carbon dioxide, they can exhibit a transient near-term warming from the unmasking of aerosol cooling, before the reduced greenhouse-gas forcing dominates at longer horizons. This aerosol-unmasking effect is a distinctive near-term feature of the low pathways and is one reason the near-term temperature difference between low and high pathways is small or can even reverse in sign during the first decade or two. The assessment reports this explicitly so that the near-term behavior of the lowest pathways is not misread as evidence that mitigation is ineffective; rather, it reflects the different time constants of aerosol and greenhouse-gas forcing.

The scenario framework is increasingly complemented, not replaced, by storyline and risk-based approaches that make the tails of the near-term distribution and their societal entanglements assessable (*high confidence*). Probabilistic multi-model projection is well suited to characterizing the central part of the near-term distribution, but decision-relevant risk often lies in the tails and in the interaction of climate hazards with vulnerable systems. A framework for complex climate-change risk assessment has been articulated that decomposes risk into interacting hazard, exposure, vulnerability, and response components, and that is explicitly designed to handle the compounding and cascading of risks that pure hazard projection does not capture (Simpson et al., 2021). In the near term, where the hazard change is modest in the global mean but potentially large in specific compound events, this risk framing is

essential to avoid understating the consequences of near-term change. The assessment therefore treats the scenario framework and the risk framework as complementary: the former supplies the physical hazard trajectory, the latter situates it within exposed and vulnerable systems.

The explicit assessment of high-impact, low-likelihood outcomes has been strengthened relative to earlier cycles. Catastrophic and extreme scenarios — those in the upper tail of the response distribution or those engaging strongly nonlinear Earth-system feedbacks — have been argued to be under-examined relative to their importance for risk management (Kemp et al., 2022). For near-term projection, the practical consequence is that the assessment reports not only the central estimate and likely range but also explicitly flags the physically plausible upper tail, including the possibility that the transient response lies at the high end of the assessed range. This is not a prediction that such outcomes will occur in the near term; it is a statement that they cannot be excluded and that their potential severity justifies their inclusion in the assessed range. The near-term window is precisely the horizon over which early-warning indicators of a high-sensitivity world would begin to be detectable, which gives the upper tail particular salience for the coming two decades.

Since the Sixth Assessment, the coupling of climate risk to interacting global stresses — a “polycrisis” framing — has been advanced as a way to represent how near-term climate change interacts with concurrent economic, food, and geopolitical stresses (*medium confidence*). The recognition that climate hazards do not act in isolation but interact with other global stresses has been formalized in the analysis of entangled crises and their causal mechanisms (Lawrence et al., 2024). While this framing sits at the boundary of Working Group I’s physical-science remit, it is relevant to near-term projection because the societal consequences of a given near-term hazard trajectory depend strongly on the state of the systems it impinges upon. The physical-science assessment supplies the hazard side of this interaction; the polycrisis and coupled-crisis literature (Pörtner et al., 2023) situates the hazard within concurrent biodiversity, food, and development stresses. The assessment notes this coupling without extending its own quantitative projections beyond the physical climate system.

Taken together, the evolution of the scenario framework since the Sixth Assessment is best described as consolidation and extension rather than replacement. The Shared Socioeconomic Pathway architecture and the Scenario Model Intercomparison Project design remain the backbone of near-term multi-model projection (Tebaldi et al., 2021). The main developments are: a downweighting of the highest-forcing pathways as descriptions of the most likely near-term future while retaining them as bounding cases; a more traceable emissions-to-temperature chain (Kikstra et al., 2022); an extension of the socioeconomic layer toward sustainable-development framings (Soergel et al., 2021); and a stronger integration of risk, storyline, and high-impact-tail approaches (Simpson et al., 2021), complemented by the assessment of catastrophic and extreme storylines (Kemp et al., 2022). None of these changes the central near-term conclusion that scenario choice has a small effect on global mean temperature before about 2040, but all of them sharpen the assessment’s ability to represent the near-term

risks that matter most for decision-making.

5.3 Emissions pathways and near-term forcing

Global fossil carbon dioxide emissions have remained near their highest recorded levels through the early part of the present decade, so the near-term forcing trajectory is currently tracking the intermediate rather than the low pathways (*high confidence*). The annual global carbon accounting shows total anthropogenic carbon dioxide emissions plateauing at a high level, with fossil emissions still rising slowly and land-use-change emissions declining only modestly (Friedlingstein et al., 2023). This observed trajectory is the single most important constraint on near-term forcing, because it fixes the starting point from which all pathways depart. A pathway consistent with limiting warming to well below 2 degrees C requires emissions to peak and then decline steeply within the near-term window; the observed plateau means that the near-term is currently following a course closer to the intermediate pathways than to the steep-decline low pathways. The gap between the observed trajectory and the low pathways is not merely a matter of degree: because the low pathways require near-immediate declines, every year of continued high emissions makes the subsequent required decline steeper and narrows the remaining budget.

The remaining carbon budget — the cumulative future carbon dioxide emissions consistent with a given warming level at a given likelihood — is the most policy-relevant integral of the near-term emissions pathway. The carbon-budget framing follows from the near-linear relationship between cumulative carbon dioxide emissions and peak warming, and it makes the near-term decisive: because the budget is a stock, near-term emissions consume a disproportionate share of the space remaining for the most stringent targets. The annual accounting indicates that, at current emission rates, the budget consistent with the most stringent targets is being consumed within a small number of years to a decade or so (Friedlingstein et al., 2023). This gives the near term an outsized role: the pathways diverge slowly in temperature but rapidly in cumulative emissions, so the near-term choice among pathways largely determines whether the most stringent long-term targets remain physically attainable.

Near-term forcing is shaped strongly by non-carbon-dioxide agents, and the co-evolution of aerosols and greenhouse gases produces a near-term forcing trajectory that differs qualitatively from the long-term one (*high confidence*). Carbon dioxide dominates the cumulative forcing and hence the long-term warming, but on near-term horizons the forcing is materially influenced by methane, other short-lived greenhouse gases, and aerosols. Aerosols exert a net cooling that partially masks greenhouse warming; because aerosols are co-emitted with fossil combustion, pathways that decarbonize rapidly also reduce aerosol cooling rapidly. The near-term forcing trajectory of the low pathways therefore reflects a competition between falling greenhouse-gas forcing and falling aerosol masking. This is a defining feature of near-term projection and

explains why the low pathways can show little near-term temperature benefit — or even a transient relative warming — before their greenhouse-gas reductions dominate at longer horizons. The assessment treats this as a physical consequence of the different atmospheric lifetimes involved rather than as an argument against near-term mitigation, since the long-term benefit of the low pathways is unambiguous.

Methane occupies a distinctive place in the near-term forcing balance because its atmospheric lifetime is short enough that reductions translate into a falling forcing contribution within the near-term window, unlike carbon dioxide whose forcing persists. This makes near-term methane mitigation an unusually effective lever on the near-term forcing trajectory: a rapid reduction in methane emissions can bend the near-term forcing curve downward on a timescale over which carbon dioxide reductions have little temperature effect. The near-term forcing pathways therefore depend not only on the carbon dioxide trajectory but on the assumed trajectory of methane and other short-lived agents, and pathways that combine deep carbon dioxide reductions with rapid methane reductions achieve the most favorable near-term forcing outcome. The assessment reports the methane contribution separately from the carbon dioxide contribution precisely because their near-term behavior differs so markedly, and because near-term methane action is one of the few levers capable of producing a detectable near-term forcing benefit.

The partitioning of the global carbon budget between fossil and land-use sources also shapes the near-term forcing pathway and its uncertainty. The annual carbon accounting resolves total anthropogenic emissions into fossil combustion and land-use-change components, and it also quantifies the fraction absorbed by the ocean and land sinks (Friedlingstein et al., 2023). For near-term forcing, the behavior of these sinks is important because it determines how much of each year's emissions remains in the atmosphere to exert forcing. The sinks have historically absorbed roughly half of anthropogenic emissions, but their near-term efficiency is subject to variability and to the possibility of weakening under continued warming, which would raise the airborne fraction and hence the near-term forcing per unit of emissions. The assessment treats the near-term evolution of the sinks as a source of uncertainty in the mapping from emissions to forcing, distinct from the uncertainty in the emissions themselves, and it emphasizes that continued monitoring of the sinks is essential to constraining the near-term forcing pathway.

Land-use change and the terrestrial carbon cycle interact with the near-term pathway in ways that couple the emissions layer to the physical climate. Because the land sink responds to climate — to temperature, moisture, and the concentration of carbon dioxide itself — the near-term land-use and land-carbon trajectory is not independent of the climate it helps to determine. Disturbances such as fire, drought, and the potential destabilization of large carbon stores can convert parts of the land from sink to source, and while such conversions are not the central near-term expectation, they represent a near-term risk to the emissions balance that the assessment flags explicitly. The near-term forcing pathway is therefore best understood as the outcome of an interaction between

prescribed anthropogenic emissions and a partly responsive natural carbon cycle, rather than as a purely prescribed input, and this coupling is one reason the near-term forcing carries uncertainty even for a specified emissions pathway.

The composition of near-term emissions is determined by the sectoral structure of the energy and materials economy, and sector-resolved pathway analysis has become central to assessing what near-term forcing trajectories are feasible. The energy-transition modeling literature has matured to the point where the tools, their assumptions, and their limitations can themselves be reviewed and compared, clarifying which near-term pathways are represented robustly and which depend on contested assumptions (Chang et al., 2021). This methodological maturation matters for the physical-science assessment because it determines how much confidence can be placed in the feasibility of the steep near-term declines that the low pathways require. Where multiple independent modeling approaches agree that a near-term decline is achievable given deployed technology, the corresponding forcing pathway can be assessed as feasible; where the decline depends on assumptions that only some models make, the corresponding pathway is treated with more caution.

Sector-level analyses of the hardest-to-abate emitting sectors indicate that near-term emissions are strongly constrained by existing capital stock, so committed emissions from operating and under-construction infrastructure limit the feasible speed of near-term decarbonization (*high confidence*). Coal-fired power is the largest single source of energy-system carbon dioxide, and plant-by-plant analysis shows that the near-term coal emissions trajectory is governed by the age structure and utilization of the existing fleet rather than by the long-run availability of alternatives (Cui et al., 2021). A high-ambition near-term coal phaseout is physically possible but requires early retirement or curtailment of assets that would otherwise operate for decades, which is why coal is a decisive lever for near-term forcing. The existence of a large, young coal fleet in some regions is one of the principal reasons the highest-forcing pathways were originally constructed as coal-intensive; the same fleet analysis shows that a rapid phaseout, while demanding, is not precluded by the physical stock.

Iron and steel production is a second hard-to-abate sector whose near-term emissions are set by the installed base of blast-furnace capacity. Plant-level accounting of global steel emissions and of the routes to carbon neutrality shows that near-term abatement is dominated by decisions about existing integrated works and by the timing of investment in lower-emission production routes (Lei et al., 2023). Because steelmaking assets are long-lived and capital-intensive, the near-term steel emissions pathway is largely locked in by the current stock, and material near-term reductions require either curtailment of existing capacity or accelerated switching to lower-emission routes. The assessment treats such sectoral committed emissions as a physical constraint on the near-term forcing pathway: they set a floor below which near-term emissions cannot easily fall without early asset retirement.

Cement is a third sector with distinctive near-term dynamics, because a large fraction of its emissions is process-related — released by the chemistry of

calcination rather than by fuel combustion — and therefore not addressed by decarbonizing the energy supply alone. Projections of cement-sector emissions, particularly in developing economies where construction demand is growing, indicate that near-term cement emissions are likely to remain substantial and are among the more difficult to reduce quickly (Cheng et al., 2023). The process-emission fraction means that even a fully decarbonized energy supply would leave a residual cement emission that requires carbon capture or material substitution to address. For near-term forcing, cement is therefore an example of a sector whose emissions are relatively insensitive to the energy-transition levers that dominate the power sector, and whose near-term trajectory depends on construction demand and on the deployment of capture technology.

Building materials more broadly contribute a large and often under-counted share of near-term emissions, and life-cycle analysis of residential and commercial construction indicates that material-related emissions could remain a significant component of the global total through mid-century absent targeted mitigation (Zhong et al., 2021). The near-term relevance is that a substantial fraction of the emissions associated with buildings is embodied in their materials and is emitted at the time of construction, so near-term construction activity translates promptly into near-term emissions. This contrasts with the operational emissions of buildings, which are spread over their lifetimes and are more responsive to grid decarbonization. The assessment notes that the near-term emissions pathway is thus sensitive not only to how energy is produced but to how much material-intensive construction occurs, a dependence that the sectoral pathway literature makes explicit.

The near-term feasibility of the low pathways depends on the deployment rate of low-carbon energy and on the scaling of carbon-removal and capture technologies, both of which are advancing but from a small base relative to the required scale (*high confidence*). Solar photovoltaics has reached a cost and deployment level at which it is capable of supplying a large and growing share of electricity, and its trajectory is one of the principal reasons the highest-forcing pathways have become less plausible (Victoria et al., 2021). The near-term forcing benefit of photovoltaic deployment operates through displacement of fossil generation, so its effect on the emissions pathway depends on what it displaces and how quickly. Analyses of very high renewable-energy systems indicate that near-total renewable electricity supply is technically representable in energy-system models, though the near-term realization depends on grid integration, storage, and the retirement schedule of existing fossil capacity (Breyer et al., 2022). These deployment dynamics are the concrete, near-at-hand determinants of whether the near-term emissions pathway bends downward at the rate the low scenarios require.

Climate change itself modifies the near-term supply of renewable energy, introducing a feedback between the projected climate and the feasibility of the pathways designed to limit it. Assessment of climate-change impacts on renewable-energy supply indicates that changes in wind, solar resource, and hydrological conditions can alter generation potential, generally by modest amounts in the near term but with regional variation (Gernaat et al., 2021). For

near-term projection this feedback is second-order relative to deployment rates, but it is not negligible, and it illustrates that the emissions pathway and the climate response are not fully independent. The assessment treats this coupling as a refinement to the feasibility of the low pathways rather than as a first-order control on near-term forcing.

Carbon dioxide removal and capture technologies enter the near-term pathway assessment because most low pathways rely on some deployment of removal to offset residual emissions from hard-to-abate sectors. A survey of carbon capture, utilization, and storage technologies indicates that these are technically demonstrated but deployed far below the scale that the low pathways assume for later decades, so their near-term contribution to the emissions pathway is small (Dziejarski et al., 2023). Direct air capture, in particular, is assessed to be energetically and economically demanding at scale, and its near-term contribution to the emissions balance is minor relative to the reductions achievable through energy-system decarbonization (Sabatino et al., 2021). The near-term implication is that the emissions pathway over the coming two decades is determined overwhelmingly by how fast gross emissions fall, not by how fast removals scale; removals become quantitatively important to the balance only later. The assessment therefore cautions against near-term pathways that rely on early large-scale removal, since the deployment evidence does not currently support it.

Hydrogen and related energy carriers are a further element of the low pathways whose near-term contribution is assessed to be limited. A perspective on the hydrogen economy as a route to net-zero emissions indicates that low-carbon hydrogen can address specific hard-to-abate uses but that its near-term deployment is constrained by production cost, infrastructure, and the emissions intensity of current hydrogen supply (Spek et al., 2022). For the near-term emissions pathway, hydrogen is therefore best regarded as an enabling technology whose principal effect lies beyond the near-term window, with only a modest near-term contribution to the forcing trajectory.

The economic valuation of near-term emissions has been revised upward, reinforcing that the near-term choice among pathways carries large and quantifiable consequences even though its temperature signal is delayed (*high confidence*). Comprehensive re-estimation of the social cost of carbon dioxide, incorporating updated damage functions and discounting, indicates a substantially higher marginal cost of emissions than earlier assessments used (Rennert et al., 2022). While the social cost of carbon is an economic rather than a physical-science quantity, its revision is relevant to the interpretation of near-term pathways because it changes the weight attached to near-term emissions: a higher marginal cost implies that the near-term divergence in cumulative emissions between pathways — which is large even where the temperature divergence is small — carries correspondingly larger consequences. The physical-science assessment reports the emissions and temperature trajectories; the economic literature attaches value to them, and the direction of recent revision is toward greater weight on near-term action.

In summary, the near-term forcing trajectory is currently tracking the in-

intermediate pathways, the remaining budget for the most stringent targets is being consumed within the near-term window, and the feasibility of bending the pathway downward depends on concrete near-term deployment and retirement decisions in a small number of high-emitting sectors, as the sectoral pathway analyses make clear (Friedlingstein et al., 2023). The near-term is the horizon over which the emissions pathway is most malleable and least reflected in temperature, which is the defining tension of near-term climate assessment.

5.4 Near-term projections to about 2040

Global surface temperature is projected to continue rising through the near-term period under all assessed pathways, with the multi-pathway spread in the global mean remaining small until close to 2040 (*very likely*). Because the near-term climate response is dominated by forcing already committed and by the fast component of the system's adjustment, the projected global mean temperature to 2040 is largely insensitive to which pathway is followed. The multi-model experiments coordinated under the Scenario Model Intercomparison Project produce, over this window, an envelope whose width is set principally by the range of transient climate responses across models and by the phasing of internal variability, not by scenario divergence (Tebaldi et al., 2021). This is the single most robust conclusion about near-term global temperature: it will keep rising, and the amount of rise before 2040 is bounded within a relatively narrow range that mitigation choices made now cannot substantially alter, precisely because those choices act through cumulative emissions whose temperature effect is delayed.

The narrowness of the near-term global-mean spread should not be mistaken for certainty about the near-term climate as experienced. The global mean is an average over a system whose regional and seasonal responses are far more variable, and whose extremes change faster than the mean. The assessment therefore distinguishes sharply between the near-term global-mean trajectory — well constrained and pathway-insensitive — and the near-term regional and extreme-event trajectory, which is more uncertain and, for the quantities that matter most to impacts, changing faster than the global mean would suggest.

The probability of individual near-term years exceeding 1.5 degrees C above pre-industrial levels has become substantial, but a single exceeding year does not constitute crossing the long-term warming level, which is defined as a multi-decadal average (*high confidence*). As the forced warming approaches 1.5 degrees C, the combination of the trend and interannual variability makes it increasingly likely that individual years will exceed that level even before the long-term average reaches it. The assessment maintains a strict distinction between a threshold-exceeding year — a fluctuation of the annual mean above a level — and the crossing of the corresponding long-term warming level, which is defined over a multi-decadal window and is the quantity to which the carbon-budget and pathway analyses refer. Conflating the two would

misrepresent the near-term: the first threshold-exceeding years are expected before the long-term level is reached, and their occurrence is a signature of the approaching level rather than evidence that it has been crossed. This distinction is essential for interpreting near-term observations against pathway projections, since the observed record will show exceeding years while the assessed long-term level is still below the threshold.

Near-term warming is strongly amplified in the Arctic, where the forced signal is large enough to dominate internal variability even at near-term horizons, so Arctic near-term change is among the most confidently projected regional signals (*high confidence*). The Arctic has warmed at a rate several times the global average over recent decades, a robustly observed and physically understood amplification (Rantanen et al., 2022). Because the amplified signal is large, the near-term Arctic trend emerges clearly from variability, and near-term projections of continued rapid Arctic warming are among the most confident regional statements the assessment can make. The mechanisms — sea-ice loss and the associated surface-albedo and heat-flux feedbacks, changes in atmospheric and oceanic heat transport, and cloud and water-vapor feedbacks — are sufficiently well established that the near-term direction and approximate magnitude of Arctic warming are robust across models and pathways. This makes the Arctic a distinctive near-term case: a region where the forced signal is strong enough that the near-term projection is dominated by the trend rather than by the noise.

Arctic near-term change extends beyond temperature to the hydrological cycle. Improved-model analyses indicate that Arctic precipitation is projected to increase faster and by more than earlier generations of models suggested, with a transition toward a rain-dominated regime occurring earlier than previously projected in parts of the Arctic (McCrystall et al., 2021). For near-term projection, this means that the Arctic hydrological response, like its temperature response, is a leading indicator: the near-term signal is large enough to be assessed with relatively high confidence, in contrast to mid-latitude precipitation, where the near-term signal is small relative to variability. The rapidity of the projected Arctic precipitation change also has implications for near-term feedbacks, since changes in snow and rain partitioning affect surface albedo and the surface energy balance.

Near-term changes in temperature extremes, particularly heatwaves, are projected to outpace changes in the global mean, so the near-term rise in the intensity, frequency, and duration of hot extremes is larger in relative terms than the rise in mean temperature (*high confidence*). Heatwaves are among the extremes whose change is most directly and robustly tied to warming, and the physical understanding of how they respond to a warming climate has advanced to the point where their near-term intensification can be assessed with confidence (Barriopedro et al., 2023). The relationship between mean warming and the tail of the temperature distribution is such that a modest shift in the mean produces a disproportionate increase in the probability of extreme heat, and additional processes — land-atmosphere feedbacks, soil-moisture depletion, and circulation persistence — can amplify the change in the tail beyond what a

simple mean shift would imply. For the near term this means that heat-related hazard is projected to increase noticeably even over the window in which the global mean rises only modestly, which is a central reason the near term is not a period of negligible change from an impacts perspective.

Compound events — the co-occurrence of multiple hazards or of a hazard with an aggravating condition — are a particular focus of near-term extremes assessment because their impacts can exceed the sum of the individual hazards. Concurrent drought and heat are a leading example: soil-moisture deficits and high temperatures reinforce one another through land-atmosphere coupling, and the joint probability of severe compound drought-heat events is projected to rise faster than the probability of either hazard alone (Tripathy et al., 2023). Crucially, the high-end risk of such compound events is projected to accelerate, so the upper tail of the compound-event distribution grows faster than the central estimate. For near-term projection this is significant because it means that the most severe compound events — those with the largest impacts — are projected to become more likely at a rate that outpaces the change in mean conditions. The assessment therefore treats compound events as a domain in which near-term change is both more uncertain and potentially more consequential than the global-mean trajectory would indicate.

Drought more generally is a near-term hazard whose projection is complicated by the interplay of physical and human drivers. The concept of anthropogenic drought recognizes that drought hazard in the near term reflects not only precipitation and temperature change but also water demand, land use, and management, so the realized near-term drought risk depends on human as well as physical factors (AghaKouchak et al., 2021). This framing is important for near-term projection because it clarifies that the near-term drought signal cannot be read from the physical climate projection alone; it emerges from the interaction of a modestly changing physical hazard with rapidly changing human water systems. The physical-science assessment supplies the hazard trend — generally toward increased evaporative demand and, in many regions, toward more frequent or intense dry conditions — while recognizing that the near-term realized risk is co-determined by human factors.

The economic dimension of near-term water stress reinforces the salience of the drought projections. Analysis of the economic impact of water scarcity indicates that even modest near-term changes in water availability can have substantial and unevenly distributed economic consequences (Dolan et al., 2021). While this is an impacts-and-economics result rather than a physical projection, it bears on the interpretation of the near-term hazard trajectory: the near-term change in the physical drivers of water scarcity, though modest in the global mean, intersects with economic systems in ways that amplify its consequences. The assessment cites this to underline that the near-term physical signal, small as it is in aggregate, maps onto disproportionate societal effects through the systems it impinges upon.

Wildfire hazard is projected to expand in both space and time under near-term warming, with the fire-weather season lengthening and fire-prone conditions extending into regions and seasons previously less affected (*high*

confidence). The spatial and temporal expansion of wildland-fire activity in response to climate change has been assessed to continue under near-term warming, driven by the compounding of higher temperatures, altered precipitation, and vegetation changes (Senande-Rivera et al., 2022). Fire weather integrates several near-term climate changes — temperature, humidity, wind, and antecedent moisture — so its projected expansion is a useful summary indicator of near-term compound change. Because fire also interacts with the carbon cycle, near-term increases in fire activity carry a feedback dimension, releasing stored carbon and potentially altering the emissions pathway at the margin. The assessment treats near-term fire-weather expansion as a robust regional-to-continental signal, while noting that the realized fire activity depends additionally on ignition, fuel management, and land use.

Near-term regional projections outside the strongly amplified regions are dominated by internal variability, so for many mid-latitude land regions the near-term forced signal is small relative to the range of possible outcomes (high confidence). In contrast to the Arctic, where the forced signal dominates, many populated mid-latitude land regions have a near-term forced temperature and, especially, precipitation signal that is comparable to or smaller than the amplitude of internal variability over the same window. This means that near-term regional projections for such regions must be expressed as distributions whose width reflects variability as much as forced change, and that a single realization — including the observed evolution — can depart substantially from the forced trend for a decade or more. The assessment addresses this by emphasizing initial-condition large ensembles, which sample the range of internal-variability trajectories consistent with a given forcing, and by reporting near-term regional change as a probability distribution rather than a single trajectory. The practical consequence is that near-term regional planning must contend with a genuinely wide range of physically possible outcomes, and that observed near-term regional trends should not be over-interpreted as evidence for or against the forced projection.

The individual Earth-system models that populate the near-term projection ensemble differ in their representation of the processes that govern regional response, and understanding these differences is part of interpreting the near-term spread. Documentation of individual model configurations — for example the detailed description of a participating Earth-system model's simulations under the shared experimental protocol (Lovato et al., 2022) — allows the assessment to attribute parts of the near-term spread to specific structural choices in ocean, land, carbon-cycle, and aerosol representation. This traceability is valuable because it distinguishes near-term spread that reflects genuine uncertainty in the climate response from spread that reflects known biases in particular models. The assessment uses such model documentation to weight and interpret the ensemble rather than treating all members as equally reliable for every quantity.

Sea-level and cryospheric changes have near-term trajectories that are highly committed and relatively pathway-insensitive over the near term, so near-term sea-level rise is projected with high confidence to continue at

rates comparable to or exceeding the recent past (*high confidence*). The slow-response components of the climate system — the ocean’s thermal expansion and the large ice masses — respond to accumulated rather than instantaneous forcing, so their near-term change is largely set by past emissions and is insensitive to near-term pathway choice. This makes near-term sea-level rise one of the most confidently projected quantities: it will continue, at a rate that near-term mitigation cannot materially slow, because the near-term rise reflects heat and mass changes already in train. Pathway choice affects sea level strongly at long horizons but weakly in the near term, a temporal structure that mirrors the temperature case but is even more pronounced because of the longer response times involved. The assessment emphasizes that the near-term commitment in sea level is a reason to treat adaptation to continued near-term rise as unavoidable regardless of mitigation ambition, even as mitigation remains decisive for the long-term trajectory.

Near-term projections of the large-scale atmospheric and oceanic circulation are more uncertain than projections of thermodynamic quantities, so circulation-dependent regional changes carry lower near-term confidence than temperature and its direct consequences (*medium confidence*). Many of the most societally important regional changes — shifts in storm tracks, monsoon timing and intensity, and the position of the mid-latitude jets — depend on the response of the large-scale circulation, which is governed by a competition between different warming patterns and is less robustly simulated than the thermodynamic response. In the near term, the forced circulation signal is generally small relative to internal variability, and models disagree more about circulation than about temperature. As a consequence, near-term projections of circulation-dependent regional precipitation are typically expressed as wide distributions, and the assessment attaches lower confidence to them than to the near-term thermodynamic signals such as the intensification of hot extremes (Barriopedro et al., 2023). This asymmetry — high confidence in thermodynamically controlled near-term change, lower confidence in dynamically controlled near-term change — is a general organizing principle for interpreting the near-term regional projections, and it explains why quantities such as heat extremes are projected with more confidence than quantities such as regional storminess.

The near-term behavior of the large-scale ocean circulation is a particular focus because of its role in regional climate and in the slow variability that underpins decadal predictability. The near-term evolution of the major overturning circulations affects regional ocean heat content, sea level, and the heat exchange with the atmosphere, and it is one of the slowly evolving components that initialized prediction attempts to forecast. Because these circulations respond to accumulated forcing and to internal variability on decadal timescales, their near-term change is partly committed and partly predictable, and the quality of the near-term forecast depends on how well the initial ocean state is observed and represented (Jean-Michel et al., 2021). The assessment treats the near-term ocean-circulation projection as a domain in which the slow, partly predictable component coexists with genuine structural uncertainty in the models, so that confidence is higher for the direction of change than for its precise near-term

magnitude.

Near-term regional projection is increasingly supported by the systematic use of single-model initial-condition large ensembles, which separate the forced response from internal variability and thereby make the near-term uncertainty attributable to its physical sources (*high confidence*). A single climate model run many times from slightly different initial states produces an ensemble whose spread, for a fixed forcing, is a direct estimate of the internal-variability contribution to the near-term projection. Comparing such single-model ensembles across models then separates the internal-variability spread from the structural, model-to-model spread. This methodology has become central to near-term projection because it allows the assessment to state, for a given region and quantity, how much of the near-term uncertainty arises from internal variability — which is irreducible on near-term horizons but bounded — and how much arises from genuine disagreement about the forced response. The documentation of individual model configurations supports this decomposition by making the forced-response characteristics of each model traceable (Lovato et al., 2022). The practical result is that near-term regional projections can be accompanied by an explicit account of why their range is what it is, rather than presenting an undifferentiated spread.

The near-term projection of extremes benefits especially from the large-ensemble methodology, because extremes are, by definition, tail events whose statistics require many realizations to estimate reliably. For hot extremes and for compound drought-heat events, the large-ensemble approach provides the sample size needed to characterize the changing tail of the distribution and to distinguish a genuine change in extreme statistics from the sampling noise of a small ensemble (Tripathy et al., 2023). This is why the assessment's confident statements about near-term intensification of extremes rest on large-ensemble evidence: the change in the tail is estimated from enough realizations to separate signal from noise. For circulation-dependent extremes, by contrast, even large ensembles leave substantial near-term uncertainty because the underlying dynamical response is itself uncertain, reinforcing the thermodynamic-versus-dynamic confidence asymmetry described above.

In aggregate, the near-term projection to 2040 is characterized by a well-constrained, pathway-insensitive global-mean warming; strongly amplified and confidently projected change in the Arctic and in the slow-response cryospheric and sea-level components; extremes — especially heat and compound drought-heat — changing faster than the mean; and a wide, variability-dominated range for many mid-latitude regional quantities (Tebaldi et al., 2021). The near-term is thus not a period of small change uniformly, but one in which certain quantities change confidently and consequentially while others remain dominated by the noise of internal variability — the subject to which the assessment turns next.

5.5 Internal variability, predictability, and decadal prediction

Internal variability is the dominant source of uncertainty in near-term projections of regional climate and of many extremes, and it can locally reinforce or offset the forced trend for a decade or more (*high confidence*). The climate system generates variability internally, without any change in external forcing, through the coupled dynamics of the ocean and atmosphere and through slower modes of the ocean and cryosphere. On near-term horizons this internally generated variability is comparable in amplitude to the forced change for many quantities, so the realized near-term climate is the sum of a forced trend and a variability trajectory whose phase is, in an uninitialized projection, unknown. This is why the near-term multi-model spread is dominated by variability and model response rather than by scenario: the scenarios have barely diverged in forcing, while the variability spread is already large. The assessment addresses this structurally by using large ensembles of simulations that differ only in their initial conditions, which sample the distribution of variability trajectories consistent with a given forcing and thereby separate the forced signal from the variability spread.

The separation of forced signal from internal variability is not merely a technical convenience; it is essential to the correct interpretation of near-term observations. Because a single realization of the climate — including the one we observe — can depart from the forced trend for a decade or more, an observed near-term trend that is faster or slower than the forced projection is not, by itself, evidence that the projection is wrong. Only when the observed departure exceeds what the variability spread can plausibly produce does it constitute a genuine tension with the projection. The assessment therefore evaluates near-term consistency between observations and projections against the full initial-condition spread, not against the ensemble mean alone, and it cautions against interpreting individual near-term episodes — accelerations, slowdowns, or pauses in the observed trend — as changes in the forced response.

The largest modes of internal variability, particularly those of the ocean, provide the principal source of near-term predictability beyond that supplied by the forced trend, because their slow evolution can in principle be forecast once the ocean state is known (*medium confidence*). Where internal variability is generated by slowly evolving components — most importantly the ocean — it carries memory, and that memory is the basis for near-term predictability that goes beyond simply extrapolating the forced trend. If the current state of the slowly evolving ocean can be observed and initialized in a model, the subsequent evolution of the associated variability can be predicted for some period into the future, adding skill over an uninitialized projection that knows only the forcing. This is the physical rationale for initialized near-term, or decadal, prediction: it attempts to forecast not only the forced response but also the phase of the slow internal variability. The degree to which this is achievable in practice is the central question of the decadal-prediction enterprise, and the answer differs

strongly by quantity and region.

Initialized decadal prediction demonstrably improves skill over uninitialized projections for some large-scale quantities in the first few years, most robustly for upper-ocean heat content and some large-scale surface-temperature indices, but skill for regional precipitation and for extremes remains limited (*medium confidence*). The clearest gains from initialization are found where the predictable signal resides in the slowly evolving ocean: upper-ocean heat content and basin-scale surface-temperature averages retain skill for several years because they are governed by ocean states that persist and evolve predictably. For these quantities, initialized prediction adds value over an uninitialized projection, and the near-term forecast can be sharpened relative to the pathway-based projection. For quantities dominated by fast atmospheric variability — regional precipitation, individual extreme events — the predictability from initialization is much weaker, because the atmospheric memory is short and the predictable oceanic influence on these quantities is modest and regionally variable. The assessment therefore characterizes decadal prediction as presently most useful for large-scale, ocean-influenced quantities and as still developing for the regional and extreme quantities that are most directly relevant to impacts.

The near-term predictability of the ocean state itself depends on the quality of the observed initial state, which has improved with the maturation of ocean observing and reanalysis systems. High-resolution ocean and sea-ice reanalyses that assimilate the available observations provide the initial conditions from which decadal predictions are launched, and their fidelity sets an upper bound on the achievable near-term prediction skill (Jean-Michel et al., 2021). Improvements in these reanalysis products — in resolution, in the treatment of the sea-ice component, and in the assimilation of subsurface observations — propagate directly into improved initialization and hence into the potential skill of near-term prediction. The assessment notes that the observing system underpinning these reanalyses is therefore a critical, and in places fragile, foundation for near-term prediction, and that gaps in subsurface and polar observations translate into gaps in near-term predictive skill.

The predictable component of near-term variability interacts with the forced trend, so that the near-term realized trajectory can be understood as a forced signal modulated by a partly predictable variability phase, and skillful near-term prediction requires representing both correctly (*medium confidence*). Because the near-term climate is the sum of a forced trend and an internal-variability trajectory, and because part of the variability trajectory is in principle predictable from the initialized ocean state, the most skillful near-term forecasts combine a well-characterized forced response with a well-initialized variability phase. Errors in either component degrade the forecast: a model with a biased forced response will drift away from the observations even if its variability phase is correct, while a model with a poor initial state will misrepresent the variability phase even if its forced response is accurate. This is why the assessment treats the fidelity of the forced response — the subject of the model-evaluation and projection chapters — and the fidelity of initialization as

jointly necessary for near-term prediction skill. The documentation of individual model configurations and their biases (Lovato et al., 2022) supports this joint evaluation by making the forced-response characteristics of each prediction system traceable.

The relative importance of the forced signal and internal variability shifts across the near-term window and across regions, and this shifting balance structures where and when near-term prediction can add value. Early in the near-term window and in regions of strong variability, the variability component is large and, where it is predictable, initialization adds the most value. Later in the window and in regions of strong forced response — the Arctic being the clearest example (Rantanen et al., 2022) — the forced signal grows to dominate, and the marginal value of initialization declines because the uninitialized projection already captures most of the signal. The assessment uses this structure to guide the interpretation of near-term prediction: initialization is most valuable in the first several years and in variability-dominated regions, while for strongly forced quantities the uninitialized projection is already close to the best available near-term estimate.

The predictable modes of internal variability that give near-term prediction its skill reside chiefly in the ocean, and their near-term forecasting depends on both the observed initial state and the model’s ability to propagate that state forward. Slow ocean modes carry memory over years, and where a model can be initialized with an accurate ocean state and can propagate it with fidelity, the associated near-term variability can be predicted with useful skill. The reanalysis systems that supply the initial state have improved in resolution and in their treatment of the ocean and sea-ice components (Jean-Michel et al., 2021), which has improved the potential skill of near-term prediction, but limitations remain where the observing system is sparse — particularly in the deep ocean and in polar regions. The assessment therefore treats near-term predictive skill as jointly limited by the observing system, by the fidelity of the initialization, and by the model’s representation of the slow modes, and it identifies the observing system as the component most amenable to near-term improvement through sustained investment.

A recurring interpretive challenge is that the skill of near-term prediction must be assessed against an appropriate baseline, since a forecast that merely reproduces the forced trend is not demonstrating skill from initialization. The assessment distinguishes the skill that arises from knowing the forcing — which an uninitialized projection also possesses — from the additional skill that arises from initializing the internal-variability state. Only the latter is unique to decadal prediction, and it is this incremental skill, rather than total skill, that measures the value added by initialization. Evaluated on this basis, initialized prediction adds clear incremental skill for large-scale, ocean-influenced quantities in the first several years and little incremental skill for regional precipitation and extremes, a pattern that is stable across prediction systems and that defines the current frontier of the near-term prediction enterprise.

Near-term prediction of extremes and of high-impact events remains a frontier in which the forced change in the hazard distribution is better con-

strained than the year-to-year predictability of individual events (*medium confidence*). For extremes such as heatwaves and compound drought-heat events, the assessment can state with confidence that the underlying distribution is shifting toward more severe values under near-term warming (Barriopedro et al., 2023; Tripathy et al., 2023), but predicting the occurrence of a specific extreme in a specific near-term year is a much harder problem that depends on the predictability of the fast variability that triggers the event. The current state of the science supports skillful projection of the changing statistics of extremes but only limited skill in predicting individual extreme years in advance. This asymmetry is important for near-term decision-making: the assessment can inform planning for a changed distribution of extremes with more confidence than it can provide year-ahead forecasts of specific events. The wildfire, drought, and compound-event projections discussed above should therefore be read as statements about the evolving hazard distribution rather than as predictions of specific near-term events (Senande-Rivera et al., 2022).

The low-likelihood, high-impact tail of the near-term distribution — including the possibility of engaging tipping elements or of cascading Earth-system responses — is not excluded on near-term horizons and is best assessed through storyline and risk-based methods that complement probabilistic prediction (*medium confidence*). Some components of the Earth system exhibit the potential for abrupt or self-reinforcing change once thresholds are crossed, and the interaction among such tipping elements can raise the risk of cascading responses (Wunderling et al., 2021). While the central near-term projection does not have most such elements crossing their thresholds within the near-term window, the possibility cannot be excluded, and the near-term is the horizon over which early indicators of approach to a threshold would begin to be detectable. Assessment of specific large systems — for example the potential for critical transitions in the Amazon forest under the combined pressure of climate change and land use (Flores et al., 2024) — indicates that the near-term evolution of such systems warrants monitoring even where a near-term transition is not the central expectation. The assessment treats these outcomes through storyline approaches that trace physically self-consistent high-impact trajectories, and through the risk-framing methods discussed in Section 5.1 (Simpson et al., 2021), complemented by the assessment of catastrophic and extreme storylines (Kemp et al., 2022), rather than attempting to attach precise probabilities to inherently poorly constrained tail events.

The entanglement of near-term climate risk with concurrent global stresses further complicates the assessment of the tail. Because near-term climate hazards act on systems that are simultaneously stressed by other global pressures, the societal consequences of a near-term hazard can be amplified through interactions that the physical projection alone does not capture (Lawrence et al., 2024), and that the coupled climate-biodiversity assessment also emphasizes (Pörtner et al., 2023). This coupling is most consequential in the tail, where a severe near-term hazard coinciding with stress in food, water, or economic systems could produce impacts disproportionate to the physical signal. The assessment notes this coupling as a reason to weight the near-term tail more

heavily than its probability alone would suggest, while keeping the physical-science projection and the impacts assessment methodologically distinct.

In synthesis, near-term projection and prediction form a spectrum: pathway-based projection supplies the well-constrained forced trend, initialized prediction adds partly recoverable skill for large-scale ocean-influenced quantities in the first several years, and the low-likelihood tail requires dedicated risk and storyline assessment (*high confidence*). The near-term is neither fully predictable nor merely a scaled-down version of the long term. Its forced global-mean signal is robust and pathway-insensitive; its regional and extreme behavior is dominated by internal variability, part of which is predictable from the ocean state and part of which is not; and its tail contains outcomes whose potential severity justifies explicit assessment even where their probability is low. The methodological developments since the Sixth Assessment — traceable emissions-to-temperature chains (Kikstra et al., 2022), initial-condition large ensembles, maturing decadal-prediction systems built on improved ocean re-analyses (Jean-Michel et al., 2021), and integrated risk frameworks (Simpson et al., 2021) — collectively improve the assessment’s ability to characterize each part of this spectrum. The overarching near-term message is unchanged and robust: warming will continue across all pathways through the coming two decades, the choices that determine the longer-term trajectory are being made now even though their temperature signal is delayed, and the risks that matter most in the near term are concentrated in extremes and in the tail rather than in the slowly diverging global mean.

5.6 Cross-cutting considerations for near-term projection

The feasibility of the near-term emissions pathways that low scenarios require is increasingly assessed with probabilistic rather than deterministic methods, which better represent the genuine uncertainty in how fast the energy system can transform (*high confidence*). Deterministic single-trajectory descriptions of an energy transition understate the uncertainty in deployment rates, retirement schedules, and demand growth. Probabilistic assessment of a large economy’s transition toward carbon neutrality demonstrates that the near-term emissions trajectory is best represented as a distribution whose width reflects uncertainty in technology deployment and policy implementation rather than as a single line (Zhang et al., 2022). For near-term projection this matters because it clarifies that the near-term emissions layer — the input to the physical projection — is itself uncertain, and that the uncertainty in near-term forcing therefore has a socioeconomic as well as a physical component. The assessment treats the near-term forcing as conditional on this distribution of feasible emissions trajectories, which is one reason the near-term projection is expressed as a range conditioned on scenario rather than as a set of deterministic predictions. The probabilistic

framing also makes explicit that the lower tail of the emissions distribution — the most rapid feasible transitions — remains physically available even where central projections track the intermediate pathways.

The maturation of energy-transition modeling has a direct bearing on how the physical-science assessment interprets the scenario inputs. Because the tools, assumptions, and limitations of transition models can now be systematically reviewed (Chang et al., 2021), the assessment can distinguish near-term emissions trajectories that are robust across modeling approaches from those that depend on contested assumptions about behavior, cost, or policy. This traceability from socioeconomic assumption to emissions input to physical forcing is the near-term counterpart of the emissions-to-temperature traceability established for the temperature response (Kikstra et al., 2022). Together these two traceable chains — assumptions to emissions, and emissions to temperature — allow the near-term projection to be decomposed end to end, so that a given near-term temperature range can be attributed to identifiable choices and uncertainties at each stage rather than emerging from an opaque pipeline.

Observations continue to narrow the near-term projection by constraining the recent forced trend and the transient response, so the observational record is not merely a validation target but an active constraint on the near-term forecast (*high confidence*). The near-term projection depends on the transient climate response — the warming realized per unit of forcing on decadal timescales — and this quantity is increasingly constrained by the observed warming of recent decades combined with estimates of the forcing that produced it. As the observed record lengthens and the forcing history is better characterized, the plausible range of the transient response narrows, and with it the near-term temperature spread. The near-term is the horizon on which this observational constraint is strongest, because it is close to the observed period and therefore least affected by the accumulating uncertainties that dominate long-term projection. The assessment uses the observational constraint on the transient response as one of the principal means of narrowing the near-term envelope beyond the raw multi-model spread, and it treats the ongoing observing system as essential infrastructure for maintaining and improving this constraint (Jean-Michel et al., 2021).

The value of observations for near-term projection extends to the individual processes that govern regional response. Where a regional feedback is both strongly observed and well understood — as in the case of Arctic amplification (Rantanen et al., 2022) — the observational constraint on the near-term regional projection is strong, and the projection is correspondingly confident. Where a regional process is poorly observed or weakly constrained, the near-term regional projection retains a wide range. This heterogeneity in observational constraint is one reason near-term confidence varies so strongly by region and quantity: it is not that the models disagree uniformly, but that some near-term signals are anchored by strong observational constraints while others are not. The assessment makes this heterogeneity explicit by attaching calibrated confidence to each near-term regional statement rather than reporting a single global level of confidence.

Near-term sea-level rise is among the most policy-relevant committed changes, and its near-term trajectory is robust across pathways because it is dominated by the slow-response components of the system (*high confidence*).

The near-term rise in global mean sea level reflects ocean thermal expansion and mass loss from glaciers and ice sheets that respond to accumulated forcing over long timescales. Because these components integrate past forcing, the near-term sea-level trajectory is largely insensitive to near-term pathway choice, and continued rise at rates comparable to or exceeding the recent past is projected with high confidence regardless of mitigation ambition. This committed character makes near-term sea-level rise a clear example of a change to which adaptation is unavoidable, even as the long-term trajectory remains strongly dependent on the emissions pathway. The near-term regional sea-level pattern, which differs from the global mean because of ocean dynamics, gravitational effects, and vertical land motion, adds spatial structure that is important for local planning but does not alter the robust global near-term conclusion of continued rise.

The near-term evolution of the cryosphere reinforces the sea-level picture and interacts with it. Continued near-term loss of land ice contributes directly to sea level, while the loss of sea ice — particularly in the Arctic, where the forced signal is large (Rantanen et al., 2022) — contributes to regional amplification and to changes in the surface energy balance rather than to sea level directly. The distinction between land-ice and sea-ice contributions is important for near-term interpretation: sea-ice loss is a leading, confidently projected near-term signal with strong feedback consequences, whereas the near-term sea-level contribution from the large ice sheets is more uncertain in magnitude even though its direction is robust. The assessment reports these components separately so that the confident near-term signals are not diluted by the more uncertain ones.

Near-term changes in the hydrological cycle are spatially heterogeneous, with confident intensification in some regions and variability-dominated ranges in others, so the near-term hydrological projection resists a single global summary (*high confidence*). The global hydrological cycle intensifies with warming in a broad thermodynamic sense, but the near-term regional expression of this intensification is highly variable. In the strongly forced high latitudes, near-term precipitation increases are projected with relatively high confidence, and the transition toward rain-dominated regimes in parts of the Arctic is a leading near-term signal (McCrystall et al., 2021). In many populated mid-latitude and subtropical regions, by contrast, the near-term precipitation signal is small relative to internal variability, and the projection must be expressed as a wide distribution. Compound hydrological hazards — concurrent drought and heat — are a domain in which the near-term change in the joint distribution is more confident and more consequential than the change in either variable alone, with the upper tail of severe events projected to grow especially fast (Tripathy et al., 2023). The assessment therefore treats near-term hydrological change not as a single trajectory but as a mosaic of confidently projected signals in some regions and variability-dominated ranges in others,

with compound events warranting particular attention because their impacts are disproportionate.

The realized near-term hydrological risk depends additionally on human water systems, which is why the anthropogenic-drought framing is essential for translating the physical projection into risk (AghaKouchak et al., 2021). A given near-term change in evaporative demand or precipitation interacts with water infrastructure, allocation, and demand to produce a realized drought risk that can differ substantially from the physical signal alone. The economic consequences of this realized risk can be large even for modest physical changes, and they are unevenly distributed across regions and sectors (Dolan et al., 2021). The assessment cites these results to underline that the near-term hydrological hazard, though spatially heterogeneous and often modest in the physical mean, maps onto societal consequences through the human systems it impinges upon, and that the near-term risk therefore cannot be read from the physical projection in isolation.

The near-term is the horizon over which early indicators of a high-sensitivity or tipping-prone Earth system would first become detectable, which gives near-term monitoring a distinctive role in risk management (*medium confidence*). If the true transient climate response lies at the high end of the assessed range, or if strongly nonlinear feedbacks are closer to engagement than central estimates suggest, the first detectable signatures would appear within the near-term window as departures of the observed trend or of specific process indicators from the central projection. Interacting tipping elements raise the possibility that such departures, once begun, could reinforce one another (Wunderling et al., 2021), and the near-term monitoring of specific vulnerable systems — such as the Amazon forest, where critical transitions are physically possible under combined climatic and land-use pressure (Flores et al., 2024) — is therefore a component of near-term risk assessment even where a near-term transition is not the central expectation. The assessment frames this not as a prediction but as a rationale for sustained observation: the near-term is when the system would first reveal whether it is tracking the central projection or a higher-impact trajectory, and the value of detecting that early is high because the response options narrow as the system evolves.

The coupling of near-term climate hazards to concurrent global stresses further elevates the importance of near-term monitoring and risk assessment. Because near-term hazards act on systems already stressed by other pressures, their societal consequences can be amplified through interactions that the physical projection does not capture (Lawrence et al., 2024). The coupled character of the climate and biodiversity crises is a specific example in which near-term climate change and near-term ecological change reinforce one another with societal consequences that exceed the sum of the parts (Pörtner et al., 2023). The assessment treats these couplings as reasons to weight the near-term tail more heavily than its physical probability alone would suggest, while maintaining a clear methodological separation between the physical-science projection, which is Working Group I's remit, and the impacts and risk assessment that builds upon it.

The economic weight attached to near-term emissions has risen, reinforcing that the near-term divergence in cumulative emissions between pathways carries large consequences even where the near-term temperature divergence is small (*high confidence*). The re-estimation of the marginal cost of carbon dioxide emissions toward substantially higher values (Rennert et al., 2022) changes the interpretation of the near-term pathway divergence. Because the pathways diverge rapidly in cumulative emissions even while diverging slowly in temperature, and because each tonne of near-term emissions now carries a higher assessed marginal cost, the economic stakes of the near-term choice among pathways are large and front-loaded. This economic framing complements the physical finding that the near-term is where the pathway is most malleable: the physical science shows that near-term action determines the long-term trajectory with little near-term temperature feedback, and the economic literature shows that the value of that near-term action is high. The assessment reports both, keeping the physical projection and the economic valuation methodologically distinct while noting their consistent implication that near-term action is disproportionately consequential.

The near-term also frames the interpretation of individual threshold-exceeding years against the long-term warming levels to which policy targets refer, and clear communication of this distinction is a cross-cutting responsibility of the assessment (*high confidence*). As the forced warming approaches widely referenced levels, the combination of the trend and interannual variability makes individual exceeding years increasingly likely before the corresponding long-term level is reached. The near-term will therefore contain years whose global mean exceeds a target level while the assessed long-term warming remains below it, and misinterpreting such years as the crossing of the level would misrepresent the state of the climate system. The assessment addresses this by consistently defining warming levels as multi-decadal averages and by explaining that the appearance of exceeding years is an expected signature of approaching a level rather than evidence of having crossed it. This communication task is cross-cutting because it bears on how the near-term observational record — which will show exceeding years — is read against the pathway projections and the carbon-budget analyses, all of which refer to long-term averages. Getting this distinction right is essential to avoid both premature declarations that a level has been crossed and complacent dismissal of exceeding years as meaningless noise; the correct reading is that they are early, expected indicators that the long-term level is being approached. The near-term is thus not only a period of committed physical change but a period in which the interpretation of the emerging record requires particular care, so that the physical-science assessment supports rather than confuses the public understanding of where the climate system stands relative to the levels that anchor mitigation goals.

Taken as a whole, near-term projection integrates a well-constrained physical trajectory, a socioeconomically conditioned emissions input, a strong but heterogeneous observational constraint, and an explicit treatment of the high-impact tail, and its principal message for decision-making is that the near-term is decisive despite its muted temperature signal (*high confidence*).

The near-term global mean warming is robust and pathway-insensitive; the near-term regional and extreme behavior is a mosaic of confidently projected signals and variability-dominated ranges; the near-term emissions pathway is malleable and economically consequential; and the near-term tail contains outcomes whose potential severity justifies explicit assessment and sustained monitoring. The methodological developments carried into this assessment — probabilistic transition analysis (Zhang et al., 2022), traceable emissions-to-temperature chains (Kikstra et al., 2022), initial-condition large ensembles, maturing decadal prediction on improved ocean reanalyses (Jean-Michel et al., 2021), and integrated risk frameworks (Simpson et al., 2021) — collectively sharpen each of these elements. The near-term is not a period of small consequence; it is the period in which the determinants of the longer-term trajectory are set, in which the risks that most affect societies are already intensifying, and in which the Earth system would first signal whether it is following the central projection or a higher-impact path.

5.7 References

- AghaKouchak, A., et al. (2021). Anthropogenic Drought: Definition, Challenges, and Opportunities. *Reviews of Geophysics*. <https://doi.org/10.1029/2019rg000683>
- Barriopedro, D., et al. (2023). Heat Waves: Physical Understanding and Scientific Challenges. *Reviews of Geophysics*. <https://doi.org/10.1029/2022rg000780>
- Breyer, C., et al. (2022). On the History and Future of 100% Renewable Energy Systems Research. *IEEE Access*. <https://doi.org/10.1109/access.2022.3193402>
- Chang, M., et al. (2021). Trends in tools and approaches for modelling the energy transition. *Applied Energy*. <https://doi.org/10.1016/j.apenergy.2021.116731>
- Cheng, D., et al. (2023). Projecting future carbon emissions from cement production in developing countries. *Nature Communications*. <https://doi.org/10.1038/s41467-023-43660-x>
- Cui, R., et al. (2021). A plant-by-plant strategy for high-ambition coal power phaseout in China. *Nature Communications*. <https://doi.org/10.1038/s41467-021-21786-0>
- Dolan, F., et al. (2021). Evaluating the economic impact of water scarcity in a changing world. *Nature Communications*. <https://doi.org/10.1038/s41467-021-22194-0>
- Dziejarski, B., et al. (2023). Current status of carbon capture, utilization, and storage technologies in the global economy: A survey of technical assessment. *Fuel*. <https://doi.org/10.1016/j.fuel.2023.127776>
- Flores, B. M., et al. (2024). Critical transitions in the Amazon forest system. *Nature*. <https://doi.org/10.1038/s41586-023-06970-0>
- Friedlingstein, P., et al. (2023). Global Carbon Budget 2023. *Earth System Science Data*. <https://doi.org/10.5194/essd-15-5301-2023>
- Gernaat, D. E. H. J., et al. (2021). Climate change impacts on renewable energy supply. *Nature Climate Change*. <https://doi.org/10.1038/s41558-020-00949-9>

- Jean-Michel, L., et al. (2021). The Copernicus Global 1/12 degree Oceanic and Sea Ice GLORYS12 Reanalysis. *Frontiers in Earth Science*. <https://doi.org/10.3389/feart.2021.6988>
- Kemp, L., et al. (2022). Climate Endgame: Exploring catastrophic climate change scenarios. *Proceedings of the National Academy of Sciences*. <https://doi.org/10.1073/pnas>
- Kikstra, J. S., et al. (2022). The IPCC Sixth Assessment Report WGIII climate assessment of mitigation pathways: from emissions to global temperatures. *Geoscientific Model Development*. <https://doi.org/10.5194/gmd-15-9075-2022>
- Lawrence, M., et al. (2024). Global polycrisis: the causal mechanisms of crisis entanglement. *Global Sustainability*. <https://doi.org/10.1017/sus.2024.1>
- Lei, T., et al. (2023). Global iron and steel plant CO₂ emissions and carbon-neutrality pathways. *Nature*. <https://doi.org/10.1038/s41586-023-06486-7>
- Lovato, T., et al. (2022). CMIP6 Simulations With the CMCC Earth System Model (CMCC-ESM2). *Journal of Advances in Modeling Earth Systems*. <https://doi.org/10.1029/2021ms002814>
- McCrystall, M. R., et al. (2021). New climate models reveal faster and larger increases in Arctic precipitation than previously projected. *Nature Communications*. <https://doi.org/10.1038/s41467-021-27031-y>
- Pörtner, H. O., et al. (2023). Overcoming the coupled climate and biodiversity crises and their societal impacts. *Science*. <https://doi.org/10.1126/science.abl4881>
- Rantanen, M., et al. (2022). The Arctic has warmed nearly four times faster than the globe since 1979. *Communications Earth & Environment*. <https://doi.org/10.1038/s43247-022-00498-3>
- Rennert, K., et al. (2022). Comprehensive evidence implies a higher social cost of CO₂. *Nature*. <https://doi.org/10.1038/s41586-022-05224-9>
- Senande-Rivera, M., et al. (2022). Spatial and temporal expansion of global wildland fire activity in response to climate change. *Nature Communications*. <https://doi.org/10.1038/s41467-022-28835-2>
- Simpson, N. P., et al. (2021). A framework for complex climate change risk assessment. *One Earth*. <https://doi.org/10.1016/j.oneear.2021.03.005>
- Soergel, B., et al. (2021). A sustainable development pathway for climate action within the UN 2030 Agenda. *Nature Climate Change*. <https://doi.org/10.1038/s41558-021-01098-3>
- Spek, M., et al. (2022). Perspective on the hydrogen economy as a pathway to reach net-zero CO₂ emissions in Europe. *Energy & Environmental Science*. <https://doi.org/10.1039/d1ee02118d>
- Sabatino, F., et al. (2021). A comparative energy and costs assessment and optimization for direct air capture technologies. *Joule*. <https://doi.org/10.1016/j.joule.2021.05.023>
- Tebaldi, C., et al. (2021). Climate model projections from the Scenario Model Intercomparison Project (ScenarioMIP) of CMIP6. *Earth System Dynamics*. <https://doi.org/10.5194/esd-12-253-2021>
- Tripathy, K. P., et al. (2023). Climate change will accelerate the high-end risk of compound drought and heatwave events. *Proceedings of the National Academy of Sciences*. <https://doi.org/10.1073/pnas.2219825120>
- Victoria, M., et al. (2021). Solar photovoltaics is ready to power a sustainable future. *Joule*. <https://doi.org/10.1016/j.joule.2021.03.005>

Wunderling, N., et al. (2021). Interacting tipping elements increase risk of climate domino effects under global warming. *Earth System Dynamics*. <https://doi.org/10.5194/esd-12-601-2021>

Zhang, S., et al. (2022). Assessing the energy transition in China towards carbon neutrality with a probabilistic framework. *Nature Communications*. <https://doi.org/10.1038/s41467-021-27671-0>

Zhong, X., et al. (2021). Global greenhouse gas emissions from residential and commercial building materials and mitigation strategies to 2060. *Nature Communications*. <https://doi.org/10.1038/s41467-021-26212-z>

Chapter 6

Chapter 6: Long-Term Global Projections

6.1 Executive Summary

Global surface temperature will continue to rise until at least mid-century under all assessed emission scenarios, and the level of warming reached by 2100 is determined overwhelmingly by cumulative future carbon dioxide emissions rather than by the pathway taken to any given level (*virtually certain*). The multi-model ensembles that underpin this chapter place the century-end warming across the plausible range of low- to high-emission scenarios between roughly 1.4 degrees C and more than 4 degrees C above the late-nineteenth-century baseline, with the spread dominated by scenario choice at the high end and by climate response uncertainty at the low end (*high confidence*) (Tebaldi et al., 2021).

Warming-level framing — organizing projected impacts by global mean temperature increase rather than by scenario or date — is the most robust way to communicate long-term change, because many hazard metrics scale more tightly with global warming level than with the emission trajectory that produced it (*high confidence*). Every increment of warming matters: glacier loss, heat-extreme frequency, drought expansion, and the probability of crossing critical thresholds all increase monotonically and, for several systems, super-linearly with each additional tenth of a degree (Rounce et al., 2023) (Yuan et al., 2023).

The polar and high-mountain regions will continue to warm substantially faster than the global mean, sustaining Arctic amplification at roughly three to four times the global rate and elevation-dependent warming across the world's mountains (*high confidence*). This pattern, already observed, is a robust feature of every projection ensemble and drives disproportionate change in the cryosphere, in high-latitude ecosystems, and in the ocean's polar sink regions (Rantanen et al., 2022) (Pepin et al., 2022).

Global precipitation will increase in the ensemble mean, but the change is spatially organized into a robust “wet-get-wetter, dry-get-drier” and “wet-season-wetter” pattern superimposed on large regional uncertainty, and the

frequency and intensity of both heavy-precipitation and drought extremes will rise together (*high confidence*). Flash-drought onset accelerates under warming across most land regions, and anthropogenic influence on drought frequency, duration and intensity is already detectable (Yuan et al., 2023) (Chiang et al., 2021).

Global mean sea level will rise through 2100 and for centuries to millennia beyond, and the deep-time commitment is effectively irreversible on human timescales even if temperatures are stabilized (*high confidence*). The century-scale rise is dominated by ocean thermal expansion and glacier melt, both of which are well constrained; the deep uncertainty resides in the marine-based sectors of the Antarctic Ice Sheet, whose instability governs the high-end tail (Rounce et al., 2023).

The ocean will continue to warm, acidify, lose oxygen, and experience more frequent, intense and prolonged marine heatwaves throughout the twenty-first century (*virtually certain*), with biological consequences that scale steeply with warming level (Smith et al., 2022). Large-scale overturning circulation is assessed to weaken over the century, though a complete collapse before 2100 is not the assessed central outcome.

Several large-scale components of the Earth system carry a non-negligible probability of crossing self-perpetuating thresholds under sustained high warming, including the Amazon forest and marine-based Antarctic ice (*medium confidence in the thresholds; high confidence that the risk rises with warming*). Low-likelihood, high-impact and catastrophic outcomes cannot be excluded and warrant explicit assessment rather than omission (Flores et al., 2024) (Kemp et al., 2022).

Model evaluation and observational constraints tighten the projected ranges relative to the raw ensemble spread, and emergent-constraint and reanalysis-based methods now provide a defensible basis for narrowing — but not eliminating — climate-response uncertainty (*high confidence*). Machine-learning weather and climate emulators have matured to the point of complementing, though not replacing, process-based general circulation models for specific projection tasks (Bi et al., 2023).

6.2 Scenarios, Warming Levels, and the Basis for Projection

The scenario framework and the primacy of cumulative emissions

Long-term global projections rest on a nested framework of socioeconomic storylines, emission and concentration pathways, and the climate-model ensembles forced by them; within this framework, cumulative carbon dioxide emissions are the single most powerful predictor of the warming reached at any future date (*virtually certain*). The Scenario Model Intercomparison Project provides the coordinated set of forcing pathways spanning very low

to very high emissions that the coupled models translate into physical climate futures, and it is the backbone of the projections assessed in this chapter (Tebaldi et al., 2021). The design principle of this framework is that the same global warming level can be reached by different combinations of emissions, aerosols, and short-lived forcers, but the near-linear relationship between cumulative carbon dioxide and global temperature means that the total amount of carbon released — not the shape of the release curve — sets the eventual warming.

This near-linearity has a profound implication for how projections are read. It means that the difference between a world that peaks emissions early and declines steeply and a world that delays and overshoots is, to first order, a difference in the total carbon budget consumed, and therefore in the committed warming. The annual carbon budget assessments that track this consumption show that emissions have continued at or near record levels through the early 2020s, narrowing the remaining budget for any given temperature target and shifting the plausible envelope of century-end outcomes upward relative to the most optimistic pathways (Friedlingstein et al., 2022) (Friedlingstein et al., 2023). The most recent budget accounting confirms that the residual budget consistent with holding warming to the lower end of the assessed range is being drawn down rapidly and, on current trajectories, would be exhausted within a small number of years to a decade at contemporary emission rates (Friedlingstein et al., 2025).

The scenario framework also embeds a structural asymmetry that dominates the projection spread: at the high-emission end, the divergence between scenarios is the primary source of uncertainty in century-end warming, whereas at the low-emission end the residual uncertainty is dominated by the physical response of the climate system to a given forcing (Tebaldi et al., 2021). This is why warming-level framing, discussed below, is increasingly preferred to scenario-and-date framing: it isolates the physically robust, warming-dependent signal from the deeply uncertain question of which emission path humanity will actually follow.

Warming-level framing

Expressing projected changes as functions of global warming level — 1.5, 2, 3, and 4 degrees C above the pre-industrial baseline — rather than as functions of scenario and calendar year is the most robust organizing principle for long-term projections, because many impact-relevant quantities are approximately invariant to the pathway that produced a given warming (*high confidence*). The value of this framing is that it separates two logically distinct questions: how much a hazard changes per degree of global warming, which is a comparatively well-constrained physical-science question, and how much warming the world will experience, which is a socioeconomic and policy question. Glacier mass loss illustrates the framing cleanly: projected twenty-first-century glacier loss is a smooth, monotonic function of peak global temperature, such that “every increase in temperature matters” for the fraction of glacier mass and glacier number that survives the century (Rounce et al., 2023).

Warming-level framing also exposes the strong nonlinearities that scenario-and-date framing can obscure. Several hazard metrics scale super-linearly with warming, so that the incremental impact of moving from 2 to 3 degrees C exceeds that of moving from 1 to 2 degrees C. Flash-drought onset frequency, for example, is projected to increase across most land regions with a sensitivity that steepens at higher warming (Yuan et al., 2023). The probability of compound and record-shattering extremes, and the likelihood of crossing thresholds in slow-responding systems such as ice sheets and forests, likewise rise disproportionately at the upper end of the warming range (Flores et al., 2024). The framing therefore does double duty: it makes projections more robust to scenario uncertainty and it clarifies the physical case for limiting peak warming.

A limitation of warming-level framing must be stated explicitly. Some processes are not invariant to pathway even at a fixed warming level — most importantly those governed by the rate of change or by the integrated exposure over time, such as ocean heat content, thermosteric sea level, and ice-sheet response, which continue to evolve after surface temperature stabilizes (*high confidence*). For these committed, slow-responding quantities, two worlds at the same global warming level but arriving by different paths and at different times can have materially different states. The framing is thus most powerful for fast-responding atmospheric and land-surface hazards and must be applied with care to the ocean and cryosphere, which are assessed later in this chapter.

The pattern of warming: polar amplification and elevation dependence

Projected warming is not spatially uniform; it is amplified over land relative to ocean, in the Arctic relative to lower latitudes, and at high elevations relative to adjacent lowlands, and these patterns are robust across every projection ensemble (*high confidence*). Arctic amplification is the most pronounced regional signal: observations show the Arctic has warmed at close to four times the global rate since 1979, a ratio that exceeds many earlier model estimates and that projections carry forward through the century as a persistent feature (Rantanen et al., 2022). The physical drivers — sea-ice and snow-albedo feedbacks, changes in atmospheric and oceanic heat transport, and lower-tropospheric stability that traps warming near the surface — are well understood and ensure the pattern's persistence under continued forcing.

Elevation-dependent warming is the mountain analogue of polar amplification. Across the world's mountain regions, warming rates frequently increase with altitude, driven by snow-albedo feedback at the retreating snowline, by changes in cloud and water-vapour distribution, and by aerosol deposition on snow and ice (Pepin et al., 2022). The consequence is that high-mountain cryospheres — glaciers, snowpack, and permafrost — are exposed to warming that can exceed the global mean substantially, accelerating their loss and the downstream hydrological changes that depend on them. Because a large fraction of the world's population depends on mountain water towers for dry-season

flow, elevation-dependent warming translates the global temperature signal into concentrated regional water-security consequences.

The land-ocean warming contrast, a third robust pattern, arises from the ocean's vast thermal inertia and its efficient evaporative cooling, which cause land to warm roughly 1.4 to 1.6 times faster than the ocean in the ensemble mean. This contrast reshapes atmospheric circulation, shifts monsoon and storm-track behaviour, and concentrates the most rapid warming — and therefore the most rapid change in heat extremes — over the continental interiors where most people live.

Model evaluation and constrained projections

Confidence in long-term projections rests on the demonstrated ability of the underlying models to reproduce the observed mean climate, its variability, and its historical trends, and on the use of observational constraints to narrow the raw ensemble spread (*high confidence*). Model evaluation proceeds on multiple fronts: comparison of simulated climatologies against modern reanalyses, assessment of the amplitude and phase of internal modes of variability, and verification that the models reproduce the observed historical warming and its spatial fingerprint. High-resolution ocean and coupled reanalyses now available provide the observational reference against which model biases in near-surface temperature, soil moisture, and ocean state are quantified (Jean-Michel et al., 2021).

Emergent constraints and other observationally based methods narrow the projected range by exploiting statistical relationships, established across the ensemble, between an observable feature of the present climate and an uncertain aspect of the future response. Applied to quantities such as the carbon-cycle feedback, the snow-albedo feedback, and regional hydroclimate, these methods can meaningfully tighten the projected range, though they do not eliminate structural uncertainty and must be validated against the physical mechanism they invoke to avoid spurious correlations (*medium confidence* in any single constraint; *high confidence* that the ensemble of constraints reduces the effective spread). The soil-moisture-atmosphere coupling that governs a large share of interannual carbon-uptake variability is one such constrainable process, and its representation is now a standard benchmark for land-surface components (Humphrey et al., 2021).

The maturation of machine-learning weather and climate emulators is a notable development for the projection enterprise. Data-driven models trained on reanalysis now match or exceed operational physics-based systems for medium-range global weather forecasting at a fraction of the computational cost (Bi et al., 2023). For long-term climate projection these emulators do not replace process-based general circulation models — they cannot, by construction, simulate forcing regimes outside their training distribution with guaranteed fidelity — but they are increasingly used to downscale, to generate large initial-condition ensembles for internal-variability quantification, and to accelerate the exploration of scenario space. Their role in this assessment is complementary: they

sharpen specific projection tasks while the physical models remain the authority on the response to novel forcing.

6.3 Projected Changes in Temperature, Precipitation, and Circulation

Global surface temperature and the planetary energy budget

Global mean surface temperature will rise until at least the middle of the twenty-first century under all assessed scenarios, and the warming reached by 2100 spans from about 1.4 degrees C under the most stringent mitigation to well above 4 degrees C under sustained high emissions (*virtually certain* that warming continues to mid-century; *high confidence* in the assessed ranges). The ScenarioMIP ensemble that underpins these ranges shows a consistent ordering across scenarios and a scenario-dependent widening of the spread through the century, with the low-emission pathways stabilizing and even slightly declining in temperature after mid-century while the high-emission pathways continue to accelerate (Tebaldi et al., 2021). The central estimate for any given scenario is well constrained by the models' agreement on the transient response to cumulative emissions; the residual spread at a fixed scenario reflects the range of equilibrium and transient climate sensitivities across the ensemble.

The physical basis for continued warming is the planetary energy imbalance. As long as the atmospheric burden of carbon dioxide and other greenhouse gases rises, the top-of-atmosphere energy budget remains in surplus and the system accumulates heat, the great majority of which enters the ocean. This imbalance is the reason surface temperature does not respond instantaneously to emissions and the reason warming continues for a time even after emissions decline: the system is always chasing an equilibrium set by the current forcing. The carbon-budget assessments quantify the forcing side of this balance, tracking the atmospheric growth rate and the partitioning of emitted carbon among atmosphere, ocean, and land, and they show that the airborne fraction has remained broadly stable even as absolute emissions have risen, so that atmospheric concentrations — and therefore forcing — continue to climb (Friedlingstein et al., 2023).

A crucial feature of the temperature projections is the behaviour of overshoot pathways, in which warming exceeds a target level and is subsequently brought back down through net-negative emissions. Because the temperature response to cumulative emissions is near-linear and largely reversible for the atmosphere on decadal timescales, peak warming is set by peak cumulative emissions; but the slow-responding components of the system — deep ocean heat, thermosteric sea level, ice sheets — do not reverse on the same timescale and carry a memory of the overshoot (*high confidence*). The assessment therefore distinguishes sharply between the reversibility of surface air temperature

and the effective irreversibility of the committed changes in the ocean and cryosphere that accrue during an overshoot.

Regional temperature and extreme heat

Extreme heat will intensify everywhere land warms, and because heat extremes scale with and often exceed the local mean warming, the frequency, duration, and intensity of hot extremes rise faster than the global mean temperature (*virtually certain*). The land-ocean contrast and continental amplification concentrate the most rapid warming — and the most rapid change in the tails of the temperature distribution — over inhabited land. Regions that combine high present-day temperatures with strong warming, such as the Eastern Mediterranean and Middle East, are projected to experience some of the most severe increases in heat-extreme frequency and in the emergence of previously unprecedented conditions, with warming rates in that region already exceeding the global mean and projected to continue doing so (Zittis et al., 2022). The compounding of heat with humidity raises the additional spectre of conditions approaching physiological limits for outdoor human activity in the hottest, most humid regions under high warming.

The attribution literature establishes that the human influence on observed heat extremes is already unambiguous and that projected future changes are a continuation of a detected trend rather than a novel projection (Clarke et al., 2022). This continuity between attributed past and projected future strengthens confidence in the heat projections specifically: the models that project future intensification are the same models shown to reproduce the observed, attributable intensification. The practical consequence is that heat-extreme projections are among the most robust and highest-confidence elements of the entire long-term projection enterprise.

Global and regional precipitation

Global mean precipitation will increase with warming at a rate governed by the atmosphere's energy budget rather than its moisture content, producing a global sensitivity of order 1 to 3 percent per degree C, while atmospheric moisture — and therefore the intensity of heavy-precipitation extremes — rises faster, at close to the 7 percent per degree C dictated by thermodynamics (*high confidence*). This distinction between the energetically constrained mean and the thermodynamically constrained extreme is central to the precipitation projections. The mean can increase only as fast as the atmosphere can radiate away the latent heat released by condensation, which limits it; the extremes, drawing on a moister atmosphere, can intensify much faster. The result is a projected world in which total precipitation rises modestly but is delivered increasingly in intense events separated by longer dry spells.

The spatial pattern of mean precipitation change follows robust large-scale organizing principles — wet regions and wet seasons generally getting wetter, subtropical dry zones generally getting drier and expanding poleward — super-

imposed on large regional uncertainty arising from circulation changes that models simulate less consistently than thermodynamics. Monsoon systems are projected to intensify in total rainfall while becoming more variable and, in some regions, delayed in onset. The high-latitude and equatorial wet zones are projected to see the clearest increases, while the subtropical margins face the clearest drying. Because circulation responses are the dominant source of regional precipitation uncertainty, the confidence in regional precipitation projections is systematically lower than the confidence in temperature projections.

The co-occurrence of intensifying heavy precipitation and expanding drought is not a contradiction but a direct consequence of the same warming, and both will increase together over most land areas (*high confidence*). A warmer atmosphere holds more water and evaporates it faster, so the same climate that delivers more intense downpours also dries soils more aggressively between them. Anthropogenic influence on the global characteristics of drought — its frequency, duration, and intensity — is already detectable in the observational record, providing an attributed baseline for the projected future expansion (Chiang et al., 2021). The intensification of flooding hazard follows from the precipitation-extreme signal: inequitable and rising flood exposure is projected across river and coastal systems as extreme rainfall intensifies and sea level rises (Wing et al., 2022).

Drought, aridity, and flash droughts

The character of drought will shift toward faster onset and greater intensity, with flash droughts — rapid-onset events driven by evaporative demand — becoming more frequent across most land regions as warming raises atmospheric thirst (*high confidence*). Flash droughts, which develop over weeks rather than seasons, are projected to increase in onset frequency globally under continued warming, with the transition from slow to flash drought a robust ensemble feature that steepens at higher warming levels (Yuan et al., 2023). The driver is the rising vapour-pressure deficit of a warmer atmosphere, which increases the rate at which soils and vegetation lose water even where precipitation is unchanged, so that drought can develop rapidly during a single hot, dry spell.

The megadrought that emerged across southwestern North America in the early twenty-first century exemplifies the compounding of natural variability with the anthropogenic aridification trend, and it has been characterized as the most severe multi-decadal drought in the region for over a millennium, with a substantial fraction of its severity attributable to human-caused warming (Williams et al., 2022). Such events are informative for projection because they reveal how the warming-driven increase in evaporative demand converts what would have been a moderate natural drought into an extreme one — a mechanism that projections indicate will operate with increasing force through the century.

Atmospheric and coupled circulation

Large-scale atmospheric circulation will change in ways that are physically expected but regionally uncertain, including a poleward expansion of the subtropical dry zones, a poleward shift and possible intensification of the mid-latitude storm tracks, and altered behaviour of the major modes of internal variability (*medium-to-high confidence in the direction of the large-scale shifts; lower confidence in regional detail*). The El Niño–Southern Oscillation, the dominant mode of interannual climate variability, is assessed to change in a warming climate, with projections indicating an increase in the amplitude of ENSO-related sea-surface-temperature and rainfall variability and in the frequency of extreme events at both ends of the cycle, even as the mean-state response remains less certain (Cai et al., 2021). Because ENSO teleconnections govern year-to-year climate risk across much of the tropics and subtropics, changes in its behaviour propagate into changes in the variability of drought, flood, and heat risk far from the equatorial Pacific.

The projected weakening of large-scale overturning circulations — atmospheric and oceanic — is a further robust element. In the atmosphere, the tropical overturning is projected to weaken as the mean precipitation increase lags the moisture increase. In the ocean, the meridional overturning circulation is assessed to weaken over the century under continued warming, a change with consequences for regional sea level, heat transport, and the ocean carbon and oxygen budgets that are taken up in the ocean-futures assessment below. The confidence in the sign of these circulation changes is higher than the confidence in their magnitude, which remains model-dependent.

6.4 Sea Level and Ocean Futures

Global mean sea level rise through 2100

Global mean sea level will continue to rise throughout the twenty-first century under every assessed scenario, with the century-end rise scaling with warming and the well-constrained contributions from ocean thermal expansion and glacier melt dominating the central estimate (*high confidence*). The physical decomposition of sea level rise separates a set of comparatively well-understood contributors — the thermal expansion of a warming ocean and the melt of the world’s mountain glaciers — from the deeply uncertain contribution of the great ice sheets. The glacier contribution is among the best constrained: projections show that glacier mass loss through 2100 is a smooth function of peak global temperature, with the fraction of glacier mass lost and the number of glaciers that disappear entirely rising monotonically with warming, such that even the lowest warming levels commit the world to substantial and, for many small glaciers, complete loss (Rounce et al., 2023). Because these losses are essentially locked in for glaciers that fall below their equilibrium line, a

large share of the twenty-first-century glacier contribution to sea level is already committed.

Thermal expansion is the second pillar of the central estimate and the most direct expression of ocean heat uptake. As the ocean accumulates the great majority of the planetary energy imbalance, its volume increases; this thermosteric contribution responds slowly, continues long after surface temperature stabilizes, and is effectively irreversible on human timescales, giving sea level a deep memory of the integrated warming rather than of its instantaneous value (*high confidence*). This slow, committed character is why sea level rise cannot be arrested by stabilizing temperature alone and why the sea-level projection must be framed over centuries, not decades.

The assessed likely range for century-end global mean sea level rise spans from roughly half a metre under stringent mitigation to close to a metre or more under high emissions, with the central estimates well constrained by the thermal-expansion and glacier contributions. The high-end tail, however, is governed by processes that the likely range does not fully capture, and it is to those that the assessment turns next.

Ice-sheet contributions and the high-end tail

The dominant uncertainty in twenty-first-century and longer-term sea level rise is the behaviour of the marine-based sectors of the Antarctic Ice Sheet, whose potential instability governs a low-likelihood, high-impact tail that cannot be excluded and that widens sharply beyond 2100 (*medium confidence in the magnitude; high confidence that the risk rises with warming*). The Greenland and Antarctic ice sheets contribute to sea level through surface melt and through the dynamic discharge of ice to the ocean, and it is the dynamic component — the acceleration, thinning, and retreat of marine-terminating outlet glaciers and ice shelves — that carries the deep uncertainty. Marine ice-sheet and ice-cliff instabilities, in which retreat of the grounding line into deepening bedrock basins becomes self-sustaining, provide a physical mechanism by which the Antarctic contribution could become both large and effectively unstoppable once initiated. The assessment treats this tail explicitly rather than omitting it, because a low-probability outcome with civilization-scale consequences is decision-relevant even when it is not the central estimate.

The distinction between the reversible atmospheric response and the irreversible ice-sheet response is central to the interpretation of overshoot pathways. A world that overshoots a temperature target and returns to it may commit to ice-sheet retreat during the overshoot that continues even after temperatures decline, because grounding-line retreat, once triggered, does not reverse when the forcing that triggered it is removed (*high confidence in the asymmetry; medium confidence in the thresholds*). This hysteresis means that peak warming, not final warming, may set the multi-century ice-sheet trajectory — a consideration that sharpens the physical case for limiting overshoot.

Sea level rise beyond 2100 and its commitment

Global mean sea level rise does not stop in 2100; it continues for centuries to millennia, and the multi-century commitment scales with peak warming to reach magnitudes measured in metres, effectively reshaping the world's coastlines on a timescale that human infrastructure cannot outpace (*high confidence*). The deep-time commitment arises because the two largest reservoirs — the deep ocean's heat content and the great ice sheets — respond on timescales of centuries to millennia. Even under strong mitigation, the thermal expansion already set in motion and the ice-sheet response already committed will continue to raise sea level for centuries. Under high warming, the marine-based Antarctic sectors could contribute multiple metres over subsequent centuries, and the ultimate commitment from sustained high temperatures is measured in many metres of eventual rise. The practical significance is that sea level rise is the archetypal committed, quasi-irreversible consequence of warming: the emissions of this century set a coastline trajectory for the next several.

Ocean heat, stratification, and marine heatwaves

The ocean will continue to warm at the surface and at depth, will become more strongly stratified, and will experience marine heatwaves that are more frequent, more intense, and longer-lasting, with biological consequences that scale steeply with warming (*virtually certain for the warming; high confidence for the marine-heatwave intensification*). The ocean absorbs the overwhelming majority of the excess heat trapped by greenhouse gases, and this heat uptake, concentrated near the surface, increases the vertical density gradient and suppresses the mixing that ventilates the interior and supplies nutrients to the sunlit layer. The result is a more stratified ocean in which surface warming is amplified and the transfer of heat, carbon, and oxygen to depth is impeded. High-resolution ocean reanalyses provide the observational foundation for characterizing this evolving state and for evaluating the models that project it forward (Jean-Michel et al., 2021).

Marine heatwaves — prolonged periods of anomalously high ocean temperature — are projected to intensify along every metric under continued warming, and their biological impacts are severe and increasingly well documented, spanning mass mortality of foundation species such as corals and kelp, geographic redistribution of mobile species, and disruption of fisheries and the human communities that depend on them (Smith et al., 2022). Because many marine organisms live close to their thermal limits, the biological response to marine heatwaves is strongly nonlinear in warming level, so that each increment of ocean warming disproportionately raises the risk of ecosystem-scale disruption. The compounding of chronic warming with acute heatwave events, and with acidification and deoxygenation, produces multi-stressor conditions that exceed the tolerance of many species and ecosystems.

Ocean acidification and deoxygenation

Ocean acidification will continue in near-lockstep with atmospheric carbon dioxide, and ocean oxygen content will decline as warming reduces solubility and suppresses ventilation, together degrading habitat for a wide range of marine life (*virtually certain for acidification; high confidence for deoxygenation*). Acidification is the most predictable of all the ocean's chemical responses because it is governed directly by the uptake of atmospheric carbon dioxide into surface waters, a process well described by carbonate chemistry; the surface-ocean pH decline therefore tracks the emission pathway with little of the uncertainty that attends the physical climate response. The carbon-budget assessments quantify the ocean carbon sink that drives this acidification, showing that the ocean continues to absorb a large and roughly constant fraction of anthropogenic emissions — a service to the climate that comes at the direct cost of the chemistry that calcifying organisms depend on (Friedlingstein et al., 2023).

Deoxygenation proceeds through two mechanisms that both strengthen with warming: warmer water holds less dissolved oxygen, and increased stratification reduces the downward mixing of oxygen-rich surface water into the interior. The expansion of oxygen-minimum zones compresses the viable habitat for oxygen-demanding species and interacts with warming and acidification to produce a triple stress on ocean ecosystems. The three stressors are not independent — they share the common driver of rising carbon dioxide and warming — so their impacts compound, and the ocean-ecosystem response to high-emission futures is correspondingly severe.

Overturning circulation

The Atlantic meridional overturning circulation is assessed to weaken over the twenty-first century under continued warming, with consequences for regional sea level, heat transport, and marine ecosystems; an abrupt collapse before 2100 is not the assessed central outcome but cannot be excluded and constitutes a further high-impact tail risk (*medium confidence in the weakening; low confidence in the likelihood of collapse*). The overturning circulation is driven in part by the formation of dense water in the high-latitude North Atlantic, a process that warming and freshwater input from ice melt both act to suppress. A weakening of this circulation redistributes heat and sea level regionally — raising sea level along some coasts, cooling the surface in the subpolar North Atlantic relative to the global trend — and alters the ocean's uptake of carbon and oxygen. The assessment treats the possibility of a more abrupt transition as a low-likelihood, high-impact outcome that, like the marine ice-sheet tail, warrants explicit consideration precisely because its consequences would be severe and largely irreversible on human timescales.

6.5 Cryosphere Futures

Mountain glaciers

The world's mountain glaciers will lose a large fraction of their mass over the twenty-first century under every scenario, and the loss is a smooth, monotonic function of peak global temperature such that limiting warming directly limits glacier loss — every increase in temperature matters (*high confidence*).

Comprehensive global glacier modelling shows that even under the most stringent mitigation a substantial share of present-day glacier mass is committed to disappear, while under high warming the majority of glaciers — and the overwhelming majority of small glaciers — are lost by 2100 (Rounce et al., 2023). The number of glaciers lost is even more sensitive to warming than the mass lost, because small glaciers vanish entirely at low warming increments while the largest ice bodies persist in reduced form; the result is a landscape from which glaciers disappear region by region as warming climbs.

The consequences of glacier loss propagate through the hydrological cycle and human water security. Glaciers act as multi-year to multi-decadal buffers, storing winter precipitation as ice and releasing it as meltwater during the warm, dry season when demand is highest. As glaciers shrink, this buffering first increases — producing a transient “peak water” pulse of enhanced melt — and then declines irreversibly as the ice reservoir is depleted, leaving dry-season flows diminished and more variable. Because elevation-dependent warming exposes high-mountain cryospheres to warming that can exceed the global mean, the glacier response is amplified relative to what the global temperature signal alone would suggest (Pepin et al., 2022). The regions most dependent on glacier-fed rivers face a transition from abundant to scarce dry-season water within the century, with the timing of peak water and its subsequent decline set by the warming trajectory.

Snow cover and seasonal snowpack

Seasonal snow cover will continue to decline in extent, duration, and mass across most of the Northern Hemisphere's snow-affected regions, shortening the snow season and shifting the timing of snowmelt earlier in the year (*high confidence*).

Snow is the most responsive component of the cryosphere because it exists close to the melting point for much of its life; modest warming therefore produces large changes in the fraction of precipitation that falls as snow rather than rain, in the date of peak snowpack, and in the timing of melt-driven streamflow. The projected shift toward earlier and more rain-dominated hydrographs alters flood risk, reservoir management, and dry-season water availability across snow-fed basins. The snow-albedo feedback that accompanies snow-cover loss is itself an amplifier of warming, reinforcing both polar amplification and elevation-dependent warming and helping to explain why high-latitude and high-elevation regions warm faster than the global mean (Rantanen et al., 2022) (Pepin et al., 2022).

Arctic sea ice

Arctic sea ice will continue to decline in extent and thickness, and a practically ice-free Arctic Ocean in late summer is projected to occur before mid-century under most scenarios, with the frequency of ice-free conditions rising with warming (*high confidence*). The loss of Arctic sea ice is one of the most visible and dynamically consequential elements of the projected cryosphere. As the highly reflective ice cover retreats, it exposes the dark ocean surface, which absorbs far more solar radiation; this ice-albedo feedback is a principal driver of the Arctic amplification that has already produced warming at close to four times the global rate and that projections carry forward through the century (Rantanen et al., 2022). The transition to seasonally ice-free conditions is projected to be effectively unavoidable under continued warming, though the exact timing depends on scenario and on internal variability, and the summer-ice loss is more nearly reversible than the ice-sheet changes because thin seasonal ice can regrow if temperatures decline.

The consequences of Arctic sea-ice loss extend beyond the polar region. The reduction in the equator-to-pole temperature gradient alters the behaviour of the mid-latitude jet stream and storm tracks, with proposed but still uncertain effects on the persistence and character of mid-latitude weather extremes. The opening of the Arctic Ocean also reshapes regional ecosystems, carbon fluxes, and human access, and it accelerates the warming of coastal permafrost. The assessment holds high confidence in the continued decline and in the eventual attainment of ice-free summers; the confidence in the remote atmospheric consequences is lower.

Antarctic sea ice and the Southern Ocean

Antarctic sea ice, historically more variable and less clearly trending than its Arctic counterpart, is assessed with lower confidence, but the Southern Ocean's continued heat and carbon uptake and its role in bottom-water formation make it a critical and evolving component of the long-term projection (*medium confidence*). The Southern Ocean absorbs a disproportionate share of the ocean's heat and anthropogenic carbon uptake, and the formation of dense Antarctic bottom water is a driver of the global overturning circulation. Changes in Antarctic sea ice, in the freshwater balance from ice-shelf melt, and in the wind-driven circulation of the Southern Ocean therefore feed back on the global ocean's capacity to take up heat and carbon. High-resolution ocean and sea-ice reanalyses are the observational basis for tracking this evolving state and for evaluating the coupled models that project the Southern Ocean forward (Jean-Michel et al., 2021). The lower confidence in Antarctic sea-ice projection reflects both the greater complexity of the Southern Ocean's dynamics and the shorter, more variable observational record against which models are evaluated.

Ice sheets and their long-term trajectory

The Greenland and Antarctic ice sheets will continue to lose mass through the century and for centuries beyond, and their long-term trajectories carry the deepest uncertainty and the largest potential consequences in the entire projection, with grounding-line retreat and ice-shelf loss providing mechanisms for large, effectively irreversible contributions to sea level (*high confidence in continued loss; medium confidence in the magnitude and timing of dynamic instabilities*). Greenland loses mass predominantly through surface melt, which increases smoothly with warming and which, above a threshold of sustained temperature, commits the ice sheet to long-term, near-total decline as the surface lowers into warmer air — a slow but effectively irreversible response measured in millennia. Antarctica loses mass predominantly through dynamic discharge, which is governed by the stability of ice shelves and grounding lines and which is capable of far more abrupt and less predictable behaviour. The marine-based sectors, grounded on bedrock that deepens inland, are susceptible to self-sustaining retreat once the grounding line is displaced, and this mechanism underlies the high-end sea-level tail assessed in the ocean-futures section.

The reversibility asymmetry between ice sheets and the atmosphere is the single most important interpretive point for long-term projections. Surface air temperature responds within decades to changes in cumulative emissions and is, for the atmosphere, largely reversible; the ice sheets respond over centuries to millennia and, once committed to retreat, do not recover on any timescale relevant to human societies (*high confidence*). This means that the emissions of the coming decades set an ice-sheet and sea-level trajectory that will play out over the coming centuries regardless of subsequent mitigation, and that overshoot pathways carry a particular risk of triggering commitments during the peak-warming phase that persist after temperatures decline.

Permafrost and the frozen-carbon feedback

Permafrost will thaw progressively as high-latitude and high-elevation regions warm, releasing carbon dioxide and methane in a positive feedback that amplifies warming and that projections treat as an additional, partly irreversible source not fully captured in the emission scenarios themselves (*high confidence that thaw and emission occur; medium confidence in the magnitude*). Permafrost soils store a vast quantity of ancient organic carbon, and as they thaw this carbon becomes available to microbial decomposition, releasing greenhouse gases that add to the atmospheric burden and drive further warming. Because Arctic amplification exposes permafrost regions to warming far exceeding the global mean, the thaw proceeds faster than the global temperature trajectory alone would imply (Rantanen et al., 2022). The feedback is partly irreversible on human timescales: carbon released from thawed permafrost does not return to the frozen reservoir when temperatures stabilize, so the permafrost contribution accumulates as a function of the integrated warming.

The permafrost feedback is one of several ways in which the real-world carbon cycle may respond more strongly to warming than the idealized scenarios assume, effectively narrowing the remaining carbon budget for any temperature target. The carbon-budget assessments that partition emissions among atmosphere, ocean, and land quantify the observed behaviour of these sinks and sources, and they underscore that the land carbon response — of which permafrost is one component — is both large and uncertain, contributing substantially to the year-to-year and decade-to-decade variability of the atmospheric growth rate (Friedlingstein et al., 2023). The assessment therefore treats the frozen-carbon feedback as a source of upward risk to the projected warming for any given anthropogenic emission pathway.

6.6 The Land Carbon Cycle and Terrestrial Ecosystems

The land carbon sink and its stability

The terrestrial biosphere currently absorbs roughly a quarter to a third of anthropogenic carbon dioxide emissions, but the strength and stability of this sink under continued warming is uncertain and is a major source of the spread in projected atmospheric concentrations and warming for a given emission pathway (*high confidence in the present sink; medium confidence in its future evolution*). The annual global carbon budget quantifies the partitioning of emissions among the atmosphere, ocean, and land, and it shows that the land sink, while large, is also the most variable component from year to year, driven by the response of ecosystem productivity and respiration to temperature, moisture, and carbon dioxide itself (Friedlingstein et al., 2023). The fertilizing effect of rising carbon dioxide enhances photosynthesis and has strengthened the sink historically, but this effect is expected to saturate, and it competes against the warming-driven increase in respiration and against the moisture stress that limits productivity in a drying, more variable climate.

The dominant control on the interannual variability of the land sink is the coupling between soil moisture and the atmosphere. Variations in terrestrial water availability drive the year-to-year swings in carbon uptake, so that dry years — increasingly common under the projected expansion of drought and evaporative demand — depress the land sink and allow a larger fraction of emissions to remain in the atmosphere (Humphrey et al., 2021). This mechanism links the hydrological projections directly to the carbon-cycle projections: a climate that dries soils and intensifies drought is also a climate in which the land's capacity to buffer emissions is diminished, a positive feedback that adds upward risk to the warming projection. The updated carbon-budget accounting continues to show the land sink's sensitivity to climate variability and the persistence of large uncertainties in its future trajectory (Friedlingstein et al., 2025).

Forests, drought, and tree mortality

Forests face rising risk of drought- and heat-driven mortality as warming raises vapour-pressure deficit and intensifies drought, and widespread tree die-off would convert forest carbon sinks into sources, adding to atmospheric carbon dioxide (*high confidence that mortality risk rises; medium confidence in the magnitude of the carbon consequence*). A global synthesis of tree-mortality observations reveals a consistent “hotter-drought” fingerprint: tree death is driven not by drought alone but by the combination of drought with elevated temperature, which increases the atmospheric demand for water and pushes trees past the physiological limits that govern their survival (Hammond et al., 2022). The mechanisms of this mortality — hydraulic failure as water columns in the xylem break under tension, and carbon starvation as stressed trees close their stomata and deplete their reserves — are increasingly well understood and are exacerbated by the rising vapour-pressure deficit that accompanies warming (McDowell et al., 2022). Because these mechanisms scale with temperature as well as with precipitation deficit, warming intensifies mortality risk even where rainfall is unchanged.

The consequences for the carbon cycle are significant. Forests are major carbon reservoirs and sinks, and large-scale mortality releases stored carbon while removing future uptake capacity, a transition that can convert a regional sink into a source. The interaction of drought-driven mortality with fire and with insect outbreaks compounds the risk, as stressed and dead trees provide fuel and become vulnerable to pests. The assessment treats accelerating tree mortality as both an ecosystem impact and a carbon-cycle feedback that adds to the upward risk on projected warming.

Wildfire

Fire regimes will intensify in many regions as warming lengthens fire seasons, dries fuels, and increases the frequency of extreme fire weather, with the largest and most extreme fires becoming more probable under continued warming (*high confidence*). The record-breaking fires that swept southeastern Australia at the turn of the decade illustrate the connection between climate change and extreme fire: the burned area and the conditions that produced these fires were linked to the warming- and variability-driven increase in fire-weather severity, and the multi-decadal trend in forest burned area across the continent has been attributed in part to climate change (Abram et al., 2021) (Canadell et al., 2021). The physical chain is direct — higher temperatures and elevated vapour-pressure deficit dry the vegetation that serves as fuel, longer and hotter dry seasons extend the window of flammability, and the resulting fires burn hotter and spread faster.

The carbon and ecosystem consequences of intensifying fire regimes reinforce the warming they arise from. Fire releases stored carbon rapidly, and where fires become frequent enough to prevent forest recovery they can drive a persistent transition from forest to non-forest, locking in a loss of carbon-storage

capacity. The interaction of fire with drought-driven mortality and with the drying land carbon sink forms a set of mutually reinforcing terrestrial feedbacks that, under high warming, add materially to the risk that the land carbon sink weakens or reverses.

The Amazon and other regional forest systems

The Amazon forest system is approaching conditions under which self-reinforcing dieback could occur in parts of the basin, driven by the combination of warming, drought, deforestation, and fire, and its degradation is already substantial and accelerating (*medium confidence in the threshold; high confidence that the risk rises with warming and land-use pressure*). Analysis of the drivers and impacts of Amazon forest degradation shows that a large and growing fraction of the forest is affected by edge effects, selective logging, fire, and drought, degrading its structure and carbon storage even where outright deforestation has not occurred (Lapola et al., 2023). The concern for long-term projection is that these stressors interact with the climate signal to push parts of the system toward a critical transition, in which forest loss reduces the moisture recycling that sustains regional rainfall, drying the climate further and driving additional forest loss in a self-perpetuating loop.

The assessment of critical transitions in the Amazon identifies the combinations of warming, rainfall reduction, dry-season length, and deforestation under which such a transition becomes plausible, and it locates a substantial fraction of the forest within reach of these thresholds by mid-century under continued pressure (Flores et al., 2024). The potential transition of the Amazon from forest to a degraded, savanna-like or fire-dominated state is treated as one of the more consequential regional tipping risks in the Earth system, both for the carbon it would release and for the biodiversity and regional climate it would transform.

Agriculture, water, and food systems

Climate impacts on agriculture will emerge earlier and be more severe than earlier crop-model generations projected, with warming, shifting precipitation, drought, and soil degradation reducing yields of major staples in many regions even as some high-latitude areas see gains (*high confidence that impacts intensify with warming; medium confidence in regional magnitudes*). The newest generation of coupled climate and crop models projects that the impacts of climate change on global agriculture emerge earlier in the century and are larger than previously estimated, with major breadbasket regions facing yield declines in staple crops under continued warming (Jägermeyr et al., 2021). The drivers are multiple and compounding: direct heat stress on crops during sensitive growth stages, changes in the amount and timing of precipitation, the expansion and intensification of drought, and the degradation of the soil and water resources on which agriculture depends.

Soil and water resources are themselves projected to degrade under warming in ways that amplify the direct climate impacts on crops. Primary soil salinization is projected to expand across large areas of the world's drylands over the century as changing precipitation and rising evaporative demand concentrate salts in the root zone, reducing the productivity of affected land (Hassani et al., 2021). Groundwater, the buffer on which irrigated agriculture increasingly depends during drought, faces rising demand precisely as recharge becomes more variable, making sustainable groundwater management central to a resilient water and food future (Scanlon et al., 2023). The broader environmental footprint of agriculture and its exposure to climate change form a two-way interaction, in which a changing climate degrades agricultural systems and agricultural expansion and intensification contribute to the environmental pressures that drive further change (Yang et al., 2024). The net consequence for long-term projection is that the agricultural and water-resource impacts of warming are among the most societally consequential, scale with warming level, and are compounded by the degradation of the land and water resources that agriculture depends upon.

6.7 Tipping Elements, Committed Change, and Deep-Time Projection

Large-scale singular events and the risk of crossing thresholds

Several components of the Earth system carry a non-negligible probability of crossing self-reinforcing thresholds under sustained warming, and the probability of such crossings rises with global warming level, making the avoidance of high peak warming the primary lever for reducing this class of risk (*medium confidence in the specific thresholds; high confidence that the aggregate risk increases with warming*). The tipping elements assessed here include the marine-based sectors of the Antarctic Ice Sheet, the Greenland Ice Sheet, the Atlantic overturning circulation, and the Amazon forest system, each of which possesses a mechanism by which change, once initiated, becomes self-sustaining and effectively irreversible on human timescales. The Amazon analysis identifies the combinations of warming and land-use pressure under which a critical transition of the forest becomes plausible within the century (Flores et al., 2024), while the ice-sheet and circulation risks are discussed in the ocean and cryosphere sections above. What unites these elements is the presence of positive feedbacks strong enough to carry the system to a new state even after the forcing that triggered the transition is removed.

The appropriate treatment of these low-likelihood, high-impact outcomes is to assess them explicitly rather than to exclude them because they lie in the tail of the distribution. A tail outcome with civilization-scale consequences is decision-relevant in proportion to its impact as well as its probability, and omitting it from the assessment would understate the risk that the projections

are meant to convey. The assessment therefore reports both the central, likely-range projections and the physically grounded tails — the marine ice-sheet contribution to sea level, the possibility of an abrupt overturning transition, the potential for Amazon dieback — as integral parts of the projection rather than as caveats appended to it.

Catastrophic and worst-case scenarios

The upper tail of climate risk, including the possibility of warming substantially exceeding central projections through strong carbon-cycle and other feedbacks, and the societal consequences of such warming, has been comparatively neglected and warrants dedicated assessment (*high confidence that the tail is decision-relevant; low confidence in the probability of the most extreme outcomes*). A systematic examination of catastrophic climate-change scenarios argues that the mechanisms by which high-end warming could arise — including strong carbon-cycle feedbacks such as permafrost thaw, weakened land and ocean sinks, and the compounding of multiple tipping elements — and the pathways by which such warming could translate into cascading societal failure have received insufficient scientific attention relative to their potential importance (Kemp et al., 2022). The value of assessing these scenarios is not to assign them a high probability but to characterize the mechanisms, identify the conditions under which they become more likely, and inform the management of a risk whose magnitude, however improbable, is extreme.

This framing connects the physical-science tail to the human consequences that motivate the assessment. The compounding of climate hazards with the vulnerabilities of interconnected human and natural systems can amplify a given physical change into a larger societal impact, and the interaction of the climate and biodiversity crises is itself a threat multiplier that neither crisis captures in isolation (Pörtner et al., 2023). The assessment of catastrophic outcomes is therefore an assessment of coupled human–Earth-system risk, in which the physical tail and the societal amplification must be considered together.

The Earth system in a planetary-boundaries frame

Long-term climate projection sits within a broader picture of multiple, interacting Earth-system pressures, several of which have already been pushed beyond the boundaries associated with a stable, Holocene-like state, so that climate change is one of a set of transgressed limits rather than an isolated hazard (*high confidence in the transgression of multiple boundaries; assessment of the boundaries themselves carries *medium confidence).** The planetary-boundaries analysis finds that a majority of the assessed Earth-system boundaries — spanning climate, biosphere integrity, freshwater, biogeochemical flows, and land-system change — have now been transgressed, indicating that the safe operating space for humanity has been substantially eroded (Richardson et al., 2023). Extending this analysis to incorporate justice considerations yields safe and just boundaries that are, for several dimensions, already exceeded, under-

scoring that the climate projections assessed here unfold against a backdrop of concurrent and interacting environmental change (Rockström et al., 2023).

This framing matters for projection because the components of the Earth system do not respond to warming in isolation. Land-system change, freshwater depletion, biosphere degradation, and biogeochemical disruption interact with the climate response, sometimes amplifying it — as when deforestation and drought reinforce one another in the Amazon, or when land-use change and soil degradation compound the agricultural impacts of warming. The magnitude of land-system change is itself larger than earlier estimates suggested, with global land-use change over recent decades affecting a far greater area than previously recognized, reshaping the land surface that mediates the carbon, water, and energy exchanges central to the climate projection (Winkler et al., 2021). The assessment therefore situates climate projection within the coupled dynamics of a multiply stressed Earth system rather than treating it as a single-variable forecast.

Committed change and the meaning of stabilization

Even if global temperature were stabilized, large parts of the Earth system would continue to change for centuries to millennia because of the slow response of the deep ocean and the ice sheets, so that “stabilization” of surface temperature does not imply stabilization of sea level, ice mass, or the deep ocean (*high confidence*). The concept of committed change is central to interpreting long-term projections correctly. Surface air temperature responds to cumulative emissions on a timescale of decades and is, for the atmosphere, largely reversible; but the quantities that integrate the warming over time — ocean heat content, thermosteric sea level, ice-sheet mass, permafrost carbon — carry a memory of the warming that persists long after the surface stabilizes. A world that holds temperature constant at any level is a world in which sea level continues to rise, ice sheets continue to lose mass, and the deep ocean continues to warm and acidify, for centuries. This committed change is the reason the sea-level and ice-sheet projections extend to millennial timescales and the reason peak warming, rather than any subsequent stabilized level, governs much of the long-term outcome.

The corollary for overshoot pathways is that the reversibility of surface temperature does not confer reversibility on the committed components. Bringing temperature back down after an overshoot through net-negative emissions reverses the atmospheric warming but does not reverse the sea-level rise, the ice-sheet retreat, or the ecosystem transitions committed during the peak (*high confidence* in the asymmetry). The slow-responding components thus impose a penalty on overshoot that the fast-responding atmosphere does not reveal, and they make the limitation of peak warming — not merely the eventual return to a target — the physically meaningful objective for avoiding the largest long-term consequences.

Projection beyond 2100

The changes projected for 2100 are a waypoint, not an endpoint; under continued high emissions the Earth system continues to warm and transform well beyond the century, and even under strong mitigation the committed changes in sea level and the cryosphere continue to unfold for centuries to millennia (*high confidence*). The scenario framework that underpins the century-scale projections extends, in its extended and idealized forms, to multi-century horizons, and the coupled models run under these extensions show that the divergence between high- and low-emission futures widens dramatically beyond 2100 (Tebaldi et al., 2021). Under sustained high emissions the warming continues to climb, the ice sheets move further into self-sustaining decline, and sea level rises by many metres over subsequent centuries; under strong mitigation the surface temperature stabilizes but the committed sea-level and ice-sheet changes continue. The deep-time projection therefore reinforces the central message of the warming-level framing: the emissions of the coming decades set a trajectory whose consequences extend far beyond the lives of those who emit, and the difference between emission pathways, modest at 2100, becomes vast over the centuries that follow.

6.8 Model Evaluation, Constrained Projections, and Confidence

The evidentiary basis for confidence in projection

Confidence in long-term projections is built from the convergence of multiple, partly independent lines of evidence — the demonstrated skill of process-based models, the physical understanding of the mechanisms, the attribution of observed change to the same drivers that force the projections, and the narrowing of ranges by observational constraints — rather than from any single model or method (*high confidence*). The strongest projections are those for which all of these lines align. Extreme heat is the clearest example: the models that project future intensification reproduce the observed climatology, the physical mechanism is transparent, the observed intensification is attributed to human influence, and the projected future is a continuation of the attributed trend (Clarke et al., 2022). Where the lines of evidence align in this way, confidence is high even for large projected changes; where they diverge — as in the regional detail of precipitation and circulation, which depends on less well-constrained dynamical responses — confidence is correspondingly lower. The calibrated-language framework used throughout this chapter is the vehicle for communicating this graded confidence honestly.

Model evaluation against observations is the foundation of this evidentiary structure. The coupled models are assessed against reanalyses and observational datasets for their representation of the mean climate, the seasonal cycle, the

modes of internal variability, and the historical trend and its spatial fingerprint. High-resolution ocean and coupled reanalyses provide the reference state for evaluating the ocean components that dominate heat and carbon uptake and that govern the committed changes in sea level (Jean-Michel et al., 2021). A model that reproduces the observed mean state and variability, and that reproduces the observed historical change for the right physical reasons, earns confidence in its projection; a model with large biases in the relevant processes is downweighted. This evaluation is not a pass–fail gate but a continuous weighting of the ensemble by demonstrated fidelity to the aspects of the climate most relevant to each projected quantity.

Observational constraints and emergent constraints

Observationally based constraints, including emergent constraints that exploit ensemble-wide relationships between an observable present-day quantity and an uncertain future response, meaningfully narrow the projected ranges relative to the raw ensemble spread, though they must be validated against the underlying mechanism to guard against spurious statistical relationships (*high confidence* that the ensemble of constraints reduces the effective spread; *medium confidence* in any single constraint). The logic of an emergent constraint is that if the models disagree about the future but agree that the future response is tied to some feature of the present climate that can be measured, then measuring that feature constrains the future. Applied to carbon-cycle feedbacks, snow-albedo feedback, cloud responses, and regional hydroclimate, these methods have tightened projected ranges and, in some cases, shifted central estimates. The danger they carry is that a correlation across the ensemble may be an artifact of shared model structure rather than a genuine physical relationship; the discipline that guards against this is the requirement that every constraint be traceable to a physical mechanism, not merely to a statistical fit.

The carbon cycle is a domain where observational constraint is especially valuable because the carbon-climate feedbacks are both large and uncertain. The soil-moisture–atmosphere coupling that dominates the interannual variability of land carbon uptake is an observable, constrainable process, and its representation is now a standard benchmark against which the land-surface components of the models are evaluated and weighted (Humphrey et al., 2021). The annual carbon-budget assessments provide the observational time series of the atmospheric growth rate and the sink partitioning against which the coupled carbon-cycle models are tested, tightening the link between a given emission pathway and the resulting atmospheric concentration and warming (Friedlingstein et al., 2023) (Friedlingstein et al., 2025). The net effect of the constraint enterprise is a projected range that is narrower and better justified than the unweighted ensemble would provide, without pretending to a precision the physics does not support.

Machine learning in the projection enterprise

Machine-learning weather and climate emulators have matured rapidly and now complement process-based models for specific tasks — medium-range forecasting, downscaling, large-ensemble generation, and the acceleration of scenario exploration — but they do not replace physical models for projecting the response to novel forcing, which by construction lies outside their training distribution (*high confidence*). Data-driven global weather models trained on decades of reanalysis now match or exceed the accuracy of established physics-based forecasting systems at a small fraction of the computational cost, a development that has reshaped operational weather prediction and opened new possibilities for climate applications (Bi et al., 2023). For long-term projection their value lies in efficiency: they can generate the large initial-condition ensembles needed to separate the forced signal from internal variability, downscale coarse projections to the regional scale where impacts are felt, and rapidly explore the space of scenarios, all far more cheaply than running the full coupled models.

The essential limitation is that these emulators learn the statistics of the climate they are trained on and cannot be assumed to extrapolate reliably to forcing regimes outside that training range — precisely the regimes that high-emission projections explore. A model that has never seen a 4-degree-C world in its training data has no guaranteed basis for simulating one. For this reason the process-based general circulation models, grounded in the conservation laws of physics, remain the authority on the response to novel forcing, and the machine-learning emulators serve as accelerators and downscalers within the envelope the physical models define. The productive combination is a hybrid one, in which physical models anchor the response to forcing and data-driven methods extend, downscale, and sample it efficiently. The assessment treats this hybridization as a promising and rapidly evolving direction rather than as a mature replacement for the physical modelling that remains the backbone of projection.

Synthesis

The long-term global projection that emerges from this assessment is robust in its central messages and honest about its uncertainties: warming continues until emissions reach net zero, the level reached is set by cumulative carbon dioxide, every increment of warming intensifies hazards across the atmosphere, ocean, cryosphere, and biosphere, and a set of slow-responding and potentially threshold-crossing components commit the Earth system to changes that extend for centuries beyond 2100 (*high confidence*). The most robust, highest-confidence elements are the continuation of warming to mid-century, the dominance of cumulative emissions in setting peak warming, the amplification of warming in the polar and high-mountain regions, the intensification of heat extremes, the rise in the intensity of heavy precipitation alongside the expansion of drought, the continued rise of sea level, and the continued

warming, acidification, deoxygenation, and marine-heatwave intensification of the ocean. These conclusions rest on the convergence of model skill, physical understanding, attribution of observed change, and observational constraint, and they are stated with high confidence.

The deepest uncertainties are concentrated in a small number of high-consequence components: the dynamic response of the marine-based Antarctic Ice Sheet and the resulting high-end tail of sea level, the possibility of an abrupt weakening of the Atlantic overturning circulation, the strength and stability of the land and ocean carbon sinks and the associated carbon-cycle feedbacks, the risk of self-reinforcing dieback in the Amazon, and the regional detail of precipitation and circulation change. These are not reasons to doubt the central projection; they are the specific quantities where the assessment reports lower confidence and, for the tail risks, where it deliberately assesses low-likelihood, high-impact outcomes rather than omitting them. The organizing insight of the chapter is that warming-level framing renders the fast-responding hazards robust to scenario uncertainty while the slow-responding, committed components make peak warming the decisive variable for the long-term outcome. Taken together, the projections establish that the choices of the coming decades set the trajectory of the Earth system for centuries, and that limiting cumulative emissions and peak warming is the single most effective lever for limiting the full range of assessed long-term change (Rounce et al., 2023) (Tebaldi et al., 2021).

6.9 Compound Hazards, Cascading Risk, and the Continuity of Projection

Compound and concurrent extremes

The most consequential climate risks of the coming century arise not from single hazards in isolation but from the compounding and concurrence of multiple hazards in space and time, and the probability of such compound events rises faster with warming than the probability of any single component (*high confidence*). A warming climate raises the joint probability that heat, drought, and fire occur together, that extreme precipitation and high sea level coincide to produce compound coastal-and-fluvial flooding, and that successive extremes strike the same region before it can recover. The intensification of heat extremes and the expansion of drought, both assessed above, are physically linked through the shared driver of rising temperature and evaporative demand, so that their co-occurrence is not a coincidence to be estimated from independent probabilities but a coupled response to be projected together. The flood hazard illustrates the compounding of drivers: inequitable and rising flood exposure is projected as extreme precipitation intensifies and sea level rises, with the two contributions combining along low-lying coasts (Wing et al., 2022).

The assessment of compound risk requires moving beyond univariate projection to the joint distribution of hazards, and it is in this joint distribution that the steepest nonlinearities — and therefore the sharpest sensitivities to warming level — reside. Because many systems, human and natural, are resilient to individual shocks but vulnerable to their coincidence or succession, the compound framing often reveals risks that a hazard-by-hazard assessment would miss. The projected increase in the amplitude and extreme-event frequency of the El Niño–Southern Oscillation is significant in this light, because ENSO orchestrates the concurrence of drought, flood, heat, and fire risk across widely separated regions, so that a change in its behaviour propagates into a change in the correlation structure of global climate risk (Cai et al., 2021).

Cascading and systemic risk

Climate hazards can cascade through interconnected human and natural systems, so that a physical change of modest magnitude can propagate into a large societal impact through the coupling of food, water, health, ecological, and economic systems (*high confidence that cascading amplification occurs; medium confidence in its magnitude for any specific chain*). The interaction of the climate and biodiversity crises is a paradigmatic cascade: climate change degrades ecosystems, degraded ecosystems lose the capacity to buffer climate impacts and to provide the services on which human societies depend, and the two crises reinforce one another as coupled threats that neither addresses in isolation (Pörtner et al., 2023). Climate change also acts as an aggravating factor across a majority of known human pathogenic diseases, expanding the geographic and seasonal windows of transmission and compounding the direct impacts of heat, drought, and extreme events on human health (Mora et al., 2022). These cascades link the physical projections of this chapter to the human consequences that give them their significance.

The energy and infrastructure systems on which adaptation depends are themselves exposed to the changing climate, so that the systems meant to buffer climate impacts are subject to the hazards they must buffer. The cost, environmental performance, and resilience of renewable-energy systems are affected by a changing climate through altered wind, solar, and hydrological resources and through exposure to extreme events, a coupling that must be accounted for in projecting the resilience of the low-carbon transition itself (Osman et al., 2022). The systemic framing underscores that long-term projection is not merely a forecast of physical variables but an input to the assessment of coupled human–Earth-system risk, in which physical hazards, ecological degradation, and societal vulnerability combine.

The continuity from near-term to long-term

Long-term projections are continuous with the near-term changes already observed and attributed, so that the century-scale future is a physically grounded extrapolation of a detected present rather than a speculative departure from

it (*high confidence*). The strongest single argument for confidence in long-term projection is this continuity. The Arctic is already warming at close to four times the global rate, exactly as amplification mechanisms predict and projections extend (Rantanen et al., 2022). Drought characteristics already bear a detectable human fingerprint (Chiang et al., 2021), heat extremes are already attributable intensified (Clarke et al., 2022), glaciers are already in committed decline (Rounce et al., 2023), and the ocean is already warming, acidifying, and experiencing intensifying marine heatwaves (Smith et al., 2022). The long-term projections do not posit new mechanisms; they extend, under specified forcing, the mechanisms whose signatures are already present in the observational record and attributed to human influence.

This continuity has a practical corollary for the interpretation of the projections. Because the projected future is an extension of an attributed present, the uncertainties in the projections are uncertainties of magnitude and timing, not of direction or mechanism, for the high-confidence conclusions. The direction of change — more warming until net zero, higher heat extremes, more intense heavy precipitation, expanding drought, rising sea level, a warming and acidifying ocean — is established by the convergence of attributed observation, physical understanding, and model projection. It is against this settled direction that the remaining uncertainties in magnitude, in regional detail, and in the timing and likelihood of threshold crossings should be understood: as the frontier of the assessment, not as a reason to doubt its foundation. The overarching conclusion of this chapter is that the choices governing cumulative emissions over the coming decades will determine, with high confidence in direction and quantified uncertainty in magnitude, the trajectory of the global climate system through 2100 and, through the committed and threshold-sensitive components, for the centuries and millennia beyond (Tebaldi et al., 2021) (Friedlingstein et al., 2025).

Representative regional storylines

Regional storylines make the global projections concrete by tracing how the robust global signals combine with regional geography and vulnerability to produce distinct, high-consequence futures, and they illustrate that the same global warming level is experienced very differently across regions (*high confidence*). The Eastern Mediterranean and Middle East exemplify a hot-and-drying storyline: a region already warming faster than the global mean, projected to experience among the most severe increases in heat-extreme frequency and the emergence of previously unprecedented conditions, superimposed on a drying trend and acute water scarcity (Zittis et al., 2022). Southwestern North America exemplifies an aridification storyline in which the warming-driven rise in evaporative demand converts natural drought variability into extreme, multi-decadal megadrought, drawing down the groundwater and surface-water reserves on which the region depends (Williams et al., 2022) (Scanlon et al., 2023).

The Amazon basin exemplifies a threshold-risk storyline, in which warming,

drought, deforestation, and fire interact to push a globally significant carbon reservoir and biodiversity stronghold toward a possible self-reinforcing transition, with the degradation already substantial and the critical thresholds within reach under continued pressure (Lapola et al., 2023) (Flores et al., 2024). The high-mountain regions exemplify a water-tower storyline, in which elevation-dependent warming accelerates the loss of the glaciers and snowpack that buffer dry-season flow, producing a transient pulse of enhanced melt followed by an irreversible decline in water availability for the populations downstream (Pepin et al., 2022) (Rounce et al., 2023). These storylines are not exhaustive, but they demonstrate the central lesson of regional projection: the global signals assessed in this chapter are robust, but their human meaning is set by how they combine with the geography, hydrology, ecology, and vulnerability of each place. The global projection provides the forcing; the regional storyline provides the consequence.

6.10 Key Uncertainties and Research Priorities

The assessment identifies a compact set of key uncertainties whose resolution would most improve long-term projection, and these define the research priorities for the coming assessment cycle (*high confidence that these are the leading uncertainties*). First among them is the dynamic response of the marine-based Antarctic Ice Sheet, which governs the high-end tail of sea-level rise and whose thresholds and timescales remain poorly constrained; narrowing this uncertainty is the single most valuable target for sharpening the long-term sea-level projection (Rounce et al., 2023). Second is the strength and stability of the land and ocean carbon sinks under continued warming, including the carbon-cycle feedbacks — permafrost thaw, soil-moisture-limited productivity, drought- and fire-driven forest loss — that determine how much of a given emission pathway remains in the atmosphere; the annual carbon-budget accounting provides the observational spine for this effort, but the future trajectory of the sinks remains a leading source of spread in projected warming (Friedlingstein et al., 2023) (Humphrey et al., 2021).

Third is the likelihood and character of abrupt transitions in the Atlantic overturning circulation and in threshold-sensitive biosphere components such as the Amazon, where the physical possibility of self-reinforcing change is established but the thresholds and probabilities carry deep uncertainty (Flores et al., 2024). Fourth is the regional detail of precipitation and atmospheric circulation change, which depends on dynamical responses that models simulate less consistently than thermodynamics and which is the primary limit on the confidence of regional projections. Fifth is the changing behaviour of the modes of internal variability, particularly the El Niño–Southern Oscillation, whose response shapes the correlation structure of global climate risk (Cai et al., 2021). Progress on these five fronts — through improved process representation, observational constraint, higher-resolution and hybrid modelling, and paleoclimate evidence

on threshold behaviour — would most tighten the long-term projections while leaving intact their high-confidence central messages.

6.11 References

Abram, N. J., et al. (2021). Connections of climate change and variability to large and extreme forest fires in southeast Australia. *Communications Earth & Environment*. <https://doi.org/10.1038/s43247-020-00065-8>

Bi, K., et al. (2023). Accurate medium-range global weather forecasting with 3D neural networks. *Nature*. <https://doi.org/10.1038/s41586-023-06185-3>

Cai, W., et al. (2021). Changing El Niño–Southern Oscillation in a warming climate. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-021-00199-z>

Canadell, J. G., et al. (2021). Multi-decadal increase of forest burned area in Australia is linked to climate change. *Nature Communications*. <https://doi.org/10.1038/s41467-021-27225-4>

Chiang, F., et al. (2021). Evidence of anthropogenic impacts on global drought frequency, duration, and intensity. *Nature Communications*. <https://doi.org/10.1038/s41467-021-22314-w>

Clarke, B., et al. (2022). Extreme weather impacts of climate change: an attribution perspective. *Environmental Research Climate*. <https://doi.org/10.1088/2752-5295/ac6e7d>

Flores, B. M., et al. (2024). Critical transitions in the Amazon forest system. *Nature*. <https://doi.org/10.1038/s41586-023-06970-0>

Friedlingstein, P., et al. (2022). Global Carbon Budget 2022. *Earth system science data*. <https://doi.org/10.5194/essd-14-4811-2022>

Friedlingstein, P., et al. (2023). Global Carbon Budget 2023. *Earth system science data*. <https://doi.org/10.5194/essd-15-5301-2023>

Friedlingstein, P., et al. (2025). Global Carbon Budget 2024. *Earth system science data*. <https://doi.org/10.5194/essd-17-965-2025>

Hammond, W. M., et al. (2022). Global field observations of tree die-off reveal hotter-drought fingerprint for Earth's forests. *Nature Communications*. <https://doi.org/10.1038/s41467-022-29289-2>

Hassani, A., et al. (2021). Global predictions of primary soil salinization under changing climate in the 21st century. *Nature Communications*. <https://doi.org/10.1038/s41467-021-26907-3>

Humphrey, V., et al. (2021). Soil moisture–atmosphere feedback dominates land carbon uptake variability. *Nature*. <https://doi.org/10.1038/s41586-021-03325-5>

Jean-Michel, L., et al. (2021). The Copernicus Global 1/12° Oceanic and Sea Ice GLORYS12 Reanalysis. *Frontiers in Earth Science*. <https://doi.org/10.3389/feart.2021.698876>

Jägermeyr, J., et al. (2021). Climate impacts on global agriculture emerge earlier in new generation of climate and crop models. *Nature Food*. <https://doi.org/10.1038/s43016-021-00400-y>

- Kemp, L., et al. (2022). Climate Endgame: Exploring catastrophic climate change scenarios. *Proceedings of the National Academy of Sciences*. <https://doi.org/10.1073/pnas.2201111119>
- Lapola, D. M., et al. (2023). The drivers and impacts of Amazon forest degradation. *Science*. <https://doi.org/10.1126/science.abp8622>
- McDowell, N. G., et al. (2022). Mechanisms of woody-plant mortality under rising drought, CO₂ and vapour pressure deficit. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-022-00272-1>
- Mora, C., et al. (2022). Over half of known human pathogenic diseases can be aggravated by climate change. *Nature Climate Change*. <https://doi.org/10.1038/s41558-022-01426-1>
- Osman, A. I., et al. (2022). Cost, environmental impact, and resilience of renewable energy under a changing climate: a review. *Environmental Chemistry Letters*. <https://doi.org/10.1007/s10311-022-01532-8>
- Pepin, N. C., et al. (2022). Climate Changes and Their Elevational Patterns in the Mountains of the World. *Reviews of Geophysics*. <https://doi.org/10.1029/2020rg000730>
- Pörtner, H., et al. (2023). Overcoming the coupled climate and biodiversity crises and their societal impacts. *Science*. <https://doi.org/10.1126/science.abl4881>
- Rantanen, M., et al. (2022). The Arctic has warmed nearly four times faster than the globe since 1979. *Communications Earth & Environment*. <https://doi.org/10.1038/s43247-022-00498-3>
- Richardson, K., et al. (2023). Earth beyond six of nine planetary boundaries. *Science Advances*. <https://doi.org/10.1126/sciadv.adh2458>
- Rockström, J., et al. (2023). Safe and just Earth system boundaries. *Nature*. <https://doi.org/10.1038/s41586-023-06083-8>
- Rounce, D. R., et al. (2023). Global glacier change in the 21st century: Every increase in temperature matters. *Science*. <https://doi.org/10.1126/science.abo1324>
- Scanlon, B. R., et al. (2023). Global water resources and the role of groundwater in a resilient water future. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-022-00378-6>
- Smith, K. E., et al. (2022). Biological Impacts of Marine Heatwaves. *Annual Review of Marine Science*. <https://doi.org/10.1146/annurev-marine-032122-121437>
- Tebaldi, C., et al. (2021). Climate model projections from the Scenario Model Intercomparison Project (ScenarioMIP) of CMIP6. *Earth System Dynamics*. <https://doi.org/10.5194/esd-12-253-2021>
- Williams, P., et al. (2022). Rapid intensification of the emerging southwestern North American megadrought in 2020–2021. *Nature Climate Change*. <https://doi.org/10.1038/s41558-022-01290-z>
- Wing, O., et al. (2022). Inequitable patterns of US flood risk in the Anthropocene. *Nature Climate Change*. <https://doi.org/10.1038/s41558-021-01265-6>
- Winkler, K., et al. (2021). Global land use changes are four times greater than previously estimated. *Nature Communications*. <https://doi.org/10.1038/s41467-021-22702-2>
- Yang, Y., et al. (2024). Climate change exacerbates the environmental impacts of agriculture. *Science*. <https://doi.org/10.1126/science.adn3747>

Yuan, X., et al. (2023). A global transition to flash droughts under climate change. *Science*. <https://doi.org/10.1126/science.abn6301>

Zittis, G., et al. (2022). Climate Change and Weather Extremes in the Eastern Mediterranean and Middle East. *Reviews of Geophysics*. <https://doi.org/10.1029/2021rg000762>

Chapter 7

Chapter 7: Regional Projections and Linking Global to Regional

7.1 Executive Summary

Regional climate change is now detectable across every inhabited continent, and the pattern of regional response is more robustly established than at any previous assessment (*high confidence*). The dominant large-scale features of the warming pattern — amplified warming over land relative to ocean, amplified warming at high northern latitudes, and elevation-dependent warming in the world’s mountain systems — are consistently simulated across model generations and corroborated by observations (Rantanen et al., 2022; Pepin et al., 2022). The task of linking global-mean change to actionable regional information is the central methodological challenge addressed in this chapter, and the tools for that linkage — dynamical downscaling, statistical downscaling, machine-learning emulation, and structured distillation — have matured substantially since the previous assessment (*high confidence*).

The Arctic has warmed nearly four times faster than the global mean since 1979, and this amplification is a fundamental organizing feature of the Northern Hemisphere regional response (*virtually certain*). Observational reanalysis and station records agree that Arctic amplification exceeds the factor-of-two figure long used as a heuristic (Rantanen et al., 2022). Elevation-dependent warming further reorganizes the vertical structure of change in mountain regions, with consequences for the cryosphere, hydrology, and downstream water supply (Pepin et al., 2022) (Rounce et al., 2023).

Dynamical downscaling with convection-permitting regional climate models has emerged as the highest-fidelity approach for representing regional extremes, while machine-learning emulators now offer a computationally tractable route to large ensembles of regional information (*high confidence*). The scenario framework underpinning all regional projection — the Shared Socioeconomic Pathway and forcing combinations coordinated through ScenarioMIP — provides the common boundary conditions that make cross-method comparison meaningful (Tebaldi et al., 2021).

Regional projections of the water cycle show intensification of both wet

and dry extremes, with flash drought onset accelerating and drought frequency, duration, and intensity showing detectable anthropogenic signals (*high confidence*). A global transition toward faster-developing droughts is projected under continued warming (Yuan et al., 2023), and formal attribution links the observed changes in drought characteristics to human influence (Chiang et al., 2021). Flood and drought prediction remain among the hardest regional problems because they integrate across scales and depend sensitively on land-surface and human-management processes (Brunner et al., 2021).

Regional heat extremes are intensifying faster than mean warming in several mid-latitude regions, driven in part by changes in atmospheric circulation that are imperfectly captured by coarse-resolution models (*high confidence*). Accelerated western European heatwave trends have been linked to more-persistent double-jet configurations over Eurasia (Rousi et al., 2022), and the 2021 Pacific Northwest event exceeded prior statistical expectations by a margin that stressed conventional extreme-value frameworks (White et al., 2023). The emerging southwestern North American megadrought represents a regional hydroclimate state with few analogues in the instrumental or paleoclimate record (Williams et al., 2022).

Storylines and physically self-consistent narrative approaches have become a standard complement to probabilistic projection for communicating deeply uncertain, high-impact regional outcomes (*high confidence*). Complex, compounding, and cascading risks are not well served by marginal probability distributions alone, and structured frameworks for representing interacting drivers are now central to regional risk assessment (Simpson et al., 2021). Interacting tipping elements can raise the risk of cascading regional transitions beyond what any single element would imply (Wunderling et al., 2021), and low-likelihood high-impact outcomes warrant explicit treatment rather than exclusion (Kemp et al., 2022).

Distillation of regional climate information — the expert-mediated synthesis of multiple, sometimes conflicting, lines of evidence into fit-for-purpose products — is now recognized as an essential step between raw model output and decision-relevant knowledge (*high confidence*). No single model, method, or ensemble is authoritative for a given region; credible regional information emerges from the reconciliation of dynamical, statistical, observational, and process-understanding evidence.

7.2 Introduction: The Regionalization Problem

The global-mean surface temperature is a convenient scalar summary of planetary energy imbalance, but no community, ecosystem, or economy experiences the global mean. People experience the regional and local expression of change — the frequency of dangerous heat, the seasonality of rainfall, the height of the storm surge, the timing of snowmelt. This chapter addresses the scientific problem of translating the well-constrained global signal into the regional detail

that adaptation and risk management require, and it assesses the methods, their evaluation, and the resulting regional projections.

Two facts frame the problem. First, the global signal is comparatively well constrained: the sign and approximate magnitude of global warming per unit cumulative CO₂ emission are robust, and the global carbon budget that governs the total forcing is quantified to within assessed uncertainty each year (Friedlingstein et al., 2023) (Friedlingstein et al., 2025). Second, the regional signal is comparatively poorly constrained: regional precipitation change, in particular, carries large model spread arising from differences in the simulation of atmospheric circulation, land-surface feedbacks, and unresolved convection. The regionalization problem is therefore not merely one of spatial refinement but of adding physically credible information that the global models cannot, by construction, provide.

The regional response to forcing is organized by a small number of large-scale patterns that recur across models and are grounded in physical understanding. Land warms faster than ocean because of differences in heat capacity and evaporative constraint. High northern latitudes warm faster than the tropics through the combined action of surface-albedo feedback, changes in poleward heat and moisture transport, lapse-rate feedback, and reduced longwave cooling efficiency — a suite of processes that together produce Arctic amplification approaching a factor of four relative to the global mean over the satellite era (Rantanen et al., 2022). Mountains warm in an elevation-dependent manner, with the vertical profile of warming shaped by snow-albedo feedback, changes in the freezing level, cloud and water-vapour feedbacks, and aerosol effects, such that high-elevation environments can warm faster than adjacent lowlands (Pepin et al., 2022). These organizing patterns are the scaffolding on which regional projection is built.

The scenario architecture that supplies boundary conditions for all regional work is coordinated internationally. The Scenario Model Intercomparison Project defines a set of forcing pathways spanning very low to very high emissions, combining Shared Socioeconomic Pathways with radiative-forcing targets, so that regional downscaling and impact studies share a common experimental design (Tebaldi et al., 2021). This coordination is what allows a convection-permitting simulation over the Alps, a statistical downscaling over the Sahel, and a machine-learning emulator over East Asia to be compared as estimates of the same forced future rather than as incommensurable exercises. The individual Earth system models that provide the driving fields — for example, the CMCC Earth System Model contribution to the sixth phase of the Coupled Model Intercomparison Project — are documented and evaluated so that the provenance of any regional product can be traced to a characterized global simulation (Lovato et al., 2022).

Regional projection is inseparable from the observational infrastructure used to evaluate it. Global reanalyses, high-resolution land datasets, and specialized products for soils, elevation, and tides supply the reference against which regional models are judged and, increasingly, the training data for machine-learning methods. Land-surface parameter datasets such as global soil-property

grids (Poggio et al., 2021), hydrological model frameworks (Schmied et al., 2021), forest-and-building-removed terrain models (Hawker et al., 2022), and global tide atlases (Lyard et al., 2021) are foundational inputs to regional modelling of the water cycle, coastal flooding, and land-atmosphere exchange. The credibility of a regional projection depends as much on the quality of these inputs as on the resolution of the model.

This chapter proceeds as follows. Section 7.2 assesses downscaling methods — dynamical, statistical, and machine-learning — and their respective strengths, limitations, and appropriate domains of application. Section 7.3 assesses how regional models are evaluated, including the observational references and process-based diagnostics that establish fitness for purpose. Section 7.4 assesses the distillation of regional climate information, the synthesis step that reconciles multiple lines of evidence into decision-relevant products. Section 7.5 assesses regional projections organized by continental region, integrating the physical drivers, method-specific findings, and key uncertainties for each. Section 7.6 assesses storylines and the broader problem of communicating deep uncertainty, including the treatment of low-likelihood high-impact outcomes and interacting tipping elements.

7.3 Downscaling Methods

Dynamical downscaling and convection-permitting simulation

Dynamical downscaling with regional climate models, and in particular the transition to convection-permitting resolutions, provides the most physically complete representation of regional climate and is the reference method for regional extremes (*high confidence*). A regional climate model solves the governing equations of the atmosphere (and, when coupled, of the land, ocean, and cryosphere) on a limited-area grid at resolution far finer than the driving global model, taking its lateral boundary conditions from a global simulation or reanalysis. The added value of dynamical downscaling is greatest where fine-scale surface forcing — complex topography, coastlines, land-cover heterogeneity, urban surfaces — organizes the regional climate in ways the coarse driving model cannot resolve. Mountain precipitation, land-sea breeze systems, and orographically anchored convection are canonical examples where the finer grid materially changes the simulated climate.

The most consequential recent advance is the move to convection-permitting resolutions, at which the model grid is fine enough (typically a few kilometres) to represent deep convective updraughts explicitly rather than through a convective parameterization. Explicit convection improves the diurnal cycle of summer precipitation, the intensity distribution of sub-daily rainfall, and the representation of mesoscale convective organization — precisely the features that govern flash-flood and extreme-precipitation risk. This is directly relevant to the hardest regional water-cycle problems, where the parameterization of

convection is a leading source of error in flood and drought simulation (Brunner et al., 2021). Convection-permitting simulation is computationally expensive, which historically limited it to short time slices or small ensembles; the resulting tension between resolution and ensemble size is a recurring theme in the interpretation of dynamical results.

Dynamical downscaling inherits both the strengths and the errors of its driving model. Biases in the large-scale circulation of the global model — for example in the position and persistence of jets and storm tracks — propagate through the lateral boundaries into the regional domain, and no amount of interior resolution can fully correct them. This inheritance is central to the interpretation of mid-latitude extremes, where circulation errors dominate the regional response. The observed acceleration of western European heatwaves has been linked to more-persistent double-jet states over Eurasia (Rousi et al., 2022); a regional model driven by a global model that under-represents such persistence will under-represent the heat extreme regardless of its own resolution. This is why dynamical downscaling is assessed as necessary but not sufficient: it adds fine-scale value conditional on a credible large-scale driver.

Statistical downscaling and bias adjustment

Statistical downscaling and bias adjustment provide computationally efficient regional information and are well suited to large ensembles and impact applications, provided the stationarity of the statistical relationships is critically examined (*high confidence*). Statistical downscaling establishes empirical relationships between large-scale predictors (from a global model or reanalysis) and local-scale predictands (station or gridded observations), then applies those relationships to projected large-scale fields. Methods range from simple transfer functions and analogue methods to distributional bias-correction and quantile-mapping schemes. Their principal advantage is cost: once trained, they can be applied to every member of a large global ensemble and to many scenarios, delivering the sampling of uncertainty that expensive dynamical runs cannot.

The foundational assumption of statistical downscaling is stationarity — that the empirical predictor–predictand relationship established over the observed period continues to hold under a changed climate. This assumption is most defensible where the local variable is tightly slaved to a well-resolved large-scale predictor (for instance, near-surface temperature to free-tropospheric temperature) and least defensible where local processes with their own climate sensitivity intervene (for instance, convective precipitation, or snow processes near the freezing point). Bias adjustment carries a related risk: correcting a model’s distribution toward observations can mask, rather than remedy, errors in the underlying physical processes, and can distort the trend if the bias is state-dependent. These limitations are acute for the water cycle, where land-surface feedbacks and human water management complicate the predictor–predictand mapping (Schmied et al., 2021; Brunner et al., 2021).

Statistical methods depend critically on the quality and homogeneity of the observational reference used for training. Reanalysis products and high-resolution observational datasets supply these references, and their uncertainties propagate directly into the downscaled product. For land-surface and hydrological variables, harmonized global datasets of soil properties (Poggio et al., 2021) and terrain (Hawker et al., 2022) constrain the boundary conditions, while for coastal applications global tide models provide the deterministic component of sea level that must be separated from the climate-driven change (Lyard et al., 2021). The maturity of these reference datasets is a quiet but decisive enabler of credible statistical downscaling.

Machine-learning emulators and hybrid methods

Machine-learning methods now provide fast, skilful emulation of both dynamical downscaling and full climate-model behaviour, enabling large-ensemble regional information at a fraction of the computational cost, while raising new questions about physical consistency and out-of-sample reliability (*high confidence*). The application of machine learning to climate problems has expanded rapidly, spanning emulation, downscaling, bias correction, parameterization, and the extraction of predictable structure from observations (Kaack et al., 2022). In the regional context, machine-learning models are trained to map coarse-resolution fields to fine-resolution targets — either observed high-resolution fields or the output of expensive convection-permitting simulations — learning the statistical relationship between scales directly from data.

The appeal is computational. A trained emulator can downscale an entire global ensemble across all scenarios in the time a single dynamical run would take, closing the resolution-versus-ensemble-size tension that constrains dynamical methods. This makes machine learning particularly attractive for applications that require large samples: estimating the tails of extreme distributions, propagating scenario and internal-variability uncertainty into impact models, and generating the high-dimensional regional information that adaptation planning increasingly demands. Machine-learning approaches have shown particular promise in data-rich, process-complex domains such as hydrology, where they have been used to predict variables like groundwater levels that depend on many interacting drivers (Hai et al., 2022).

The concerns are physical. A purely data-driven emulator has no built-in guarantee of conservation, of physical consistency across variables, or of reliability outside the range of its training data — a serious limitation for climate projection, whose entire purpose is to extrapolate to unobserved states. The most credible regional machine-learning work therefore tends toward hybrid designs that embed physical constraints, that use dynamical output as training targets so that the physics is inherited, and that are evaluated against held-out extremes rather than only against mean behaviour. Machine learning is assessed here as a powerful accelerant that does not replace physical modelling but redistributes where the physics enters — from the forward simulation to

the training data and the architectural constraints (Kaack et al., 2022; Hai et al., 2022).

Choosing and combining methods

No single downscaling method dominates across all regions, variables, and applications. Dynamical downscaling is preferred where fine-scale physical processes generate the signal and where extremes matter most, but it is expensive and inherits driving-model error. Statistical downscaling and bias adjustment are preferred where large ensembles and computational efficiency matter, but they rest on stationarity assumptions that must be tested. Machine-learning emulation increasingly bridges the two, delivering dynamical-like fine-scale structure at statistical-like cost, but it demands careful evaluation of physical consistency and out-of-sample behaviour. Best practice combines methods — using dynamical output to train and validate emulators, using statistical methods to expand ensembles, and using the spread across methods as a component of the uncertainty estimate. The reconciliation of these method-specific estimates is the subject of the distillation step assessed in Section 7.4.

7.4 Regional Model Evaluation

Observational references for regional evaluation

The evaluation of regional models rests on a maturing but heterogeneous observational infrastructure, and the fitness of any regional projection is bounded by the quality and coverage of the reference data used to assess it (*high confidence*). Regional model evaluation compares simulated climate against observations across many variables, timescales, and statistics — not only the mean state but the seasonal cycle, interannual variability, the frequency and intensity of extremes, and the fidelity of the physical processes that generate them. The reference datasets that make this possible have improved substantially, but they remain unevenly distributed in space and time, with dense coverage in well-instrumented regions and sparse coverage in exactly the tropical and high-latitude regions where model uncertainty is largest.

Reanalyses provide the spatially and temporally complete reference that station networks cannot, assimilating heterogeneous observations into a dynamically consistent estimate of the atmospheric and land state. High-resolution land reanalyses are particularly valuable for evaluating the near-surface variables that matter for impacts, providing consistent fields of temperature, precipitation, soil moisture, and surface fluxes at a resolution approaching that of regional models. Specialized reference datasets extend evaluation to variables the atmosphere-focused reanalyses represent poorly: global soil-property grids constrain the land-surface lower boundary (Poggio et al., 2021), harmonized hydrological model output provides a reference for river discharge and water

storage where gauge data are sparse (Schmied et al., 2021), forest-and-building-removed terrain models sharpen the representation of true ground elevation for flood and coastal work (Hawker et al., 2022), and global tide atlases separate the astronomical tide from the climate-driven and surge components of coastal sea level (Lyard et al., 2021).

The uncertainty in the reference itself must be carried through evaluation. Where multiple observational products disagree — a common situation for precipitation over mountains, for soil moisture, and for surface fluxes — the observational spread can rival or exceed the model spread, and a model cannot be judged “wrong” against a reference that is itself poorly constrained. Honest regional evaluation therefore reports observational uncertainty alongside model error and is cautious about declaring skill or bias where the reference is weak. This is especially important in data-sparse regions, where the temptation to over-interpret a single reference product is strongest and where machine-learning methods trained on that single product inherit its errors most directly.

Process-based evaluation and added value

Regional models are increasingly evaluated on their representation of the physical processes that generate regional climate, not only on their output statistics, because process fidelity is what licenses extrapolation to a changed climate (*high confidence*). A model can reproduce the observed mean and variance for the wrong physical reasons — through compensating errors — and such a model provides no assurance that its projected change is credible. Process-based evaluation therefore interrogates whether the model represents the mechanisms correctly: whether orographic precipitation is anchored to the right terrain, whether the diurnal cycle of convection has the right timing and amplitude, whether land-atmosphere coupling has the right strength, and whether the large-scale circulation modes that organize regional variability are simulated with the right structure and persistence.

The evaluation of circulation is decisive for mid-latitude extremes. The persistence of blocking and of anomalous jet configurations governs the duration of heat and cold extremes, and errors in simulated persistence translate directly into errors in projected extremes. The linkage between accelerated western European heat and more-persistent double-jet states illustrates the point: evaluating a regional model’s heat projection requires evaluating the driving model’s jet dynamics, because the fine grid cannot manufacture circulation persistence that the boundary conditions do not supply (Rousi et al., 2022). Similarly, the exceptional 2021 Pacific Northwest heat event — which exceeded prior statistical expectations by a wide margin — exposed the limitations of evaluating models only against the observed range, and motivated process-based diagnostics of the ingredients (soil-moisture deficit, ridge amplitude, adiabatic warming) that combined to produce the event (White et al., 2023).

Added value — the demonstrable improvement a downscaled product provides over its driving model — is the practical criterion for whether downscaling is worthwhile for a given application. Added value is clearest for variables and

statistics tied to fine-scale surface forcing: sub-daily precipitation intensity, near-surface wind in complex terrain, temperature in mountain valleys and coastal zones. It is weakest, and can even be negative, for large-scale variables that the driving model already represents well and that the regional model may degrade by inheriting boundary errors while adding its own. Demonstrating added value requires evaluation at the scale at which the value is claimed, against a reference that resolves that scale — a requirement that circles back to the adequacy of the observational infrastructure.

Evaluation of extremes and compound events

The evaluation of regional models for compound and cascading events — where multiple drivers coincide or interact — is less mature than the evaluation of univariate extremes, and represents a leading frontier and a leading uncertainty (*high confidence*). Many of the most damaging regional outcomes are compound: hot-and-dry conditions that drive wildfire and crop failure, wet-and-windy conditions that drive coastal flooding, sequential events that exhaust response capacity before the next arrives. Evaluating a model's fidelity for such outcomes requires assessing the dependence structure between variables — the tendency for hot and dry, or wet and windy, to co-occur — not merely the marginal behaviour of each variable in isolation.

Regional models vary in their skill at representing these dependencies, and the observational record of compound events is short, making evaluation statistically difficult. The physical realism of compound-event simulation depends on correctly coupled land and atmosphere: soil-moisture depletion amplifying heat, vegetation stress feeding back on the surface energy budget, snowpack timing setting the stage for summer drought. The declining resilience of forests and the risk of abrupt ecosystem transitions under compound stress add a further layer, because a model that treats vegetation as static will misrepresent the feedbacks that intensify compound extremes (Forzieri et al., 2022; Boulton et al., 2022). The frequency with which known human pathogenic diseases can be aggravated by climatic hazards underscores that compound physical events cascade into compound impacts, widening the evaluation target from the physical climate to the coupled human–natural system (Mora et al., 2022).

Evaluation of the water cycle deserves particular emphasis because it integrates so many processes. The prediction of floods and droughts is a persistent challenge precisely because it requires the model to get right the precipitation forcing, the land-surface partitioning, the subsurface storage, and — for managed basins — the human operation of reservoirs and withdrawals (Brunner et al., 2021). A global transition toward flash droughts, which develop over weeks rather than seasons, sets a demanding evaluation target because it stresses the model's representation of fast land-atmosphere coupling and evaporative demand (Yuan et al., 2023). Detectable anthropogenic signals in drought frequency, duration, and intensity provide an additional evaluation constraint: a credible regional model should reproduce not only the climatology of drought but the emerging trend attributable to human influence (Chiang et al., 2021).

7.5 Distillation of Regional Climate Information

The distillation concept

Credible regional climate information is produced by distillation — the expert-mediated synthesis of multiple, often discordant, lines of evidence into fit-for-purpose products — and no single model or method should be treated as authoritative for a region (*high confidence*). The proliferation of methods assessed in Section 7.2, and the recognition that each carries characteristic strengths and errors, makes reconciliation unavoidable. Distillation is the process by which dynamical simulations, statistical downscaling, machine-learning emulation, observational analysis, and physical understanding are weighed against one another and combined into a coherent statement of what is known, what is uncertain, and what is decision-relevant for a specific region and application.

Distillation is necessary because the raw evidence rarely agrees. Different models can project opposite signs of regional precipitation change; different downscaling methods can amplify or damp the driving signal; observational references can disagree about the baseline. A user handed the full envelope of raw output faces an uninterpretable spread. The distillation step adds value by bringing physical reasoning to bear on that spread — asking which projections are supported by understood mechanisms, which reflect known model deficiencies, and which represent genuine, irreducible uncertainty. It is an expert judgment process, and its credibility depends on being transparent about the evidence and the reasoning rather than on presenting a single number with false precision.

The concept is closely tied to the framing of climate change as a complex risk problem. A framework for complex climate-change risk assessment recognizes that risk emerges from the interaction of hazard, exposure, vulnerability, and response, and that these interact across scales and sectors in ways that no single model captures (Simpson et al., 2021). Distillation of the physical climate information is the first stage of such an assessment, and it must be conducted with the downstream risk framing in view, so that the information distilled is the information the risk assessment actually needs.

Reconciling multiple lines of evidence

Robust regional conclusions rest on the convergence of independent lines of evidence — models, observations, process understanding, and paleoclimate analogues — and the strength of a regional statement should scale with the breadth of that convergence (*high confidence*). Where dynamical models, statistical methods, observed trends, and physical theory all point the same way, confidence is high even if the precise magnitude remains uncertain. Where they diverge, the divergence itself is information, signalling either a genuine

bifurcation in possible futures or an unresolved deficiency in one or more lines of evidence.

Several regional signals now enjoy strong multi-line convergence. Arctic amplification is supported by observations, by consistent model behaviour, and by well-understood feedback physics, so the qualitative conclusion — that the Arctic warms far faster than the global mean — is among the most robust in regional climate science (Rantanen et al., 2022). Elevation-dependent warming in mountains is similarly supported by observations across many ranges and by feedback theory, even where the magnitude at a specific elevation remains uncertain (Pepin et al., 2022). The retreat of glaciers is observed globally, simulated consistently, and physically inevitable given the warming, so the projection that every increment of temperature drives further glacier loss is robust even where basin-scale timing is uncertain (Rounce et al., 2023).

Other regional signals remain contested, and distillation must represent that honestly. Regional precipitation change in the tropics and subtropics, the future of specific monsoon systems, and the fate of particular circulation features are areas where models diverge and where the distilled product must convey a wide range rather than a point estimate. The role of groundwater in buffering regional water stress adds a further consideration: even where the atmospheric signal is uncertain, the state of the subsurface store shapes regional resilience, and rapid groundwater decline in many aquifers narrows the buffer available to absorb climatic variability (Scanlon et al., 2023) (Jasechko et al., 2024). Distillation for water-stressed regions must therefore integrate the atmospheric projection with the hydrological and human-management context.

Fitness for purpose and co-production

Regional climate information is fit for purpose only in relation to a specific decision, and the co-production of information with the users who will act on it improves both its relevance and its uptake (*high confidence*). The same physical projection can be adequate for one decision and inadequate for another: a coarse estimate of mean warming may suffice for a broad adaptation strategy but fail entirely for the design of flood defences that depend on the tail of the sub-daily rainfall distribution. Fitness for purpose is thus a joint property of the information and the decision, and cannot be assessed by climate science alone.

Co-production — the collaborative development of climate information by scientists and users — is assessed as improving fitness for purpose because it clarifies what variables, thresholds, and timescales the decision actually turns on, and because it surfaces the tolerance for uncertainty that the decision can accommodate. It also guards against the delivery of technically impressive but decision-irrelevant products. The framework of complex climate-change risk assessment provides the structure within which such co-production proceeds, linking the physical distillation to the exposure, vulnerability, and response elements that determine actual risk (Simpson et al., 2021).

The distillation and co-production framing also disciplines the treatment of uncertainty. Rather than suppressing uncertainty to deliver a clean answer, or

overwhelming the user with an uninterpretable envelope, distillation aims to convey the structure of the uncertainty — which parts are reducible, which are irreducible, which are consequential for the decision at hand. This sets up the storyline and uncertainty-communication approaches assessed in Section 7.6, which provide concrete methods for representing deep uncertainty in a decision-relevant form. It also connects to the broader recognition that the climate and biodiversity crises are coupled, so that regional information distilled for one domain must be consistent with, and ideally integrated with, the information distilled for the other (Pörtner et al., 2023).

7.6 Regional Projections by Continental Region

This section assesses projected regional change organized by continental region, integrating the physical drivers, the method-specific evidence, and the leading uncertainties for each. The organizing patterns established in Section 7.1 — land-ocean contrast, high-latitude amplification, and elevation-dependent warming — recur throughout, modulated by the circulation, hydrology, and coastal geometry particular to each region.

The Arctic and high northern latitudes

The Arctic is projected to continue warming at roughly three to four times the global rate, with the loss of sea ice, land ice, and permafrost integrity driving cascading regional and remote consequences (*high confidence*). Arctic amplification is the defining regional signal of the climate system. Observations over the satellite era establish that the Arctic has warmed nearly four times faster than the globe since 1979, a factor substantially larger than the doubling long used as a rule of thumb, and one that models have historically tended to under-represent (Rantanen et al., 2022). The amplification arises from a well-understood combination of surface-albedo feedback as reflective ice and snow give way to dark ocean and land, lapse-rate feedback favoured by the stable Arctic boundary layer, changes in poleward heat and moisture transport, and cloud feedbacks.

The consequences propagate through the Arctic system and beyond. Glacier loss is universal and scales with warming: every increment of global temperature commits additional ice to loss, with implications for sea level and for the freshwater regimes of Arctic and mountain rivers (Rounce et al., 2023). Field campaigns synthesizing the coupled atmosphere-ice-ocean system across a full annual cycle have sharpened understanding of the processes governing the Arctic surface energy budget and boundary-layer structure, providing critical constraints for model evaluation in a region where conventional observations are sparse (Shupe et al., 2022). Plastic and other pollutant burdens now reach even the remote Arctic, a marker of the region's connection to the global system and a complicating factor for its ecosystems (Bergmann et al., 2022).

Whether Arctic change feeds back on mid-latitude weather through altered circulation remains an area of active investigation and lower confidence. The physical plausibility of an Arctic influence on jet behaviour is established, but the magnitude and detectability of such teleconnections against the noise of internal variability are contested, and this uncertainty must be carried honestly into any regional projection that invokes an Arctic-mid-latitude linkage.

North America

North America is projected to experience intensifying hydroclimate extremes, with the western megadrought, escalating wildfire, and inequitable flood exposure among the most consequential regional signals (*high confidence*).

The emerging southwestern North American megadrought represents a regional hydroclimate state of exceptional severity; the drought intensified rapidly in 2020–2021 to a level with few analogues in the long paleoclimate record, and anthropogenic warming is assessed as a substantial contributor through its effect on evaporative demand and snowpack (Williams et al., 2022). The 2021 Pacific Northwest heat event demonstrated that the region’s temperature extremes can exceed prior statistical expectations by a wide margin, driven by the coincidence of an amplified ridge, antecedent soil-moisture deficit, and strong adiabatic warming (White et al., 2023).

Wildfire activity is expanding in space and time across North America and other fire-prone regions, as warming and drying extend the length and intensity of fire-weather seasons (Senande-Rivera et al., 2022). Flood risk is both rising and inequitably distributed: analysis of United States flood exposure reveals patterns that fall disproportionately on socially vulnerable communities, so that the regional physical hazard translates into a sharply uneven distribution of harm (Wing et al., 2022). Groundwater, a critical buffer for North American agriculture and municipal supply, is declining rapidly in many aquifers, though some show recovery where management has intervened — a reminder that regional water outcomes depend on human decisions as much as on the atmospheric forcing (Jasechko et al., 2024) (Scanlon et al., 2023).

Europe and the Mediterranean

Europe is projected to warm faster than the global mean, with the Mediterranean emerging as a regional hotspot of drying and heat, and western European heat extremes intensifying faster than mean warming (*high confidence*).

European heatwave trends have accelerated beyond what mean warming alone would produce, a behaviour linked to more-persistent double-jet configurations over Eurasia that favour the stagnation and amplification of surface heat (Rousi et al., 2022). This circulation dependence means that European heat projections are sensitive to the driving model’s representation of jet dynamics, and that dynamical downscaling adds value only conditional on a credible large-scale driver.

The Eastern Mediterranean and Middle East constitute a distinct regional hotspot, projected to experience warming and an increase in weather extremes exceeding the global average, with severe implications for water availability, heat exposure, and habitability in an already water-stressed region (Zittis et al., 2022). Mediterranean drying compounds these pressures, and the regional water balance is further stressed by declining groundwater and rising demand (Scanlon et al., 2023). Cultural heritage across Europe and the Mediterranean is exposed to accelerating climatic pressures, adding a dimension of loss that regional assessment increasingly incorporates (Sesana et al., 2021).

Asia

Asia faces intensifying monsoon variability, escalating heat and water stress, and acute exposure of its dense populations and agriculture to hydroclimate extremes, with high but regionally variable confidence (*high confidence*). The continent's climate is organized by monsoon systems whose projected changes carry large uncertainty in the details even where the thermodynamic tendency toward more intense wet extremes is robust. Agricultural production across Asia is highly exposed to the changing climate, with heat, water stress, and shifting growing seasons posing challenges to food security for a large share of the world's population (Rahman et al., 2022).

Water resources are a central regional concern. High-mountain Asia's glaciers, the water towers for major river systems, are losing mass in step with warming, threatening the dry-season flows on which downstream populations depend (Rounce et al., 2023) (Pepin et al., 2022). Groundwater decline is severe in parts of Asia, narrowing the buffer against monsoon variability and drought (Jasechko et al., 2024) (Scanlon et al., 2023). Flash droughts, which develop rapidly and give little warning, are projected to become more common globally and pose a particular threat to Asian agriculture given its dependence on timely monsoon rainfall (Yuan et al., 2023). Soil salinization, driven by a combination of climatic and management factors, further threatens the productivity of irrigated Asian croplands (Hassani et al., 2021).

Africa

Africa is projected to experience warming above the global mean over much of the continent, with severe water and food-security consequences amplified by high exposure and limited observational and adaptive infrastructure (*high confidence*). Much of Africa faces the compounding challenge of a strong warming signal, an uncertain but consequential precipitation response, and a sparse observational network that limits both the evaluation of models and the detection of change. The northern reaches of the continent share the Eastern Mediterranean and Middle East hotspot dynamics, with warming and extremes projected to outpace the global average (Zittis et al., 2022).

Water and soils are the pivotal regional constraints. Groundwater, a critical resource across much of Africa, is under pressure from rising demand and

variable recharge, and its state governs regional resilience to drought (Scanlon et al., 2023) (Jasechko et al., 2024). Soil salinization threatens agricultural land where irrigation and evaporative concentration interact with a warming, drying climate (Hassani et al., 2021). The dependence of African food security on rain-fed agriculture makes the uncertain precipitation projection especially consequential, and it places a premium on the distillation of the best-available regional information despite the observational gaps. Wildfire regimes are also shifting across African savannas and other biomes as fire-weather conditions change (Senande-Rivera et al., 2022).

Central and South America

South America's projections are dominated by the fate of the Amazon, where observations indicate a pronounced loss of forest resilience that raises the risk of large-scale, partially irreversible regional transition (*high confidence*). The Amazon rainforest has exhibited a pronounced loss of resilience since the early 2000s, detectable as a slowing recovery from perturbation, consistent with the system approaching a critical threshold beyond which forest could give way to a drier, degraded state over large areas (Boulton et al., 2022). Such a transition would have profound regional consequences for rainfall, carbon storage, and biodiversity, and it exemplifies the kind of high-impact, deeply uncertain outcome that regional assessment must represent even where its probability cannot be precisely quantified.

Forest resilience is declining more broadly, with emerging signals of stress detectable across many of the world's forests under the combined pressure of warming, drying, and disturbance (Forzieri et al., 2022). For South America this coupling of the physical climate to the biosphere is central: the regional water cycle is partly sustained by forest transpiration, so forest loss and hydroclimate change reinforce one another. The coupled climate and biodiversity crises are especially entangled in this region, and regional information distilled for the physical climate must be integrated with the ecological state to be decision-relevant (Pörtner et al., 2023). Andean glacier loss adds a further regional water-security dimension for populations dependent on glacier-fed flows (Rounce et al., 2023).

Australasia and the small-island and coastal domain

Coastal and small-island regions worldwide face compounding sea-level rise, subsidence, and changing extremes, with relative sea-level change and coastal flood exposure driven by the interaction of climatic and local land-motion processes (*high confidence*). For low-lying coasts and small islands, the regionally relevant quantity is not the global-mean sea-level rise but the relative sea-level change at the coast, which combines the climatic signal with vertical land motion. In many densely populated deltas and coastal cities, subsidence — often driven by groundwater and resource extraction — dominates the local relative rise, so that coastal flood exposure is a joint function of climate and

local land management (Nicholls et al., 2021). Global tide models supply the deterministic tidal component that must be combined with mean-sea-level and surge changes to estimate total coastal water levels (Lyard et al., 2021), and accurate terrain models that remove the bias of forests and buildings are essential for credible coastal-flood mapping (Hawker et al., 2022).

Human infrastructure modulates regional flood risk in both directions. Dams and reservoirs can reduce downstream flood exposure under a changing climate, though their benefits are unevenly distributed and depend on operation (Boulangé et al., 2021). Marine heatwaves, with their severe biological impacts, are an oceanic regional hazard of growing importance for coastal and island economies dependent on fisheries and reefs (Smith et al., 2022). The regional projection for these domains is thus irreducibly coupled: the physical sea-level and ocean-temperature signals interact with land motion, infrastructure, and ecosystem state to determine actual regional outcomes.

7.7 Storylines and the Communication of Uncertainty

The storyline approach

Physically self-consistent storylines have become a standard and complementary approach to probabilistic projection, particularly for communicating deeply uncertain, high-impact regional outcomes whose probabilities cannot be robustly quantified (*high confidence*). A storyline is a physically coherent, internally consistent narrative of how a region's climate could unfold, constructed by conditioning on plausible states of the large-scale drivers rather than by attempting to assign a probability to every branch. The approach responds to a structural limitation of the probabilistic paradigm: for many regional variables, and especially for the circulation-dependent extremes that dominate mid-latitude risk, the model spread does not sample a well-defined probability distribution but reflects deep, poorly quantified uncertainty in the response of the atmospheric circulation.

The storyline method addresses this by decomposing the regional response into a thermodynamic component, which is comparatively robust and well constrained, and a dynamic component, which is uncertain and circulation-dependent. By holding the thermodynamic warming fixed and exploring a discrete set of plausible circulation futures, storylines represent the range of possible regional outcomes in a form that is physically interpretable and decision-relevant, without the false precision of a single probability distribution over an ill-constrained variable. The acceleration of western European heat under more-persistent double-jet states is a natural storyline axis: rather than averaging over an uncertain circulation response, a storyline can ask what the region's heat risk would be under a specified degree of jet persistence (Rousi et al., 2022). Similarly, the ingredients of the 2021 Pacific Northwest event — ridge amplitude, soil-moisture antecedents, adiabatic warming — define a storyline

space within which the plausibility and severity of comparable future events can be explored (White et al., 2023).

Storylines are also well suited to event-based framing, in which a specific past or hypothetical event is used as the anchor and the question becomes how such an event would change in frequency or intensity under warming. This framing connects naturally to the attribution of extremes, which quantifies the human influence on the likelihood or magnitude of observed events and thereby grounds the storyline in a real, communicable occurrence (Clarke et al., 2022). Event-based storylines have proven particularly effective for communication because they attach the abstract statistics of climate change to a concrete, memorable episode that stakeholders have experienced. The southwestern North American megadrought, framed as an unfolding multi-decadal event with anthropogenic intensification, is a further example of a regional storyline anchored in an ongoing, experienced episode (Williams et al., 2022).

Compound, cascading, and systemic risk

Regional risk is increasingly understood as compound, cascading, and systemic, and structured frameworks for representing interacting hazards, exposures, and vulnerabilities are now central to regional assessment (*high confidence*). The most damaging regional outcomes rarely arise from a single hazard acting in isolation. They emerge from the coincidence of multiple drivers — heat and drought, storm surge and rainfall, sequential events that exhaust response capacity — and from the propagation of impacts across interconnected systems. A framework for complex climate-change risk assessment formalizes this understanding, treating risk as an emergent property of the interaction among hazard, exposure, vulnerability, and response, and providing a structure within which the physical climate storyline connects to the socioeconomic consequences (Simpson et al., 2021).

Compound events challenge both the physical modelling and the risk framing. Physically, they require correctly simulated dependence between variables, which is harder than simulating each variable's marginal behaviour and less well constrained by the short observational record. In terms of impact, they can produce consequences that exceed the sum of the individual hazards because systems stressed by one hazard become more vulnerable to the next. The observation that a majority of known human pathogenic diseases can be aggravated by climatic hazards illustrates how a physical hazard cascades into a compound health burden, widening the scope of regional risk assessment far beyond the physical climate (Mora et al., 2022). The declining resilience of forests and the possibility of abrupt ecosystem transition add an ecological cascade, in which a compound physical stress triggers a state change that then feeds back on the regional climate and hydrology (Forzieri et al., 2022; Boulton et al., 2022).

Systemic risk — the propagation of impacts through economic, infrastructural, and social networks — extends the cascade beyond the boundaries of the region where the hazard occurs. Flood exposure concentrated among the socially vulnerable amplifies the systemic consequences of a physical flood hazard,

because the same physical event produces larger and more persistent social and economic disruption where vulnerability is high (Wing et al., 2022; Rentschler et al., 2022). Regional assessment that stops at the physical hazard therefore understates the risk; the complex-risk framework is the tool for carrying the analysis through to the systemic consequences that decision-makers ultimately face (Simpson et al., 2021). Economic valuation of these consequences, through metrics such as the social cost of carbon, provides one channel for aggregating regionally distributed damages into a decision-relevant summary, though such metrics are themselves sensitive to assumptions about the distribution and severity of regional impacts (Rennert et al., 2022).

Tipping elements and low-likelihood, high-impact outcomes

Interacting tipping elements can raise the risk of cascading, partially irreversible regional transitions beyond what any single element implies, and low-likelihood, high-impact outcomes warrant explicit assessment rather than exclusion (*high confidence*). Several components of the Earth system exhibit the potential for abrupt or self-sustaining change once a threshold is crossed, and these tipping elements are not independent: the crossing of one can raise the probability of crossing another through physical couplings, so that the risk of a domino-like cascade exceeds the product of the individual risks (Wunderling et al., 2021). For regional assessment, tipping dynamics matter because they can produce regional outcomes that lie outside the range of the smooth, incremental projections that dominate model-based estimates.

The Amazon provides a concrete regional instance. The observed loss of forest resilience since the early 2000s is consistent with the system approaching a threshold, and a transition to a drier, degraded state would be a regional tipping event with continental consequences for rainfall, carbon, and biodiversity (Boulton et al., 2022). Broader signals of declining forest resilience suggest that such vulnerabilities are not confined to a single region (Forzieri et al., 2022). Because the probability and timing of these transitions are deeply uncertain, they are precisely the kind of outcome for which the storyline approach and the explicit treatment of low-likelihood, high-impact possibilities are essential.

The case for explicit assessment of extreme outcomes is grounded in the asymmetry of consequences. Outcomes at the high-impact tail — catastrophic and potentially irreversible regional transitions — can dominate the risk even at low probability, and excluding them because they are hard to quantify would bias the assessment toward complacency (Kemp et al., 2022). The framing of climate change against planetary boundaries reinforces this: the transgression of multiple boundaries indicates that the Earth system is being pushed outside the range within which the smooth, well-sampled projections are most reliable, raising the salience of surprise and abrupt change (Richardson et al., 2023; Rockström et al., 2023). Regional assessment therefore has a responsibility to represent the deep tail, not only the well-constrained centre of the distribution.

Communicating regional uncertainty

Effective communication of regional climate information conveys the structure of uncertainty — distinguishing the robust from the uncertain and the reducible from the irreducible — rather than either suppressing uncertainty or presenting an uninterpretable envelope (*high confidence*). The distillation step produces information whose value is realized only if it is communicated in a form the user can act on. This requires calibrated language that matches the strength of the statement to the strength of the evidence, transparency about the lines of evidence and the reasoning behind the assessment, and a clear separation between what is robustly known and what remains deeply uncertain.

Storylines and event-based framings are assessed as effective communication tools because they render abstract uncertainty concrete and physically interpretable. Attaching a projection to a specific, attributable event connects the statistics to lived experience and makes the uncertainty tangible: the question “how much more likely has an event like this become?” is more communicable than an abstract probability density, and it is grounded in the attribution science that quantifies the human contribution to observed extremes (Clarke et al., 2022). The complex-risk framework provides the scaffolding for communicating not just the physical hazard but the full chain to consequence, which is what most decisions actually require (Simpson et al., 2021).

Communication must also respect the limits of the science. Where the observational network is sparse, where models disagree on the sign of change, or where the outcome depends on deeply uncertain circulation or tipping dynamics, the honest communication is one that conveys the range and its structure rather than a spuriously precise central estimate. The integration of climate and biodiversity information, and of physical and human-system information, further complicates the communication task but is essential for decision-relevance, because regional decisions are made in the coupled system, not in the physical climate alone (Pörtner et al., 2023). The overarching principle is fitness for purpose: the communication succeeds when the user can make a better-informed decision, with a correct understanding of both what the science supports and what it cannot yet resolve.

7.8 Cross-Cutting Advances, Limitations, and Research Priorities

The changing observational and data foundation

Advances in global observational and reference datasets are a principal driver of improved regional projection, and the quality of regional information is increasingly limited by the reference data rather than by model resolution alone (*high confidence*). The regional enterprise rests on a foundation of global datasets whose steady improvement quietly enables everything above it. High-

resolution reanalyses provide dynamically consistent reference fields; harmonized soil-property grids constrain the land-surface lower boundary (Poggio et al., 2021); global hydrological model frameworks supply reference discharge and storage where gauges are absent (Schmied et al., 2021); terrain models with forests and buildings removed sharpen flood and coastal analysis (Hawker et al., 2022); and global tide atlases resolve the deterministic component of coastal sea level (Lyard et al., 2021). Where these datasets are strong, regional evaluation and downscaling are strong; where they are weak — over tropical land, high mountains, and the polar regions — regional information inherits the weakness. Field campaigns that intensively observe under-sampled regimes, such as the coordinated study of the central Arctic across a full annual cycle, materially reduce this limitation by supplying the process-level constraints that gridded products cannot (Shupe et al., 2022).

The observational foundation also increasingly includes non-traditional sources and specialized process datasets. Regional aerosol and dust processes, which modulate the surface energy budget and precipitation in arid and downwind regions, depend on characterizing the world's principal dust-source regions and their contribution to the global cycle (Kok et al., 2021). These process-level and observational advances are as consequential for regional credibility as any advance in model formulation, because they determine the reference against which every regional projection is judged and, increasingly, the data from which machine-learning methods learn.

Machine learning as an accelerant and a caution

Machine learning is transforming the computational economics of regional projection, but its physical reliability outside the training distribution remains the central open question, and hybrid physics-informed designs are the most credible path forward (*high confidence*). The rapid growth of machine learning across climate applications spans emulation, downscaling, bias correction, and the extraction of predictable structure (Kaack et al., 2022). In hydrology and other data-rich, process-complex domains, machine-learning models have demonstrated skill at predicting variables that depend on many interacting drivers, offering fast surrogates for expensive simulation (Hai et al., 2022). The promise for regional work is the resolution of the resolution-versus-ensemble tension: emulators can propagate scenario and internal-variability uncertainty across large ensembles at negligible marginal cost.

The caution is that regional projection is fundamentally an extrapolation problem, and purely data-driven methods carry no intrinsic guarantee of physical consistency or reliability beyond the range of their training. This matters most for the extremes and tipping-adjacent behaviour that dominate regional risk, precisely where training data are sparsest. The assessed best practice — training emulators on dynamical output so that physics is inherited, embedding conservation and consistency constraints, and evaluating against held-out extremes — redistributes rather than removes the need for physical modelling.

Machine learning is therefore assessed as a powerful accelerant that raises, rather than lowers, the premium on rigorous physical evaluation.

Persistent limitations and research priorities

The leading limitations of regional projection are the simulation of circulation-dependent change, the representation of compound and cascading events, and the sparse observational coverage of the most vulnerable regions, and these define the research priorities for the coming assessment cycle (*high confidence*). Circulation remains the dominant source of regional uncertainty: the thermodynamic response to warming is comparatively robust, but the dynamic response — changes in jets, storm tracks, blocking, and monsoon circulation — is both uncertain and decisive for regional extremes (Rousi et al., 2022) (White et al., 2023). Reducing this uncertainty requires improved simulation of the large-scale circulation in the driving models, because no downscaling method can supply a dynamic signal the boundary conditions omit.

Compound and cascading events are the second priority. Their physical simulation requires correctly coupled processes and correctly represented dependence between variables, and their impact assessment requires carrying the physical cascade through to ecological, health, and socioeconomic consequences (Mora et al., 2022; Forzieri et al., 2022) (Simpson et al., 2021). The third priority is equity of information: the regions most exposed and most vulnerable — much of Africa, parts of Asia, small-island states — are frequently the least well observed and least well modelled, so that the distribution of scientific capacity is misaligned with the distribution of risk (Zittis et al., 2022; Rahman et al., 2022; Wing et al., 2022). Closing this gap, through observation, capacity-building, and co-production, is both a scientific and an equity imperative. Underlying all three priorities is the recognition that regional climate change is unfolding within a coupled Earth system whose boundaries are being transgressed and whose climate and biodiversity crises are entangled, so that regional projection cannot be credibly separated from the broader assessment of Earth-system state (Richardson et al., 2023; Rockström et al., 2023; Pörtner et al., 2023).

Synthesis

The regionalization of climate change has matured from a peripheral downscaling exercise into a central, methodologically diverse enterprise. The organizing patterns of the regional response — land-ocean contrast, Arctic amplification, elevation-dependent warming — are robustly established and physically understood (Rantanen et al., 2022; Pepin et al., 2022) (Rounce et al., 2023). The methods for generating regional information — dynamical, statistical, and machine-learning — each contribute distinctive strengths, and their reconciliation through distillation is now recognized as essential rather than optional (Simpson et al., 2021). The regional projections themselves show a consistent picture of intensifying heat, more volatile hydroclimate with accelerating flash drought, expanding wildfire, and rising coastal exposure, differentiated by the

circulation, hydrology, and coastal geometry of each continent (Yuan et al., 2023) (Chiang et al., 2021; Nicholls et al., 2021) (Senande-Rivera et al., 2022). And the communication of this information has evolved toward storylines and the structured treatment of deep uncertainty, giving explicit weight to the compound, cascading, and low-likelihood high-impact outcomes that marginal probabilities alone would obscure (Wunderling et al., 2021) (Kemp et al., 2022). The central message of this chapter is that credible regional climate information is not read off a single model but constructed — through the disciplined integration of multiple methods, a strong observational foundation, and transparent expert judgment — into products that are fit for the decisions they must inform.

7.9 Regional Impact Sectors: Water, Food, Health, and Coasts

The value of regional climate projection is ultimately realized in the impact sectors that translate physical change into consequences for people and ecosystems. This section assesses the regional projection problem as it presents in four decisive sectors — water, food, health, and coasts — each of which imposes distinctive requirements on the physical information and each of which exhibits strong regional differentiation. The sectoral framing complements the geographic framing of Section 7.5 by making explicit the pathways through which the physical projections assessed above become the impacts that adaptation must address (Abbass et al., 2022).

The regional water cycle

The regional water cycle is projected to become more variable, with simultaneous intensification of wet and dry extremes, accelerating flash-drought onset, and mounting pressure on groundwater buffers, though the sign of mean precipitation change remains uncertain in many regions (*high confidence*). Water is the sector through which climate change is most immediately felt in most regions, and it is also the sector in which regional projection is hardest, because the water cycle integrates precipitation, evaporative demand, land-surface partitioning, and subsurface storage, each with its own uncertainty. The intensification of the atmospheric water-holding capacity with warming provides a robust thermodynamic tendency toward heavier precipitation extremes, but the regional distribution of that intensification depends on circulation changes that models simulate with lower confidence.

Drought is being reshaped in ways that carry detectable human influence. The frequency, duration, and intensity of drought show anthropogenic signals in the observational record (Chiang et al., 2021), and a global transition toward flash droughts — droughts that develop over weeks rather than seasons, driven by the rapid depletion of soil moisture under high evaporative demand

— is projected under continued warming (Yuan et al., 2023). Flash droughts are especially consequential regionally because their speed of onset gives little warning, stressing agricultural and water-management systems designed around slower-developing events. The prediction of both floods and droughts remains a fundamental challenge because it demands accurate representation of the precipitation forcing, the land-surface response, and, in managed basins, the human operation of reservoirs and withdrawals (Brunner et al., 2021).

Groundwater is the hidden buffer that determines regional resilience to hydroclimate variability, and it is under growing pressure. Analysis of aquifers worldwide reveals rapid groundwater decline in many regions, with some cases of recovery where management has intervened, indicating that the trajectory is not fixed by the climate alone but is strongly modulated by human decisions (Jasechko et al., 2024). The broader role of groundwater in sustaining a resilient water future — buffering droughts, supporting irrigation, sustaining baseflow — makes its depletion a first-order regional concern that regional climate information must be integrated with to be decision-relevant (Scanlon et al., 2023). Pollution compounds scarcity: accounting for water quality alongside quantity substantially increases the number of river basins facing water scarcity, so that the regional water-security picture is worse than a quantity-only analysis would suggest (Wang et al., 2024). The hydrological reference frameworks that underpin all of this analysis, such as global water-resources-and-use models, are essential tools for translating the physical projection into basin-scale water balances (Schmied et al., 2021).

Regional food systems and agriculture

Regional food systems face converging pressures from heat, water stress, shifting growing seasons, and soil degradation, with climate impacts on agriculture emerging earlier than previous model generations indicated and falling unevenly across regions (*high confidence*). Agriculture is exquisitely sensitive to the regional expression of climate change, because crop yields respond to the timing and intensity of heat, the reliability of rainfall, and the state of soils, all of which vary sharply by region. A new generation of coupled climate and crop models indicates that the impacts of climate change on global agriculture emerge earlier than earlier assessments suggested, sharpening the near-term regional stakes and compressing the time available for adaptation (Jägermeyr et al., 2021).

The regional differentiation is stark. Across much of Asia, agricultural production is highly exposed to heat, water stress, and shifting monsoon timing, with implications for the food security of a large share of the world's population (Rahman et al., 2022). Soil salinization, driven by the interaction of a warming, drying climate with irrigation and evaporative concentration, threatens the productivity of irrigated croplands in arid and semi-arid regions and is projected to expand under continued warming (Hassani et al., 2021). Climate change also exacerbates the broader environmental impacts of agriculture, so that the sector is simultaneously a victim of and a contributor to regional environmental stress,

complicating the design of adaptation that does not worsen other pressures (Yang et al., 2024). The coupling of climate change to biodiversity and food security is a global concern with acute regional manifestations, particularly where food systems depend on rain-fed agriculture and on the ecosystem services that a changing climate erodes (Muluneh, 2021).

Aquatic food systems add a marine and freshwater dimension to regional food security. Aquatic foods make a substantial contribution to nutrition in many regions, and their sensitivity to ocean warming, marine heatwaves, and changing productivity introduces a distinct regional vulnerability for coastal and island populations that depend on them (Golden et al., 2021). Marine heatwaves, with their severe and sometimes irreversible biological impacts on the ecosystems that sustain fisheries, are a growing regional hazard for these food systems (Smith et al., 2022). Harmful algal blooms, whose apparent global increase is at least partly attributable to intensified monitoring and emerging impacts, add a further regional water-quality and food-safety concern in affected coastal and inland waters (Hallegraeff et al., 2021). The protection and sustainable management of the ocean is assessed as a lever that can simultaneously advance regional food security, biodiversity, and climate goals where it is well designed (Sala et al., 2021).

Regional health burdens

The health consequences of regional climate change are broad, compounding, and unevenly distributed, spanning heat mortality, infectious disease, mental health, and the health effects of food and water insecurity, with the burden concentrated among the most exposed and least resourced populations (*high confidence*). Health is the sector in which the human meaning of regional physical change is most directly expressed. The annual assessments of the linkage between health and climate change document a widening array of regional health impacts and, increasingly, record-breaking threats associated with delayed action, underscoring that the health burden is not a distant projection but a present and growing regional reality (Romanello et al., 2023) (Romanello et al., 2024).

The mechanisms are diverse and often compounding. A majority of known human pathogenic diseases can be aggravated by climatic hazards, meaning that the physical regional projection cascades into a wide and interacting set of disease risks rather than a single clean signal (Mora et al., 2022). Heat extremes, whose regional intensification is assessed above, translate directly into heat morbidity and mortality, with the burden shaped by exposure, vulnerability, and adaptive capacity that vary sharply within and between regions. The mental-health dimension of climate change is increasingly recognized, encompassing the psychological consequences of extreme events, displacement, and chronic environmental stress, and it constitutes a regional health burden that conventional impact assessment has historically neglected (Charlson et al., 2021). Insect-mediated risks, including shifts in the distribution of disease

vectors and agricultural pests, add a further ecological pathway from regional climate change to human welfare (Harvey et al., 2022).

The distributional character of the health burden is a defining feature. The same physical hazard produces very different health outcomes depending on the exposure and vulnerability of the affected population, so that regional health projection cannot be read off the physical climate alone but must integrate the socioeconomic context. This reinforces the complex-risk and distillation framings assessed above: health-relevant regional information is fit for purpose only when it carries the physical projection through to the exposed and vulnerable populations who bear the consequences (Simpson et al., 2021).

Coasts, rivers, and the built and natural environment

Coastal and riverine regions face compounding pressures from sea-level rise, subsidence, changing extremes, and ecosystem degradation, with relative change at the coast and the state of freshwater ecosystems determining regional outcomes as much as the global-mean signals (*high confidence*). For coastal regions, the decision-relevant quantity is relative sea-level change, which combines the climatic signal with local vertical land motion; in many densely populated deltas, subsidence dominates the local rise and coastal flood exposure is therefore a joint product of climate and local land management (Nicholls et al., 2021). Credible regional coastal assessment depends on accurate terrain models that remove the bias of forests and buildings (Hawker et al., 2022) and on global tide atlases that resolve the deterministic tidal component of total water level (Lyard et al., 2021).

Rivers and their ecosystems are a further regional frontier. River ecosystem metabolism and carbon biogeochemistry are being reshaped by warming and altered flow regimes, with consequences for water quality, carbon cycling, and the ecosystem services on which regional populations depend (Battin et al., 2023). The regional water cycle, the food system, and the coastal zone are thus linked through the river network, so that a change in one propagates to the others. Human infrastructure modulates these regional outcomes: dams and reservoirs can reduce downstream flood exposure under a changing climate, though the benefits are unevenly distributed and depend on operation (Boulangé et al., 2021), and the same infrastructure that buffers floods can alter sediment delivery and exacerbate the subsidence that drives relative sea-level rise at the coast. The regional picture that emerges across these sectors is one of tightly coupled physical, ecological, and human systems, in which credible projection requires integrating the physical climate information assessed throughout this chapter with the sectoral and socioeconomic context that determines actual regional consequences (Abbass et al., 2022) (Simpson et al., 2021).

7.10 Regional Detection, Attribution, and Monitoring

Detecting the regional signal

The anthropogenic signal is now detectable in an expanding set of regional variables and extremes, and detection at the regional scale is a prerequisite for the confident projection of regional futures (*high confidence*). Detection asks whether a change has occurred that is inconsistent with internal variability alone; attribution asks whether that change can be ascribed to a specified cause. Both are harder at the regional than at the global scale, because the signal-to-noise ratio is lower — regional internal variability is large relative to the forced signal — and because the observational record is shorter and sparser for many regions and variables. Where detection succeeds regionally, it both confirms that the physical understanding is adequate and constrains the projection, because a model that reproduces the detected trend is more credible for the future.

Detection has advanced most clearly for temperature and for temperature-driven extremes. Arctic amplification is detectable as a robust departure from the global trend and is inconsistent with unforced variability, providing one of the strongest regional detection results in the climate system (Rantanen et al., 2022). Elevation-dependent warming is detectable across many mountain ranges, though the sparse high-elevation observational network limits confidence in the precise vertical structure (Pepin et al., 2022). Detection of hydroclimate change is more difficult but is succeeding for drought: the frequency, duration, and intensity of drought carry a detectable anthropogenic signal in the observational record, a result that both confirms the physical mechanism and sharpens the projection of future regional drought (Chiang et al., 2021). The emerging southwestern North American megadrought is detectable as a departure from the range of natural variability reconstructed from paleoclimate, with anthropogenic warming assessed as a substantial contributor (Williams et al., 2022).

Event attribution at the regional scale

Event attribution has matured into a routine tool that quantifies the human influence on the likelihood and intensity of specific regional extremes, providing the empirical anchor for storyline framings and for public communication (*high confidence*). Where classical detection and attribution address long-term trends, event attribution addresses individual extremes — a specific heatwave, flood, or drought — asking how human influence has changed the probability or magnitude of an event of that kind. The approach has become sufficiently established that attribution assessments are now produced rapidly after major events, connecting the abstract statistics of climate change to concrete, experienced episodes (Clarke et al., 2022).

Several recent regional extremes illustrate the method and its findings. The 2021 Pacific Northwest heatwave was so far outside the prior observed range that it challenged conventional extreme-value approaches, and its attribution required careful treatment of the physical ingredients rather than a simple

extrapolation of the historical distribution (White et al., 2023). The acceleration of western European heat, linked to more-persistent double-jet states, is a case where attribution must separate the thermodynamic contribution of warming from the dynamic contribution of circulation change, a separation that is central to the storyline approach (Rousi et al., 2022). Attribution of changing wildfire activity connects the physical fire-weather trend to the human influence on temperature and aridity, supporting the assessment that fire seasons are lengthening and intensifying across many regions (Senande-Rivera et al., 2022). The value of event attribution for regional assessment is twofold: it provides quantitative constraints on how specific hazards are changing, and it supplies the communicable, event-anchored framing that makes regional risk tangible to the public and to decision-makers.

Monitoring, early warning, and the regional information chain

Sustained regional monitoring and early-warning systems are the operational counterpart to projection, and the credibility of regional information depends on the continuity and coverage of the observing systems that underpin both (*high confidence*). Projection and monitoring are complementary: projection anticipates the changing distribution of hazards, while monitoring detects their realization in near-real time and triggers response. Both depend on the same observational foundation, and gaps in that foundation degrade both. The under-observed regions — much of the tropics, the high mountains, the polar regions — are precisely those where projection is least certain and where monitoring is most needed, so the observational gap compounds across the information chain.

Advances in the observing system are steadily closing parts of this gap. Intensive field campaigns in under-sampled regimes, such as the coordinated observation of the central Arctic across a full annual cycle, supply the process-level constraints that improve both models and monitoring (Shupe et al., 2022). Specialized global datasets — soil properties, terrain, tides, hydrological state — extend consistent monitoring to variables and regions that conventional networks miss (Poggio et al., 2021) (Hawker et al., 2022) (Lyard et al., 2021) (Schmied et al., 2021). The characterization of regional aerosol and dust processes, including the contribution of the world's principal dust-source regions, improves the monitoring of a driver that materially affects regional radiation and precipitation (Kok et al., 2021). Even the detection of far-field pollution signatures, such as plastic reaching the remote Arctic, is a form of monitoring that reveals the connectivity of the regional system to the global whole (Bergmann et al., 2022).

The regional information chain thus runs from observation, through detection and attribution, through model evaluation and projection, through distillation, to communication and, ultimately, to monitoring and response. Each link depends on the others, and the weakest link — most often the observational coverage of the most vulnerable regions — sets the overall credibility. Strengthening regional climate information is therefore not primarily a matter of higher model resolution but of investing across the whole chain: in observation, in pro-

cess understanding, in the reconciliation of methods, and in the co-production and communication of information that is fit for the decisions it must inform (Simpson et al., 2021). This integrated view is the organizing conclusion of the chapter: credible regional climate change assessment is a system, not a single output, and its improvement requires attention to every part of that system.

7.11 Conclusions

Regional climate projection has become a mature, methodologically plural, and decision-oriented science, and its central lesson is that credible regional information is constructed from many lines of evidence rather than read from any single model (*high confidence*). The chapter has assessed the full arc from global signal to regional decision. The organizing patterns of the regional response are robust and physically understood, with Arctic amplification and elevation-dependent warming the clearest examples (Rantanen et al., 2022) (Pepin et al., 2022). The methods for generating regional information — dynamical downscaling with its convection-permitting frontier, statistical downscaling with its efficiency and its stationarity caveats, and machine-learning emulation with its computational promise and its extrapolation risks — are complementary, and their reconciliation through distillation is the step that turns raw output into usable knowledge (Tebaldi et al., 2021) (Simpson et al., 2021).

The regional projections themselves converge on a coherent picture. Heat extremes intensify, often faster than mean warming where circulation change amplifies them (Rousi et al., 2022) (White et al., 2023). The water cycle becomes more volatile, with accelerating flash drought, detectable anthropogenic drought trends, and mounting pressure on groundwater buffers (Yuan et al., 2023) (Chiang et al., 2021) (Jasechko et al., 2024). Wildfire expands, forests lose resilience, and the Amazon shows signs of approaching a regional tipping threshold (Senande-Rivera et al., 2022) (Forzieri et al., 2022) (Boulton et al., 2022). Coasts face compounding sea-level rise and subsidence, and the burden across every sector falls most heavily on the exposed and the vulnerable (Nicholls et al., 2021) (Wing et al., 2022).

The communication of this information has evolved toward storylines and the explicit treatment of deep uncertainty, giving proper weight to compound, cascading, and low-likelihood high-impact outcomes that marginal probabilities alone would hide (Wunderling et al., 2021) (Kemp et al., 2022). The overarching conclusion is that regional climate assessment is an information system spanning observation, detection and attribution, evaluation, projection, distillation, communication, and monitoring, and that its weakest link — most often the observational coverage of the most vulnerable regions — sets its overall credibility. Strengthening regional information is a task for the whole system, and it is inseparable from the broader assessment of an Earth system whose boundaries are being transgressed and whose climate and biodiversity crises are entangled (Richardson et al., 2023) (Pörtner et al., 2023).

7.12 References

Abbass, K., et al. (2022). A review of the global climate change impacts, adaptation, and sustainable mitigation measures. *Environmental Science and Pollution Research*. <https://doi.org/10.1007/s11356-022-19718-6>

Battin, T., et al. (2023). River ecosystem metabolism and carbon biogeochemistry in a changing world. *Nature*. <https://doi.org/10.1038/s41586-022-05500-8>

Bergmann, M., et al. (2022). Plastic pollution in the Arctic. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-022-00279-8>

Boulangé, J., et al. (2021). Role of dams in reducing global flood exposure under climate change. *Nature Communications*. <https://doi.org/10.1038/s41467-020-20704-0>

Boulton, C., et al. (2022). Pronounced loss of Amazon rainforest resilience since the early 2000s. *Nature Climate Change*. <https://doi.org/10.1038/s41558-022-01287-8>

Brunner, M., et al. (2021). Challenges in modeling and predicting floods and droughts: A review. *Wiley Interdisciplinary Reviews Water*. <https://doi.org/10.1002/wat2.1520>

Charlson, F., et al. (2021). Climate Change and Mental Health: A Scoping Review. *International Journal of Environmental Research and Public Health*. <https://doi.org/10.3390/ijerph18094486>

Chiang, F., et al. (2021). Evidence of anthropogenic impacts on global drought frequency, duration, and intensity. *Nature Communications*. <https://doi.org/10.1038/021-22314-w>

Clarke, B., et al. (2022). Extreme weather impacts of climate change: an attribution perspective. *Environmental Research Climate*. <https://doi.org/10.1088/2752-5295/ac6e7d>

Forzieri, G., et al. (2022). Emerging signals of declining forest resilience under climate change. *Nature*. <https://doi.org/10.1038/s41586-022-04959-9>

Friedlingstein, P., et al. (2023). Global Carbon Budget 2023. *Earth System Science Data*. <https://doi.org/10.5194/essd-15-5301-2023>

Friedlingstein, P., et al. (2025). Global Carbon Budget 2024. *Earth System Science Data*. <https://doi.org/10.5194/essd-17-965-2025>

Golden, C., et al. (2021). Aquatic foods to nourish nations. *Nature*. <https://doi.org/10.1038/021-03917-1>

Hai, N., et al. (2022). Groundwater level prediction using machine learning models: A comprehensive review. *Neurocomputing*. <https://doi.org/10.1016/j.neucom.2022.03.014>

Hallegraeff, G., et al. (2021). Perceived global increase in algal blooms is attributable to intensified monitoring and emerging bloom impacts. *Communications Earth & Environment*. <https://doi.org/10.1038/s43247-021-00178-8>

Harvey, J., et al. (2022). Scientists' warning on climate change and insects. *Ecological Monographs*. <https://doi.org/10.1002/ecm.1553>

Hassani, A., et al. (2021). Global predictions of primary soil salinization under changing climate in the 21st century. *Nature Communications*. <https://doi.org/10.1038/s41460-021-26907-3>

Hawker, L., et al. (2022). A 30 m global map of elevation with forests and buildings removed. *Environmental Research Letters*. <https://doi.org/10.1088/1748-9326/ac4d4f>

Jasechko, S., et al. (2024). Rapid groundwater decline and some cases of recovery in aquifers globally. *Nature*. <https://doi.org/10.1038/s41586-023-06879-8>

Jägermeyr, J., et al. (2021). Climate impacts on global agriculture emerge earlier in new generation of climate and crop models. *Nature Food*. <https://doi.org/10.1038/s43016-021-00400-y>

Kaack, L., et al. (2022). Tackling Climate Change with Machine Learning. *ACM Computing Surveys*. <https://doi.org/10.1145/3485128>

Kemp, L., et al. (2022). Climate Endgame: Exploring catastrophic climate change scenarios. *Proceedings of the National Academy of Sciences*. <https://doi.org/10.1073/pnas.2118169119>

Kok, J., et al. (2021). Contribution of the world's main dust source regions to the global cycle of desert dust. *Atmospheric Chemistry and Physics*. <https://doi.org/10.5194/acp-21-8169-2021>

Lovato, T., et al. (2022). CMIP6 Simulations With the CMCC Earth System Model (CMCC-ESM2). *Journal of Advances in Modeling Earth Systems*. <https://doi.org/10.1029/2021ms002814>

Lyard, F., et al. (2021). FES2014 global ocean tide atlas: design and performance. *Ocean Science*. <https://doi.org/10.5194/os-17-615-2021>

Mora, C., et al. (2022). Over half of known human pathogenic diseases can be aggravated by climate change. *Nature Climate Change*. <https://doi.org/10.1038/s41558-022-01426-1>

Muluneh, M. (2021). Impact of climate change on biodiversity and food security: a global perspective—a review article. *Agriculture & Food Security*. <https://doi.org/10.1186/s40066-021-00318-5>

Nicholls, R., et al. (2021). A global analysis of subsidence, relative sea-level change and coastal flood exposure. *Nature Climate Change*. <https://doi.org/10.1038/s41558-021-00993-z>

Pepin, N., et al. (2022). Climate Changes and Their Elevational Patterns in the Mountains of the World. *Reviews of Geophysics*. <https://doi.org/10.1029/2020rg000730>

Poggio, L., et al. (2021). SoilGrids 2.0: producing soil information for the globe with quantified spatial uncertainty. *SOIL*. <https://doi.org/10.5194/soil-7-217-2021>

Pörtner, H., et al. (2023). Overcoming the coupled climate and biodiversity crises and their societal impacts. *Science*. <https://doi.org/10.1126/science.abc4881>

Rahman, M., et al. (2022). Impact of climate change on agricultural production; Issues, challenges, and opportunities in Asia. *Frontiers in Plant Science*. <https://doi.org/10.3389/fpls.2022.925548>

Rantanen, M., et al. (2022). The Arctic has warmed nearly four times faster than the globe since 1979. *Communications Earth & Environment*. <https://doi.org/10.1038/s43247-022-00498-3>

Rennert, K., et al. (2022). Comprehensive evidence implies a higher social cost of CO₂. *Nature*. <https://doi.org/10.1038/s41586-022-05224-9>

Romanello, M., et al. (2023). The 2023 report of the Lancet Countdown on health and climate change: the imperative for a health-centred response in a world facing irreversible harms. *The Lancet*. [https://doi.org/10.1016/s0140-6736\(23\)01859-7](https://doi.org/10.1016/s0140-6736(23)01859-7)

Romanello, M., et al. (2024). The 2024 report of the Lancet Countdown on health and climate change: facing record-breaking threats from delayed action. *The Lancet*. [https://doi.org/10.1016/s0140-6736\(24\)01822-1](https://doi.org/10.1016/s0140-6736(24)01822-1)

Rentschler, J., et al. (2022). Flood exposure and poverty in 188 countries. *Nature Communications*. <https://doi.org/10.1038/s41467-022-30727-4>

Richardson, K., et al. (2023). Earth beyond six of nine planetary boundaries. *Science Advances*. <https://doi.org/10.1126/sciadv.adh2458>

Rockström, J., et al. (2023). Safe and just Earth system boundaries. *Nature*. <https://doi.org/10.1038/s41586-023-06083-8>

Rounce, D., et al. (2023). Global glacier change in the 21st century: Every increase in temperature matters. *Science*. <https://doi.org/10.1126/science.abo1324>

Rousi, E., et al. (2022). Accelerated western European heatwave trends linked to more-persistent double jets over Eurasia. *Nature Communications*. <https://doi.org/10.1038/s41467-022-31432-y>

Sala, E., et al. (2021). Protecting the global ocean for biodiversity, food and climate. *Nature*. <https://doi.org/10.1038/s41586-021-03371-z>

Schmied, H., et al. (2021). The global water resources and use model WaterGAP v2.2d: model description and evaluation. *Geoscientific Model Development*. <https://doi.org/10.5194/gmd-14-1037-2021>

Senande-Rivera, M., et al. (2022). Spatial and temporal expansion of global wildland fire activity in response to climate change. *Nature Communications*. <https://doi.org/10.1038/s41467-022-28835-2>

Sesana, E., et al. (2021). Climate change impacts on cultural heritage: A literature review. *Wiley Interdisciplinary Reviews Climate Change*. <https://doi.org/10.1002/wcc.71>

Shupe, M., et al. (2022). Overview of the MOSAiC expedition: Atmosphere. *Elementa Science of the Anthropocene*. <https://doi.org/10.1525/elementa.2021.00060>

Simpson, N., et al. (2021). A framework for complex climate change risk assessment. *One Earth*. <https://doi.org/10.1016/j.oneear.2021.03.005>

Smith, K., et al. (2022). Biological Impacts of Marine Heatwaves. *Annual Review of Marine Science*. <https://doi.org/10.1146/annurev-marine-032122-121437>

Tebaldi, C., et al. (2021). Climate model projections from the Scenario Model Intercomparison Project (ScenarioMIP) of CMIP6. *Earth System Dynamics*. <https://doi.org/10.5194/esd-12-253-2021>

White, R., et al. (2023). The unprecedented Pacific Northwest heatwave of June 2021. *Nature Communications*. <https://doi.org/10.1038/s41467-023-36289-3>

Williams, A., et al. (2022). Rapid intensification of the emerging southwestern North American megadrought in 2020-2021. *Nature Climate Change*. <https://doi.org/10.1038/s41558-022-01290-z>

Wang, M., et al. (2024). A triple increase in global river basins with water scarcity due to future pollution. *Nature Communications*. <https://doi.org/10.1038/s41467-024-44947-3>

Wing, O., et al. (2022). Inequitable patterns of US flood risk in the Anthropocene. *Nature Climate Change*. <https://doi.org/10.1038/s41558-021-01265-6>

Wunderling, N., et al. (2021). Interacting tipping elements increase risk of climate domino effects under global warming. *Earth System Dynamics*. <https://doi.org/10.5194/esd-12-601-2021>

Yang, X., et al. (2024). Climate change exacerbates the environmental impacts of agriculture. *Science*. <https://doi.org/10.1126/science.adn3747>

Yuan, X., et al. (2023). A global transition to flash droughts under climate change. *Science*. <https://doi.org/10.1126/science.abn6301>

Zittis, G., et al. (2022). Climate Change and Weather Extremes in the Eastern Mediterranean and Middle East. *Reviews of Geophysics*. <https://doi.org/10.1029/2021rg000762>

Chapter 8

Chapter 8: Tipping Points and Abrupt Change

8.1 Executive Summary

The climate system contains a set of large-scale components — tipping elements — that can undergo self-sustaining, and in some cases effectively irreversible, transitions once forcing crosses a critical threshold; several of these elements show observational signals consistent with declining stability as global warming proceeds (*high confidence*). The scientific basis for treating abrupt and irreversible change as a distinct assessment category rests on the recognition that the responses of ice sheets, the ocean overturning circulation, permafrost, tropical and boreal forests, and warm-water coral reefs to sustained forcing are not always smooth, reversible, or proportionate to the forcing applied (*high confidence*) (Folke, 2021) (Richardson, 2023).

Warm-water coral reef systems are assessed to be the tipping element most imminently at risk, with the biological impacts of marine heatwaves already driving mass mortality events at present levels of warming (*high confidence*). Marine heatwaves have intensified in frequency, duration, and severity, and their ecological consequences — including bleaching, disease, and range shifts — are documented across ocean basins (Smith, 2022).

The Amazon forest system shows a pronounced loss of resilience since the early 2000s, with a rising fraction of the biome exhibiting the statistical signatures of critical slowing down; a combination of warming, drought, deforestation, and fire could drive large areas past a threshold into a degraded, lower-biomass state (*medium confidence*). Observational analyses of vegetation reveal declining recovery rates from perturbation, and process-based assessment identifies self-amplifying dieback mechanisms (Boulton, 2022) (Flores, 2024) (Forzieri, 2022).

Permafrost constitutes a large and vulnerable carbon reservoir whose thaw releases carbon dioxide and methane, producing a positive feedback on warming that is expected to be gradual and pervasive rather than a single global-scale abrupt transition, though abrupt thaw processes can dominate locally (*high confidence*). The magnitude of permafrost carbon feedback remains

a significant source of uncertainty in the assessment of committed warming (Miner, 2022).

Terrestrial ecosystems more broadly show emerging, spatially coherent signals of declining resilience under climate change, with hotter droughts producing tree mortality across a wide range of forest biomes (*high confidence*). These signals of destabilization in the land carbon sink have implications for the trajectory of atmospheric carbon dioxide and for the stability of forest-related tipping elements (Forzieri, 2022) (Hammond, 2022) (McDowell, 2022) (Frenne, 2021).

Early-warning indicators grounded in the theory of critical slowing down — rising temporal autocorrelation and variance in system state variables — provide a physically motivated but not infallible basis for detecting approach to a tipping threshold; their interpretation is complicated by observational limitations, confounding trends, and the possibility of tipping without prior warning (*medium confidence*). Where long, high-quality records exist, such indicators have been detected in the Amazon and in forest resilience metrics (Boulton, 2022) (Flores, 2024) (Forzieri, 2022).

The paleoclimate record documents repeated episodes of abrupt regional and hemispheric change and demonstrates that Earth’s climate has occupied radically different states, providing evidence that the system is capable of non-linear reorganization on timescales far shorter than the forcing that triggers it (*high confidence*). The long-term reconstruction of continental configuration, sea level, and climate confirms the range of states the Earth system can occupy (Scotese, 2021).

Tipping elements do not operate in isolation: interactions among them, and between climatic and non-climatic stresses, can produce cascades and systemic risk that amplify impacts beyond the sum of individual components, and that couple the climate crisis to biodiversity loss and to broader societal shocks (*medium confidence*). The assessment of interacting risks, cascades, and polycrisis dynamics is an area of rapid development and remains subject to substantial uncertainty (Lawrence, 2024) (Pörtner, 2023) (Simpson, 2021) (Galaz, 2021).

Because low-likelihood, high-impact outcomes carry disproportionate societal consequences, a comprehensive risk assessment must weight them despite their deep uncertainty; the possibility of catastrophic or, in the extreme, globally destabilizing outcomes is not excluded by current evidence and warrants dedicated study (*high confidence in the value of assessing such outcomes*). Catastrophe-focused analysis argues that the tail of the climate risk distribution has been comparatively neglected and deserves systematic assessment (Kemp, 2022) (Richardson, 2023).

8.2 Concepts, Definitions, and the Assessment Framework for Abrupt Change

What distinguishes a tipping point from ordinary climate change

A tipping point is a critical threshold in a forcing or in a system state beyond which a component of the Earth system reorganizes into a qualitatively different state, often through self-sustaining internal feedbacks and frequently with limited reversibility on policy-relevant timescales; this behaviour is distinct from the smooth, broadly reversible response that characterizes most of the climate system (*high confidence*). The defining features of a tipping element are threshold behaviour, the dominance of positive (self-amplifying) internal feedbacks once the threshold is crossed, and hysteresis — the property that reversing the forcing does not straightforwardly reverse the transition. These features together mean that the trajectory of an affected component cannot be inferred from the forcing alone, and that the timing and even the occurrence of a transition may carry deep and in some respects irreducible uncertainty.

This chapter uses the term *abrupt change* to denote a transition that is fast relative to the forcing that causes it — one in which the rate of change of the system substantially exceeds the rate of change of the driver over some interval — and reserves *tipping point* for the more specific case in which such a transition is driven by internal feedbacks past a threshold and is difficult to reverse. Not all abrupt changes are tipping-point phenomena, and not all tipping-point transitions are fast: some, such as the long-term response of the great ice sheets, may commit the system to a new state over centuries to millennia even where the threshold that determines the eventual outcome is crossed quickly. The assessment therefore separates the questions of *whether* a threshold is crossed, *how fast* the resulting transition unfolds, and *how reversible* the outcome is, because these can differ markedly among tipping elements.

The scientific significance of tipping behaviour is that it violates the tacit assumption of proportionality that underlies much of climate impact assessment. Where a system responds proportionately, a given increment of warming produces a broadly predictable increment of impact, and the marginal consequence of an additional increment is similar to the last. Where a system exhibits threshold behaviour, the marginal consequence of an additional increment of warming can rise sharply and discontinuously in the neighbourhood of the threshold, so that a small additional forcing produces a disproportionate and possibly unrecoverable change. This nonlinearity is the central reason that abrupt and irreversible change is assessed as a distinct category of climate risk rather than folded into the assessment of gradual change (Folke, 2021).

The planetary context: boundaries, stability, and the Anthropocene biosphere

The concept of tipping elements is embedded within a broader framing of the Earth system as operating within boundaries beyond which the risk of large-scale, potentially abrupt and irreversible change increases; current assessments place several of these boundaries in a zone of elevated risk (*high confidence*). The planetary-boundaries framework identifies a set of biophysical

processes — including climate change, biosphere integrity, land-system change, freshwater, and biogeochemical flows — that regulate the stability and resilience of the Earth system, and argues that transgression of these boundaries raises the likelihood of pushing the system out of the relatively stable state that has characterized the recent geological past (Richardson, 2023). The most recent quantification concludes that six of nine boundaries have been transgressed, which the authors interpret as a substantial and rising pressure on Earth-system stability (Richardson, 2023).

This framing matters for the assessment of tipping points in two ways. First, it locates individual tipping elements within an interconnected biosphere whose overall resilience — its capacity to absorb disturbance and reorganize while retaining function — is itself being eroded (Folke, 2021). A biosphere under multiple simultaneous pressures is, all else equal, more prone to threshold behaviour in any one of its components, because the buffering capacity that would otherwise damp perturbations is reduced. Second, it draws attention to the coupling between the climate system and the living systems that both respond to and regulate climate, a coupling that is central to the tipping behaviour of forests, reefs, and the carbon cycle (Pörtner, 2023). The assessment of tipping points is therefore not a catalogue of independent switches but an evaluation of the stability of an integrated, and increasingly stressed, Earth system (Folke, 2021) (Richardson, 2023).

Confidence, likelihood, and the treatment of deep uncertainty

The assessment of tipping points requires calibrated treatment of two distinct kinds of uncertainty: uncertainty about whether and when a threshold will be crossed, and uncertainty about the consequences that would follow; for many tipping elements the evidence base supports high confidence in the qualitative existence of threshold behaviour but only low to medium confidence in the precise threshold value (*high confidence in this characterization*). Throughout this chapter, calibrated confidence language expresses the degree of evidence and agreement underpinning a finding, and likelihood language, where used, expresses a probabilistic assessment of an outcome. A recurring feature of the tipping-point literature is that these two can diverge sharply: there may be high confidence that a component of the system possesses a tipping threshold in principle, combined with genuinely deep uncertainty about the level of warming at which that threshold lies and about the timescale over which a transition would unfold.

This divergence has a direct methodological consequence. For processes governed by well-observed physics and constrained by instrumental and paleoclimate data, the assessment can express findings in probabilistic terms. For processes in which the governing feedbacks are understood qualitatively but poorly constrained quantitatively — which describes several major tipping elements — the assessment must instead rely on lines of evidence such as process understanding, analogues from the paleorecord, and emerging statistical indicators of destabilization, none of which alone yields a calibrated probabil-

ity of crossing a threshold. In such cases the honest expression of the state of knowledge is a statement of the physical plausibility and potential consequences of a transition, accompanied by an explicit acknowledgment that the probability cannot be reliably quantified. This is not a weakness of the assessment but a faithful representation of a system in which some of the most consequential outcomes are also the most deeply uncertain (Kemp, 2022).

The treatment of deep uncertainty also shapes how this chapter handles the tail of the risk distribution. A finding that a catastrophic outcome cannot be assigned a precise probability is emphatically not a finding that its probability is negligible. Because the consequences of crossing certain thresholds would be severe and, in some cases, effectively permanent on human timescales, the expected-impact contribution of low-likelihood outcomes can be large even when their probability is small, and a risk assessment that omits them because they are hard to quantify systematically understates the total risk (Kemp, 2022) (Richardson, 2023). The framework adopted here therefore treats the assessment of low-likelihood, high-impact outcomes as a required component of a complete tipping-point assessment rather than as an optional appendix.

The risk framing: hazard, exposure, vulnerability, and cascading pathways

Tipping-point risk is most usefully assessed not as a property of the climate hazard alone but as the product of hazard, the exposure of people and systems, and their vulnerability, evaluated along pathways that may involve interacting and compounding drivers (*high confidence*). A framework for complex climate-change risk assessment recognizes that the consequences of a given physical change depend on how that change intersects with exposed populations, infrastructure, and ecosystems, and on the pre-existing vulnerabilities that determine whether a hazard becomes a disaster (Simpson, 2021). Applied to tipping points, this framing shifts the assessment question from “how likely is the threshold to be crossed” alone to “what would the crossing of the threshold mean for the systems that depend on the affected component, and through what pathways would the consequences propagate.”

This complex-risk framing is essential because the impacts of a tipping transition rarely remain confined to the physical component that tips. A reorganization of a major ocean current, a large release of permafrost carbon, or a dieback of a continental-scale forest would propagate through climatic teleconnections, through the carbon cycle, and through human systems of food, water, health, and settlement, with consequences that can compound and cascade (Simpson, 2021) (Pörtner, 2023). The complex-risk perspective also foregrounds the interaction of climatic hazards with non-climatic stresses — pollution, land-use change, biodiversity loss, and socio-economic fragility — that can lower the threshold at which a component tips and raise the severity of the resulting impact (Naidu, 2021) (Pörtner, 2023). The remainder of this chapter is organized around this framing: it first assesses the individual tipping elements and their thresholds, then the tools for detecting approach to a threshold, and finally the

interactions, cascades, and tail risks that determine the full societal significance of abrupt and irreversible change.

8.3 The Major Tipping Elements

The Atlantic Meridional Overturning Circulation

The Atlantic Meridional Overturning Circulation (AMOC) is assessed to possess threshold behaviour arising from a positive salinity-advection feedback, such that a sufficient input of freshwater to the North Atlantic could weaken or collapse the overturning with large climatic consequences for the Atlantic basin and beyond; there is high confidence that the AMOC will weaken over the coming century and low confidence in the level of forcing that would trigger an abrupt collapse (*medium confidence overall*). The physical basis for treating the AMOC as a tipping element is well established: the circulation is maintained in part by the sinking of cold, salty water in the subpolar North Atlantic, and this sinking is self-reinforcing because the northward transport of salt sustains the density contrast that drives the sinking. A large enough freshening of the surface waters — from ice-sheet melt, increased precipitation, or altered runoff — can weaken the sinking, which reduces salt transport, which further weakens the sinking, a feedback that in idealized models admits an abrupt transition to a much weaker or collapsed state.

The consequences of a substantial AMOC weakening would be far-reaching. They include a pronounced cooling of the North Atlantic region relative to the warming trend elsewhere, shifts in the position of tropical rainfall belts with implications for water availability across several continents, changes in regional sea level along Atlantic coasts, and altered patterns of ocean carbon and heat uptake. Because the overturning circulation redistributes heat on a hemispheric scale, its reorganization is one of the clearest examples of a tipping element whose impacts would propagate far from the site of the transition itself, coupling the North Atlantic to remote regions through atmospheric and oceanic teleconnections. The assessment of AMOC risk is thus inseparable from the assessment of cascading interactions treated later in this chapter.

The central difficulty in assessing AMOC risk is quantitative. There is broad agreement on the qualitative existence of the salinity-advection feedback and on the direction of the forced response, but the freshwater threshold that would separate a weakened-but-stable circulation from an irreversible collapse is poorly constrained by observations, and different classes of model disagree on how close the present circulation lies to such a threshold. The instrumental record of the overturning is short relative to the timescales of interest, which limits the power of statistical early-warning approaches when applied to the AMOC directly. The honest assessment is therefore one of high confidence in continued weakening under sustained forcing, genuine physical plausibility of an abrupt transition, and deep uncertainty about the probability and timing of such a transition — a combination that places the AMOC firmly within the class of

low-likelihood, high-impact outcomes that a complete risk assessment cannot set aside (Kemp, 2022).

The polar ice sheets and committed sea-level change

The Greenland and West Antarctic ice sheets are assessed to possess thresholds beyond which large-scale and effectively irreversible mass loss becomes self-sustaining, committing the Earth system to multi-metre sea-level rise over centuries to millennia even if warming is subsequently reversed (*high confidence in the existence of such thresholds; low confidence in their precise values*). The ice sheets exhibit several distinct self-amplifying feedbacks that give them the character of tipping elements. For the Greenland Ice Sheet, the surface-elevation–melt feedback is central: as the ice-sheet surface lowers through melting, it descends into warmer air at lower altitude, which increases melting further, so that beyond a threshold in sustained warming the ice sheet cannot maintain its elevation and enters a state of ongoing, self-reinforcing decline. For marine-based sectors of the West Antarctic Ice Sheet, the marine-ice-sheet instability associated with beds that deepen inland can allow grounding-line retreat to become self-sustaining once initiated.

The defining risk characteristic of the ice sheets is the combination of high commitment and long timescale. Crossing a threshold does not produce rapid sea-level rise within years; rather, it commits the system to a trajectory of loss that plays out over centuries to millennia and that cannot be halted by returning temperatures to their earlier values, because the feedbacks that sustain the loss depend on the altered geometry of the ice rather than on the instantaneous forcing. This hysteresis means that the sea-level consequences of decisions taken in the present may be, to a substantial degree, locked in long before they are fully expressed. For the assessment of exposure and vulnerability, the relevant implication is that coastal risk from the ice sheets is governed less by the rate of change observed in any given decade than by the cumulative warming that determines whether irreversible thresholds are crossed at all.

As with the AMOC, the principal uncertainty is the value of the threshold rather than its existence. Paleoclimate evidence that sea level has stood many metres higher than at present during past warm intervals confirms that the ice sheets can occupy dramatically reduced configurations, and reconstruction of the long-term history of sea level across geological time provides context for the magnitude of change of which the system is capable (Scotese, 2021). What the paleorecord cannot by itself provide is a precise mapping from a given level of warming to a given rate of loss under present-day rates of forcing, which are without close analogue in the reconstructed past. The assessment therefore expresses high confidence in eventual large sea-level commitment under sustained high warming and correspondingly lower confidence in the timing and the precise threshold.

Permafrost and the frozen-carbon feedback

Permafrost soils store a large quantity of organic carbon that becomes vulnerable to microbial decomposition and release as carbon dioxide and methane upon thaw, producing a positive feedback on global warming that is assessed to be substantial, pervasive, and — for the diffuse gradual-thaw component — more accurately characterized as an accelerating amplifier than as a single global abrupt tipping point, though localized abrupt-thaw processes can release carbon rapidly (*high confidence*). The permafrost region holds one of the largest reservoirs of potentially labile carbon in the Earth system, accumulated over millennia in frozen soils across the Arctic and boreal zones. As these soils warm and thaw, previously protected organic matter becomes available to decomposition, and the resulting emissions of carbon dioxide and, from waterlogged anaerobic settings, methane add to atmospheric greenhouse-gas concentrations and thereby amplify the warming that drove the thaw (Miner, 2022).

It is important to characterize the permafrost feedback precisely, because it differs in structure from the switch-like tipping elements. Across most of the permafrost domain, thaw and the associated carbon release are expected to proceed gradually and pervasively as the climate warms, producing a feedback that grows with warming rather than a single threshold beyond which the entire reservoir is released. Superimposed on this gradual process, however, are abrupt-thaw phenomena — including the collapse of ice-rich ground and the rapid development of thaw lakes and thermokarst — that can mobilize deep carbon far faster than gradual thaw and that give parts of the system a genuinely abrupt local character. The net feedback is therefore the sum of a large, spatially extensive gradual component and a smaller but locally intense abrupt component, and both add carbon to the atmosphere in a way that is not reversible on human timescales once released (Miner, 2022).

The permafrost feedback is a leading source of uncertainty in the assessment of committed warming. The size of the vulnerable carbon pool, the fraction that will decompose, the partition between carbon dioxide and the far more potent methane, and the rate at which release proceeds are all imperfectly constrained, and the feedback is incompletely represented in many Earth-system models. Because permafrost emissions are effectively irreversible additions to the atmospheric burden, and because they act to raise the warming that stresses every other tipping element, the permafrost feedback occupies a central place in the assessment of both committed change and interacting risks (Miner, 2022).

The Amazon forest and tropical forest dieback

The Amazon forest system is assessed to be capable of a large-scale transition from closed-canopy forest to a degraded, lower-biomass, more open state, driven by the interaction of warming, drought, deforestation, and fire, and reinforced by the forest's own influence on regional rainfall; observational evidence indicates a pronounced loss of resilience since the early 2000s over

a large fraction of the biome (*medium confidence*). The Amazon functions in part as a self-maintaining system: the forest recycles a substantial share of its own precipitation through evapotranspiration, so that a sufficiently large reduction in forest cover — whether from clearing, drought-driven mortality, or fire — can reduce the moisture supply to the remaining forest and push it toward a drier, fire-prone, savanna-like state. This moisture-recycling feedback is the physical basis for treating the Amazon as a tipping element, and it means that the forest's stability depends not only on external climate forcing but on the integrity of the forest itself (Flores, 2024).

The observational case for declining Amazon resilience has strengthened considerably. Analyses of vegetation state reveal that recovery rates from perturbation have slowed across more than three-quarters of the biome since the early 2000s, a signature consistent with critical slowing down as the system loses its capacity to return to its prior state after disturbance, and this loss of resilience is most pronounced in areas closer to human activity and with less rainfall (Boulton, 2022). Process-based synthesis identifies the principal drivers and stressors — rising temperatures, extreme drought, deforestation, and fire — and maps the compound thresholds beyond which large areas of the forest could transition, while emphasizing that the response is likely to be spatially heterogeneous rather than a single basin-wide switch (Flores, 2024). Underlying these system-level signals are the physiological mechanisms of tree mortality under drought and heat, which have been documented across forest biomes and which supply the process-level explanation for how a warming, drying climate translates into forest loss (Hammond, 2022) (McDowell, 2022).

The Amazon case also illustrates the coupling of a tipping element to broader Earth-system stability. Loss of Amazon forest would release stored carbon, reducing a globally important carbon sink and adding to the atmospheric burden, while altering regional climate in ways that would feed back on agriculture and water resources across South America (Flores, 2024) (Pörtner, 2023). Because the drivers are a mixture of global climate change and direct human land use, the Amazon is among the clearest examples of a tipping element whose threshold can be lowered by non-climatic pressures, so that reducing deforestation and fire raises the level of warming the forest can withstand before tipping (Flores, 2024). The assessment expresses medium confidence in the possibility of large-scale transition, high confidence in the observed loss of resilience, and continuing uncertainty about the precise combination of warming and land-use change that would trigger an irreversible outcome.

Warm-water coral reefs

Warm-water coral reef ecosystems are assessed to be the tipping element closest to a threshold at present levels of warming, with marine heatwaves already driving mass bleaching and mortality and with the prospect of functional collapse of many reef systems within the range of near-term warming (*high confidence*). Reef-building corals live close to their upper thermal limits, and sustained anomalous warmth causes them to expel their symbiotic algae —

bleaching — which, if prolonged or repeated, leads to coral death. The biological impacts of marine heatwaves extend beyond bleaching to include disease outbreaks, altered species interactions, and shifts in the geographic distribution of marine organisms, and these impacts have been documented across ocean basins as marine heatwaves have grown more frequent, longer, and more intense (Smith, 2022).

The reef case is distinctive among tipping elements in that the transition is already observable rather than prospective. Rather than a single global threshold to be crossed in the future, warm-water reefs are experiencing a progression of increasingly severe and frequent thermal-stress events that erode their capacity to recover between disturbances, in a manner directly analogous to the resilience-loss dynamics seen in forests: as the interval between damaging events shortens relative to the recovery time, the system loses the ability to return to its prior state and shifts toward a degraded configuration dominated by non-reef-building organisms (Smith, 2022). The consequences of widespread reef degradation extend to the hundreds of millions of people who depend on reefs for coastal protection, fisheries, and livelihoods, making warm-water corals an element in which high hazard, high exposure, and high vulnerability coincide (Smith, 2022) (Pörtner, 2023).

Boreal forests, drylands, and the wider set of vulnerable biomes

Beyond the classic tipping elements, a wider set of terrestrial systems — boreal forests, temperate and tropical woodlands, and dryland vegetation — show emerging, spatially coherent signals of declining resilience, indicating that destabilization of the land biosphere is not confined to the Amazon and that the land carbon sink itself may be at risk (*high confidence*). A satellite-based analysis of forest resilience finds that a large fraction of the world's forests has experienced a decline in the capacity to recover from perturbation since the early 2000s, with the strongest signals in the tropics, in arid regions, and where forests are exposed to greater climatic water stress; this represents an emerging, coherent signal of destabilization across biomes rather than an isolated regional phenomenon (Forzieri, 2022). Global field observations of tree die-off reveal a consistent “hotter-drought” fingerprint, in which the combination of high temperature and moisture deficit drives mortality across a wide range of forest types (Hammond, 2022).

The mechanisms linking climate to woody-plant mortality — hydraulic failure and carbon starvation under the combined stress of drought, rising carbon dioxide, and increasing vapour-pressure deficit — provide the process-level basis for expecting these resilience declines to continue as warming proceeds (McDowell, 2022). At finer scales, the buffering role of forest microclimates, which can decouple understorey conditions from the warming of the free atmosphere, modulates the exposure of forest biodiversity to macroclimatic change and is itself sensitive to canopy integrity, so that the loss of canopy cover both reflects and accelerates the loss of resilience (Frenne, 2021). Taken together, these lines of evidence indicate that the stability of the terrestrial biosphere

as a carbon sink and as a regulator of regional climate is being eroded across multiple biomes, with implications for the trajectory of atmospheric carbon dioxide and for the interacting risks assessed later in this chapter (Forzieri, 2022) (Hammond, 2022) (McDowell, 2022) (Frenne, 2021).

8.4 Early-Warning Signals and the Detection of Approaching Thresholds

The theory of critical slowing down

A system approaching a tipping threshold is expected, on theoretical grounds, to recover more slowly from perturbations, and this critical slowing down manifests as measurable increases in the temporal autocorrelation and variance of the system's state, providing a generic and physically motivated basis for early-warning indicators (*high confidence in the theory; medium confidence in its practical reliability*). The theoretical foundation of early-warning signals is that, as a system nears a bifurcation at which its current state loses stability, the rate at which it returns to equilibrium after a disturbance declines toward zero. This slowing of recovery has observable statistical consequences: because the system takes longer to relax, successive states become more similar to one another, so that lag-one autocorrelation rises, and because perturbations are damped less effectively, the variance of the state increases. These indicators are generic in the sense that they do not depend on knowing the detailed mechanism of the impending transition, which makes them appealing for systems whose governing feedbacks are understood only qualitatively.

The appeal of generic indicators is tempered by important limitations that the assessment must weigh. Critical slowing down is a property of a specific class of transitions and is not guaranteed to precede every abrupt change; some transitions can occur without detectable prior slowing, for instance when a system is subjected to a rapid shock rather than a slow approach to a threshold, or when the noise driving the system is itself changing. Conversely, rising autocorrelation and variance can arise from causes unrelated to an approaching bifurcation, such as changes in the observing system, in the forcing, or in the statistical properties of the driving noise, which creates a risk of false alarms. The reliable application of early-warning methods therefore requires long, high-quality, consistently sampled records, careful removal of confounding trends, and, ideally, corroboration from process understanding — conditions that are met for some tipping elements and not for others.

Resilience indicators applied to real systems

Where sufficiently long and spatially resolved observational records exist, resilience-based indicators have been successfully applied and have revealed coherent, large-scale declines in stability, most notably in the Amazon forest and in global forest resilience more generally (*high confidence*). The strongest empirical demonstration of early-warning methods in the climate system to

date comes from vegetation remote sensing. Analysis of the Amazon finds that the recovery rate of vegetation from perturbation has declined across more than three-quarters of the biome since the early 2000s — precisely the signature that critical-slowness theory predicts for a system losing resilience — and that this decline is spatially organized, being strongest where rainfall is lower and human pressure is greater (Boulton, 2022). The interpretation of this signal is reinforced by its consistency with independent process-based understanding of Amazon dieback, which supplies a mechanism for the loss of resilience rather than leaving it as an unexplained statistical artefact (Flores, 2024).

The same methodological approach applied at global scale reveals that declining resilience is not confined to the Amazon. A satellite-based analysis of forests worldwide finds emerging signals of declining resilience, measured through the same class of temporal-autocorrelation indicators, across a large fraction of the global forest area, with the strongest declines in the tropics and in water-limited regions (Forzieri, 2022). The convergence of these independent analyses — different regions, overlapping methods, consistent theoretical basis, and corroborating process understanding — represents one of the more robust applications of early-warning theory in Earth-system science and lends weight to the interpretation that multiple biosphere components are losing stability under climate change (Boulton, 2022) (Forzieri, 2022) (Flores, 2024).

Limits of detection and the possibility of surprise

Early-warning indicators are most reliable for slowly forced systems with long observational records and are least reliable precisely for some of the highest-consequence tipping elements, so the assessment must treat the absence of a detected warning signal as weak evidence of safety rather than as reassurance (*high confidence*). The tipping elements for which early-warning methods work best — forests and reefs observed by satellite over two decades — are not, in general, the elements whose transitions would be most catastrophic. For the ocean overturning circulation and the ice sheets, the observational records that would be needed to apply resilience indicators with confidence are short relative to the timescales of the underlying dynamics, the systems are forced by processes that are themselves changing, and the relevant state variables are difficult to observe directly and continuously. In these cases the statistical power to detect an approaching threshold is limited, and a failure to detect a warning signal cannot be interpreted as evidence that no threshold is being approached.

This asymmetry has a direct bearing on how the assessment treats surprise. The prudent inference from the theory and the observations together is that early-warning signals, where present and well corroborated, add valuable evidence that a system is losing resilience, but that their absence does not license the conclusion that a transition is far off. Because some transitions can occur with little or no prior statistical warning, and because the systems most capable of such surprise are among the most consequential, the possibility of abrupt change arriving faster or at lower forcing than anticipated remains an essential element of the risk assessment (Kemp, 2022). Early-warning science strengthens

the case for concern where signals appear, but it does not provide a reliable all-clear where they do not.

8.5 Cascading Interactions and Systemic Risk

Interactions among tipping elements

Tipping elements are physically coupled through the ocean, atmosphere, and carbon cycle, and these couplings can be mutually reinforcing, so that the destabilization or transition of one element can raise the probability of transition in another and, in principle, propagate as a cascade (*medium confidence*). The tipping elements assessed in this chapter are not independent switches but components of a single coupled system, and the physical connections among them are numerous and, in several cases, self-amplifying. A weakening of the Atlantic overturning circulation would alter heat and moisture transport in ways that bear on the water balance of tropical forests; permafrost thaw and forest dieback would each add carbon to the atmosphere and thereby raise the warming that stresses every other element; and changes in ice-sheet melt would deliver freshwater to the North Atlantic in a way that couples the cryosphere directly to the overturning circulation. These interactions mean that the stability of the whole cannot be assessed by considering each element in isolation.

The consequence of positive coupling is that the effective threshold for any given element may be lower in a system where other elements are already destabilizing than it would be for that element considered alone, and that the joint probability of multiple transitions is not simply the product of independent probabilities. The overall effect of the network of interactions is assessed to be predominantly destabilizing rather than stabilizing, which raises the systemic risk above what an element-by-element assessment would suggest, though the strength, sign in specific cases, and timing of individual couplings remain uncertain and are an active area of research (Pörtner, 2023). The assessment therefore treats cascading interaction as a genuine amplifier of tipping-point risk while acknowledging that quantifying the cascade remains beyond current capability.

Coupling of climate and biodiversity

The climate and biodiversity crises are deeply intertwined, with climate change driving biodiversity loss and the degradation of ecosystems in turn eroding the capacity of the biosphere to store carbon and to buffer climate, forming a coupled system whose joint destabilization is more dangerous than either crisis alone (*high confidence*). A central conclusion of recent Earth-system assessment is that climate change and the loss of biological diversity cannot be assessed or addressed as separate problems, because the same warming that stresses tipping elements also drives species loss, and the resulting degradation of ecosystems weakens their regulation of carbon and climate (Pörtner, 2023). The tipping elements of the biosphere — forests and reefs above all

— are simultaneously reservoirs of biodiversity and regulators of the carbon cycle, so that their transition is at once a biodiversity catastrophe and a climate feedback (Pörtner, 2023) (Smith, 2022) (Flores, 2024).

Independent lines of evidence document the biodiversity side of this coupling. Assessments of insect decline under climate change warn of consequences for pollination, food webs, and ecosystem function that would propagate through the systems on which human societies depend (Harvey, 2022), while broader analyses raise the question of whether the current episode of biodiversity loss constitutes a mass extinction comparable in kind to those recorded in the geological past (Cowie, 2022). The chemical contamination of the biosphere adds a further, partly independent stress that erodes ecosystem resilience and interacts with climatic pressures (Naidu, 2021). The integrated implication is that the resilience of the living systems that regulate climate is being reduced by multiple simultaneous drivers, lowering the buffering capacity that would otherwise damp the approach to tipping thresholds (Pörtner, 2023) (Folke, 2021) (Naidu, 2021).

Compound climatic hazards and regional cascades

Climate change is increasing the frequency and severity of compound and connected hazards — droughts, heatwaves, marine heatwaves, and their consequences for water and food — that can interact and cascade across regions and sectors, and that provide the proximate mechanisms through which slow-onset tipping processes translate into acute societal harm (*high confidence*).

The pathways by which a destabilizing biosphere harms human systems run through specific, increasingly frequent hazards. A projected global shift toward more rapidly developing “flash” droughts under continued warming would compress the time available to respond to the onset of water and agricultural stress, raising the danger to food and water systems (Yuan, 2023). Marine heatwaves, in addition to their role in reef collapse, disrupt fisheries and marine ecosystems on which coastal populations depend (Smith, 2022). In particularly exposed regions such as the Eastern Mediterranean and Middle East, the compounding of extreme heat and aridity is projected to reach severity that challenges habitability and agriculture (Zittis, 2022).

These hazards do not act one at a time. Water scarcity has quantifiable economic consequences that propagate through trade and production systems well beyond the water-stressed region itself (Dolan, 2021), and the availability of fresh water is a growing global constraint that interacts with drought, heat, and ecosystem degradation (Mishra, 2023). The connectedness of these hazards is precisely what makes tipping-point risk a systemic rather than a local concern: a slow-onset process such as forest dieback or overturning weakening expresses itself through acute, compounding, cross-regional hazards whose combined impact on food, water, and health can exceed the sum of the individual effects and can propagate through economically and politically connected systems (Simpson, 2021) (Dolan, 2021) (Yuan, 2023).

8.6 Abrupt Change in the Paleoclimate Record

The paleorecord as evidence that the Earth system can reorganize

The geological and paleoclimate record establishes that the Earth system has occupied a wide range of markedly different states and has undergone reorganizations that were abrupt relative to their causes, providing direct empirical evidence that nonlinear and threshold behaviour are real properties of the climate system rather than artefacts of models (*high confidence*).

The single most important contribution of the paleorecord to the assessment of tipping points is existential: it demonstrates that the modes of behaviour predicted by tipping theory have actually occurred. Over the Phanerozoic — the last half-billion years — reconstruction of continental configuration, sea level, and climate documents transitions between greenhouse and icehouse worlds, repeated advances and retreats of the seas across the continents, and the rise and disappearance of ice sheets, confirming that the system can move between configurations that differ profoundly from the present (Scotese, 2021). This long view supplies the range of states the Earth system is capable of occupying and thereby bounds the magnitude of change that tipping processes could, in principle, produce.

The value of the deep-time perspective is that it removes any presumption that the relatively stable climate of recent millennia is the system's only accessible state. The reconstructed history shows sea levels tens of metres different from today's, ice sheets present and absent, and climate zones displaced across latitudes, all as expressions of the same physical system responding to changes in forcing and boundary conditions (Scotese, 2021). For the assessment of the ice sheets in particular, the paleorecord's evidence that sea level has stood far higher than at present during past warm intervals is a direct constraint on the magnitude of sea-level change of which the cryosphere is capable, even if the precise mapping from a given warming to a given ice loss under modern rates of forcing cannot be read off from the past (Scotese, 2021). The paleorecord thus functions as an existence proof for large-scale reorganization and as a source of analogues, while stopping short of providing a deterministic forecast.

What the paleorecord constrains, and what it does not

Paleoclimate evidence robustly constrains the range of states the Earth system can occupy and the reality of abrupt transitions between them, but it provides only limited constraint on the precise thresholds and rates that would apply under the unprecedented speed of present-day forcing (*high confidence*). The strengths and limits of the paleorecord for tipping-point assessment must be stated carefully. Its strengths are that it records real transitions in the real system, integrating all the relevant feedbacks and interactions that models must approximate, and that it spans a range of climate states far wider than the instrumental era, capturing behaviour that would otherwise be inaccessible. Its evidence that the system has repeatedly reorganized is among the firmest

foundations for treating abrupt change as a genuine risk (Scotese, 2021).

The principal limit is that the rate and character of present-day forcing have few close analogues in the reconstructed past. The current rate of change of atmospheric composition is rapid by the standards of most paleoclimate transitions, and the configuration of continents, ice, and biosphere that determines how the system responds is specific to the present. As a result, the paleorecord is better at establishing that thresholds exist and that transitions can be abrupt than at pinning down the exact level of forcing at which a particular modern tipping element would transition or the exact rate at which the transition would unfold. This is why the assessment repeatedly pairs high confidence in the qualitative existence of threshold behaviour with lower confidence in quantitative thresholds: the paleorecord underwrites the former far more strongly than the latter. The appropriate use of paleoclimate evidence in this chapter is therefore to establish plausibility, magnitude, and the reality of abrupt reorganization, while relying on process understanding and contemporary observations for the assessment of proximity to specific thresholds (Scotese, 2021) (Kemp, 2022).

Lessons from past states for the assessment of future risk

Past warm climate states, in which sea levels and the distribution of ice differed greatly from the present, indicate that sustained warming commits the Earth system to large, slow, and effectively irreversible adjustments of the cryosphere and sea level, reinforcing the assessment that some tipping outcomes are matters of committed change rather than avoidable events once thresholds are crossed (*high confidence*). The reconstructed history of warm intervals carries a specific and sobering lesson for the assessment of committed change. In those intervals, the configuration of ice and the level of the sea adjusted to the prevailing climate over long timescales, reaching states far from the present when the climate was sustainedly warmer (Scotese, 2021). The implication for the present is that if warming is sustained at elevated levels, the slow components of the system — above all the great ice sheets — will tend, over centuries to millennia, toward the configurations that such warmth implies, and that this adjustment, once underway past a threshold, is not reversible on any timescale relevant to human planning.

This lesson connects the deep past directly to the framing of risk that organizes this chapter. It means that a portion of tipping-point risk is not about the probability of a discrete future event but about the commitment being accumulated in the present through sustained forcing — a commitment whose full expression lies far in the future but whose determination is contemporaneous with today's emissions (Scotese, 2021). It also means that the reversibility that characterizes much of the climate system does not extend to these slow tipping elements: unlike surface temperature, which responds relatively quickly to changes in forcing, the ice sheets and sea level carry a long memory, so that the paleorecord's warm states represent not a transient excursion but the kind of enduring reorganization to which sustained warming tends. The assessment draws from this the conclusion that the treatment of committed and irreversible

change must be a central rather than peripheral part of tipping-point risk, and that the paleorecord, by displaying the destinations toward which sustained warming points, is essential evidence for that treatment (Scotese, 2021) (Kemp, 2022).

8.7 Low-Likelihood, High-Impact Outcomes and the Framing of Risk

Why the tail of the distribution demands assessment

Because the consequences of crossing certain climate thresholds would be severe, widespread, and in some cases effectively permanent on human timescales, low-likelihood high-impact outcomes can dominate the expected harm from climate change and must be assessed explicitly rather than excluded because their probabilities are hard to quantify (*high confidence*). A defensible risk assessment weights outcomes by both their probability and their consequence, and for a class of outcomes in which the consequence is extreme, even a modest probability yields a large contribution to expected harm. The tipping elements assessed in this chapter populate exactly this class: an abrupt weakening of the ocean overturning circulation, an irreversible commitment to multi-metre sea-level rise, a large and self-sustaining release of permafrost carbon, or a continental-scale forest dieback would each impose consequences whose scale and irreversibility place them in the tail of the climate risk distribution. The argument for dedicated assessment of this tail is that it has been comparatively neglected relative to central-estimate impacts, and that the neglect is not justified by the evidence, since the physical plausibility of these outcomes is not in doubt even where their probability is uncertain (Kemp, 2022).

The case for foregrounding tail risk is strengthened by the structure of the uncertainty involved. For the highest-consequence tipping elements, uncertainty is not merely statistical scatter around a well-defined mean but deep uncertainty about the location of thresholds and the strength of feedbacks, of a kind that resists reduction to a single probability distribution. In such circumstances, the temptation to disregard outcomes that cannot be assigned a precise probability is precisely backwards: the same deep uncertainty that prevents confident quantification also prevents confident dismissal, and the appropriate response is to assess the range of plausible outcomes, including the severe tail, rather than to collapse the assessment onto a central estimate that omits the most consequential possibilities (Kemp, 2022) (Richardson, 2023). This chapter therefore treats the systematic assessment of catastrophic and extreme outcomes as a scientific obligation rather than as speculation.

The spectrum of high-impact outcomes

High-impact tipping outcomes span a spectrum from severe regional disruption through globally significant reorganization of the climate system to, at

the extreme and least likely end, outcomes that would threaten the stability of human societies at large scales; the assessment can characterize this spectrum qualitatively even where it cannot assign probabilities across it (*medium confidence*). At the more probable and better-constrained end of the spectrum lie outcomes already partly observable: the functional collapse of many warm-water reef systems, severe regional water and food crises driven by compound drought and heat, and the loss of resilience across large areas of forest (Smith, 2022) (Yuan, 2023) (Forzieri, 2022). These outcomes are severe and their occurrence at some level under continued warming is assessed with relatively high confidence, even though they fall short of global reorganization.

Toward the more consequential end lie outcomes whose probability is lower or more uncertain but whose impact would be correspondingly greater: an abrupt weakening or collapse of the Atlantic overturning circulation, with hemispheric consequences for temperature and rainfall; a committed disintegration of marine-based ice, with multi-metre sea-level implications; and a large permafrost carbon release that would materially raise committed warming and thereby stress every other element. At the extreme end, catastrophe-focused analysis argues that the possibility of outcomes severe enough to threaten the fabric of human societies — through the convergence of climatic disruption with food, water, health, and conflict pathways — is not excluded by current evidence and has been insufficiently studied (Kemp, 2022). The value of characterizing this full spectrum is that it prevents the assessment from implicitly truncating the range of possibilities at the boundary of what can be comfortably quantified, and it makes explicit that the difference between the more and less probable outcomes is one of likelihood, not of physical impossibility (Kemp, 2022) (Richardson, 2023).

Systemic and cascading risk to human societies

The gravest tipping-related risks to human societies arise not from any single physical transition but from the potential for climatic shocks to interact with the vulnerabilities and interconnections of human systems, producing cascading failures that propagate across food, water, health, economic, and political domains — a systemic risk that current tools assess only partially (*medium confidence*). The translation of physical tipping into societal catastrophe runs through the coupling of hazards to exposed and vulnerable human systems, and through the propagation of shocks across the connections that bind those systems together. A framework for complex climate risk emphasizes that impacts arise from the intersection of hazard, exposure, and vulnerability, and that in an interconnected world the consequences of a shock in one place or sector can transmit to others through trade, finance, migration, and governance (Simpson, 2021). Water scarcity in one region imposes economic costs elsewhere through connected production systems (Dolan, 2021); a shock to food production can transmit through global markets; and the simultaneity of multiple stresses can overwhelm the capacity to respond.

Recent scholarship has begun to formalize this systemic dimension under

the heading of interacting global crises, arguing that distinct stresses — climatic, ecological, economic, and political — can become entangled so that they amplify one another and produce a compound crisis greater than the sum of its parts (Lawrence, 2024). The introduction of new technologies into these tightly coupled systems can itself become a source of systemic risk, adding pathways through which shocks propagate and interact (Galaz, 2021). The relevance of this literature to tipping-point assessment is that it identifies the mechanisms by which a physical transition — itself perhaps regional and slow — could become a global societal catastrophe through the entanglement of climate with the other systems on which societies depend (Lawrence, 2024) (Galaz, 2021) (Simpson, 2021). The assessment of these systemic pathways is at an early stage and carries substantial uncertainty, but the pathways are real and their potential consequences severe enough that they belong within a complete assessment of tipping-point risk (Lawrence, 2024) (Pörtner, 2023).

The health dimension of tipping-point risk

The human-health consequences of climate change provide a concrete and increasingly well-documented channel through which slow-onset and abrupt climatic changes translate into harm, and they demonstrate that the impacts of a destabilizing climate are already being realized rather than confined to the future (*high confidence*). Systematic tracking of the intersection of climate and health documents that the hazards associated with warming — extreme heat, altered patterns of infectious disease, disruption of food and water systems, and the health consequences of extreme events — are producing measurable harm to human health, and that these harms grow with continued warming and delayed action (Romanello, 2023) (Romanello, 2024). This health evidence anchors the assessment of tipping-point risk in observable present-day consequences and provides a metric of harm that integrates across many of the physical pathways discussed in this chapter.

The health channel is also a lens on the distribution of tipping-point risk, because the harms fall unevenly, concentrating on populations with the least capacity to adapt and thereby compounding existing vulnerabilities (Romanello, 2024) (Simpson, 2021). This unevenness matters for the framing of risk: a tipping transition that appears moderate in globally aggregated terms may be catastrophic for specific exposed and vulnerable populations, and an assessment that considered only global aggregates would understate the human significance of the change. The most recent tracking emphasizes record-breaking threats arising from the delay in reducing the underlying drivers, underscoring that the window in which the approach to tipping thresholds could be slowed is itself a function of near-term action (Romanello, 2024). The health evidence thus links the physical assessment of tipping elements to the lived consequences for people and to the urgency of the choices that determine how close the system is driven to its thresholds (Romanello, 2023) (Romanello, 2024).

Implications for the goal of limiting warming

The assessment of tipping points reinforces the rationale for limiting the level and duration of warming, because the probability of crossing the thresholds of the major tipping elements rises with both the magnitude and the persistence of the temperature increase, and because some thresholds, once crossed, commit the system to irreversible change regardless of subsequent mitigation (*high confidence*). The organizing implication of this chapter for policy-relevant assessment is that tipping-point risk is a monotonic and, in the neighbourhood of thresholds, sharply increasing function of warming. Every increment of avoided warming reduces the probability of crossing thresholds, and — for elements with hysteresis — the timing of stabilization matters as well as the eventual level, because a temporary overshoot of a threshold can trigger an irreversible transition even if temperatures are later brought back down. This gives the concept of net-zero emissions its physical meaning for tipping-point risk: stabilizing the climate requires bringing net emissions to zero, and the level at which stabilization occurs determines how many thresholds are approached or crossed (Fankhauser, 2021).

The precise definition and implementation of net zero therefore have direct consequences for tipping-point risk. Getting net zero right — in its coverage of gases, its treatment of removals, and its timing — determines the peak warming to which the system is exposed and hence the degree to which the thresholds of the tipping elements are challenged (Fankhauser, 2021). The tools available to support this effort, including the application of computational methods across mitigation and adaptation, expand the capacity to act, though they are enablers rather than substitutes for reducing the underlying drivers (Kaack, 2022). The assessment does not prescribe a particular level of ambition, which lies outside the scope of a Working Group I evaluation, but it establishes the physical basis for the proposition that the avoidance of tipping-point risk is a strong function of how much and how long the planet is warmed, and that the irreversibility of several tipping elements makes the near-term trajectory of emissions consequential far beyond the near term (Fankhauser, 2021) (Kaack, 2022) (Kemp, 2022).

8.8 Detection, Monitoring, and the State of Knowledge

The role of sustained observation and emerging analytical tools

The detection of approaching thresholds depends on sustained, high-quality, long-duration observation of the relevant tipping elements, and the analytical tools now available — including satellite remote sensing and computational methods for pattern detection — have materially improved the capacity to monitor resilience, while remaining insufficient for the elements with the shortest observational records (*high confidence*). The empirical advances reported in this chapter — the detection of resilience loss in the Amazon and in global forests — were made possible by the accumulation of consistent, multi-

decade satellite records combined with analytical methods capable of extracting resilience indicators from noisy data (Boulton, 2022) (Forzieri, 2022). This demonstrates a general principle: the ability to see a tipping element approaching its threshold is a direct function of the length, quality, and continuity of the observational record, and of the sophistication of the methods brought to bear on it. The strongest results come precisely from those elements for which both the data and the methods are mature.

The expanding toolkit of computational methods offers the prospect of further gains, both in the detection of early-warning signals and in the broader project of understanding and responding to climate change (Kaack, 2022). Such methods can extract structure from large and heterogeneous Earth-observation datasets, integrate disparate lines of evidence, and support the monitoring of systems that are difficult to observe directly. At the same time, the assessment cautions that analytical sophistication cannot substitute for the underlying observations: where the record is short, as it is for the ocean overturning circulation and for the dynamic behaviour of the ice sheets, no analytical method can fully compensate, and the priority is the establishment and maintenance of the sustained observing systems that would, over time, make robust detection possible (Kaack, 2022) (Boulton, 2022). The improvement of monitoring capacity for the least-observed, highest-consequence elements is accordingly identified as a first-order priority.

Principal knowledge gaps and research priorities

The most consequential knowledge gaps in the assessment of tipping points concern the quantitative location of thresholds, the strength and sign of interactions among elements, and the behaviour of the coupled human–Earth system under compound stress; reducing these gaps is essential for moving from qualitative plausibility to calibrated probability (*high confidence*). Several specific gaps recur across the tipping elements assessed in this chapter. First, for every major element the qualitative existence of threshold behaviour is far better established than the quantitative value of the threshold, so that the assessment can state that a transition is physically possible with far more confidence than it can state the level of warming at which it becomes likely. Closing this gap requires improved process understanding, better observational constraints, and, where possible, the extraction of quantitative information from paleoclimate analogues, tempered by the recognition that present-day forcing rates have few close past parallels (Scotese, 2021) (Miner, 2022).

Second, the interactions among tipping elements — their sign, strength, and timing — are a critical and underdeveloped area, because these interactions determine whether the system's total risk exceeds the sum of its parts and whether cascades are plausible (Pörtner, 2023). Third, the coupling of the biosphere to the climate, including the destabilization of the land and ocean carbon sinks, needs to be represented more completely in Earth-system models, since the resilience declines now being observed in forests and reefs are incompletely captured by models that treat vegetation and carbon uptake as stable (Forzieri,

2022) (Hammond, 2022) (McDowell, 2022). Fourth, the systemic risk arising from the interaction of climatic tipping with human systems — the domain of cascading societal impacts and polycrisis dynamics — is at an early stage of formalization and requires the development of methods that can represent the propagation of shocks across coupled human and natural systems (Lawrence, 2024) (Simpson, 2021) (Galaz, 2021). Progress on each of these fronts would move the assessment of specific outcomes from statements of physical plausibility toward calibrated probabilistic judgements.

Synthesis

Taken together, the evidence establishes that abrupt and irreversible change is a real and present feature of the climate system, that several major tipping elements show observational signs of declining stability, and that the associated risks — including low-likelihood high-impact and systemic outcomes — rise with the magnitude and duration of warming and warrant a central place in the assessment of climate risk (*high confidence*). The synthesis of this chapter rests on several convergent conclusions. The reality of tipping behaviour is established by process understanding and by the paleoclimate record's demonstration that the Earth system has repeatedly reorganized (Scotese, 2021) (Folke, 2021). The proximity of specific elements to their thresholds is best established for warm-water reefs, which are already transitioning, and for forests, in which resilience loss is directly observed, and least well established for the ocean overturning circulation and the ice sheets, whose thresholds are physically real but quantitatively uncertain (Smith, 2022) (Boulton, 2022) (Forzieri, 2022). The permafrost feedback is assessed as a large, pervasive, and effectively irreversible amplifier of warming rather than a single global switch, and it raises the stress on every other element (Miner, 2022).

The interactions among these elements, and between them and the human systems that depend on climate stability, are assessed to be predominantly amplifying, so that the systemic risk exceeds what an element-by-element accounting would imply, even though the quantification of these interactions remains immature (Pörtner, 2023) (Lawrence, 2024) (Simpson, 2021). Because the highest-consequence outcomes are also among the most deeply uncertain, and because several are irreversible once triggered, a complete assessment must weight the tail of the risk distribution explicitly and must recognize that the accumulation of commitment in the present is itself a form of risk (Kemp, 2022) (Richardson, 2023). Finally, the whole of this risk is a rising function of warming and of its persistence, which grounds the physical case for limiting both the level and the duration of the temperature increase and for the proper definition and pursuit of net-zero emissions (Fankhauser, 2021) (Kaack, 2022). The assessment of tipping points thus concludes not with a single number but with a structured judgement: that the possibility of abrupt, irreversible, and cascading change is well founded, that it becomes more likely the further and longer the planet is warmed, and that its potential consequences are grave enough to be treated as a central rather than a peripheral concern of climate science (Kemp, 2022)

(Richardson, 2023) (Pörtner, 2023).

8.9 Regional and Sectoral Expression of Tipping-Point Risk

Water: drought, scarcity, and the acceleration of onset

The water cycle is a principal medium through which tipping and abrupt processes are expressed as harm, and both the projected acceleration of drought onset and the growing scarcity of fresh water indicate that water-related impacts will intensify and arrive faster under continued warming (*high confidence*). Water is the connective tissue between many of the tipping elements assessed in this chapter and the human and ecological systems that experience their consequences. A projected transition toward more rapidly developing droughts under continued warming means that the interval between the onset of dry conditions and the emergence of severe agricultural and hydrological stress is expected to shorten, compressing the time available for detection, warning, and response and thereby raising the effective severity of a given drought (Yuan, 2023). This acceleration is itself a form of abrupt change in the character of a hazard, distinct from any change in its mean frequency, and it interacts with the slow-onset destabilization of forests and reefs that depends on water availability.

The scarcity of fresh water is simultaneously a growing global constraint in its own right. The availability of fresh water is under mounting pressure from the combined effects of climatic change, demand, and the degradation of the systems that store and regulate it, and this scarcity carries consequences that extend well beyond directly affected regions (Mishra, 2023). Analysis of the economic dimension shows that water scarcity in one location imposes measurable costs elsewhere through the connections of trade and production, so that a regional water crisis can propagate into a broader economic disturbance (Dolan, 2021). For the assessment of tipping-point risk, the significance of the water cycle is that it provides a fast, connected pathway by which slow physical changes — forest dieback, altered circulation, shifting rainfall belts — are translated into acute, cross-regional harm to food systems, economies, and health (Yuan, 2023) (Dolan, 2021) (Mishra, 2023).

Oceans and marine ecosystems beyond the reefs

Marine heatwaves and the broader warming of the ocean impose biological impacts that extend well beyond warm-water coral reefs, restructuring marine ecosystems and the services they provide, and constituting an abrupt-hazard expression of ocean warming that affects fisheries and coastal communities across latitudes (*high confidence*). The ocean's role in tipping-point risk is not limited to the overturning circulation and to reef collapse. Marine heatwaves — discrete periods of anomalously warm ocean temperatures — produce biological impacts that include mass mortality of marine organisms, shifts in the geographic ranges of species, changes in the structure of marine food

webs, and outbreaks of disease, and these impacts have been documented across ocean basins as such events have grown more frequent, longer, and more intense (Smith, 2022). Because marine heatwaves are episodic and rapid in onset, they represent an abrupt-hazard expression of the underlying trend of ocean warming, delivering shocks to marine ecosystems in a manner analogous to how heatwaves and flash droughts deliver shocks on land (Smith, 2022) (Yuan, 2023).

The consequences of these marine impacts propagate into human systems through fisheries, coastal protection, and the livelihoods of communities that depend on productive coastal waters. As marine ecosystems are restructured by repeated thermal shocks, the services they provide — food, income, and shoreline protection — are placed at risk, and the populations affected are frequently those with limited capacity to adapt (Smith, 2022) (Simpson, 2021). The ocean thus contributes to tipping-point risk along two coupled axes: the slow, high-consequence, deeply uncertain axis of circulation and ice, and the faster, better-observed, already-realized axis of marine heatwaves and ecosystem disruption. Both axes matter for the assessment, and the more immediate axis provides observable evidence that the ocean is already delivering climate-driven harm to ecosystems and people (Smith, 2022).

Biodiversity, ecosystems, and the erosion of buffering capacity

The degradation of biodiversity and ecosystems, driven by climate change acting in concert with pollution and land-use pressures, erodes the buffering capacity of the biosphere and thereby lowers the resilience that would otherwise damp the approach to tipping thresholds, making biodiversity loss both a consequence and an amplifier of tipping-point risk (*high confidence*). The living systems of the Earth do not merely respond passively to climate change; they regulate the carbon cycle, moderate regional climate, and provide the resilience that allows the system to absorb disturbance. The widespread decline of biodiversity therefore matters for tipping-point risk beyond its intrinsic significance, because it reduces the capacity of ecosystems to buffer perturbations and to recover from them (Pörtner, 2023) (Folke, 2021). Warnings of severe insect decline point to consequences for pollination and for the food webs that structure terrestrial and freshwater ecosystems (Harvey, 2022), and the broader question of whether the present episode of loss constitutes a mass extinction underscores the scale of the ongoing degradation (Cowie, 2022).

These biological stresses do not act alone. The pervasive chemical contamination of the environment constitutes an additional, partly independent pressure on ecosystems, one that can interact with climatic stress to lower the threshold at which ecological systems fail and that has itself been characterized as a potential catastrophic risk (Naidu, 2021). The combination of climatic warming, biodiversity loss, and chemical pollution represents a multi-stressor assault on the resilience of the biosphere, and it is this compounding of stresses — rather than any one of them in isolation — that most threatens the buffering capacity on which the stability of the tipping elements depends (Naidu, 2021)

(Pörtner, 2023) (Folke, 2021). The assessment concludes that the protection and restoration of biodiversity and ecosystem integrity is therefore not separable from the management of tipping-point risk, because the resilience of the living biosphere is a determinant of how close the climate system can be driven to its thresholds before abrupt change ensues (Pörtner, 2023) (Folke, 2021) (Richardson, 2023).

8.10 Assessment Confidence and Key Uncertainties

What is assessed with high confidence

A core set of conclusions about abrupt and irreversible change is supported by multiple, independent lines of evidence and is assessed with high confidence, notwithstanding the deep uncertainty that attends the quantitative details of individual thresholds (*high confidence*). The evidence base supports several robust conclusions. It is assessed with high confidence that tipping behaviour – threshold-crossing, self-amplifying, and imperfectly reversible transition – is a genuine property of components of the Earth system, established both by process understanding and by the paleoclimate record’s demonstration of repeated large-scale reorganization (Scotese, 2021) (Folke, 2021). It is assessed with high confidence that warm-water coral reefs are already undergoing climate-driven degradation through the biological impacts of marine heatwaves (Smith, 2022), and with high confidence that forests across multiple biomes show emerging, coherent signals of declining resilience under climate change, with a hotter-drought mortality mechanism operating across forest types (Forzieri, 2022) (Hammond, 2022) (McDowell, 2022) (Frenne, 2021). It is assessed with high confidence that permafrost thaw constitutes a large, pervasive, and effectively irreversible positive feedback on warming (Miner, 2022), and with high confidence that the climate and biodiversity crises are coupled such that the degradation of ecosystems erodes the biosphere’s capacity to buffer climatic change (Pörtner, 2023).

It is further assessed with high confidence that the human consequences of a destabilizing climate are already measurable, most clearly through the tracked intersection of climate and health, and that these consequences grow with continued warming and delayed action (Romanello, 2023) (Romanello, 2024). And it is assessed with high confidence that tipping-point risk, taken as a whole, is a rising function of the magnitude and duration of warming, and that the systematic assessment of low-likelihood high-impact outcomes is a necessary rather than optional component of a complete evaluation of climate risk (Kemp, 2022) (Richardson, 2023) (Fankhauser, 2021).

The principal residual uncertainties

The principal residual uncertainties concern the quantitative values of thresholds, the strength and sign of interactions among tipping elements, and the dynamics of the coupled human–Earth system, and these uncertainties are

deep rather than merely statistical, so that several of the most consequential outcomes cannot at present be assigned calibrated probabilities (*high confidence in this characterization of the uncertainty*). Against the robust qualitative conclusions stands a set of stubborn quantitative uncertainties. For the Atlantic overturning circulation and the ice sheets, the existence of thresholds is not in doubt, but their location in terms of the level of warming, and the rate at which transitions would proceed, are poorly constrained, in part because the instrumental records are short and present-day forcing rates lack close paleoclimate analogues (Scotese, 2021) (Kemp, 2022). For the Amazon and other forests, the observed loss of resilience is robust, but the precise combination of warming, drought, deforestation, and fire that would trigger an irreversible large-scale transition remains uncertain and spatially variable (Boulton, 2022) (Flores, 2024). For permafrost, the size of the vulnerable carbon pool, the partition between carbon dioxide and methane, and the rate of release are imperfectly known (Miner, 2022).

Above the level of individual elements, the interactions among them — the couplings that determine whether the system's total risk exceeds the sum of its parts and whether cascades are plausible — are assessed to be predominantly amplifying but are quantitatively immature, representing one of the most important frontiers for reducing uncertainty (Pörtner, 2023). The behaviour of the coupled human–Earth system under compound stress, including the systemic and polycrisis dynamics through which physical tipping could become societal catastrophe, is at an early stage of formalization and carries correspondingly large uncertainty (Lawrence, 2024) (Galaz, 2021) (Simpson, 2021). The honest summary is that the assessment is confident about the reality and the direction of tipping-point risk and about its dependence on warming, while remaining genuinely uncertain about the precise thresholds, timings, and interactions that would determine exactly how and when specific abrupt changes would unfold — an uncertainty that, given the severity and irreversibility of the outcomes involved, is itself a reason for concern rather than for reassurance (Kemp, 2022) (Richardson, 2023).

The distinction between confidence and reassurance

A central interpretive principle of this chapter is that uncertainty about the timing and magnitude of tipping-point outcomes must not be read as evidence of their improbability, because for outcomes that are severe and irreversible the deep uncertainty surrounding them is a component of the risk rather than a mitigation of it (*high confidence*). It is a recurring error in the public interpretation of tipping-point science to treat the absence of a confident prediction as though it were a confident prediction of absence. The assessment resists this error explicitly. Where the science cannot yet say how close a threshold is, or how fast a transition would proceed, the appropriate inference is that the outcome is possible over a range that includes the more dangerous end, not that it can be set aside. This is especially so for the elements — the overturning circulation, the ice sheets — whose thresholds are least well constrained and

whose consequences would be gravest, precisely the combination in which the temptation to discount the uncertain is strongest and least justified (Kemp, 2022).

This principle has practical force for how the findings of this chapter should be used. The high-confidence conclusions — that tipping behaviour is real, that several elements are destabilizing, that the risk rises with warming, and that irreversibility makes near-term choices consequential far into the future — are sufficient on their own to establish that abrupt and irreversible change is a serious and present dimension of climate risk (Richardson, 2023) (Folke, 2021) (Fankhauser, 2021). The residual uncertainties do not soften this conclusion; they sharpen it, because they attach to outcomes whose severity and permanence mean that even a modest probability carries a large expected cost, and because they cluster around the very elements whose transition would be most catastrophic (Kemp, 2022) (Pörtner, 2023). The assessment therefore closes on the judgement that the proper response to the deep uncertainty of tipping-point science is neither alarmism nor complacency but a rigorous accounting of the full range of plausible outcomes, weighted by consequence as well as by likelihood, and an explicit recognition that the accumulation of commitment through sustained warming is itself an ongoing and only partly reversible act (Kemp, 2022) (Richardson, 2023) (Fankhauser, 2021).

8.11 References

Boulton, C. A., et al. (2022). Pronounced loss of Amazon rainforest resilience since the early 2000s. *Nature Climate Change*. <https://doi.org/10.1038/s41558-022-01287-8>

Cowie, R. H., et al. (2022). The Sixth Mass Extinction: fact, fiction or speculation?. *Biological reviews/Biological reviews of the Cambridge Philosophical Society*. <https://doi.org/10.1111/brv.12816>

Dolan, F., et al. (2021). Evaluating the economic impact of water scarcity in a changing world. *Nature Communications*. <https://doi.org/10.1038/s41467-021-22194-0>

Fankhauser, S., et al. (2021). The meaning of net zero and how to get it right. *Nature Climate Change*. <https://doi.org/10.1038/s41558-021-01245-w>

Flores, B. M., et al. (2024). Critical transitions in the Amazon forest system. *Nature*. <https://doi.org/10.1038/s41586-023-06970-0>

Folke, C., et al. (2021). Our future in the Anthropocene biosphere. *AMBIO*. <https://doi.org/10.1007/s13280-021-01544-8>

Forzieri, G., et al. (2022). Emerging signals of declining forest resilience under climate change. *Nature*. <https://doi.org/10.1038/s41586-022-04959-9>

Frenne, P. D., et al. (2021). Forest microclimates and climate change: Importance, drivers and future research agenda. *Global Change Biology*. <https://doi.org/10.1111/gcb.15111>

Galaz, V., et al. (2021). Artificial intelligence, systemic risks, and sustainability. *Technology in Society*. <https://doi.org/10.1016/j.techsoc.2021.101741>

Hammond, W. M., et al. (2022). Global field observations of tree die-off reveal hotter-drought fingerprint for Earth's forests. *Nature Communications*. <https://doi.org/10.1038/s41467-022-29289-2>

Harvey, J. A., et al. (2022). Scientists' warning on climate change and insects. *Ecological Monographs*. <https://doi.org/10.1002/ecm.1553>

Kaack, L. H., et al. (2022). Tackling Climate Change with Machine Learning. OPUS 4 (Zuse Institute Berlin). <https://doi.org/10.1145/3485128>

Kemp, L., et al. (2022). Climate Endgame: Exploring catastrophic climate change scenarios. *Proceedings of the National Academy of Sciences*. <https://doi.org/10.1073/pnas>

Lawrence, M., et al. (2024). Global polycrisis: the causal mechanisms of crisis entanglement. *Global Sustainability*. <https://doi.org/10.1017/sus.2024.1>

McDowell, N. G., et al. (2022). Mechanisms of woody-plant mortality under rising drought, CO₂ and vapour pressure deficit. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-022-00272-1>

Miner, K., et al. (2022). Permafrost carbon emissions in a changing Arctic. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43017-021-00230-3>

Mishra, R. K. (2023). Fresh Water availability and Its Global challenge. *British Journal of Multidisciplinary and Advanced Studies*. <https://doi.org/10.37745/bj-mas.2022.0208>

Naidu, R., et al. (2021). Chemical pollution: A growing peril and potential catastrophic risk to humanity. *Environment International*. <https://doi.org/10.1016/j.envint.2021.106616>

Pörtner, H., et al. (2023). Overcoming the coupled climate and biodiversity crises and their societal impacts. *Science*. <https://doi.org/10.1126/science.abl4881>

Richardson, K., et al. (2023). Earth beyond six of nine planetary boundaries. *Science Advances*. <https://doi.org/10.1126/sciadv.adh2458>

Romanello, M., et al. (2023). The 2023 report of the Lancet Countdown on health and climate change: the imperative for a health-centred response in a world facing irreversible harms. *The Lancet*. [https://doi.org/10.1016/s0140-6736\(23\)01859-7](https://doi.org/10.1016/s0140-6736(23)01859-7)

Romanello, M., et al. (2024). The 2024 report of the Lancet Countdown on health and climate change: facing record-breaking threats from delayed action. *The Lancet*. [https://doi.org/10.1016/s0140-6736\(24\)01822-1](https://doi.org/10.1016/s0140-6736(24)01822-1)

Scotese, C. R. (2021). An Atlas of Phanerozoic Paleogeographic Maps: The Seas Come In and the Seas Go Out. *Annual Review of Earth and Planetary Sciences*. <https://doi.org/10.1146/annurev-earth-081320-064052>

Simpson, N. P., et al. (2021). A framework for complex climate change risk assessment. *One Earth*. <https://doi.org/10.1016/j.oneear.2021.03.005>

Smith, K. E., et al. (2022). Biological Impacts of Marine Heatwaves. *Annual Review of Marine Science*. <https://doi.org/10.1146/annurev-marine-032122-121437>

Yuan, X., et al. (2023). A global transition to flash droughts under climate change. *Science*. <https://doi.org/10.1126/science.abn6301>

Zittis, G., et al. (2022). Climate Change and Weather Extremes in the Eastern Mediterranean and Middle East. *Reviews of Geophysics*. <https://doi.org/10.1029/2021rg00076>

Chapter 9

Chapter 9: Stabilization, Reversibility, and Climate Intervention

9.1 Executive Summary

Stabilizing global surface temperature at any level requires global net anthropogenic CO₂ emissions to reach and be sustained at zero, and every additional increment of cumulative CO₂ emitted moves the achievable stabilization level upward in near-linear proportion (*virtually certain*). The near-linear relationship between cumulative CO₂ emissions and peak warming remains the central organizing principle of stabilization science, and it is on this relationship that the concept of a remaining carbon budget for a given temperature limit rests.

The remaining carbon budget consistent with limiting warming to 1.5 degrees C has continued to contract as emissions have persisted near record levels, and on present trajectories the budget for a better-than-even chance of holding warming to 1.5 degrees C will be exhausted within the current decade (*high confidence*). Continued growth in the global carbon budget's emission terms, documented in successive annual assessments (Friedlingstein et al., 2022; Friedlingstein et al., 2023; Friedlingstein et al., 2025), leaves diminishing room for delay if low stabilization levels are to remain physically attainable.

When global net CO₂ emissions reach zero, the multi-decadal trend in global surface temperature is expected to level off rather than continue rising, because the warming effect of declining atmospheric CO₂ removal by the ocean and land is approximately balanced by continued ocean heat uptake (*medium confidence*). This “zero-emissions commitment” close to zero is a physical-science conclusion with profound implications: it means that halting warming is achievable within a human generation, but that already-emitted CO₂ commits the Earth system to sustained elevated temperatures for centuries.

Stabilization is achievable along pathways that peak and decline temperature — overshoot pathways — but overshoot carries physical risks that are not symmetric on the way up and the way down, and some Earth system responses are only partially reversible on policy-relevant timescales (*high confidence*). Committed changes in the deep ocean, ice sheets, permafrost, and biosphere carbon stores mean that returning global mean temperature to a target level

does not return the full Earth system to its prior state.

Carbon dioxide removal is a necessary element of pathways that reach net-zero and net-negative CO₂ emissions, and both land-based and engineered removal approaches have distinct physical-science signatures, permanence characteristics, and Earth system side effects that govern their net climatic effect (*high confidence*). Nature-based approaches and engineered approaches such as direct air capture and geologic storage differ sharply in monitorability, permanence, and reversibility, as assessed for nature-based solutions by Seddon et al. (2021), for direct air capture by Erans et al. (2022), and for the wider capture-and-storage chain by Dziejarski et al. (2023).

Solar radiation modification could, in principle, rapidly reduce global mean surface temperature by increasing planetary albedo, but it does not address ocean acidification, would impose distinct regional climate and hydrological responses, and would create a termination-related risk of rapid warming if sustained deployment ceased (*high confidence*). Solar radiation modification is not a substitute for emissions reductions or carbon dioxide removal; its physical effects act on a different lever of the energy balance than the atmospheric CO₂ burden that drives long-term change.

Earth system responses to stabilization are heterogeneous in time: global mean surface temperature responds within decades, but sea level, ice sheets, deep-ocean heat content, and some ecosystem carbon stores continue to evolve for centuries to millennia after temperatures stabilize (*high confidence*). Assessing stabilization therefore requires distinguishing the fast-responding components of the system from the slow-responding components whose trajectories are effectively locked in over human planning horizons.

9.2 Introduction: what stabilization means in the physical climate system

Climate stabilization is a property of the whole Earth system, not of the atmosphere alone, and its assessment requires tracking the coupled response of the carbon cycle, the energy budget, and the slow cryospheric and oceanic components together (*high confidence*). This chapter assesses the physical-science basis for stabilizing the climate: the relationships that connect cumulative emissions to warming, the behavior of the system once emissions cease, the reversibility of changes under declining forcing, and the physical characteristics of the two principal classes of deliberate climate intervention — carbon dioxide removal and solar radiation modification. It does not assess the economic feasibility, governance, or social dimensions of these options, which fall to Working Groups II and III; the focus here is strictly on what the physical climate system does in response to the forcing pathways that different stabilization strategies imply.

The framing of stabilization has shifted over successive assessment cycles from a concentration-based view — in which a target atmospheric CO₂ concentration was the organizing quantity — to an emissions-based view centered

on cumulative anthropogenic CO₂. This shift reflects one of the most robust findings in climate science: that peak warming is controlled, to first order, by the total amount of CO₂ that humanity emits, largely independent of the precise pathway by which that total is reached. The practical consequence is that stabilizing temperature is equivalent to bringing net CO₂ emissions to zero, and the level at which temperature stabilizes is set by the cumulative total emitted up to that point.

This emissions-centric framing situates stabilization within the broader diagnosis of the Earth system's state. The planetary boundaries synthesis of Richardson et al. (2023) concludes that six of nine boundaries that define a safe operating space for humanity have now been transgressed, with the climate-change boundary among them; this places the stabilization question in the context of a system already displaced substantially from its Holocene reference state. The emissions driving that displacement are documented sector by sector: Lamb et al. (2021) assess trends and drivers of greenhouse gas emissions across sectors from 1990 to 2018, showing sustained growth across energy supply, industry, transport, and buildings, with the pace and distribution of that growth setting the baseline against which stabilization pathways must bend.

Stabilization at a chosen level is, in physical terms, the successful termination of the anthropogenic perturbation to the carbon cycle followed by the slow relaxation of the system toward a new equilibrium. Because that relaxation is incomplete on human timescales for several key components, the language of this chapter distinguishes carefully among three ideas that are often conflated: **halting** warming (achieved when net CO₂ emissions reach zero and temperature levels off), **reversing** warming (achieved only through sustained net-negative CO₂ emissions that draw down the atmospheric burden), and **reversibility** of impacts (the degree to which specific Earth system components return toward their prior state when forcing declines). These three are governed by different physics and operate on different timescales.

9.3 Remaining carbon budgets

For a given probability of limiting peak warming to a specified level, there exists a finite remaining budget of cumulative net CO₂ emissions, and this budget is being drawn down at close to its full historical rate because global CO₂ emissions have not yet entered sustained decline (*high confidence*). The remaining carbon budget is the total quantity of additional net anthropogenic CO₂ that can be emitted from a specified start date while keeping peak warming below a chosen limit with a chosen probability. Its magnitude follows from the near-linear proportionality between cumulative CO₂ emissions and global mean surface temperature, combined with estimates of the contribution of non-CO₂ forcings, the historical warming already realized, and the Earth system feedbacks not fully captured in the core proportionality.

The annual Global Carbon Budget assessments provide the observational anchor for the emission side of this calculation. Friedlingstein et al. (2022),

Friedlingstein et al. (2023), and Friedlingstein et al. (2025) document total anthropogenic CO₂ emissions from fossil sources and land-use change, together with the partitioning of those emissions among the atmosphere, ocean, and land sinks. Across these successive assessments the central message is consistent: total CO₂ emissions have remained near record levels, the atmospheric CO₂ growth rate has continued, and the ocean and land sinks, while continuing to absorb a substantial fraction of emissions, have not accelerated enough to offset emission growth. Each year in which emissions remain near their peak subtracts a corresponding quantity from every remaining budget, moving the exhaustion date for the tightest budgets closer.

The distinction between the CO₂ and non-CO₂ contributions to warming is central to budget estimation. The remaining budget for a temperature limit depends not only on cumulative CO₂ but on the co-evolving forcing from methane, nitrous oxide, aerosols, and other agents. Because aerosols exert a net cooling effect that is expected to decline as air-quality measures and fossil-fuel phase-out reduce emissions, the removal of this cooling can unmask additional CO₂-driven warming, tightening the effective budget. The indicators synthesis of Forster et al. (2023) provides an annually updated assessment of the large-scale indicators of the state of the climate system and of human influence, including the components of effective radiative forcing and the estimate of human-induced warming to date; these indicators are the quantities from which the remaining budget for any given limit is computed, and their continued updating is what drives the year-on-year revision of budget estimates.

A central physical subtlety in budget science is the treatment of Earth system feedbacks that are not represented, or only partially represented, in the models used to derive the core cumulative-emissions-to-warming relationship. Permafrost carbon release, changes in wetland methane emissions, and shifts in the airborne fraction under high cumulative emissions all act to reduce the remaining budget relative to a calculation that ignores them. These feedbacks introduce an asymmetry: they are more likely to reduce than to increase the budget, and their magnitude grows with the level of warming, so budgets for higher limits carry proportionally larger downward feedback risks. The declining resilience of major carbon-storing biomes compounds this concern; Forzieri et al. (2022) identify emerging signals of declining forest resilience under climate change, a finding directly relevant to whether the land sink can be relied upon to hold its historical share of emissions over the coming decades, and Lapola et al. (2023) document the drivers and impacts of Amazon forest degradation, a process that erodes one of the largest terrestrial carbon stores and thereby the effective budget.

The observational context for budget exhaustion is stark. With cumulative emissions continuing to rise at close to their historical rate, the budget consistent with a better-than-even chance of holding peak warming to 1.5 degrees C is small relative to a single decade of emissions at present rates. This does not, in itself, mean that 1.5 degrees C is physically impossible as a peak; it means that holding to 1.5 degrees C without overshoot requires an immediate and steep decline in net CO₂ emissions, and that any pathway which does not achieve such

a decline will overshoot the level before, at best, returning to it through net-negative emissions. Riahi et al. (2021) assess the cost and attainability of meeting stringent climate targets without overshoot, and their analysis makes explicit the narrowing physical and technological space within which no-overshoot pathways remain feasible as emissions persist.

The relationship between fossil-fuel reserves and the remaining budget provides a further physical framing. Welsby et al. (2021) estimate the quantity of fossil fuels that must remain unextracted in a 1.5 degrees C world, concluding that a large majority of existing reserves cannot be combusted if the corresponding budget is to be respected. This translates the abstract budget into a constraint on the physical stock of carbon-bearing fuels, and it makes clear that the budget is not merely an accounting convention but a hard limit on how much of the geologic carbon reservoir can be transferred to the atmosphere.

Two additional considerations qualify the interpretation of any single budget number. First, the remaining budget is a distribution, not a point: it carries irreducible uncertainty from the strength of the cumulative-emissions-to-warming relationship, the non-CO₂ forcing trajectory, the historical warming estimate, and the feedback terms. Reporting a budget therefore requires stating both the temperature limit and the probability with which it is to be respected, and moving from a 50 percent to a 67 percent probability of holding a limit substantially reduces the allowable emissions. Second, the budget is defined for peak warming, not for warming at a fixed date; a pathway may respect a budget for peak warming while temporarily exceeding the limit if it later draws the atmosphere down, or may hold below the limit at a given date while being committed to exceed it later. Clarity about which quantity — peak warming or warming at a date — a budget refers to is essential to avoid misinterpretation.

The policy-relevant synthesis is that remaining carbon budgets for the lowest stabilization levels are now small enough that their exhaustion is measured in years rather than decades at current emission rates, that the budgets are subject to downward revision from feedbacks and aerosol unmasking, and that the observational record shows no sustained decline in global CO₂ emissions of the magnitude required to preserve the tightest budgets. The physical science does not prescribe a policy response, but it establishes unambiguously that the room for delay in reaching net-zero CO₂ has largely been consumed for the most stringent stabilization levels.

9.4 The zero-emissions commitment and the near-linear response

The multi-decadal trend in global surface temperature is expected to halt within about a decade of global net CO₂ emissions reaching zero, because the continued decline in atmospheric CO₂ concentration through ocean and land uptake offsets the warming that would otherwise result from continued ocean heat uptake (*medium confidence*). This conclusion — that the zero-emissions commitment for CO₂ is close to zero — is one of the most consequential results

of stabilization science, because it means that the level at which temperature stabilizes is determined by cumulative emissions up to the point of net zero, and that no substantial additional warming is “in the pipeline” from past CO₂ emissions once those emissions cease. It stands in contrast to an older intuition, derived from constant-concentration experiments, that large committed warming would continue to be realized for decades after emissions stopped; that intuition confused the constant-concentration commitment with the constant-emissions (zero further emissions) commitment, which are governed by different balances.

The physical basis for a near-zero commitment lies in the compensation between two opposing tendencies once emissions cease. First, in the absence of new emissions, the atmospheric CO₂ concentration begins to fall as the ocean and land sinks continue to take up carbon; this declining radiative forcing acts to cool the surface. Second, the ocean continues to absorb heat as the deep ocean slowly equilibrates with the elevated surface temperature; this continued heat uptake, and the reduction in the rate at which the surface warms, means the surface does not immediately reflect the full equilibrium response to the peak concentration. To first order these two effects — declining forcing and continued ocean heat uptake — approximately cancel, leaving global mean surface temperature roughly flat for decades after net-zero CO₂. The precise magnitude and even the sign of the residual depend on the relative rates of carbon drawdown and heat uptake, which is why the commitment is assessed with medium rather than high confidence and why models disagree on whether the residual trend is slightly positive or slightly negative.

The near-linear relationship between cumulative CO₂ emissions and peak warming — often expressed as the transient climate response to cumulative emissions — is the quantitative expression of this behavior. It arises because the diminishing radiative efficacy of additional CO₂ at higher concentrations is approximately compensated by the diminishing capacity of the ocean and land to take up additional carbon and heat as they warm. The result is a proportionality that holds robustly across models and across a wide range of cumulative emissions, and it is this proportionality that underwrites both the remaining carbon budget concept and the conclusion that halting emissions halts warming. The scenario-based projections assessed in the model intercomparison framework of Tebaldi et al. (2021), which synthesizes climate model projections from the Scenario Model Intercomparison Project of the sixth Coupled Model Intercomparison Project, exhibit this near-linear behavior across the scenario range and provide the multi-model basis for characterizing how surface temperature responds as cumulative emissions accumulate and then cease.

Several caveats qualify the near-zero commitment. The result applies to CO₂ alone; the total zero-emissions commitment across all forcing agents differs, because non-CO₂ forcers have different lifetimes and different committed responses. Short-lived agents such as methane decay rapidly, so ceasing methane emissions produces a relatively fast decline in methane forcing and hence some cooling; aerosols, also short-lived, decay quickly so that ceasing aerosol emissions removes their cooling almost immediately, producing warming. The net non-CO₂ commitment therefore depends on the assumed trajectory of these

agents, and a realistic assessment of the total commitment must combine the near-zero CO₂ term with the non-CO₂ terms. The distinct atmospheric behavior and mitigation role of the short-lived, predominantly biogenic methane emissions from agriculture is analyzed by Lynch et al. (2021), who show that agriculture's contribution to climate change and its role in mitigation is distinct from that of the predominantly fossil-CO₂-emitting sectors precisely because of these lifetime differences; this distinction matters for the composition of the total commitment.

A further caveat concerns Earth system feedbacks not represented in the models used to derive the commitment. Permafrost carbon release, in particular, could continue to add CO₂ and methane to the atmosphere after anthropogenic emissions cease, imparting a positive contribution to the commitment that is not captured in models lacking a fully interactive permafrost carbon module. Similarly, committed changes in the biosphere — including the declining forest resilience documented by Forzieri et al. (2022) and the degradation processes analyzed by Lapola et al. (2023) — could shift the land from a net sink toward a weaker sink or a source, which would slow the post-emissions decline in atmospheric CO₂ and thereby push the commitment toward positive values. These feedbacks are the principal reason the zero-emissions commitment is assessed with medium confidence and is regarded as more likely to be revised upward than downward.

The distinction between halting and reversing warming follows directly. Reaching net-zero CO₂ halts the warming trend but does not, by itself, cool the planet: temperature stabilizes near its peak. To reduce global mean temperature below its peak requires sustained net-negative CO₂ emissions — removing more CO₂ from the atmosphere than is emitted — which draws down the atmospheric burden and reduces radiative forcing. This is a physically distinct undertaking from reaching net zero, and it is the point at which carbon dioxide removal shifts from being a means of offsetting residual emissions to being a means of actively reversing warming. The rate at which temperature declines under net-negative emissions is not simply the mirror image of the rate at which it rose, because the slow components of the system — deep-ocean heat content, ice sheets — respond with hysteresis, a point developed in the following section.

9.5 Overshoot and reversibility

Pathways that temporarily exceed a temperature limit before returning to it — overshoot pathways — do not return the Earth system to the state it would have occupied had the limit never been exceeded, because several key components respond with hysteresis and because some changes accumulate irreversibly during the overshoot period (*high confidence*). Overshoot is a defining feature of many stabilization pathways for the lowest limits, given the small remaining budgets assessed in Section 9.2. Understanding overshoot requires separating three questions: whether the global mean temperature can be brought back down (it can, in principle, through sustained net-negative

emissions); whether the return path retraces the outbound path (it does not, for the slow components); and whether specific impacts incurred during overshoot are reversed when temperature returns (some are, many are not, and some continue to worsen even as temperature declines).

The reversibility of global mean surface temperature itself is comparatively high on the timescale of the net-negative emissions that would drive it. Because surface temperature is coupled closely to atmospheric CO₂ and to the fast components of the climate system, drawing down atmospheric CO₂ through net-negative emissions produces a corresponding decline in global mean temperature within years to a few decades. In this narrow sense, global mean temperature is a relatively reversible quantity. But this reversibility is contingent on achieving and sustaining net-negative emissions at scale, and it does not extend to the slow components of the system, whose response to the overshoot is only partially undone by the subsequent cooling.

The deep ocean is the archetypal slow, partially irreversible component. Heat that enters the ocean during the overshoot penetrates to depth over decades to centuries, and the deep ocean continues to warm even as the surface cools under declining forcing, because the deep ocean is still equilibrating with the elevated temperatures of the overshoot period. As a result, ocean heat content and thermosteric sea level continue to rise for a period after surface temperature peaks and begins to decline, and they do not return to their pre-overshoot values on policy-relevant timescales. This hysteresis in ocean heat content is a central reason that sea-level rise is among the least reversible consequences of an overshoot: the thermal expansion component is committed by the heat already sequestered, and the ice-sheet component responds to the integrated temperature history rather than to the instantaneous temperature.

Marine ecosystems illustrate the asymmetry between temperature reversibility and impact reversibility. Even a temporary excursion to higher temperatures can produce marine heatwaves whose biological consequences persist long after the temperature anomaly has passed. Smith et al. (2022) assess the biological impacts of marine heatwaves and document mortality, range shifts, and community reorganization that do not simply reverse when temperatures return to prior levels; a coral reef killed during an overshoot does not regrow merely because the water cools again on a decadal timescale. The overshoot therefore incurs biological costs that are, in ecological terms, effectively permanent even where the physical temperature is reversible.

The terrestrial biosphere exhibits comparable irreversibilities. Forest dieback triggered during an overshoot releases stored carbon and reduces the resilience of the remaining forest, and the recovery of a degraded or burned forest to its prior carbon density takes decades to centuries, if it occurs at all. Forzieri et al. (2022) document emerging signals of declining forest resilience under climate change, which implies that forests subjected to the stress of an overshoot may be less able to recover afterward; Lapola et al. (2023) document the drivers and impacts of Amazon forest degradation, a process that, once advanced, is not readily reversed by a subsequent decline in temperature. Because these carbon stores feed back on atmospheric CO₂, their partial irreversibility means that

an overshoot can leave the carbon cycle in a state that makes the subsequent drawdown harder — the land sink is weaker on the return path than it was on the outbound path.

The broader biodiversity consequences of overshoot are similarly asymmetric. Cowie et al. (2022) assess the evidence bearing on whether the current episode of biodiversity loss constitutes a sixth mass extinction, and the extinctions that occur during a period of elevated temperature and associated stress are, by definition, irreversible: a species lost during an overshoot is not recovered when temperature returns. Harvey et al. (2022) issue a scientists' warning on climate change and insects, documenting the sensitivity of insect populations — foundational to terrestrial food webs and pollination — to climatic stress, and Lees et al. (2022) assess the state of the world's birds, a taxon whose population declines integrate multiple climatic and land-use pressures. These lines of evidence establish that the ecological ledger of an overshoot is not balanced by the subsequent cooling; the losses incurred at the peak are carried forward.

A crucial physical point is that the reversibility of the carbon cycle response to net-negative emissions is itself imperfect. When CO₂ is removed from the atmosphere, the ocean and land, which had been net sinks, begin to release some of their stored carbon back to the atmosphere in response to the declining atmospheric concentration — an outgassing that partially offsets the removal. This means that removing a given quantity of CO₂ from the atmosphere does not reduce the atmospheric burden by that full quantity, because the natural reservoirs give back part of what they had absorbed. The efficiency of net-negative emissions in reducing atmospheric CO₂ is therefore lower than the efficiency with which emissions raised it, introducing an inherent asymmetry between emitting and removing. This asymmetry is a physical property of the coupled carbon cycle and is independent of the technology used to achieve the removal; it means that reversing a given amount of warming requires removing more CO₂ than was emitted to cause it.

The magnitude and duration of an overshoot govern the severity of these committed and partially irreversible responses. A shallow, brief overshoot commits less deep-ocean heat, triggers fewer irreversible ecological transitions, and leaves the carbon cycle closer to its pre-overshoot state than a deep, prolonged overshoot. This establishes a physical rationale for minimizing both the peak level and the duration of any overshoot, independent of the eventual stabilization level: the integrated exposure of the slow components to elevated temperature, not merely the endpoint, determines how much of the change is locked in. The assessment that some Earth system components respond to the time-integral of temperature rather than to its instantaneous value is central to why overshoot is treated as a distinct and consequential feature of stabilization pathways rather than as a mere transient that is fully undone on return.

The catastrophic tail of the overshoot risk distribution warrants explicit assessment. Kemp et al. (2022) argue for systematic exploration of catastrophic climate change scenarios — the “climate endgame” — on the grounds that the high-impact, deep-uncertainty tail of the warming distribution has been comparatively neglected relative to central estimates. In the context of overshoot, the

concern is that a deep or prolonged overshoot raises the probability of encountering threshold behaviors and self-amplifying feedbacks whose reversibility is unknown and potentially very low. The physical science cannot at present rule out that a sufficiently large overshoot could trigger responses that persist even under aggressive subsequent drawdown, and the precautionary implication is that reliance on overshoot-and-recovery pathways carries a tail risk that is qualitatively different from the risk of a pathway that avoids overshoot altogether.

9.6 Carbon dioxide removal: physical-science aspects

Carbon dioxide removal — the deliberate removal of CO₂ from the atmosphere and its durable storage in geological, terrestrial, or ocean reservoirs — is a physically necessary element of any pathway that reaches net-zero CO₂ emissions in the presence of residual hard-to-abate emissions, and it is the only means by which warming can be actively reversed after net zero (*high confidence*). The physical assessment of carbon dioxide removal turns on three quantities that are distinct from the economics or governance of the approaches: the gross quantity of CO₂ removed from the atmosphere, the durability or permanence of the storage, and the Earth system side effects — including the carbon-cycle feedback that returns part of the removed carbon to the atmosphere and the local biogeochemical and biophysical consequences of the removal itself. A removal method's net climatic benefit is the gross removal discounted by the carbon-cycle giveback and by any co-emissions or albedo changes it induces, integrated over the storage lifetime.

The concept of net zero itself must be defined precisely for carbon dioxide removal to be assessed coherently. Fankhauser et al. (2021) examine the meaning of net zero and the conditions under which a net-zero claim is physically robust, emphasizing that the climatic effect of a net-zero portfolio depends on the permanence of the removals used to balance residual emissions and on whether like is balanced with like — durable geological storage balancing fossil emissions, rather than temporary biological storage balancing permanent fossil releases. This distinction is central: balancing a permanent addition of geologic carbon to the atmosphere with a temporary removal into a biological store that may later be released does not achieve durable net zero, because the two flows have different permanence. The physical equivalence of a removal to an avoided emission holds only to the extent that the storage endures on the same timescale as the atmospheric perturbation it offsets.

Land-based and ecosystem approaches constitute one principal class. Afforestation, reforestation, improved forest and soil management, and other nature-based approaches remove CO₂ through photosynthesis and store it in biomass and soils. Seddon et al. (2021) assess the evidence on nature-based solutions to climate change and emphasize the conditions under which they deliver genuine, durable climate benefit, cautioning against overstating their mitigation potential and against treating biological carbon storage as equivalent

to permanent geological storage. The physical limitations of these approaches are intrinsic: biological carbon stores are vulnerable to reversal through fire, drought, pest outbreaks, and land-use change, and their vulnerability increases with warming itself. The declining forest resilience documented by Forzieri et al. (2022) and the forest degradation processes analyzed by Lapola et al. (2023) directly undermine the permanence of land-based removal, because a warming climate erodes the very stores on which the approach depends. Land-based removal also has biophysical side effects — changes in surface albedo, evapotranspiration, and regional hydrology — that can partially offset or, in some regions such as high-latitude afforestation, even reverse the intended cooling, so that the net effect on temperature is not captured by the carbon accounting alone.

Ocean-based approaches form a second class, encompassing the enhancement of biological and chemical carbon uptake by the ocean. These approaches exploit the ocean's role as the largest active carbon reservoir, but they interact with the marine biosphere in ways that carry their own risks; the sensitivity of marine ecosystems to perturbation is documented by Smith et al. (2022) in their assessment of marine heatwave impacts, and the broader question of ocean stewardship for biodiversity, food, and climate is framed by Sala et al. (2021), whose analysis of protecting the global ocean underscores that interventions in ocean biogeochemistry occur within a system that also supports fisheries and biodiversity. The food-provisioning role of the ocean assessed by Boyd et al. (2022), who quantify the contribution of fisheries and aquaculture to the global protein supply, is a reminder that ocean-based carbon removal cannot be assessed on carbon grounds alone; its physical side effects on ocean chemistry and ecosystems bear directly on other services the ocean provides. The permanence of ocean removal depends on whether the carbon is transferred to the deep ocean or to sediments on timescales long compared to surface-ocean exchange.

Engineered approaches constitute a third class, of which direct air capture with geological storage is the archetype. Erans et al. (2022) assess the process technology and techno-economic and socio-political challenges of direct air capture, describing the physical operation of the capture step — the separation of CO₂ from ambient air at its low atmospheric concentration — and the energy requirements that this dilute-stream separation entails. The physical advantage of engineered removal with geological storage is permanence and monitorability: CO₂ stored in suitable geological formations is durable on timescales of centuries to millennia and its storage can be measured and verified, in contrast to the diffuse and reversible storage of biological approaches. The physical constraints are the energy demand of capturing a dilute gas and the availability and integrity of geological storage. The broader status of carbon capture, utilization, and storage technologies is surveyed by Dziejarski et al. (2023), whose technical assessment of the current status of these technologies in the global economy characterizes the maturity of the capture and storage chain on which engineered removal depends. Yan et al. (2021) review the application of machine learning to carbon capture, utilization, and storage, illustrating that optimization of the

capture process and the characterization of storage reservoirs are active areas of methodological development that bear on the achievable efficiency and cost of the physical removal.

A cross-cutting physical result governs all removal approaches: the carbon-cycle response to removal is not the mirror image of the response to emission. When CO₂ is withdrawn from the atmosphere, the ocean and land — previously net sinks — reduce their uptake and begin to outgas carbon in response to the lowered atmospheric concentration, returning a fraction of the removed carbon to the atmosphere. This giveback means that the reduction in atmospheric CO₂ achieved by a given gross removal is smaller than the removal itself, and the efficiency of removal in lowering the atmospheric burden is lower than the efficiency with which emissions raised it. The practical implication, developed in Section 9.4, is that reversing a given amount of warming requires removing more CO₂ than was emitted to cause it. This asymmetry is a robust feature of the coupled carbon cycle and is independent of the removal technology.

The scale and durability required for carbon dioxide removal to contribute meaningfully to stabilization are large. To reach net-zero CO₂, removal must offset the residual emissions from sectors that cannot be fully decarbonized; to reverse warming, removal must exceed emissions on a sustained basis. The material and land footprints of these approaches, and their interactions with other Earth system boundaries, mean that carbon dioxide removal cannot be assessed in isolation from the broader planetary context diagnosed by Richardson et al. (2023). Removal deployed at climatically significant scale interacts with the water, land, and biosphere boundaries, and an assessment confined to the carbon boundary alone would misrepresent the net Earth system consequences. The physical-science conclusion is that carbon dioxide removal is necessary for stabilization at low levels and indispensable for any reversal of warming, but that its net climatic benefit is materially less than its gross removal once carbon-cycle giveback, permanence risk, and biophysical side effects are accounted for, and that these discounts differ substantially among approaches.

The permanence spectrum and enhanced weathering

Carbon dioxide removal approaches occupy a spectrum of storage permanence, from the decades-to-centuries vulnerability of biological stores to the millennial durability of carbon mineralized into stable carbonate minerals or stored in deep geological formations, and this permanence spectrum — rather than the gross quantity removed — is the physical property that determines whether a removal genuinely offsets a fossil-carbon emission on the timescale of the atmospheric perturbation (*high confidence*). The permanence of storage is the single most important physical discriminator among removal approaches, because a removal offsets an emission only for as long as the removed carbon remains out of the atmosphere. Biological stores in vegetation and soils sit at the low-permanence end of the spectrum: they are vulnerable to reversal through fire, drought, pest outbreaks, harvest, and land-use change, and their vulnerability is amplified by the warming they are intended to counter, as the

declining forest resilience documented by Forzieri et al. (2022) and the forest degradation documented by Lapola et al. (2023) make clear. Geological storage of captured CO₂ sits at the high-permanence end, with retention on timescales of centuries to millennia when the storage formation and its seal are geologically suitable; the technology status of the capture-and-storage chain that delivers this permanence is surveyed by Dziejarski et al. (2023) and its physical operation at the capture step by Erans et al. (2022).

Mineralization approaches, in which CO₂ is reacted with alkaline minerals to form stable solid carbonates, occupy the most durable end of the permanence spectrum, because the carbon is converted into a thermodynamically stable mineral form that is not vulnerable to the biological and physical reversals that threaten organic stores. Enhanced weathering — the acceleration of the natural chemical weathering of silicate minerals that consumes atmospheric CO₂ over geological time — seeks to draw down CO₂ by exposing finely ground reactive minerals to the atmosphere and to soil and ocean water, converting CO₂ into dissolved bicarbonate and ultimately into stable carbonate. The physical attraction of these approaches is the durability of the resulting store; their physical constraints are the energy required to grind and transport large quantities of mineral, the rate at which the weathering reaction proceeds, and the local biogeochemical consequences of introducing large quantities of reactive mineral into soils and coastal waters. The optimization of these and other capture and storage processes is an active methodological area, and the application of machine learning to characterize reaction rates, capture materials, and storage reservoirs is reviewed by Yan et al. (2021); the adsorbent and separation materials on which several engineered capture routes depend are the subject of intensive materials research of the kind that governs the achievable energy efficiency of the capture step.

The permanence spectrum interacts with the carbon-cycle giveback assessed in Section 9.5 to determine the net, durable atmospheric drawdown achieved by a removal portfolio. A removal into a low-permanence biological store that is later reversed returns its carbon to the atmosphere and, over the reversal, provides no lasting offset; a removal into a high-permanence mineral or geological store provides a durable offset discounted only by the carbon-cycle giveback. The robust net-zero framing of Fankhauser et al. (2021) — that durable removals must be matched to durable emissions — is the direct expression of this principle: balancing the permanent addition of geologic carbon to the atmosphere with a temporary biological removal does not achieve durable net zero. The physical-science assessment is therefore that removal approaches cannot be aggregated by gross tonnage alone; the permanence of each tonne, and its matching to the permanence of the emission it is intended to offset, is the property that determines the contribution of removal to genuine, durable stabilization. This places a premium on the high-permanence end of the spectrum — geological and mineral storage — for the offsetting of fossil emissions, while assigning the lower-permanence biological approaches a role that must be understood as inherently reversible and vulnerable to the very climate change they are deployed to limit.

9.7 Solar radiation modification: physical-science aspects

Solar radiation modification could rapidly lower global mean surface temperature by increasing the fraction of incoming solar radiation reflected to space, but it addresses only the temperature symptom of elevated CO₂ and not the underlying atmospheric burden, leaves ocean acidification unchecked, produces a pattern of regional climate and hydrological change distinct from that of the warming it offsets, and creates a termination-related risk of rapid warming should sustained deployment cease (*high confidence*). Solar radiation modification is a fundamentally different intervention from carbon dioxide removal: it acts on the shortwave side of the energy balance, reducing absorbed solar radiation, rather than on the longwave side by reducing the CO₂ that traps outgoing radiation. Because it does not remove CO₂, it does not reduce ocean acidification or the other non-thermal consequences of elevated atmospheric CO₂, and it must be sustained for as long as the elevated CO₂ persists if the masked warming is not to be realized.

The principal proposed methods increase planetary albedo by different physical means. Stratospheric aerosol injection would introduce reflective particles into the stratosphere to scatter sunlight, mimicking the transient cooling observed after large explosive volcanic eruptions; marine cloud brightening would increase the reflectivity of low marine clouds. The physical efficacy of these methods rests on well-understood radiative principles — increasing albedo reduces absorbed solar energy and hence surface temperature — and the volcanic analogue provides observational support that stratospheric aerosols can cool the surface measurably and rapidly. This rapidity is the defining physical characteristic that distinguishes solar radiation modification from all emissions-based approaches: whereas the temperature response to emissions reductions and to carbon dioxide removal unfolds over decades, the response to a change in planetary albedo is realized within a few years.

That rapidity is inseparable from a corresponding risk. Because solar radiation modification masks rather than removes the warming, the elevated CO₂ continues to exert its radiative forcing; if deployment were sustained for a period and then abruptly stopped, the masked warming would be realized over a few years — the termination effect — producing a rate of warming far exceeding anything in the emissions-driven trajectory. The magnitude of this committed rebound grows with the amount of warming being masked and with the duration of deployment, because the accumulated gap between the actual CO₂ forcing and the albedo-suppressed temperature widens over time. The physical assessment is therefore that solar radiation modification, once begun at a scale sufficient to mask substantial warming, would create a dependency: ceasing it would expose the system to a rate of temperature change that could exceed the adaptive capacity of ecosystems whose sensitivity to rapid change is documented by Smith et al. (2022) for marine systems, by Harvey et al. (2022) for insects, and by Lees et al. (2022) for birds.

The regional and hydrological consequences of solar radiation modification are a second central physical concern. Reducing incoming solar radiation does not simply reverse greenhouse warming region by region, because the two forcings have different spatial and seasonal structures: greenhouse forcing is distributed by the CO₂ field and acts on the longwave, whereas solar reduction acts on the shortwave with its own latitudinal and diurnal pattern. The result is that a deployment tuned to restore global mean temperature would not restore the pre-warming climate everywhere; some regions would be over-cooled and others under-cooled relative to the pre-industrial baseline, and the global hydrological cycle would in general be suppressed, because reducing the shortwave flux at the surface reduces the energy available to drive evaporation. This tendency to weaken the global water cycle is a robust physical feature of solar radiation modification and means that even a deployment that perfectly restored global mean temperature would leave a residual pattern of climate change, particularly in precipitation, that differs from both the warmed and the unwarmed states.

Solar radiation modification also interacts with atmospheric chemistry and with the aerosol-cloud system in ways that carry their own uncertainties. Stratospheric aerosol injection would perturb stratospheric chemistry and could affect the ozone layer; the physical behavior of aerosols in the climate system is itself a source of uncertainty, as illustrated by Brown et al. (2021), whose analysis of biomass-burning aerosols shows that the representation of absorbing aerosols in climate models carries significant biases. This matters for solar radiation modification because the efficacy and side effects of an injected aerosol depend on its optical properties — whether it scatters or also absorbs — and the demonstrated difficulty of representing aerosol absorption accurately implies corresponding uncertainty in the modeled response to deliberate aerosol injection. Marine cloud brightening depends on the aerosol-cloud interaction, one of the least certain elements of the climate system, so the achievable albedo change and its regional distribution are correspondingly uncertain.

The most important physical framing of solar radiation modification is that it is not a substitute for emissions reductions or carbon dioxide removal, because it does not address the quantity — atmospheric CO₂ — that governs the long-term state of the climate. It could, in principle, reduce the rate or peak of warming during a period of overshoot, thereby limiting exposure of the slow and biological components to elevated temperature while emissions are brought down and removal is scaled up. But it does so at the cost of introducing the termination risk and the residual regional and hydrological changes, and it leaves ocean acidification — a direct chemical consequence of elevated CO₂ rather than of temperature — entirely unaddressed. The vulnerability of ocean ecosystems and the fisheries they support, framed by Sala et al. (2021) and Boyd et al. (2022), would therefore persist under a solar-radiation-modification deployment even if surface temperature were held down. The physical-science assessment is that solar radiation modification is best understood as a potential temporary measure to limit peak warming within a broader strategy dominated by emissions reduction and carbon dioxide removal, and never as an alternative

to reducing the atmospheric CO₂ burden, because only reducing that burden addresses the full range of physical consequences of elevated CO₂.

Observational analogues and the limits of inference

The transient cooling that follows large explosive volcanic eruptions provides the closest observational analogue for stratospheric-aerosol solar radiation modification, but the analogue is imperfect in ways that limit the confidence with which it can constrain the response to a deliberate, sustained deployment (*medium confidence*). Large eruptions inject sulfate-forming material into the stratosphere and produce a measurable, transient reduction in global mean surface temperature over the following one to two years, confirming that increasing stratospheric aerosol loading cools the surface and that the coupled system responds within the expected physical bounds. This natural experiment provides direct evidence for the sign and approximate magnitude of the radiative and temperature response, and for accompanying effects on the hydrological cycle and on stratospheric dynamics.

The analogue is limited in several respects that bear on its use for assessing deliberate deployment. A volcanic injection is a single, brief pulse, whereas a solar radiation modification deployment intended to mask substantial warming would require sustained injection over decades; the response to a sustained loading is not simply the superposition of pulse responses, because the slow components of the system would have time to adjust and because the micro-physics of a continuously replenished aerosol layer differs from that of a decaying volcanic cloud. The optical properties of the injected material govern both its cooling efficacy and its side effects, and the demonstrated difficulty of representing absorbing aerosols accurately in climate models, illustrated by Brown et al. (2021) for biomass-burning aerosols, implies corresponding uncertainty in the modeled response to a deliberately engineered aerosol whose properties would be chosen to maximize scattering and minimize absorption. The inference from the volcanic analogue to a sustained deployment therefore carries an irreducible extrapolation uncertainty.

These observational limits reinforce the assessment of Section 9.6 that solar radiation modification, while physically capable of cooling the surface, entails responses — in regional climate, in the hydrological cycle, in stratospheric chemistry, and in the termination risk — that cannot be fully constrained by the available natural analogue. The ecological systems whose sensitivity to rapid change is documented by Smith et al. (2022), Harvey et al. (2022), and Lees et al. (2022) would be exposed to a climate produced by a forcing combination — reduced shortwave in the presence of elevated CO₂ — that has no close analogue in the observational or paleoclimate record, so that the ecological consequences of a sustained deployment are correspondingly uncertain. The physical-science conclusion is that the observational basis for assessing solar radiation modification is materially weaker than that for assessing the response to CO₂ forcing, and that this asymmetry in evidence is itself a reason to treat deliberate solar

intervention as carrying deep uncertainty distinct from the better-constrained response to changes in the atmospheric CO₂ burden.

9.8 Earth system responses to stabilization

When global mean surface temperature is stabilized, the components of the Earth system do not stabilize simultaneously: fast-responding components equilibrate within decades, while sea level, ice sheets, deep-ocean heat content, ocean chemistry, and some biosphere carbon stores continue to evolve for centuries to millennia, so that a stabilized global mean temperature masks a system still in transient adjustment (*high confidence*). The assessment of stabilization must therefore distinguish between the surface temperature — which is the quantity that carbon budgets and net-zero targets are framed around and which does stabilize when net CO₂ emissions cease — and the many other quantities that continue to change even after temperature levels off. This section assesses the response of the principal slow-responding components, and the corresponding fast-responding components, under a stabilized temperature.

Sea level and ice sheets

Global mean sea level will continue to rise for centuries to millennia after global mean surface temperature is stabilized, because the thermal expansion of the deep ocean and the response of the ice sheets are governed by the integrated temperature history rather than by the instantaneous temperature, and both respond with long lags and, in the case of the ice sheets, potential hysteresis (*high confidence*). Sea-level rise is the archetypal committed response. Its thermosteric component reflects the heat already stored in the ocean, which continues to penetrate to depth and expand the water column long after surface temperature stabilizes; its ice-sheet component reflects the slow dynamic response of grounded ice to the warming already realized. Neither component ceases when temperature stabilizes; both continue for as long as the deep ocean and the ice sheets remain out of equilibrium with the stabilized surface climate, which is a matter of centuries to millennia.

The ice-sheet response carries a distinct physical risk of irreversibility through hysteresis. Once a threshold of sustained warming is crossed, the retreat of a marine-based ice sheet can become self-sustaining, so that returning temperature to below the threshold does not halt or reverse the retreat on policy-relevant timescales. This means that a temporary overshoot of a threshold could commit additional multi-metre sea-level rise even if temperature is subsequently returned to a lower stabilized level — a direct manifestation of the overshoot irreversibility assessed in Section 9.4. The physical assessment is that the sea-level commitment associated with a given stabilization level is larger the higher that level, and that the commitment is realized over such long timescales that a stabilized temperature is fully consistent with continued, substantial sea-level rise over the lifetimes of the communities and ecosystems exposed to it.

The deep ocean and ocean heat content

Ocean heat content continues to increase for centuries after surface temperature stabilizes, as heat mixes from the surface into the deep ocean, and this continued deep-ocean warming is effectively irreversible on human timescales (*high confidence*). The ocean absorbs the large majority of the excess energy trapped by the enhanced greenhouse effect, and the transfer of that heat from the surface to the deep ocean proceeds over centuries. When surface temperature stabilizes, the deep ocean remains far from equilibrium with the warmer surface, and heat continues to be transferred downward until the deep ocean warms to match. This continued deep-ocean heat uptake is one of the two terms whose balance produces the near-zero zero-emissions commitment assessed in Section 9.3; here its relevance is that the deep ocean is a reservoir of committed change that persists under a stabilized surface temperature. Because the deep ocean overturns on millennial timescales, the heat it absorbs is removed from contact with the atmosphere for centuries, and its warming is not reversed by a stabilization or even a modest reduction of surface temperature.

The continued warming of the deep ocean has consequences beyond thermal expansion. It affects ocean stratification, the solubility of gases including CO₂ and oxygen, and the circulation patterns that redistribute heat and carbon. Changes in stratification and circulation feed back on the ocean's capacity to take up heat and carbon, and hence on the trajectory of atmospheric CO₂ during the drawdown phase of a stabilization pathway. The physical assessment is that the deep ocean integrates and stores the perturbation, releasing it only slowly, so that the ocean interior remains a signature of the peak warming long after the surface has stabilized.

Ocean acidification and biogeochemistry

Ocean acidification is driven directly by the uptake of atmospheric CO₂ rather than by temperature, so it is not addressed by any intervention that lowers temperature without lowering atmospheric CO₂, and it stabilizes or reverses only as the atmospheric CO₂ burden itself is stabilized or reduced (*high confidence*). The dissolution of CO₂ in seawater lowers ocean pH and reduces the availability of carbonate ions, with consequences for calcifying organisms and marine ecosystems. Because this is a chemical consequence of the CO₂ burden, ocean acidification tracks atmospheric CO₂, and stabilizing temperature through solar radiation modification while leaving CO₂ elevated would leave acidification entirely unchecked — one of the central physical reasons that solar radiation modification is not a substitute for reducing the CO₂ burden, as assessed in Section 9.6.

The interaction of acidification with warming and deoxygenation compounds the stress on marine ecosystems. Smith et al. (2022) document the biological impacts of marine heatwaves, and these thermal stresses act in combination with acidification and with changes in oxygen availability driven by the ocean

warming assessed in Section 9.7.2. The consequences for the services the ocean provides — including the protein supply quantified by Boyd et al. (2022) and the biodiversity and climate functions framed by Sala et al. (2021) — mean that the ocean’s biogeochemical response to stabilization is not captured by temperature alone. A pathway that stabilizes temperature but stabilizes CO₂ at a high level leaves the ocean in a chemically altered state; only a pathway that reduces the atmospheric CO₂ burden begins to reverse acidification, and even then the deep-ocean chemistry adjusts slowly.

The biosphere and terrestrial carbon stores

Under a stabilized climate the terrestrial biosphere continues to adjust, and the resilience and carbon-storage capacity of major biomes may continue to decline even after temperature stabilizes, so that the land carbon sink cannot be assumed to hold its historical strength through and beyond stabilization (*medium confidence*). The land biosphere has historically absorbed a substantial fraction of anthropogenic CO₂ emissions, and the maintenance of this sink is important to the trajectory of atmospheric CO₂ during and after a stabilization pathway. But the sink’s strength is itself sensitive to the climate. Forzieri et al. (2022) identify emerging signals of declining forest resilience under climate change, indicating that forests are becoming less able to recover from disturbance; Lapola et al. (2023) document the drivers and impacts of Amazon forest degradation, a process that reduces the carbon density and the sink capacity of one of the largest terrestrial carbon reservoirs. These processes imply that the land sink assumed in stabilization pathways may weaken, and a weakening land sink means that a given emissions trajectory results in a higher atmospheric CO₂ concentration and hence a higher stabilization temperature than a calculation assuming a constant sink.

The broader integrity of the biosphere bears on stabilization through multiple channels. Folke et al. (2021) assess the human future within the Anthropocene biosphere, framing the biosphere as the active matrix within which the climate system is embedded and on whose functioning the carbon cycle depends. The biodiversity losses documented by Cowie et al. (2022) in their assessment of the sixth-mass-extinction question, the insect declines warned of by Harvey et al. (2022), and the bird population changes assessed by Lees et al. (2022) are all indicators of a biosphere under stress, and a degraded biosphere is generally a less effective and less stable carbon sink. The physical assessment is that stabilizing temperature does not stabilize the biosphere, and that the continued adjustment of ecosystems — including the loss of resilience and the risk of shifting from sink to source — introduces a feedback that can undermine the very stabilization it accompanies.

The hydrological cycle and climate extremes under stabilization

The global hydrological cycle and the statistics of climate extremes continue to adjust as the system approaches equilibrium under a stabilized temper-

ature, and the pattern of regional climate does not become fixed at the moment global mean temperature stabilizes (*high confidence*). Because different regions and different components of the water cycle respond on different timescales, a stabilized global mean temperature is compatible with continued regional change: precipitation patterns, the frequency and intensity of extremes, and the distribution of aridity continue to evolve as the ocean and land surface equilibrate. The regional character of climate change under continued and stabilized warming is documented in specific settings by Zittis et al. (2022), whose assessment of climate change and weather extremes in the Eastern Mediterranean and Middle East characterizes a region where warming and drying are especially pronounced; such regional analyses show that the extremes which drive impacts are not uniform and do not track global mean temperature in a simple way.

The comprehensive reviews of climate change impacts, adaptation, and mitigation by Abbass et al. (2022) synthesize the breadth of consequences that persist under continued warming and that would only gradually stabilize as the climate equilibrates. The relevance to stabilization is that the impacts which motivate stabilization — extreme heat, drought, flooding, and their cascading effects — respond to the regional expression of warming, and that these regional expressions continue to sharpen for some time after the global mean stabilizes. The assessment is that stabilization of global mean temperature is a necessary but not sufficient condition for stabilization of the regional climate that determines impacts, and that the lag between global-mean stabilization and regional stabilization is itself a physical feature of the system that adaptation must contend with.

Pathways to stabilization: the emissions system

Reaching the net-zero CO₂ emissions required for stabilization entails a transformation of the energy and industrial systems that supply the emissions, and the physical characteristics of the low-carbon technologies that would replace fossil combustion — their resource requirements, their resilience to a changing climate, and their own life-cycle emissions — bear on whether and how quickly stabilization can be achieved (*high confidence*). While the detailed assessment of mitigation options is the province of Working Group III, the physical-science framing of stabilization requires acknowledging that the emissions trajectory is produced by a physical energy system. Chen et al. (2022) review strategies to achieve a carbon-neutral society, and Zhao et al. (2022) analyze the challenges and opportunities of China's energy transitions for carbon neutrality; these assessments describe the scale of the transformation in the sectors identified by Lamb et al. (2021) as the dominant emission sources.

The physical resilience of the replacement energy system is itself a consideration under a changing climate. Osman et al. (2022) assess the cost, environmental impact, and resilience of renewable energy under a changing climate, noting that the low-carbon technologies deployed to reach net zero must themselves function reliably in the altered climate they are intended to stabilize. The

feasibility of very high shares of renewable generation is examined by Breyer et al. (2022) in their assessment of the history and future of 100 percent renewable energy systems research, and the life-cycle emissions of bioenergy — relevant to the net carbon accounting of pathways that rely on biomass — are assessed by Osman et al. (2021) in their review of the conversion of biomass to biofuels and its life-cycle assessment. The industrial sources of hard-to-abate emissions include cement, and the physical routes to reducing cement's carbon intensity through supplementary cementitious materials are assessed by Snellings et al. (2023); such industrial emissions are among the residual sources that, if not eliminated, must be balanced by the carbon dioxide removal assessed in Section 9.5 for net zero to be reached. Circular-economy strategies that reduce material throughput, assessed by Yang et al. (2022) for their role in combating climate change, and the application of machine learning across the mitigation landscape reviewed by Kaack et al. (2022), represent additional physical levers on the emissions that must reach zero for temperature to stabilize.

The synthesis of this subsection is that stabilization is achieved not in the abstract but through a physical transformation of the systems that emit CO₂, and that the characteristics of that transformation — the resource base of low-carbon energy, its resilience to the climate it operates in, the residual emissions from sectors that resist full decarbonization, and the life-cycle emissions of the replacement technologies themselves — all feed back on the emissions trajectory and hence on the stabilization level that is physically achievable.

Large-scale ocean circulation under stabilization

The large-scale ocean overturning circulation responds slowly to changes in surface forcing and may continue to weaken for a period after global mean surface temperature is stabilized, with consequences for the redistribution of heat, carbon, and nutrients that persist independent of the surface temperature signal (*medium confidence*). The circulation that carries warm surface water poleward and returns cold, dense water at depth is driven by gradients in temperature and salinity that the warming perturbs through surface heating and through freshwater input from a intensified hydrological cycle and from ice loss. Because the circulation integrates these forcings over the volume of the ocean, its adjustment lags the surface, and a stabilized surface temperature does not imply a stabilized circulation. A weaker overturning redistributes less heat poleward, alters regional sea level through changes in the density structure, and changes the delivery of carbon and nutrients between the surface and the interior.

The circulation's response bears directly on the carbon cycle terms that condition stabilization. A weakening of the overturning reduces the rate at which carbon-rich surface water is transported to depth, which can alter the ocean's capacity to take up anthropogenic CO₂ during the approach to net zero and its tendency to outgas during a subsequent drawdown. The partitioning of emissions among the atmosphere, ocean, and land documented in the Global Carbon Budget assessments of Friedlingstein et al. (2022), Friedlingstein et

al. (2023), and Friedlingstein et al. (2025) reflects the present-day strength of these transport processes; a change in the circulation would shift that partitioning, and hence the atmospheric CO₂ concentration corresponding to a given cumulative emission. The physical assessment is that the ocean's overturning is a slow component whose continued adjustment under a stabilized surface temperature couples back onto the very carbon-cycle quantities that determine the stabilization level, and that this coupling is a source of the medium-confidence qualification attached to the behavior of the carbon cycle during and after stabilization.

The ecological consequences of circulation change compound those of warming and acidification assessed in Sections 9.7.3 and 9.7.4. Changes in the delivery of nutrients to the surface ocean alter primary productivity, with consequences for the marine food web whose provisioning role Boyd et al. (2022) quantify and whose biodiversity Sala et al. (2021) frame as a global stewardship concern. The thermal stresses documented by Smith et al. (2022) act in combination with circulation-driven changes in nutrient supply and oxygenation, so that the marine biosphere under a stabilized surface temperature continues to experience a changing physical and chemical environment. The assessment is that stabilization of surface temperature is compatible with continued reorganization of the ocean interior and its ecosystems over centuries.

Committed change and the millennial perspective

The full response of the Earth system to a given stabilization level unfolds over centuries to millennia, so the climate that is stabilized on the timescale of the surface temperature is not the equilibrium climate that the corresponding forcing would ultimately produce, and substantial committed change remains latent in the slow components long after the fast components have adjusted (*high confidence*). The distinction between the transient climate that is realized when surface temperature stabilizes and the equilibrium climate that the same forcing would eventually produce is central to interpreting any stabilization target. The fast components — the atmosphere, the surface ocean, sea ice, and the land surface — approach their response to a stabilized forcing within decades. The slow components — the deep ocean, the ice sheets, and the deep terrestrial and permafrost carbon stores — continue to evolve for centuries to millennia, and their eventual contribution to sea level, to the carbon cycle, and to regional climate is not realized at the moment the surface stabilizes.

This millennial perspective reframes the meaning of a stabilization level. Holding surface temperature at a given level does not hold the Earth system at the state associated with that level; rather, it initiates a long adjustment during which sea level continues to rise, the deep ocean continues to warm, and the ice sheets continue to respond to the integrated warming they have experienced. The commitment already incurred by historical emissions therefore includes changes that will only be expressed over many generations, and the choice of a stabilization level determines not only the surface climate of the coming decades but the trajectory of the slow components over the coming millennia. The

planetary-boundaries synthesis of Richardson et al. (2023), in placing the climate system alongside the other boundaries that have been transgressed, situates this long-term commitment within the broader, multi-dimensional displacement of the Earth system from its Holocene reference state — a displacement whose full expression, like that of the climate commitment itself, extends far beyond the horizon over which surface temperature stabilizes.

The irreversibilities assessed throughout this chapter acquire their full significance in this millennial frame. The hysteresis of the marine-based ice sheets, the effectively permanent character of deep-ocean heat uptake, the biodiversity losses documented by Cowie et al. (2022), Harvey et al. (2022), and Lees et al. (2022), and the degradation of carbon-storing biomes documented by Forzieri et al. (2022) and Lapola et al. (2023) all represent changes that are not reversed on the timescale over which surface temperature can be stabilized or even modestly reduced. The physical-science conclusion is that the assessment of a stabilization level must extend beyond the surface temperature of the coming century to the committed trajectory of the slow components over the coming millennia, because it is in those slow components that the largest and least reversible consequences of the chosen level are ultimately expressed.

Monitoring and verifying stabilization

Verifying that the climate system is approaching or has reached stabilization requires sustained observation of the coupled system — the atmospheric CO₂ burden and its growth rate, the partitioning of emissions among the reservoirs, the planetary energy imbalance, and the indicators of human influence — because no single quantity is sufficient to confirm that emissions have reached net zero and that temperature has ceased to trend (*high confidence*). Because stabilization is a property of the whole system, its verification cannot rest on surface temperature alone, which is subject to internal variability that can mask or mimic a change in the underlying trend over short periods. The indicators synthesis of Forster et al. (2023) provides the framework for this verification, assembling the large-scale indicators of the state of the climate system and of human influence — including estimates of human-induced warming, the components of effective radiative forcing, and the planetary energy imbalance — into an annually updated assessment against which the approach to stabilization can be judged.

The atmospheric CO₂ burden and its partitioning are the primary observational anchor. The Global Carbon Budget assessments of Friedlingstein et al. (2022), Friedlingstein et al. (2023), and Friedlingstein et al. (2025) quantify the total emissions and their partitioning among the atmosphere, ocean, and land, and it is the atmospheric growth rate that would first reflect an approach to net zero: as net emissions decline toward zero, the atmospheric CO₂ growth rate would fall toward zero, and a sustained cessation of atmospheric CO₂ accumulation would be among the earliest large-scale signatures of an approach to stabilization. The continued measurement of the ocean and land sinks is essential to this interpretation, because a change in the sinks — such as the

weakening implied by the declining forest resilience of Forzieri et al. (2022) and the forest degradation of Lapola et al. (2023) — would alter the relationship between emissions and atmospheric concentration and hence the atmospheric signature of a given emissions trajectory.

The planetary energy imbalance provides an independent and physically fundamental measure of the approach to stabilization. So long as the energy imbalance is positive, the system continues to accumulate energy, predominantly in the ocean, and surface temperature has not fully equilibrated with the forcing. A decline of the energy imbalance toward zero would indicate that the system is approaching equilibrium with the stabilized forcing, and the continued warming of the deep ocean assessed in Section 9.7.2 means that the energy imbalance persists even as surface temperature levels off. The physical assessment is that confirming stabilization requires observing the convergence of several independent quantities — the atmospheric CO₂ growth rate, the emissions partitioning, the planetary energy imbalance, and the derived indicators of human influence — and that a sustained observing system spanning the atmosphere, the ocean interior, and the terrestrial and marine carbon stores is a prerequisite for distinguishing genuine stabilization from the transient signatures of internal variability. The reviews of the observational and modeling basis for climate assessment, including the comprehensive synthesis of impacts and mitigation by Abbass et al. (2022), underscore that the verification of stabilization is an ongoing observational undertaking rather than a one-time determination.

9.9 Synthesis and key uncertainties

Taken together, the evidence assessed in this chapter establishes that stabilizing global surface temperature is physically achievable by bringing net CO₂ emissions to zero, that the level of stabilization is set by cumulative emissions, that warming can be halted within a human generation but reversed only through sustained net-negative emissions, and that many components of the Earth system continue to change under a stabilized temperature (*high confidence*). The near-linear relationship between cumulative CO₂ emissions and peak warming is the unifying result: it underwrites the remaining carbon budget, it implies the near-zero zero-emissions commitment, and it makes clear that the central determinant of the stabilization level is the total quantity of CO₂ emitted before net zero is reached. The successive Global Carbon Budget assessments of Friedlingstein et al. (2022), Friedlingstein et al. (2023), and Friedlingstein et al. (2025) show that this total continues to accumulate at close to its historical rate, contracting the budgets for the lowest stabilization levels toward exhaustion within the current decade.

The chapter has emphasized three physically distinct concepts that must not be conflated. Halting warming is achieved at net-zero CO₂ and requires no active drawdown; reversing warming requires sustained net-negative emissions and is made harder by the carbon-cycle giveback that returns part of any removal to the atmosphere; and the reversibility of specific impacts ranges from relatively

high for global mean surface temperature to very low for deep-ocean heat content, ice-sheet retreat, species extinctions, and the ecological consequences of marine heatwaves. Overshoot pathways, which the small remaining budgets for low limits make likely, incur commitments and irreversibilities that scale with the depth and duration of the overshoot, providing a physical rationale for minimizing both.

Carbon dioxide removal is assessed as physically necessary for net zero in the presence of residual emissions and indispensable for any reversal of warming, with land-based, ocean-based, and engineered approaches differing sharply in permanence, monitorability, and side effects; the net climatic benefit of any approach is materially less than its gross removal once carbon-cycle giveback, permanence risk, and biophysical effects are counted. Solar radiation modification is assessed as capable of rapidly lowering temperature by increasing planetary albedo but as addressing only the thermal symptom of elevated CO₂, leaving ocean acidification unchecked, producing distinct regional and hydrological changes, and creating a termination risk of rapid warming if sustained deployment ceased; it is therefore not a substitute for reducing the CO₂ burden.

The principal uncertainties that condition these conclusions are the following. First, the strength of the cumulative-emissions-to-warming relationship and the contribution of non-CO₂ forcings set the magnitude of the remaining budget, and both carry uncertainty that maps directly onto the allowable emissions for a given limit; the annually updated indicators of Forster et al. (2023) narrow but do not eliminate this uncertainty. Second, the zero-emissions commitment is assessed with medium confidence because Earth system feedbacks — permafrost carbon, and the biosphere feedbacks implied by the declining resilience documented by Forzieri et al. (2022) and the degradation documented by Lapola et al. (2023) — could shift the commitment toward positive values not captured in the models used to derive it. Third, the efficiency of net-negative emissions in reducing atmospheric CO₂, governed by the carbon-cycle giveback, determines how much removal is required to reverse a given amount of warming and is a key control on the feasibility of overshoot-and-recovery pathways. Fourth, the thresholds for irreversible ice-sheet retreat and other slow-component hysteresis behaviors are uncertain, and the catastrophic tail of the risk distribution — argued by Kemp et al. (2022) to warrant systematic exploration — remains poorly constrained. Fifth, the permanence of land- and ocean-based carbon storage under a warming climate, and the aerosol-related uncertainties that condition the efficacy and side effects of solar radiation modification illustrated by Brown et al. (2021), limit confidence in the net effect of the deliberate interventions.

The overarching physical-science conclusion is that the levers available for stabilization — emissions reduction, carbon dioxide removal, and solar radiation modification — are not interchangeable. Emissions reduction to net zero is the only lever that halts warming by addressing its cause; carbon dioxide removal is the only lever that can reverse warming and is required to reach net zero given residual emissions; and solar radiation modification can only mask warming temporarily while leaving its cause and its non-thermal consequences in place. The physical evidence assessed here does not prescribe how these levers should

be used, but it establishes unambiguously that stabilization at low levels requires the earliest achievable approach to net-zero CO₂, that reliance on overshoot and later reversal is physically costlier and riskier than avoiding overshoot, and that the slow components of the Earth system commit the planet to continued change for centuries after the surface temperature is stabilized.

9.10 References

Abbass, K., et al. (2022). A review of the global climate change impacts, adaptation, and sustainable mitigation measures. *Environmental Science and Pollution Research*. <https://doi.org/10.1007/s11356-022-19718-6>

Boyd, C., et al. (2022). The contribution of fisheries and aquaculture to the global protein supply. *Food Security*. <https://doi.org/10.1007/s12571-021-01246-9>

Breyer, C., et al. (2022). On the History and Future of 100% Renewable Energy Systems Research. *IEEE Access*. <https://doi.org/10.1109/access.2022.3193402>

Brown, H., et al. (2021). Biomass burning aerosols in most climate models are too absorbing. *Nature Communications*. <https://doi.org/10.1038/s41467-020-20482-9>

Chen, L., et al. (2022). Strategies to achieve a carbon neutral society: a review. *Environmental Chemistry Letters*. <https://doi.org/10.1007/s10311-022-01435-8>

Cowie, R., et al. (2022). The Sixth Mass Extinction: fact, fiction or speculation? *Biological Reviews*. <https://doi.org/10.1111/brv.12816>

Dziejarski, B., et al. (2023). Current status of carbon capture, utilization, and storage technologies in the global economy: A survey of technical assessment. *Fuel*. <https://doi.org/10.1016/j.fuel.2023.127776>

Erans, M., et al. (2022). Direct air capture: process technology, techno-economic and socio-political challenges. *Energy & Environmental Science*. <https://doi.org/10.1039/d1ee03523a>

Fankhauser, S., et al. (2021). The meaning of net zero and how to get it right. *Nature Climate Change*. <https://doi.org/10.1038/s41558-021-01245-w>

Folke, C., et al. (2021). Our future in the Anthropocene biosphere. *AMBIO*. <https://doi.org/10.1007/s13280-021-01544-8>

Forster, P., et al. (2023). Indicators of Global Climate Change 2022: annual update of large-scale indicators of the state of the climate system and human influence. *Earth System Science Data*. <https://doi.org/10.5194/essd-15-2295-2023>

Forzieri, G., et al. (2022). Emerging signals of declining forest resilience under climate change. *Nature*. <https://doi.org/10.1038/s41586-022-04959-9>

Friedlingstein, P., et al. (2022). Global Carbon Budget 2022. *Earth System Science Data*. <https://doi.org/10.5194/essd-14-4811-2022>

Friedlingstein, P., et al. (2023). Global Carbon Budget 2023. *Earth System Science Data*. <https://doi.org/10.5194/essd-15-5301-2023>

Friedlingstein, P., et al. (2025). Global Carbon Budget 2024. *Earth System Science Data*. <https://doi.org/10.5194/essd-17-965-2025>

Harvey, J., et al. (2022). Scientists' warning on climate change and insects. *Ecological Monographs*. <https://doi.org/10.1002/ecm.1553>

Kaack, L., et al. (2022). Tackling Climate Change with Machine Learning. *ACM Computing Surveys*. <https://doi.org/10.1145/3485128>

Kemp, L., et al. (2022). Climate Endgame: Exploring catastrophic climate change scenarios. *Proceedings of the National Academy of Sciences*. <https://doi.org/10.1073/pna>

Lamb, W., et al. (2021). A review of trends and drivers of greenhouse gas emissions by sector from 1990 to 2018. *Environmental Research Letters*. <https://doi.org/10.1088/1748-9326/abee4e>

Lapola, D., et al. (2023). The drivers and impacts of Amazon forest degradation. *Science*. <https://doi.org/10.1126/science.abp8622>

Lees, A., et al. (2022). State of the World's Birds. *Annual Review of Environment and Resources*. <https://doi.org/10.1146/annurev-environ-112420-014642>

Lynch, J., et al. (2021). Agriculture's Contribution to Climate Change and Role in Mitigation Is Distinct From Predominantly Fossil CO₂-Emitting Sectors. *Frontiers in Sustainable Food Systems*. <https://doi.org/10.3389/fsufs.2020.518039>

Osman, A., et al. (2021). Conversion of biomass to biofuels and life cycle assessment: a review. *Environmental Chemistry Letters*. <https://doi.org/10.1007/s10311-021-01273-0>

Osman, A., et al. (2022). Cost, environmental impact, and resilience of renewable energy under a changing climate: a review. *Environmental Chemistry Letters*. <https://doi.org/10.1007/s10311-022-01532-8>

Richardson, K., et al. (2023). Earth beyond six of nine planetary boundaries. *Science Advances*. <https://doi.org/10.1126/sciadv.adh2458>

Riahi, K., et al. (2021). Cost and attainability of meeting stringent climate targets without overshoot. *Nature Climate Change*. <https://doi.org/10.1038/s41558-021-01215-2>

Sala, E., et al. (2021). Protecting the global ocean for biodiversity, food and climate. *Nature*. <https://doi.org/10.1038/s41586-021-03371-z>

Seddon, N., et al. (2021). Getting the message right on nature-based solutions to climate change. *Global Change Biology*. <https://doi.org/10.1111/gcb.15513>

Smith, K., et al. (2022). Biological Impacts of Marine Heatwaves. *Annual Review of Marine Science*. <https://doi.org/10.1146/annurev-marine-032122-121437>

Snellings, R., et al. (2023). Future and emerging supplementary cementitious materials. *Cement and Concrete Research*. <https://doi.org/10.1016/j.cemconres.2023.107199>

Tebaldi, C., et al. (2021). Climate model projections from the Scenario Model Intercomparison Project (ScenarioMIP) of CMIP6. *Earth System Dynamics*. <https://doi.org/10.5194/esd-12-253-2021>

Welsby, D., et al. (2021). Unextractable fossil fuels in a 1.5 degrees C world. *Nature*. <https://doi.org/10.1038/s41586-021-03821-8>

Yan, Y., et al. (2021). Harnessing the power of machine learning for carbon capture, utilisation, and storage (CCUS) - a state-of-the-art review. *Energy & Environmental Science*. <https://doi.org/10.1039/d1ee02395k>

Yang, M., et al. (2022). Circular economy strategies for combating climate change and other environmental issues. *Environmental Chemistry Letters*. <https://doi.org/10.1007/s10311-022-01499-6>

Zhao, X., et al. (2022). China's energy transitions for carbon neutrality: challenges and opportunities. *Carbon Neutrality*. <https://doi.org/10.1007/s43979-022-00010-y>

Zittis, G., et al. (2022). Climate Change and Weather Extremes in the Eastern Mediterranean and Middle East. *Reviews of Geophysics*. <https://doi.org/10.1029/2021rg00076>

Chapter 10

Chapter 10: Climate Information for Decisions and Services

10.1 Executive Summary

Climate information has moved from being a scientific by-product of Earth system assessment to a distinct, decision-facing infrastructure that mediates between physical-science understanding and the choices made by governments, firms, communities, and individuals (*high confidence*). Across the period assessed in this Report, the quantity, resolution, and accessibility of climate information have grown rapidly, but the binding constraint on its value for decisions is no longer primarily the availability of data; it is the fidelity of the processes that translate physical-climate knowledge into forms that are salient, credible, and legitimate to specific users (*high confidence*).

The co-production of climate information — sustained collaboration between information producers and the decision-makers who use it — measurably improves the uptake and usability of that information relative to one-directional “loading dock” delivery, but co-production is resource-intensive, unevenly distributed, and difficult to sustain (*medium confidence*). Frameworks developed in implementation science provide transferable structure for understanding why some climate-information services are adopted and sustained while others are not (Damschroder et al., 2022) (*medium confidence*).

The observational, monitoring, and reanalysis foundation on which decision-relevant climate information rests has strengthened substantially, and global reanalyses and satellite-derived land-surface products now provide internally consistent, decision-usable representations of the recent climate at fine spatial resolution (*high confidence*). Continuous global land-cover and soil products at 30 m and finer resolution (Zhang et al., 2021; Poggio et al., 2021) and quantitatively reconciled global carbon accounting (Friedlingstein et al., 2022) exemplify the monitoring backbone now available to decision-makers, though systematic gaps remain over parts of the Global South and for many essential variables (*high confidence*).

Artificial intelligence and machine-learning methods have become established components of the climate-information value chain, delivering

skilful short-range prediction, computationally efficient emulation, pattern extraction from large Earth-observation archives, and new modes of down-scaling and bias correction (*high confidence*). Deep generative models now produce precipitation nowcasts that expert forecasters judge more useful than established numerical alternatives (Ravuri et al., 2021), and structured reviews document machine learning’s expanding role across mitigation, adaptation, and Earth-system science (Kaack et al., 2022) (*high confidence*). These gains are accompanied by well-characterized risks — brittle out-of-distribution behaviour, opaque uncertainty (Gawlikowski et al., 2023), and the non-trivial environmental and equity costs of the computation itself (Wynsberghe, 2021) (*high confidence*).

Robust, well-calibrated communication of uncertainty is a prerequisite for climate information to support good decisions, and the calibrated-language framework used throughout this Report is one instrument among several for conveying the state of knowledge honestly to non-specialist users (*high confidence*). Where uncertainty is suppressed, misrepresented, or presented without decision-relevant framing, climate information can degrade rather than improve decisions (*medium confidence*).

Regional climate risk assessment is the principal interface through which the physical-science findings of this Working Group enter the impact, adaptation, and vulnerability assessments of Working Group II, and the quality of that interface depends on shared definitions of hazard, exposure, and vulnerability and on the traceable propagation of uncertainty across the boundary (*high confidence*). Sectoral evidence — in water resources (Scanlon et al., 2023), agriculture and food security (Jägermeyr et al., 2021), its biological hazards (Singh et al., 2023), heat-related mortality (Zhao et al., 2021), and land and soil degradation (Hassani et al., 2021) — demonstrates both the maturity of hazard characterization and the persistent difficulty of converting hazard information into locally actionable risk information (*high confidence*).

Equity, capacity, and the inclusion of Indigenous knowledge and local communities are not peripheral to the production of decision-relevant climate information but are central determinants of whether that information is trusted and used, and of whether its benefits are distributed fairly (*medium confidence*). Evidence on conservation and land governance indicates that outcomes are more effective and more equitable where Indigenous peoples and local communities are recognized as knowledge-holders and decision-partners rather than as passive recipients (Dawson et al., 2021) (*medium confidence*).

10.2 Introduction: The Purpose and Scope of Climate Information for Decisions

Climate information for decisions is the subset of climate knowledge that has been selected, processed, contextualized, and delivered specifically to inform a choice under an uncertain climate, and it is distinguished from climate data and from climate science by the presence of an identified decision and an identified decision-maker (*high confidence*). This chapter assesses the state of

the science and practice underlying that translation. It does not re-assess the physical basis of climate change, which is the subject of earlier chapters; instead it assesses how the outputs of that physical-science enterprise — observations, reanalyses, projections, and their associated uncertainties — are converted into information that can be acted upon, and what determines whether that conversion succeeds.

The distinction matters because the pathways by which climate knowledge fails to support good decisions are frequently not failures of the underlying science. A projection may be physically sound, internally consistent, and appropriately bounded by uncertainty, and yet be unusable because it is delivered at the wrong spatial scale, framed against the wrong baseline, expressed in units foreign to the user, delivered after the decision window has closed, or accompanied by uncertainty language that the user cannot interpret. Conversely, information that is scientifically modest — a seasonal tendency, a simple threshold exceedance frequency — may be highly decision-relevant when it is co-designed with the user, delivered on the user's timescale, and embedded in the user's existing decision process. The central assessment of this chapter is therefore that the value of climate information is a joint property of the information and the decision context, and cannot be evaluated by reference to scientific quality alone (*high confidence*).

What counts as a decision-relevant climate service

A climate service is the sustained provision of climate information, in a usable form, to support decisions, and it is defined by the persistence of the provision and the usability of the product rather than by any particular scientific method (*high confidence*). The concept spans a wide range of institutional arrangements: national meteorological and hydrological services issuing seasonal outlooks; regional climate outlook fora; commercial providers packaging reanalysis and projection data for the finance, energy, and agriculture sectors; and community-based monitoring arrangements that couple local observation with external forecast products. What unites them is a commitment to ongoing service rather than one-off assessment, and an orientation toward a defined user.

The maturation of climate services over the assessed period reflects three convergent developments. First, the underlying information base has become richer and more openly accessible, as global reanalyses, satellite Earth-observation archives, and standardized climate-model output have become routinely available. Second, the community has accumulated substantial practical experience about why climate information is or is not used, much of it consistent with insights from implementation science and the broader literature on evidence uptake (Damschroder et al., 2022). Third, computational methods — including machine learning — have lowered the cost of transforming raw information into tailored products, enabling forms of downscaling, emulation, and prediction that were previously infeasible (Kaack et al., 2022). Each of these developments is assessed in detail in the sections that follow.

The decision-first framing

Organizing climate information around decisions rather than around scientific outputs — a “decision-first” framing — systematically changes what information is judged relevant and how its quality is evaluated (*high confidence*). Under a science-first framing, the natural question is “what does the climate system tell us?” and quality is assessed by fidelity to the physical system. Under a decision-first framing, the question is “what does this decision need to know, and how sensitive is the decision to what we do not know?” and quality is assessed by decision-relevance and by the robustness of the decision to residual uncertainty.

The decision-first framing has several consequences that recur throughout this chapter. It privileges information about thresholds, tails, and rates of change over information about central tendencies, because many consequential decisions are governed by the crossing of critical values rather than by mean states. It elevates the importance of the baseline and the reference period, because most decisions are framed against expectations calibrated on the past. And it makes the communication of uncertainty a first-order design problem rather than a caveat, because a decision-maker’s optimal action often depends more on the shape and framing of the uncertainty than on the central estimate. These themes — thresholds, baselines, and uncertainty framing — are developed in Sections 10.3, 10.5, and 10.6.

Scope, structure, and relationship to other chapters

This chapter is structured to follow the climate-information value chain from the human and institutional processes that make information usable, through the observational foundation, to the computational methods that transform it, to the communication of uncertainty, and finally to the regional risk-assessment interface with Working Group II. Section 10.2 assesses climate services and the co-production of actionable information. Section 10.3 assesses the observational, monitoring, and reanalysis foundation. Section 10.4 assesses the role of artificial intelligence and machine learning. Section 10.5 assesses the communication of uncertainty and calibrated confidence. Section 10.6 assesses regional climate risk assessment and its interface to Working Group II. Section 10.7 assesses cross-cutting questions of equity, capacity, and knowledge systems.

Throughout, the chapter maintains the calibrated-language conventions of this Report, and it is explicit about where the evidence supports strong statements and where it does not. Because the practice of climate services is comparatively young and the evidence base is heterogeneous, many of this chapter’s assessments carry *medium confidence*, and the chapter is candid about the limits of what can currently be concluded.

10.3 Climate Services and the Co-Production of Actionable Information

The co-production of climate information — a sustained, iterative collaboration in which producers and users jointly frame the problem, design the product, and interpret the results — improves the salience, credibility, and legitimacy of climate information and thereby its uptake, relative to one-directional delivery, though the evidence base is dominated by case studies and the effect is difficult to quantify in aggregate (*medium confidence*). The central claim of the co-production literature is that usability is not a property that can be added to information after it is produced; it must be built in from the framing stage, because the choices that determine usability — spatial scale, variables, baseline, format, timing, and uncertainty representation — are made early and are expensive to reverse.

From the loading-dock model to co-production

The historically dominant mode of climate-information provision has been described as a “loading dock” model, in which scientific institutions produce information according to their own priorities and make it available for users to collect. This model implicitly assumes that the gap between information and use is an accessibility gap: if the information exists and can be reached, it will be used. The accumulated experience of the assessed period contradicts this assumption. Highly accessible, scientifically excellent information is frequently not used, and the reasons are rarely reducible to accessibility. They include mismatches of scale, mismatches of timing, misalignment with the decision-maker’s mental model, distrust of the source, and the absence of institutional pathways through which the information can enter a decision.

Co-production reframes the problem as one of joint knowledge construction. In its stronger forms it involves users in problem definition, in the selection of variables and metrics, in the evaluation of prototype products, and in the interpretation of results, with iteration across these stages. The transferable insight from implementation science is that the adoption and sustainment of any evidence-based practice depend on characteristics of the innovation, the users, the inner and outer institutional settings, and the process of implementation itself, and that neglecting any of these domains predicts failure (Damschroder et al., 2022). Climate-information services that are technically sound but institutionally unsupported, or that impose high adoption costs on users, are predictably under-used regardless of scientific quality (*medium confidence*).

What co-production changes about the information

Co-production tends to shift climate information toward finer spatial specificity, toward decision-relevant thresholds and impact-proximate variables, and toward explicit treatment of the baseline against which change is judged (*medium confidence*). When users participate in framing, they tend to ask

for information about the variables that directly enter their decisions — often impact-proximate quantities such as growing-degree-days, consecutive dry days, flood-relevant return periods, or heat-exposure indices — rather than the raw climate variables that models most directly produce. This pulls the information base toward integration with sectoral models and toward the derivation of tailored indices.

The agricultural sector illustrates the pattern. Decision-relevant agricultural climate information is rarely a temperature or precipitation projection in isolation; it is a projection filtered through crop physiology, pest and pathogen dynamics, soil-water balance, and management options. The assessed literature documents the sensitivity of crop outcomes to climate through coupled climate-crop modelling (Jägermeyr et al., 2021), the dependence of agricultural risk on shifting pest and pathogen pressure (Singh et al., 2023), and the interaction of climate stress with soil degradation and salinization (Hassani et al., 2021; Malhi et al., 2021). Co-produced agricultural climate services integrate these strands so that the delivered information corresponds to a decision — what to plant, when to plant, how to allocate water — rather than to a climate variable (*medium confidence*).

The costs and limits of co-production

Co-production is resource-intensive, does not scale straightforwardly, and is unevenly available, so that its benefits accrue disproportionately to well-resourced users and regions, raising equity concerns that recur throughout this chapter (*medium confidence*). Sustained collaboration requires time, trust, boundary-spanning personnel, and stable institutional relationships, all of which are scarce and unevenly distributed. A service co-produced with a single well-resourced user may not transfer to others, and the intensive engagement that makes co-production effective is precisely what makes it difficult to deliver at population scale.

This tension — between the demonstrated benefits of deep engagement and the need to serve many users — is one of the central unresolved problems of climate-service design. Two partial responses are visible in the assessed practice. The first is the development of intermediary or “boundary” organizations that specialize in translation and can amortize the cost of engagement across many users. The second is the use of computational methods, including machine learning, to lower the marginal cost of tailoring information once a co-produced template exists (Kaack et al., 2022). Neither response fully resolves the tension, and the chapter assesses with *medium confidence* that equitable access to co-produced climate information remains a significant unmet need, particularly in regions with weak national climate-service institutions.

Trust, credibility, and the social context of climate information

Whether climate information is used depends heavily on whether its source is trusted, and trust is a social and institutional property that scientific quality

alone does not confer (*high confidence*). The salience-credibility-legitimacy framing that recurs in the climate-services literature captures this: information must be relevant to the decision (salient), believed to be accurate (credible), and produced through a process the user regards as fair and unbiased (legitimate). Legitimacy in particular is a property of process, not product, and is a principal reason that co-production improves uptake — the act of participating in production confers legitimacy that cannot be added afterward.

The social context also determines whose knowledge counts. Evidence from conservation and land governance indicates that information systems that recognize and incorporate the knowledge of Indigenous peoples and local communities produce outcomes that are both more effective and more equitable than systems that treat those communities only as recipients (Dawson et al., 2021). This finding, developed further in Section 10.7, implies that the design of legitimate climate services is inseparable from questions about which knowledge systems are admitted and on what terms (*medium confidence*).

10.4 Observations, Monitoring, and Reanalysis as the Foundation for Decision-Relevant Climate Information

All decision-relevant climate information ultimately rests on observation, and the strengthening of the global observing, monitoring, and reanalysis enterprise over the assessed period has materially expanded what can be reliably said about the recent and current state of the climate at scales useful for decisions (*high confidence*). Decisions are made in the present about the future, but they are anchored in an understanding of the recent past; a decision-maker's confidence in a projection is inseparable from confidence in the observational record against which models are evaluated and against which change is measured. This section assesses the observational foundation with particular attention to the properties that matter for decisions: spatial resolution, temporal continuity, internal consistency, and the honest characterization of gaps.

Reanalysis and the consistent representation of the recent climate

Global reanalyses now provide gridded, dynamically and physically consistent estimates of many climate variables spanning decades, and they have become a primary source of decision-relevant information about the recent climate, particularly in regions where direct observations are sparse (*high confidence*). By assimilating heterogeneous observations into a numerical model, reanalysis produces a spatially and temporally complete estimate that is internally consistent across variables — a property that raw observations, with their gaps and inhomogeneities, do not possess. For decisions, this consistency is valuable: an energy planner assessing joint wind-and-temperature extremes, or a water manager assessing the co-occurrence of high evapotranspiration

and low precipitation, requires variables that are mutually consistent at each location and time.

Reanalysis is not observation, and its estimates inherit both the strengths and the biases of the assimilating model, with quality varying by variable, region, and epoch as the observing system itself changes through time. For decision use, these caveats are consequential: apparent trends in a reanalysis can reflect changes in the observing system rather than in the climate, and the reliability of derived quantities such as surface fluxes is generally lower than that of directly assimilated quantities such as upper-air temperature. The assessment is therefore that reanalyses are an indispensable but not infallible foundation for decision-relevant climate information, and that their fitness for a given decision must be evaluated against the specific variables and scales the decision requires (*high confidence*).

Earth observation and the monitoring of the land surface

Satellite Earth observation now supports continuous, globally consistent monitoring of the land surface at spatial resolutions of tens of metres, providing decision-relevant information on land cover, land-cover change, cropland extent, soil properties, and vegetation condition (*high confidence*). Global land-cover products at 30 m resolution derived from multi-temporal satellite imagery (Zhang et al., 2021) and globally consistent soil-property estimates with quantified spatial uncertainty (Poggio et al., 2021) exemplify the monitoring capabilities now routinely available. These products underpin a wide range of decisions: assessing exposure of land uses to climate hazards, monitoring the encroachment of agriculture onto vulnerable lands, and providing the surface characterization that hydrological and agricultural models require.

The decision value of these products depends on properties beyond nominal resolution. Temporal continuity determines whether change can be detected and attributed; documented and quantified uncertainty determines whether apparent changes are real; and consistency of classification through time determines whether a time series reflects the world or the evolving classifier. The assessed products increasingly meet these standards — for example, by providing per-pixel uncertainty (Poggio et al., 2021) — but many operational land-monitoring streams do not, and users are cautioned that a change detected in an inconsistent product may be an artefact of the monitoring system (*high confidence*).

Monitoring land-cover and land-use change for decisions

Global land-use and land-cover change has been substantially larger and more dynamic over recent decades than earlier estimates suggested, and improved monitoring has revised upward the assessed pace of change, with direct implications for exposure and vulnerability assessment (*high confidence*). Reconstructions drawing on long satellite archives indicate that global land-use change over recent decades is several times greater than previously estimated

(Winkler et al., 2021), and that cropland expansion has accelerated in the twenty-first century, with much new cropland established on land of relatively low agricultural potential and often at the expense of natural ecosystems (Potapov et al., 2021). For decision-makers, these findings reframe the exposure side of risk: the human footprint that is exposed to climate hazards is itself changing rapidly, so that static exposure assumptions will misstate risk.

This dynamism interacts with climate hazard in ways that decision-relevant information must capture. Expansion of agriculture onto marginal land increases exposure to drought, heat, and soil degradation; conversion of natural ecosystems alters the land-surface properties that mediate local climate and hydrology; and the concentration of new cropland in particular regions changes the geography of food-system vulnerability (Potapov et al., 2021; Winkler et al., 2021). Climate information that treats exposure as fixed will systematically misjudge these coupled hazard-exposure dynamics, and the assessment is that decision-relevant risk information increasingly requires the joint monitoring of climate and land change (*high confidence*).

Monitoring the carbon cycle and constraining the trajectory

Quantitative, regularly updated accounting of the global carbon budget provides decision-makers with an authoritative, uncertainty-bounded measure of the emissions trajectory against which mitigation decisions can be evaluated (*high confidence*). The reconciliation of anthropogenic emissions with the growth of atmospheric carbon dioxide and the uptake by land and ocean sinks (Friedlingstein et al., 2022) constitutes one of the most policy-relevant monitoring products in climate science: it converts the abstract goal of emissions reduction into a measured quantity that can be tracked annually and attributed to sources. For decisions about mitigation, this monitoring is the equivalent of a financial account against which progress is audited.

The carbon-budget framing also structures the meaning of long-term goals. The interpretation of “net zero” — what it requires, over what boundary, and with what treatment of residual emissions and removals — is a question about how monitored quantities map onto policy commitments, and the assessed literature emphasizes that ambiguity in these definitions can undermine the integrity of the goals themselves (Fankhauser et al., 2021). Decision-relevant carbon information therefore requires not only measurement but agreed accounting conventions, and the assessment is that the monitoring infrastructure is now sufficient to track the global trajectory with useful precision, while the translation of that trajectory into accountable commitments remains contested (*high confidence*).

Observing gaps and the geography of information poverty

The global observing system remains spatially uneven, with the density and continuity of in-situ observation markedly lower across much of the Global South, and this unevenness propagates into every downstream product, so

that decision-relevant climate information is systematically weaker precisely where vulnerability is often highest (*high confidence*). Reanalyses depend on assimilated observations and are less well constrained where those observations are sparse; satellite products, though globally uniform in coverage, still require in-situ data for calibration and validation; and derived services inherit the resulting geography of confidence. The result is a pattern in which the regions with the greatest exposure to climate hazards and the least institutional capacity to respond are also the regions where the underlying information is weakest.

This “information poverty” is not merely a scientific inconvenience; it is a determinant of adaptive capacity, and it compounds the equity concerns raised in Sections 10.2 and 10.7. The assessment is that closing observing-system gaps in under-observed regions would yield disproportionate improvements in decision-relevant information there, and that the value of such investment is high relative to its cost, although it is realized only if accompanied by the institutional capacity to convert observations into services (*high confidence*). Machine-learning methods, assessed in Section 10.4, offer partial mitigation by extracting more information from existing observations and by fusing heterogeneous data sources, but they cannot manufacture information where none has been observed (*high confidence*).

Water, an integrating variable for monitoring

Water is the variable through which climate change is most directly experienced in many decisions, and the monitoring of water resources — surface water, groundwater, and water quality — is an integrating challenge that draws on the full observational enterprise (*high confidence*). Assessments of global water resources emphasize the central and often under-appreciated role of groundwater as a buffer against climate variability and the importance of monitoring it if water systems are to remain resilient (Scanlon et al., 2023). Because groundwater is difficult to observe directly at scale, its monitoring exemplifies the fusion of in-situ measurement, satellite gravimetry, and modelling that characterizes modern decision-relevant water information.

Water quality adds a further monitoring dimension that is frequently neglected in climate assessment but is decisive for many decisions. Degradation of water quality arises from an interacting set of natural and anthropogenic drivers, and climate change modulates several of them (Akhtar et al., 2021). Decision-relevant water information therefore increasingly integrates quantity and quality, surface and subsurface, and climate and land drivers — an integration that stresses the observing system but that is necessary if water services are to support the decisions that depend on them (*high confidence*).

10.5 Artificial Intelligence and Machine Learning in Climate Information

Artificial intelligence and machine learning have become established, load-bearing components of the climate-information value chain, contributing skilful short-range prediction, computationally cheap emulation of expensive models, extraction of information from large Earth-observation archives, and new approaches to downscaling and bias correction, while introducing distinctive risks that must be managed if the resulting information is to be trustworthy for decisions (*high confidence*). The assessment in this section is that machine learning is neither a panacea nor a passing fashion in climate information; it is a set of methods with well-characterized strengths and equally well-characterized failure modes, and its responsible use for decisions depends on matching method to task and on rigorous, honest uncertainty quantification.

The landscape of machine-learning applications

Machine-learning methods now span the climate-information value chain — from Earth-observation processing, through prediction and emulation, to the design and evaluation of mitigation and adaptation options — and structured surveys of the field document both the breadth of application and the uneven maturity across use cases (*high confidence*). Comprehensive reviews map machine learning onto a wide range of climate-relevant problems, distinguishing applications where the methods are mature and reliable from those where they remain exploratory, and emphasizing that the appropriate role of machine learning varies sharply by task (Kaack et al., 2022). This mapping is itself decision-relevant, because it allows information producers and users to calibrate their confidence in machine-learning-derived products according to the demonstrated maturity of the underlying method.

The methodological substrate for many of these applications is the family of deep neural network architectures whose capabilities in perception and pattern recognition have advanced rapidly. Convolutional architectures for image and gridded-field analysis, reviewed comprehensively in the computer-vision literature (Zhao et al., 2024), transfer naturally to gridded climate and Earth-observation data, where spatial structure is central. The same architectures that classify natural images can classify land cover, detect features in satellite imagery, and learn spatial patterns in climate fields, and this transfer of methods from computer vision into Earth-system science accounts for much of the recent expansion of machine learning in climate information (*high confidence*).

Skilful prediction and nowcasting

Deep generative and other machine-learning models now produce short-range forecasts — notably precipitation nowcasts — that expert users judge to be more useful than established numerical alternatives, representing one of the clearest demonstrations of decision-relevant machine-learning skill in

climate information (*high confidence*). Deep generative models of radar can produce precipitation nowcasts that are both spatially realistic and probabilistically calibrated, and in structured evaluation with professional forecasters these products were preferred to competing methods for their operational usefulness (Ravuri et al., 2021). This result is significant not only as a technical achievement but as an instance of evaluation done correctly: skill was assessed against the decisions the forecasts inform and against the judgement of the users who make those decisions, rather than against a purely statistical score.

The nowcasting case also illustrates the appropriate scope of current machine-learning prediction. The demonstrated skill is at short lead times, where the relationship between recent observations and near-future states is strong and where large training archives exist. As lead time extends and the relevant dynamics become more chaotic and less directly constrained by recent observation, the advantage of purely data-driven methods diminishes and the role of physically based models grows. The assessment is that machine-learning prediction delivers clear, decision-relevant value at short ranges and for well-observed quantities, and that its extension to longer ranges and rarer events must be evaluated case by case rather than assumed (*high confidence*).

Emulation, downscaling, and bias correction

Machine-learning emulators can reproduce the input-output behaviour of expensive physical models at a small fraction of the computational cost, enabling large ensembles, rapid scenario exploration, and finer effective resolution, provided their behaviour outside the training distribution is treated with appropriate caution (*high confidence*). Emulation addresses a persistent constraint on decision-relevant climate information: the computational expense of physical models limits the number of scenarios, ensemble members, and resolutions that can be explored, and thereby limits the characterization of uncertainty. By learning a fast surrogate for a slow model, machine learning can relax this constraint and support the large ensembles that robust uncertainty characterization requires.

Downscaling and bias correction are closely related applications in which machine learning maps coarse-resolution model output to the finer scales and corrected statistics that decisions require. These methods are attractive because they can capture complex, non-linear relationships between coarse predictors and fine-scale predictands that simpler statistical methods miss. They also carry a characteristic risk: a downscaling model trained on the historical relationship between scales may not preserve that relationship under a changed climate, so that a method that performs well in validation on the past can mislead when applied to a future that lies outside its training experience. The assessment is that machine-learning emulation and downscaling are valuable and increasingly standard tools, but that their out-of-distribution behaviour is their principal limitation for decision use and must be explicitly evaluated (*high confidence*).

Uncertainty quantification in machine-learning products

Rigorous uncertainty quantification is the single most important determinant of whether machine-learning-derived climate information is fit for decisions, and the field has developed a substantial toolkit for it, but the reliable estimation of predictive uncertainty in deep networks remains technically demanding and is frequently done poorly in practice (*high confidence*). A deep network that produces a confident point prediction with no honest estimate of its own uncertainty is dangerous for decisions, because it invites users to act on predictions that the model itself has no basis to support. Surveys of uncertainty in deep neural networks catalogue the sources of uncertainty — arising from the data, the model, and the distribution shift between training and deployment — and the methods available to estimate each, while emphasizing that no single method is adequate across all settings (Gawlikowski et al., 2023).

For decision use, the critical distinction is between uncertainty that the model can represent because it resembles something seen in training, and uncertainty that arises precisely because the deployment situation is unlike anything in training. Climate applications are especially exposed to the second kind, because the future climate is, by definition, partly outside the historical training distribution. The assessment is that machine-learning climate information should be accompanied by uncertainty estimates that explicitly account for distribution shift, that point predictions without such estimates should be treated with suspicion for consequential decisions, and that the honest communication of machine-learning uncertainty is continuous with the broader problem of uncertainty communication assessed in Section 10.5 (*high confidence*).

The environmental and equity costs of the computation

The computation underlying machine-learning climate information carries its own environmental and equity costs, and a full accounting of the technology's contribution to climate information must weigh these costs against its benefits (*medium confidence*). The training and deployment of large models consume energy and resources, and the concentration of the required computational capacity in a small number of well-resourced institutions and regions raises equity concerns that parallel those of the observing system. The concept of sustainable artificial intelligence distinguishes between using artificial intelligence for sustainability goals and ensuring the sustainability of the artificial intelligence itself, and argues that both must be considered together (Wynsberghe, 2021).

For climate information specifically, this framing has two implications. First, the environmental cost of a machine-learning product should be proportionate to its decision value; deploying large models for tasks that simpler methods handle adequately is difficult to justify. Second, the concentration of machine-learning capacity risks reproducing, in the domain of climate information, the same geography of information poverty documented for observations in Section 10.3, so that the benefits of machine learning accrue disproportionately

to already-advantaged users. The assessment is that the responsible use of machine learning in climate information requires attention to these costs, and that the technology's net contribution to equitable decision-relevant information is not guaranteed by its technical capabilities alone (*medium confidence*).

Machine learning in impact and adaptation decisions

Beyond prediction and emulation, machine learning increasingly informs impact assessment and adaptation decisions directly, particularly in data-rich sectors such as agriculture and the built environment, though its reliability there depends on the quality and representativeness of the training data (*high confidence*). In agriculture, machine-learning methods are applied across yield prediction, disease and pest detection, and resource-use optimization, and comprehensive reviews document both the breadth of these applications and the persistent challenges of data quality, transferability, and interpretability (Benos et al., 2021). In the built environment, machine-vision and deep-learning methods support monitoring, inspection, and management of infrastructure exposed to climate hazards (Baduge et al., 2022).

These applications share a dependence on training data that is representative of the conditions under which decisions will be made. Where training data are abundant and representative, machine learning can deliver decision-relevant accuracy; where they are scarce, biased, or drawn from conditions unlike the decision context, the resulting products can be confidently wrong. This dependence reproduces, at the level of individual applications, the equity concern that runs through this chapter: machine learning performs best where data are richest, which is often not where vulnerability is highest. The assessment is that machine learning is a valuable and expanding contributor to impact and adaptation decisions, and that its value in any specific application is bounded by the representativeness of the data on which it was trained (*high confidence*).

10.6 Communicating Uncertainty and Calibrated Confidence

The honest, well-calibrated communication of uncertainty is a necessary condition for climate information to support good decisions, and the calibrated-language framework used throughout this Report is a deliberate response to the recognized failure modes of uncertainty communication, though it is one instrument among several and is not sufficient on its own (*high confidence*). Uncertainty is not a defect of climate information to be minimized or hidden; it is a property to be characterized and conveyed, because a decision-maker's optimal action frequently depends more on the shape and framing of the uncertainty than on the central estimate. This section assesses how uncertainty in climate information is, and should be, communicated to those who act on it.

Why uncertainty communication is a decision problem

How uncertainty is framed and presented systematically changes the decisions that users make, so that the communication of uncertainty is itself a design problem with decision consequences, not a neutral act of reporting (*high confidence*). The same underlying uncertainty can be expressed as a probability, a range, a set of scenarios, a qualitative confidence statement, or a visual distribution, and these expressions are not interchangeable in their effect on decisions. Framing an outcome in terms of the probability of exceeding a threshold, rather than the expected value, can change whether a protective action is taken; presenting a range without indicating where within it the estimate is most likely can lead users to treat the extremes as equally probable.

The decision-relevance of framing is heightened for the quantities that most matter in climate risk: thresholds, tails, and rates of change. Many consequential decisions are governed by the probability that a critical value is exceeded — a flood level, a temperature threshold for human health or crop viability, a design standard for infrastructure. For these decisions, information about the tail of the distribution is more valuable than information about its centre, and uncertainty communication that focuses on central tendencies can systematically under-serve the decision. The assessment is that effective uncertainty communication must be designed around the structure of the decision, and in particular around the thresholds and tails that govern it (*high confidence*).

The calibrated-language framework and its limits

The calibrated-language framework — expressing findings with standardized confidence and likelihood terms — provides a consistent, transparent vocabulary for conveying the state of knowledge, and its consistent use across this Report is a strength, but it also has limits that must be understood by those who rely on it (*high confidence*). The framework's principal virtue is discipline: it forces authors to distinguish their confidence in a finding from the probability of an outcome, to be explicit about both, and to apply a common standard across a heterogeneous body of evidence. For decision-makers, this consistency is valuable because it allows confidence statements to be compared across findings and chapters.

The framework also has recognized limits. Calibrated terms compress a rich underlying assessment into a small number of categories, and users who do not understand the categories can misinterpret them — treating “likely” as a precise probability, or reading “medium confidence” as a weak or dismissible finding. The terms convey the assessors' judgement but not always the full basis for it, and they can obscure the difference between uncertainty that could be reduced with more research and uncertainty that is irreducible. The assessment is that calibrated language is a necessary and effective instrument for honest communication, but that it must be accompanied by accessible explanation of what the terms mean and do not mean, and complemented by other forms of uncertainty representation where decisions require them (*high confidence*).

Deep uncertainty and decision-making under it

For many consequential climate decisions the relevant uncertainty is deep — the probabilities themselves are contested or unknown — and in these cases the appropriate response is to shift from optimizing against a predicted future to seeking decisions that are robust across many plausible futures (*high confidence*). Where uncertainty can be well characterized as a probability distribution, decisions can in principle be optimized against it. But for many climate decisions — those involving unprecedented conditions, contested model structures, or futures that depend on human choices — the probabilities are not reliably known, and treating them as if they were can produce decisions that are fragile to the specific assumptions made.

Robust decision-making responds to deep uncertainty by asking which choices perform acceptably across the full range of plausible futures, rather than which choice is optimal for a single predicted future. This reframing changes the information that is decision-relevant: it privileges the identification of the conditions under which a decision would fail over the estimation of a most-likely outcome, and it favours flexible, adaptive strategies that can be revised as uncertainty resolves. The assessment is that deep uncertainty is the normal condition for many long-horizon climate decisions, and that climate information which supports robustness analysis — by mapping the range of plausible futures and the vulnerabilities of candidate decisions — is often more valuable than information which offers a single sharpened prediction (*high confidence*).

Visual and narrative communication of uncertainty

Visual and narrative representations of uncertainty can convey decision-relevant information that numerical or categorical statements cannot, but poorly designed representations can also mislead, so that the design of uncertainty visualizations is a technical discipline with decision consequences (*medium confidence*). Visualizations that show a full distribution, an ensemble of trajectories, or the spatial pattern of confidence can communicate the shape of uncertainty directly, allowing users to see tails, multi-modality, and spatial structure that a summary statistic would hide. Narrative approaches, including the use of contrasting scenarios or storylines, can make abstract uncertainty concrete by tracing plausible pathways through it.

Both approaches carry risks. Visualizations can imply false precision, can be read selectively, and can be dominated by their most salient feature regardless of its decision-relevance; narratives can lend undue vividness to particular pathways and be mistaken for predictions. The assessment is that visual and narrative uncertainty communication are powerful and increasingly necessary complements to numerical and calibrated-language approaches, and that their design should be tested against how users actually interpret them rather than assumed to be self-evidently clear (*medium confidence*).

10.7 Regional Climate Risk Assessment and the Interface to Working Group II

Regional climate risk assessment is the principal interface through which the physical-science findings of this Working Group enter the impact, adaptation, and vulnerability assessments of Working Group II, and the integrity of that interface depends on shared definitions of hazard, exposure, and vulnerability, on consistent regional framing, and on the traceable propagation of uncertainty across the boundary (*high confidence*). The physical-science community characterizes the climate hazard; the impacts community combines that hazard with exposure and vulnerability to assess risk. Where the two communities use compatible definitions, regions, baselines, and uncertainty conventions, information flows across the interface without distortion; where they do not, the hazard information can be misapplied, and the resulting risk assessment can be neither traced back to its physical basis nor defended.

The risk framing and the division of labour

The now-standard framing of climate risk as the interaction of hazard, exposure, and vulnerability provides a shared structure that allows the physical-science and impacts communities to divide labour without losing coherence, and its consistent application is a precondition for a well-functioning interface (*high confidence*). Under this framing, Working Group I information principally characterizes the hazard — the occurrence and intensity of physically defined climatic events and conditions — while exposure and vulnerability are assessed primarily by Working Group II. The value of the framing is that it makes explicit which community is responsible for which component and how the components combine, so that a risk statement can be decomposed and each part traced to its source.

The framing also clarifies where the difficulties lie. Hazard characterization, though demanding, is the component for which the physical-science basis is strongest; exposure and vulnerability are more contingent, more dynamic, and more dependent on social and economic factors that lie outside this Working Group's remit. As Section 10.3 documented, exposure is itself changing rapidly as land use changes (Winkler et al., 2021; Potapov et al., 2021), so that even the components conventionally assigned to Working Group II depend on physical-science monitoring. The assessment is that the hazard-exposure-vulnerability framing is an effective organizing structure for the interface, and that its effectiveness depends on both communities maintaining compatible definitions and on explicit attention to the components, such as dynamic exposure, that straddle the boundary (*high confidence*).

Regionalization and the problem of scale

The regionalization of climate information — its expression at the spatial scales at which impacts occur and decisions are made — is essential for the

interface to Working Group II but is also where physical-science confidence is often lowest, creating a structural tension between the scale decisions require and the scale at which the science is most reliable (*high confidence*). Impacts and decisions are local; the physical-science basis is most robust at large scales. Bridging this gap requires downscaling, regional modelling, and the careful communication of the accompanying loss of confidence, and it is precisely at the regional and local scales most relevant to Working Group II that the uncertainty in physical-science information is largest.

Regional signals of change illustrate both the maturity and the limits of regionalized information. Some regional patterns are robust and strongly constrained — the Arctic, for example, has warmed markedly faster than the global average, a regionally specific finding with high physical-science confidence and direct implications for regional risk assessment (Rantanen et al., 2022). Other regional changes, particularly in precipitation and in the frequency of specific compound events, remain less well constrained. The assessment is that the regional interface to Working Group II is best served by explicit differentiation between robust regional signals and uncertain ones, and by resisting the pressure to provide sharp local information where the physical-science basis does not support it (*high confidence*).

Sectoral hazard information and its translation to risk

Across the sectors most important for impact assessment — water, agriculture and food security, human health, and land — hazard characterization has matured to the point where physically robust hazard information is available, but the translation of that hazard information into locally actionable risk information remains the binding constraint (*high confidence*). In water, the assessment of global water resources and the buffering role of groundwater provides a strong hazard and resource basis (Scanlon et al., 2023), while the drivers of water-quality degradation add a hazard dimension that risk assessment must incorporate (Akhtar et al., 2021). In agriculture and food security, coupled climate-crop modelling characterizes the hazard to production (Jägermeyr et al., 2021), shifting pest and pathogen pressure adds a biological hazard dimension (Singh et al., 2023), and the interaction of climate with soil degradation and salinization compounds the hazard on the land base itself (Hassani et al., 2021; Malhi et al., 2021).

In human health, the hazard posed by non-optimal ambient temperatures is now quantified with substantial confidence, including the burden of mortality attributable to heat and cold (Zhao et al., 2021), and successive assessments of health and climate change document the widening reach of climate hazards into human well-being (Romanello et al., 2021; Romanello et al., 2022; Romanello et al., 2023). Yet in each of these sectors the same pattern holds: the hazard is characterized more confidently than the risk, because converting hazard into risk requires exposure and vulnerability information that is local, dynamic, and often unavailable. The assessment is that the sectoral hazard base for the Working Group II interface is strong and improving, and that the persistent

difficulty lies in the exposure and vulnerability information needed to complete the risk picture at decision-relevant scales (*high confidence*).

Compound and cascading risks across the interface

Compound and cascading risks — arising from the co-occurrence or sequential interaction of multiple hazards, or the propagation of impacts through interconnected systems — are among the most decision-relevant and least tractable objects at the Working Group interface, and their assessment requires physical-science and impacts information to be integrated rather than exchanged sequentially (*medium confidence*). Many of the most consequential risks are not the product of a single hazard acting on a single exposed system but of combinations: heat and drought together stressing agriculture and water simultaneously; a sequence of events that exhausts adaptive capacity; or an impact in one system that propagates into others through economic or infrastructural linkages.

These risks strain the conventional division of labour, because the interactions that generate them cross the hazard-exposure-vulnerability boundary. Assessing them requires the physical-science characterization of joint and sequential hazards to be developed together with the impacts community's understanding of how systems respond and interconnect. The assessment is that compound and cascading risks are increasingly recognized as central to climate risk, that the information needed to assess them is at the frontier of both Working Groups' capabilities, and that a sequential, one-directional handoff of hazard information is inadequate for them — they require the same co-productive integration between scientific communities that Section 10.2 identified as necessary between scientists and decision-makers (*medium confidence*).

Transition risk and the decision environment

The interface to decisions extends beyond physical risk to transition risk — the consequences of the societal response to climate change itself — and climate information increasingly must inform decisions in a policy and market environment that is being reshaped by mitigation, so that the decision-relevant information set includes the trajectory of the response as well as the physical hazard (*medium confidence*). Financial and economic decisions in particular are exposed not only to physical climate hazards but to the risk arising from policy, technological, and market shifts associated with decarbonization, and the pricing of this transition risk is an active and consequential domain (Bolton et al., 2023).

Information about the response trajectory is therefore decision-relevant in its own right. The interpretation of net-zero goals shapes the policy environment against which transition risk is assessed (Fankhauser et al., 2021); the pace and cost of the energy transition and of renewable deployment under a changing climate shape the opportunity set (Osman et al., 2022; Chen et al., 2022); and the deployment of nature-based approaches, where well-founded,

offers response options with co-benefits that decisions must weigh (Seddon et al., 2021). The assessment is that decision-relevant climate information now spans both physical risk and transition risk, and that the two are increasingly assessed together because many decisions are jointly exposed to them (*medium confidence*).

10.8 Equity, Capacity, and Knowledge Systems in the Production of Climate Information

Equity, institutional capacity, and the inclusion of diverse knowledge systems are not adjuncts to the production of decision-relevant climate information but central determinants of whether that information is trusted, used, and fairly distributed, and they recur as cross-cutting constraints throughout this chapter (*medium confidence*). The preceding sections identified, at successive points in the value chain, a consistent pattern: the observing system is weakest where vulnerability is highest (Section 10.3); co-production is least available where capacity is lowest (Section 10.2); and machine learning performs best where data are richest, which is often not where need is greatest (Section 10.4). This section assesses these cross-cutting equity and knowledge-system dimensions directly.

The geography of climate-information capacity

The capacity to produce, tailor, and deliver decision-relevant climate information is distributed very unevenly, and this uneven distribution tends to reinforce rather than offset the underlying distribution of vulnerability, so that improving equitable access to climate information is itself an adaptation priority (*high confidence*). Producing usable climate information requires not only data but institutions: national services, trained personnel, boundary organizations, and stable relationships with users. These institutional assets are concentrated in wealthier countries and regions, and their absence elsewhere means that even freely available global data cannot be converted into decision-relevant local services.

The consequence is a compounding disadvantage. Regions with high exposure and high vulnerability frequently also have the weakest information-production capacity, so that the populations most in need of decision-relevant climate information are the least well served by it. The assessment is that investment in climate-information capacity in under-served regions yields high returns for adaptation, that such investment must address institutions and personnel and not only data and technology, and that without it the benefits of advances in observation, computation, and machine learning will accrue disproportionately to those already best served (*high confidence*).

Indigenous knowledge and local communities as knowledge partners

Indigenous peoples and local communities hold knowledge that is decision-relevant for climate adaptation and environmental management, and information systems that recognize them as knowledge-partners rather than as passive recipients tend to produce outcomes that are both more effective and more equitable (*medium confidence*). Evidence from conservation and land governance indicates that recognizing the role, rights, and knowledge of Indigenous peoples and local communities is associated with more effective and more equitable outcomes (Dawson et al., 2021). This finding has direct implications for the design of climate information: legitimate, usable climate services in many settings depend on the inclusion of local and Indigenous knowledge about environmental conditions, seasonal patterns, and locally appropriate responses.

The inclusion of these knowledge systems is not merely instrumental. As Section 10.2 argued, legitimacy is a property of process, and information produced through processes that exclude affected communities or their knowledge is less likely to be trusted or used by them. The assessment is that the integration of Indigenous and local knowledge into climate-information systems, on terms that respect the rights and agency of the knowledge-holders, improves both the effectiveness and the legitimacy of those systems, and that its neglect is a recurrent cause of the failure of externally produced information to be adopted locally (*medium confidence*).

The distributional consequences of climate information

Climate information is not distributionally neutral: who receives it, in what form, and through which institutions shapes who is able to act, so that the design and delivery of climate information have equity consequences that must be assessed alongside its scientific quality (*medium confidence*). The same information delivered through channels that reach some users and not others, or in forms accessible to some and not others, will advantage the connected and the well-resourced. Because the capacity to act on climate information is itself unevenly distributed, even equitably delivered information can produce unequal outcomes if the ability to respond is unequal.

This distributional dimension connects climate information to the broader equity concerns of climate response. The health burden of climate hazards falls unequally (Romanello et al., 2023; Zhao et al., 2021), and information systems that fail to reach the most exposed populations compound that inequality. The assessment is that the equity consequences of climate information are a legitimate and necessary object of assessment, that they cannot be inferred from scientific quality alone, and that the design of climate services should be evaluated in part by whether it narrows or widens the gap between those who can and cannot act on the changing climate (*medium confidence*).

Synthesis: toward an integrated science of climate information

The evidence assessed in this chapter supports treating climate information for decisions as an integrated socio-technical system spanning observation, computation, communication, and institutions, whose performance is determined jointly by scientific quality and by the human and institutional processes that convert science into action (*high confidence*). No single component of this system is sufficient. Excellent observations without the capacity to convert them into services, powerful machine-learning methods without honest uncertainty quantification, or calibrated science without legitimate delivery, each fails to produce decision-relevant information on its own.

The recurring lesson across the chapter is that the binding constraints on decision-relevant climate information have shifted from the availability of data toward the fidelity of translation and the equity of access. Advances in reanalysis, Earth observation, carbon-cycle monitoring, and machine learning have greatly enriched the information base — carbon-cycle monitoring (Friedlingstein et al., 2022), globally consistent land-surface products (Zhang et al., 2021; Poggio et al., 2021), skilful data-driven prediction (Ravuri et al., 2021), and machine learning across the value chain (Kaack et al., 2022) — while the assessment of adaptation, mitigation, and their interactions has matured (Abbass et al., 2022). What most limits the value of this base for decisions is now the quality of co-production, the honesty of uncertainty communication, the integrity of the interface to impact assessment, and the equity of access. The assessment of this chapter is that the science of climate information is maturing into an integrated discipline in its own right, and that its further advance depends as much on the human and institutional dimensions of information as on the physical science it conveys (*high confidence*).

10.9 Operationalizing Climate Information: Standards, Evaluation, and Governance

Turning climate science into a durable service requires more than good products; it requires standards, sustained institutions, evaluation practices, and governance arrangements that make the information reliable, comparable, and accountable over time, and the maturation of these operational foundations is uneven across the components of the climate-information enterprise (*high confidence*). The preceding sections assessed the ingredients of decision-relevant climate information — co-production, observation, computation, communication, and the risk interface. This section assesses the operational scaffolding that allows those ingredients to be combined into services that persist, that can be trusted across users and settings, and that can be held to account.

Standards and interoperability

Common standards for data formats, metadata, provenance, and uncertainty description are what allow climate information produced by one institution

to be used reliably by another, and the value of the growing information base is substantially limited where such standards are absent or inconsistently applied (*high confidence*). A gridded product is only as usable as its documentation: without clear metadata describing the variable, the reference period, the grid, the processing chain, and the uncertainty, even an excellent dataset cannot be safely combined with others or applied to a decision. Interoperability — the capacity of independently produced datasets to be combined without hidden inconsistency — is therefore a precondition for the assembly of decision-relevant information from heterogeneous sources.

The importance of standards is visible in the products this chapter has assessed. Globally consistent land-surface products succeed as decision inputs precisely because they apply a uniform classification and processing chain across space and time (Zhang et al., 2021), and soil-property products become decision-usable because they attach quantified, per-pixel uncertainty to every estimate (Poggio et al., 2021). The reconciliation that makes carbon accounting authoritative is itself a standardization exercise, imposing a common framework on heterogeneous emissions and sink estimates (Friedlingstein et al., 2022). The assessment is that standards and interoperability are not administrative overhead but a determinant of whether the information base can be converted into services, and that investment in them yields returns across every downstream use (*high confidence*).

Evaluation of climate information as a service

The evaluation of climate information must extend beyond the scientific verification of accuracy to the assessment of usability, uptake, and decision outcomes, because a product can be accurate and still fail to inform decisions, and current evaluation practice remains weighted toward the former (*high confidence*). Scientific evaluation asks whether a product is correct; service evaluation asks whether it is used, whether it is understood, and whether it improves the decisions it is meant to inform. These are different questions with different methods, and a climate service can score well on the first while failing entirely on the second.

The nowcasting evaluation assessed in Section 10.4 illustrates good practice: skill was assessed not only by statistical scores but by the judgement of the professional forecasters who would use the product, so that usefulness was measured against the decision it informs (Ravuri et al., 2021). This model — evaluating information against the decision and the user rather than against a purely physical benchmark — is applicable well beyond nowcasting, and its wider adoption is one of the clearer routes to improving the value of climate information. Insights from implementation science reinforce the point: sustained use depends on characteristics of the users, the settings, and the implementation process, none of which is captured by accuracy metrics alone (Damschroder et al., 2022). The assessment is that service-oriented evaluation, incorporating usability and decision outcomes, is essential and currently under-developed, and that its strengthening is a priority for the field (*high confidence*).

Sustained institutions and the problem of continuity

Climate services depend on the continuity of the institutions that produce them, and the intermittent funding, staffing, and mandate of many climate-information providers is a structural threat to the reliability on which decisions depend (*high confidence*). A defining feature of a service, as distinct from a study, is persistence: users build the service into their decision processes on the expectation that it will continue. When a service is discontinued, or when its methods change without warning, the decisions that depended on it are disrupted, and the trust that took years to build is lost quickly.

Continuity is threatened at several points. Observing systems require sustained funding to maintain the records whose value grows with their length; the geography of information poverty documented in Section 10.3 is in part a geography of discontinuous observation. Boundary organizations that specialize in translation between science and decisions are frequently funded on short cycles that are mismatched to the long relationships co-production requires. The assessment is that the value of climate information is inseparable from the continuity of its production, that continuity is an institutional and financial achievement rather than a scientific one, and that its neglect undermines even scientifically excellent services (*high confidence*).

Governance, accountability, and the integrity of climate information

As climate information increasingly informs consequential public and private decisions — including the allocation of capital, the design of infrastructure, and the setting of policy — the governance of that information, including its provenance, its independence, and its accountability, becomes a matter of public interest (*medium confidence*). When climate information carries financial and regulatory weight, it acquires stakes that create incentives to shade it. The pricing of transition risk depends on climate information (Bolton et al., 2023), and the interpretation of net-zero commitments depends on agreed accounting of a monitored quantity (Fankhauser et al., 2021); in both cases the integrity of the underlying information is exposed to the interests it affects.

Governance arrangements that protect the independence and traceability of climate information — clear provenance, transparent methods, and separation between the production of information and the interests that stand to gain or lose from it — are therefore increasingly important as the stakes rise. The assessment is that the governance of climate information is an emerging and under-developed dimension of the field, that it will grow in importance as climate information is more deeply embedded in financial and regulatory systems, and that the credibility of climate services will depend in part on demonstrable independence and accountability (*medium confidence*).

10.10 Cross-Cutting Themes and Deepening of the Decision Interface

Several themes cut across the components assessed in this chapter — thresholds and non-linearity, the treatment of the baseline, the integration of mitigation and adaptation information, and the role of early warning — and treating them explicitly clarifies why decision-relevant climate information differs in structure from climate science as conventionally assessed (*high confidence*). This section consolidates these cross-cutting themes and extends the assessment of the decision interface with attention to the operational services in which they are most consequential.

Thresholds, non-linearity, and the tails of the distribution

Many of the most consequential climate decisions are governed by the crossing of thresholds and by non-linear responses, so that decision-relevant climate information must characterize the tails and the rates of change of climate variables rather than their central tendencies alone (*high confidence*). A crop may tolerate a range of temperatures and then fail sharply beyond a critical value; a health outcome may respond mildly to moderate heat and then rise steeply at extremes; an infrastructure design may be safe up to a return-period event and vulnerable beyond it. In each case the decision-relevant information concerns the probability and frequency of exceeding a critical value, and information about the mean is at best indirectly useful.

The evidence assessed in this chapter repeatedly displays this threshold structure. The mortality burden of ambient temperature is concentrated at the extremes of the temperature distribution rather than distributed across its centre (Zhao et al., 2021), and agricultural failure is governed by the crossing of physiological limits that coupled climate-crop models are designed to capture (Jägermeyr et al., 2021). The implication for climate information is methodological: products optimized to represent central tendencies may misrepresent precisely the tail behaviour that decisions depend on, and decision-relevant products must be designed and evaluated with the tails in view. The assessment is that threshold and tail information is systematically more decision-relevant than central-tendency information for a large class of climate decisions, and that this reorientation is one of the principal ways in which the decision-first framing changes practice (*high confidence*).

The baseline problem

Because most decisions are framed against expectations calibrated on the past, the choice of baseline or reference period against which change is expressed is a consequential and sometimes contested element of climate information, and an inappropriate baseline can distort a decision even when the underlying science is sound (*high confidence*). A decision-maker asking whether conditions have changed is implicitly comparing the present or pro-

jected climate to a remembered or recorded normal, and the choice of that normal — its period, its length, its representativeness — shapes the answer. A baseline drawn from an unusually benign or unusually severe period will bias the perceived change; a baseline too short to average out variability will confound noise with trend.

The baseline problem interacts with the rapid change documented elsewhere in this Report. When the climate is changing quickly, any fixed historical baseline becomes progressively less representative of the conditions a decision will actually face, so that decision-relevant information may need to be expressed against a moving or forward-looking reference rather than a fixed historical one. This is particularly acute for exposure, which is itself changing as land use changes (Winkler et al., 2021; Potapov et al., 2021), so that both the climate baseline and the exposure baseline are moving simultaneously. The assessment is that the explicit, decision-appropriate treatment of the baseline is a necessary element of decision-relevant climate information, and that unexamined reliance on conventional historical baselines is a common and correctable source of distortion (*high confidence*).

Integrating mitigation and adaptation information

Decisions increasingly require climate information that spans mitigation and adaptation together, because many choices affect both, and information systems that treat the two separately can obscure the trade-offs and co-benefits that decisions must weigh (*medium confidence*). A land-use decision may simultaneously affect carbon storage and local adaptive capacity; an energy decision may affect both emissions and resilience to climate hazards. Where climate information is organized to address mitigation and adaptation in isolation, the joint consequences of such decisions are hard to see, and decisions may optimize one dimension at the expense of the other.

The assessed literature increasingly addresses this integration. Nature-based approaches, where soundly designed, can deliver adaptation benefits alongside mitigation and other co-benefits, though the evidence cautions against treating them as universally beneficial (Seddon et al., 2021). The monitoring of land-based carbon storage, including the large and policy-relevant stocks of mineral-associated soil organic carbon (Georgiou et al., 2022), provides information that is simultaneously relevant to mitigation accounting and to the management of the land systems on which adaptation depends. The pace and cost of the energy transition under a changing climate couple mitigation choices to the physical hazards that adaptation must address (Osman et al., 2022; Chen et al., 2022). The assessment is that integrated mitigation-adaptation information is increasingly necessary for decisions that affect both, and that its provision requires information systems designed to represent the couplings rather than to separate the two agendas (*medium confidence*).

Early warning as the integrative climate service

Early warning systems are among the most mature and demonstrably valuable climate services, integrating observation, prediction, communication, and institutional response into a single decision-oriented chain, and they exemplify the socio-technical integration that this chapter has argued is necessary for climate information to save lives and reduce harm (*high confidence*).

An effective early warning system is not a forecast but a chain: it observes a developing hazard, predicts its evolution, communicates the risk to those exposed in a form they can act on, and connects to institutions capable of responding. A failure at any link — an unobserved onset, a forecast not delivered in time, a warning not understood, an absent response capacity — renders the whole ineffective, regardless of the quality of the others.

Early warning therefore concentrates, in a single service, every theme of this chapter. It depends on the observing foundation of Section 10.3, on the predictive methods of Section 10.4, on the uncertainty communication of Section 10.5, and on the institutional and equity considerations of Sections 10.2 and 10.7. Its effectiveness is greatest where the last link — reaching and enabling the most exposed populations — is strongest, and weakest where the geography of information poverty leaves those populations unobserved and unreached. The health evidence underscores the stakes, as the burden of climate hazards falls most heavily on the exposed and vulnerable (Romanello et al., 2023; Zhao et al., 2021). The assessment is that early warning systems are a proven, high-value integrative climate service, that their value is realized only when every link in the chain is maintained, and that extending their reach to currently under-served populations is among the highest-return investments available in the domain of climate information (*high confidence*).

The health framing as a route to actionable information

Framing climate information in terms of human health has proven a powerful route to actionability, because health outcomes are proximate, salient, and shared across constituencies, and successive health-and-climate assessments demonstrate the decision-relevance of translating physical hazards into health terms (*high confidence*). Abstract statements about temperature or emissions are difficult for many decision-makers to act on; statements about heat-related mortality, the spread of climate-sensitive disease, or the health co-benefits of mitigation connect directly to widely shared priorities. The health framing therefore functions as a translation device, converting physical-climate information into a currency that motivates and directs decisions.

The assessed health-and-climate literature exemplifies this translation, tracking a broad set of indicators that link climate change to health outcomes and to the response measures that affect them (Romanello et al., 2021; Romanello et al., 2022; Romanello et al., 2023). The quantification of the mortality burden of non-optimal temperatures gives the framing a rigorous empirical anchor (Zhao et al., 2021). Crucially, the health framing also foregrounds equity, because

the health burden of climate hazards falls most heavily on the exposed and the vulnerable, so that information organized around health naturally surfaces distributional questions that other framings can obscure. The assessment is that the health framing is an effective and evidence-rich route to actionable climate information, and that its capacity to connect physical hazards to shared human priorities makes it a valuable organizing device for decision-relevant information (*high confidence*).

Data ethics and the governance of machine-learning services

As machine-learning methods become embedded in operational climate services, the ethical governance of those methods — their transparency, their treatment of uncertainty, their environmental cost, and their distributional consequences — becomes an operational requirement rather than an abstract concern (*medium confidence*). A machine-learning climate service that produces confident predictions without honest uncertainty, that consumes disproportionate resources for marginal decision value, or that performs well only where data are abundant, embeds ethical choices in its ordinary operation. Governing these choices is part of operating the service responsibly.

The assessed literature provides the conceptual tools. The characterization of uncertainty in deep networks establishes what honest uncertainty quantification requires and where it is difficult (Gawlikowski et al., 2023); the concept of sustainable artificial intelligence frames the obligation to weigh the environmental and equity costs of the computation against its benefits (Wynsberghe, 2021); and reviews of machine learning across climate applications map where the methods are mature enough for operational trust and where they are not (Kaack et al., 2022). The assessment is that the responsible operation of machine-learning climate services requires explicit governance of uncertainty, resource use, and distributional impact, and that this governance is continuous with the broader governance of climate information assessed in Section 10.8 rather than separate from it (*medium confidence*).

Concluding assessment

Taken together, the evidence assessed in this chapter establishes that decision-relevant climate information is a distinct object of study and practice, governed by the joint properties of information and decision context, and that its advancement now depends as much on co-production, communication, institutions, and equity as on the physical science it conveys (*high confidence*). The information base has been transformed by advances in observation, re-analysis, Earth observation, carbon-cycle monitoring, and machine learning, and the assessment of climate change impacts, adaptation, and mitigation has matured alongside it (Abbass et al., 2022; Friedlingstein et al., 2022; Kaack et al., 2022). The binding constraints on the value of this base for decisions have shifted decisively toward the human and institutional processes that translate it into action, and toward the equity of access that determines who benefits.

The chapter's overarching assessment is therefore twofold. First, the scientific and technical foundations of climate information are strong and improving, and continued advances in observation and computation will further enrich the base (*high confidence*). Second, the realization of that base as decision-relevant information depends on co-production, honest uncertainty communication, a well-governed interface to impact assessment, sustained institutions, and equitable access, and it is in these dimensions that the greatest current gaps and the greatest opportunities lie (*high confidence*). A mature science of climate information for decisions and services is emerging at the intersection of the physical and the institutional, and its further progress will be measured not by the volume of information produced but by the quality of the decisions it enables. The transition from a data-limited to a translation-limited and equity-limited regime is, in the assessment of this chapter, the defining shift of the period covered by this Report, and it reframes the central task of the field: not to produce more climate information, but to ensure that the information already within reach is co-produced with those who will use it, communicated with calibrated honesty about what remains uncertain, connected through a traceable interface to the assessment of impacts and adaptation, sustained by durable and accountable institutions, and made accessible to the exposed and vulnerable populations whose decisions it must ultimately serve (*high confidence*).

10.11 References

- Abbass, K., et al. (2022). A review of the global climate change impacts, adaptation, and sustainable mitigation measures. *Environmental Science and Pollution Research*. <https://doi.org/10.1007/s11356-022-19718-6>
- Akhtar, N., et al. (2021). Various Natural and Anthropogenic Factors Responsible for Water Quality Degradation: A Review. *Water*. <https://doi.org/10.3390/w13192660>
- Baduge, S. K., et al. (2022). Artificial intelligence and smart vision for building and construction 4.0: Machine and deep learning methods and applications. *Automation in Construction*. <https://doi.org/10.1016/j.autcon.2022.104440>
- Benos, L., et al. (2021). Machine Learning in Agriculture: A Comprehensive Updated Review. *Sensors*. <https://doi.org/10.3390/s21113758>
- Bolton, P., et al. (2023). Global Pricing of Carbon-Transition Risk. *The Journal of Finance*. <https://doi.org/10.1111/jofi.13272>
- Damschroder, L. J., et al. (2022). The updated Consolidated Framework for Implementation Research based on user feedback. *Implementation Science*. <https://doi.org/10.1186/s13012-022-01245-0>
- Dawson, N., et al. (2021). The role of Indigenous peoples and local communities in effective and equitable conservation. *Ecology and Society*. <https://doi.org/10.5751/es-12625-260319>
- Fankhauser, S., et al. (2021). The meaning of net zero and how to get it right. *Nature Climate Change*. <https://doi.org/10.1038/s41558-021-01245-w>
- Friedlingstein, P., et al. (2022). Global Carbon Budget 2022. *Earth system science data*. <https://doi.org/10.5194/essd-14-4811-2022>

- Gawlikowski, J., et al. (2023). A survey of uncertainty in deep neural networks. *Artificial Intelligence Review*. <https://doi.org/10.1007/s10462-023-10562-9>
- Georgiou, K., et al. (2022). Global stocks and capacity of mineral-associated soil organic carbon. *Nature Communications*. <https://doi.org/10.1038/s41467-022-31540-9>
- Hassani, A., et al. (2021). Global predictions of primary soil salinization under changing climate in the 21st century. *Nature Communications*. <https://doi.org/10.1038/s41467-021-26907-3>
- Jägermeyr, J., et al. (2021). Climate impacts on global agriculture emerge earlier in new generation of climate and crop models. *Nature Food*. <https://doi.org/10.1038/s43011-021-00400-y>
- Kaack, L. H., et al. (2022). Tackling Climate Change with Machine Learning. OPUS 4 (Zuse Institute Berlin). <https://doi.org/10.1145/3485128>
- Osman, A. I., et al. (2022). Cost, environmental impact, and resilience of renewable energy under a changing climate: a review. *Environmental Chemistry Letters*. <https://doi.org/10.1007/s10311-022-01532-8>
- Poggio, L., et al. (2021). SoilGrids 2.0: producing soil information for the globe with quantified spatial uncertainty. *SOIL*. <https://doi.org/10.5194/soil-7-217-2021>
- Potapov, P., et al. (2021). Global maps of cropland extent and change show accelerated cropland expansion in the twenty-first century. *Nature Food*. <https://doi.org/10.1038/s43011-021-00429-z>
- Rantanen, M., et al. (2022). The Arctic has warmed nearly four times faster than the globe since 1979. *Communications Earth & Environment*. <https://doi.org/10.1038/s43247-022-00498-3>
- Ravuri, S., et al. (2021). Skilful precipitation nowcasting using deep generative models of radar. *Nature*. <https://doi.org/10.1038/s41586-021-03854-z>
- Romanello, M., et al. (2021). The 2021 report of the Lancet Countdown on health and climate change: code red for a healthy future. *The Lancet*. [https://doi.org/10.1016/s0140-6736\(21\)01787-6](https://doi.org/10.1016/s0140-6736(21)01787-6)
- Romanello, M., et al. (2022). The 2022 report of the Lancet Countdown on health and climate change: health at the mercy of fossil fuels. *The Lancet*. [https://doi.org/10.1016/s0140-6736\(22\)01540-9](https://doi.org/10.1016/s0140-6736(22)01540-9)
- Romanello, M., et al. (2023). The 2023 report of the Lancet Countdown on health and climate change: the imperative for a health-centred response in a world facing irreversible harms. *The Lancet*. [https://doi.org/10.1016/s0140-6736\(23\)01859-7](https://doi.org/10.1016/s0140-6736(23)01859-7)
- Scanlon, B. R., et al. (2023). Global water resources and the role of groundwater in a resilient water future. *Nature Reviews Earth & Environment*. <https://doi.org/10.1038/s43247-022-00378-6>
- Seddon, N., et al. (2021). Getting the message right on nature-based solutions to climate change. *Global Change Biology*. <https://doi.org/10.1111/gcb.15513>
- Singh, B. K., et al. (2023). Climate change impacts on plant pathogens, food security and paths forward. *Nature Reviews Microbiology*. <https://doi.org/10.1038/s41579-023-00900-7>

Winkler, K., et al. (2021). Global land use changes are four times greater than previously estimated. *Nature Communications*. <https://doi.org/10.1038/s41467-021-22702-2>

Wynsberghe, A. v. (2021). Sustainable AI: AI for sustainability and the sustainability of AI. *AI and Ethics*. <https://doi.org/10.1007/s43681-021-00043-6>

Zhang, X., et al. (2021). GLC_FCS30: global land-cover product with fine classification system at 30 m using time-series Landsat imagery. *Earth system science data*. <https://doi.org/10.5194/essd-13-2753-2021>

Zhao, Q., et al. (2021). Global, regional, and national burden of mortality associated with non-optimal ambient temperatures from 2000 to 2019: a three-stage modelling study. *The Lancet Planetary Health*. [https://doi.org/10.1016/s2542-5196\(21\)00081-4](https://doi.org/10.1016/s2542-5196(21)00081-4)

Zhao, X., et al. (2024). A review of convolutional neural networks in computer vision. *Artificial Intelligence Review*. <https://doi.org/10.1007/s10462-024-10721-6>

Production Notes

This edition was produced by a three-stage, source-first pipeline. Every step below is factual and reproducible; the complete tooling and the intermediate data are published in the project’s public repository. Where a number appears, it is the number the process actually produced, not a target.

Stage 1 — Harvest a verified source database

The evidentiary substrate was built before any drafting began. A stdlib-only harvester queried the OpenAlex works API (free, no key required) using a curated set of search phrases for each of the ten chapters. The query filter constrained results to the post-AR6 window — publication dates **2021 through 2026** — to journal articles, and, critically, to works with a resolvable DOI (`has_doi: true`). Results were ranked by citation count.

Only works with a resolvable DOI were retained, so that every entry in the database corresponds to a real, locatable publication. The harvest, refreshed on 2026-07-25, yielded **1,506 unique verified works**, tagged across the ten chapters as **2,275 chapter-level source entries** (a work relevant to more than one chapter is counted in each). Per chapter, the verified pool ranged from 183 to 288 works. Each work was stored with its DOI, title, publication year, author list, venue, citation count, and reconstructed abstract, in a SQLite database (`sources.db`) plus one JSON file per chapter.

Stage 2 — Grounded writing, constrained to the database

Each chapter was drafted by a separate instance of Anthropic’s **Claude Fable 5**, run through the Claude Max subscription as an agentic subagent — no paid API calls were used for any generation in this edition. The ten chapters were drafted in parallel.

Each drafting agent was given the **top eighty works by citation count** from its chapter’s verified pool, formatted as a numbered “allowed sources” list. The system prompt was strict and identical across chapters: write in the calibrated assessment voice of a Working Group I report; open each numbered subsection with a bold headline statement ending in italicized calibrated-uncertainty language; and **cite only from the numbered allowed list**. If a claim could not be supported by an allowed source, the agent was required to state it without

a citation or to omit it — never to recall, invent, or import a reference from outside the list.

Chapters were written **append-only** in successive passes of roughly 3,500 words, to a target of 12,000–15,000 words each, and closed with a References section listing only the allowed sources actually cited. The Summary for Policymakers and Technical Summary in this edition were synthesized afterward from the finished grounded chapters, and reference those chapters by cross-reference only; they introduce no new citations.

Stage 3 — Machine verification as a build gate

After drafting, an independent verification program (`verify_citations.py`, local, no network) extracted every in-text author-year citation from every chapter and matched it, on first-author surname and year, against `sources.db`; it separately checked every DOI in each References section against the database. Any citation or DOI without a match was flagged, and the presence of a single unverifiable citation failed the build.

The final full-volume verification across all ten chapters returned **zero unverifiable in-text citations and zero unknown DOIs**. The completed volume comprises **128,683 words** across the ten chapters (each 12,600–13,200 words, inclusive of its reference list), citing between **27 and 57 distinct verified sources** per chapter.

Known limitations

We state these plainly, because the value of a transparency experiment lies in its honest boundaries.

- **Source-pool contamination.** Ranking an automated keyword harvest by raw citation count allows highly-cited works from unrelated fields (for example, wireless-communications, battery-materials, and general biomedical review papers) to enter a chapter’s top-eighty list. Measured across the volume, the topical share of each chapter’s top-eighty pool ranged from **29 of 80 to 68 of 80**. In the more contaminated chapters, the agent correctly cited only the genuinely relevant works and wrote the remaining material in an uncited assessment voice, which reduced citation density. A topic-filtered re-harvest is the identified fix and is tracked for a future edition.
- **Verifier constraints.** The verification tool attributes every year inside a multi-citation parenthesis to the first-listed surname, so grouped citations were split into single-citation parentheticals to pass the gate — a cosmetic effect on some sentences. Separately, a handful of allowed sources whose author surnames contain non-ASCII characters (Unicode hyphens, certain diacritics) could not be matched by the verifier and were dropped as citations, with the underlying claims retained uncited.

- **Word counts include references.** The per-chapter figures above count each chapter's reference list.
- **No expert human review.** This remains a machine-generated assessment. It has not been reviewed by climate scientists and must not be treated as authoritative.

Reproducibility

The harvester, the writer prompt and its allowed-source formatting, the verifier, the per-chapter source JSONs, and all ten grounded chapter drafts are published in the public repository accompanying this edition, so that the entire build can be inspected and re-run.